
Linux Trace Documentation

The kernel development community

Jan 15, 2023

CONTENTS

1	Function Tracer Design	1
1.1	Introduction	1
1.2	Prerequisites	1
1.3	HAVE_FUNCTION_TRACER	1
1.4	HAVE_FUNCTION_GRAPH_TRACER	3
1.5	HAVE_FUNCTION_GRAPH_FP_TEST	4
1.6	HAVE_FUNCTION_GRAPH_RET_ADDR_PTR	5
1.7	HAVE_SYSCALL_TRACEPOINTS	5
1.8	HAVE_FTRACE_MCOUNT_RECORD	5
1.9	HAVE_DYNAMIC_FTRACE	5
1.10	HAVE_DYNAMIC_FTRACE + HAVE_FUNCTION_GRAPH_TRACER	8
2	Notes on Analysing Behaviour Using Events and Tracepoints	9
2.1	1. Introduction	9
2.2	2. Listing Available Events	9
2.3	3. Enabling Events	10
2.4	4. Event Filtering	11
2.5	5. Analysing Event Variances with PCL	11
2.6	6. Higher-Level Analysis with Helper Scripts	12
2.7	7. Lower-Level Analysis with PCL	13
3	ftrace - Function Tracer	17
3.1	Introduction	17
3.2	Implementation Details	18
3.3	The File System	18
3.4	The Tracers	28
3.5	Error conditions	29
3.6	Examples of using the tracer	30
3.7	Output format:	30
3.8	Latency trace format	31
3.9	trace_options	33
3.10	irqsoff	38
3.11	preemptoff	42
3.12	preemptirqsoff	44
3.13	wakeup	48
3.14	wakeup_rt	49
3.15	Latency tracing and events	52
3.16	Hardware Latency Detector	53
3.17	function	55

3.18	Single thread tracing	56
3.19	function graph tracer	59
3.20	dynamic ftrace	63
3.21	Selecting function filters via index	68
3.22	Dynamic ftrace with the function graph tracer	68
3.23	ftrace_enabled	69
3.24	Filter commands	70
3.25	trace_pipe	72
3.26	trace entries	73
3.27	Snapshot	74
3.28	Instances	75
3.29	Stack trace	78
3.30	More	79
4	Using ftrace to hook to functions	81
4.1	Introduction	81
4.2	The ftrace context	81
4.3	The ftrace_ops structure	81
4.4	The callback function	82
4.5	Protect your callback	83
4.6	The ftrace FLAGS	83
4.7	Filtering which functions to trace	84
5	Fprobe - Function entry/exit probe	87
5.1	Introduction	87
5.2	The usage of fprobe	87
5.3	The fprobe entry/exit handler	88
5.4	Share the callbacks with kprobes	89
5.5	The missed counter	89
5.6	Functions and structures	89
6	Kernel Probes (Kprobes)	93
6.1	Concepts: Kprobes and Return Probes	93
6.2	Architectures Supported	98
6.3	Configuring Kprobes	98
6.4	API Reference	99
6.5	Kprobes Features and Limitations	102
6.6	Probe Overhead	103
6.7	TODO	104
6.8	Kprobes Example	104
6.9	Kretprobes Example	105
6.10	Deprecated Features	105
6.11	The kprobes debugfs interface	105
6.12	The kprobes sysctl interface	106
6.13	References	106
7	Kprobe-based Event Tracing	107
7.1	Overview	107
7.2	Synopsis of kprobe_events	107
7.3	Types	108
7.4	User Memory Access	109
7.5	Per-Probe Event Filtering	109

7.6	Event Profiling	109
7.7	Kernel Boot Parameter	109
7.8	Usage examples	110
8	Uprobe-tracer: Uprobe-based Event Tracing	113
8.1	Overview	113
8.2	Synopsis of uprobe_tracer	113
8.3	Types	114
8.4	Event Profiling	114
8.5	Usage examples	114
9	Using the Linux Kernel Tracepoints	117
9.1	Purpose of tracepoints	117
9.2	Usage	117
10	Event Tracing	121
10.1	1. Introduction	121
10.2	2. Using Event Tracing	121
10.3	3. Defining an event-enabled tracepoint	123
10.4	4. Event formats	123
10.5	5. Event filtering	124
10.6	6. Event triggers	127
10.7	7. In-kernel trace event API	130
11	Subsystem Trace Points: kmem	141
11.1	1. Slab allocation of small objects of unknown type	141
11.2	2. Slab allocation of small objects of known type	141
11.3	3. Page allocation	142
11.4	4. Per-CPU Allocator Activity	142
11.5	5. External Fragmentation	143
12	Subsystem Trace Points: power	145
12.1	1. Power state switch events	145
12.2	2. Clocks events	146
12.3	3. Power domains events	146
12.4	4. PM QoS events	146
13	NMI Trace Events	147
13.1	nmi_handler	147
14	MSR Trace Events	149
15	In-kernel memory-mapped I/O tracing	151
15.1	Preparation	151
15.2	Usage Quick Reference	151
15.3	Usage	152
15.4	How Mmiotrace Works	153
15.5	Trace Log Format	153
15.6	Explanation Keyword Space-separated arguments	153
15.7	Tools for Developers	154
16	Event Histograms	155
16.1	1. Introduction	155

16.2	2. Histogram Trigger Command	155
17	Histogram Design Notes	217
17.1	'hist_debug' trace event files	217
17.2	Basic histograms	217
17.3	Variables	223
17.4	Actions and Handlers	232
17.5	A couple special cases	247
18	Boot-time tracing	257
18.1	Overview	257
18.2	Options in the Boot Config	257
18.3	When to Start	260
18.4	Examples	260
19	Hardware Latency Detector	263
19.1	Introduction	263
19.2	Usage	263
20	OSNOISE Tracer	265
20.1	Usage	266
20.2	Tracer options	267
20.3	Additional Tracing	267
21	Timerlat tracer	269
21.1	Usage	269
21.2	Tracer options	270
21.3	timerlat and osnoise	270
21.4	IRQ stacktrace	272
22	Intel(R) Trace Hub (TH)	273
22.1	Overview	273
22.2	Bus and Subdevices	274
22.3	Quick example	274
22.4	Host Debugger Mode	275
22.5	Software Sinks	275
23	Lockless Ring Buffer Design	277
23.1	Terminology used in this Document	277
23.2	The Generic Ring Buffer	278
23.3	Making the Ring Buffer Lockless:	284
23.4	Nested writes	289
24	System Trace Module	295
24.1	stm_source	296
24.2	stm_console	297
24.3	stm_fttrace	297
25	MIPI SyS-T over STP	299
26	CoreSight - ARM Hardware Trace	301
26.1	Coresight - HW Assisted Tracing on ARM	301
26.2	CoreSight System Configuration Manager	313

26.3	Coresight CPU Debug Module	318
26.4	CoreSight Embedded Cross Trigger (CTI & CTM).	321
26.5	ETMv4 sysfs linux driver programming reference.	325
26.6	Trace Buffer Extension (TRBE).	337
27	user_events: User-based Event Tracing	339
27.1	Overview	339
27.2	Registering	339
27.3	Deleting	341
27.4	Status	341
27.5	Writing Data	342
27.6	Example Code	343
Index		345

FUNCTION TRACER DESIGN

Author Mike Frysinger

Caution: This document is out of date. Some of the description below doesn't match current implementation now.

1.1 Introduction

Here we will cover the architecture pieces that the common function tracing code relies on for proper functioning. Things are broken down into increasing complexity so that you can start simple and at least get basic functionality.

Note that this focuses on architecture implementation details only. If you want more explanation of a feature in terms of common code, review the common *ftrace - Function Tracer* file.

Ideally, everyone who wishes to retain performance while supporting tracing in their kernel should make it all the way to dynamic ftrace support.

1.2 Prerequisites

Ftrace relies on these features being implemented:

- STACKTRACE_SUPPORT - implement `save_stack_trace()`
- TRACE_IRQFLAGS_SUPPORT - implement `include/asm/irqflags.h`

1.3 HAVE_FUNCTION_TRACER

You will need to implement the `mcount` and the `ftrace_stub` functions.

The exact `mcount` symbol name will depend on your toolchain. Some call it “`mcount`”, “`_mcount`”, or even “`__mcount`”. You can probably figure it out by running something like:

```
$ echo 'main(){}' | gcc -x c -S -o - -pg | grep mcount
      call     mcount
```

We'll make the assumption below that the symbol is "mcount" just to keep things nice and simple in the examples.

Keep in mind that the ABI that is in effect inside of the mcount function is *highly* architecture/toolchain specific. We cannot help you in this regard, sorry. Dig up some old documentation and/or find someone more familiar than you to bang ideas off of. Typically, register usage (argument/scratch/etc...) is a major issue at this point, especially in relation to the location of the mcount call (before/after function prologue). You might also want to look at how glibc has implemented the mcount function for your architecture. It might be (semi-)relevant.

The mcount function should check the function pointer ftrace_trace_function to see if it is set to ftrace_stub. If it is, there is nothing for you to do, so return immediately. If it isn't, then call that function in the same way the mcount function normally calls __mcount_internal - the first argument is the "frompc" while the second argument is the "selfpc" (adjusted to remove the size of the mcount call that is embedded in the function).

For example, if the function foo() calls bar(), when the bar() function calls mcount(), the arguments mcount() will pass to the tracer are:

- "frompc" - the address bar() will use to return to foo()
- "selfpc" - the address bar() (with mcount() size adjustment)

Also keep in mind that this mcount function will be called *a lot*, so optimizing for the default case of no tracer will help the smooth running of your system when tracing is disabled. So the start of the mcount function is typically the bare minimum with checking things before returning. That also means the code flow should usually be kept linear (i.e. no branching in the nop case). This is of course an optimization and not a hard requirement.

Here is some pseudo code that should help (these functions should actually be implemented in assembly):

```
void ftrace_stub(void)
{
    return;
}

void mcount(void)
{
    /* save any bare state needed in order to do initial checking */

    extern void (*ftrace_trace_function)(unsigned long, unsigned long);
    if (ftrace_trace_function != ftrace_stub)
        goto do_trace;

    /* restore any bare state */

    return;
do_trace:

    /* save all state needed by the ABI (see paragraph above) */

    unsigned long frompc = ...;
    unsigned long selfpc = <return address> - MCOUNT_INSN_SIZE;
```

```

    ftrace_trace_function(frompc, selfpc);

    /* restore all state needed by the ABI */
}

```

Don't forget to export mcount for modules !

```

extern void mcount(void);
EXPORT_SYMBOL(mcount);

```

1.4 HAVE_FUNCTION_GRAPH_TRACER

Deep breath ... time to do some real work. Here you will need to update the mcount function to check ftrace graph function pointers, as well as implement some functions to save (hijack) and restore the return address.

The mcount function should check the function pointers `ftrace_graph_return` (compare to `ftrace_stub`) and `ftrace_graph_entry` (compare to `ftrace_graph_entry_stub`). If either of those is not set to the relevant stub function, call the arch-specific function `ftrace_graph_caller` which in turn calls the arch-specific function `prepare_ftrace_return`. Neither of these function names is strictly required, but you should use them anyway to stay consistent across the architecture ports - easier to compare & contrast things.

The arguments to `prepare_ftrace_return` are slightly different than what are passed to `ftrace_trace_function`. The second argument "selfpc" is the same, but the first argument should be a pointer to the "frompc". Typically this is located on the stack. This allows the function to hijack the return address temporarily to have it point to the arch-specific function `return_to_handler`. That function will simply call the common `ftrace_return_to_handler` function and that will return the original return address with which you can return to the original call site.

Here is the updated mcount pseudo code:

```

void mcount(void)
{
    ...
        if (ftrace_trace_function != ftrace_stub)
            goto do_trace;

#ifdef CONFIG_FUNCTION_GRAPH_TRACER
+   extern void (*ftrace_graph_return)(...);
+   extern void (*ftrace_graph_entry)(...);
+   if (ftrace_graph_return != ftrace_stub ||
+       ftrace_graph_entry != ftrace_graph_entry_stub)
+       ftrace_graph_caller();
#endif

    /* restore any bare state */
    ...

```

Here is the pseudo code for the new `ftrace_graph_caller` assembly function:

```
#ifdef CONFIG_FUNCTION_GRAPH_TRACER
void ftrace_graph_caller(void)
{
    /* save all state needed by the ABI */

    unsigned long *frompc = &...;
    unsigned long selfpc = <return address> - MCOUNT_INSN_SIZE;
    /* passing frame pointer up is optional -- see below */
    prepare_ftrace_return(frompc, selfpc, frame_pointer);

    /* restore all state needed by the ABI */
}
#endif
```

For information on how to implement `prepare_ftrace_return()`, simply look at the x86 version (the frame pointer passing is optional; see the next section for more information). The only architecture-specific piece in it is the setup of the fault recovery table (the `asm(...)` code). The rest should be the same across architectures.

Here is the pseudo code for the new `return_to_handler` assembly function. Note that the ABI that applies here is different from what applies to the `mcount` code. Since you are returning from a function (after the epilogue), you might be able to skimp on things saved/restored (usually just registers used to pass return values).

```
#ifdef CONFIG_FUNCTION_GRAPH_TRACER
void return_to_handler(void)
{
    /* save all state needed by the ABI (see paragraph above) */

    void (*original_return_point)(void) = ftrace_return_to_handler();

    /* restore all state needed by the ABI */

    /* this is usually either a return or a jump */
    original_return_point();
}
#endif
```

1.5 HAVE_FUNCTION_GRAPH_FP_TEST

An arch may pass in a unique value (frame pointer) to both the entering and exiting of a function. On exit, the value is compared and if it does not match, then it will panic the kernel. This is largely a sanity check for bad code generation with gcc. If gcc for your port sanely updates the frame pointer under different optimization levels, then ignore this option.

However, adding support for it isn't terribly difficult. In your assembly code that calls `prepare_ftrace_return()`, pass the frame pointer as the 3rd argument. Then in the C version of that function, do what the x86 port does and pass it along to `ftrace_push_return_trace()` instead of a stub value of 0.

Similarly, when you call `ftrace_return_to_handler()`, pass it the frame pointer.

1.6 HAVE_FUNCTION_GRAPH_RET_ADDR_PTR

An arch may pass in a pointer to the return address on the stack. This prevents potential stack unwinding issues where the unwinder gets out of sync with `ret_stack` and the wrong addresses are reported by `ftrace_graph_ret_addr()`.

Adding support for it is easy: just define the macro in `asm/ftrace.h` and pass the return address pointer as the `'retp'` argument to `ftrace_push_return_trace()`.

1.7 HAVE_SYSCALL_TRACEPOINTS

You need very few things to get the syscalls tracing in an arch.

- Support `HAVE_ARCH_TRACEHOOK` (see `arch/Kconfig`).
- Have a `NR_syscalls` variable in `<asm/unistd.h>` that provides the number of syscalls supported by the arch.
- Support the `TIF_SYSCALL_TRACEPOINT` thread flags.
- Put the `trace_sys_enter()` and `trace_sys_exit()` tracepoints calls from `ptrace` in the `ptrace` syscalls tracing path.
- If the system call table on this arch is more complicated than a simple array of addresses of the system calls, implement an `arch_syscall_addr` to return the address of a given system call.
- If the symbol names of the system calls do not match the function names on this arch, define `ARCH_HAS_SYSCALL_MATCH_SYM_NAME` in `asm/ftrace.h` and implement `arch_syscall_match_sym_name` with the appropriate logic to return true if the function name corresponds with the symbol name.
- Tag this arch as `HAVE_SYSCALL_TRACEPOINTS`.

1.8 HAVE_FTRACE_MCOUNT_RECORD

See `scripts/recordmcount.pl` for more info. Just fill in the arch-specific details for how to locate the addresses of `mcount` call sites via `objdump`. This option doesn't make much sense without also implementing dynamic `ftrace`.

1.9 HAVE_DYNAMIC_FTRACE

You will first need `HAVE_FTRACE_MCOUNT_RECORD` and `HAVE_FUNCTION_TRACER`, so scroll your reader back up if you got over eager.

Once those are out of the way, you will need to implement:

- **asm/ftrace.h:**
 - `MCOUNT_ADDR`
 - `ftrace_call_adjust()`

- struct dyn_arch_ftrace{}
- **asm code:**
 - mcount() (new stub)
 - ftrace_caller()
 - ftrace_call()
 - ftrace_stub()
- **C code:**
 - ftrace_dyn_arch_init()
 - ftrace_make_nop()
 - ftrace_make_call()
 - ftrace_update_ftrace_func()

First you will need to fill out some arch details in your asm/ftrace.h.

Define MCOUNT_ADDR as the address of your mcount symbol similar to:

```
#define MCOUNT_ADDR ((unsigned long)mcount)
```

Since no one else will have a decl for that function, you will need to:

```
extern void mcount(void);
```

You will also need the helper function ftrace_call_adjust(). Most people will be able to stub it out like so:

```
static inline unsigned long ftrace_call_adjust(unsigned long addr)
{
    return addr;
}
```

<details to be filled>

Lastly you will need the custom dyn_arch_ftrace structure. If you need some extra state when runtime patching arbitrary call sites, this is the place. For now though, create an empty struct:

```
struct dyn_arch_ftrace {
    /* No extra data needed */
};
```

With the header out of the way, we can fill out the assembly code. While we did already create a mcount() function earlier, dynamic ftrace only wants a stub function. This is because the mcount() will only be used during boot and then all references to it will be patched out never to return. Instead, the guts of the old mcount() will be used to create a new ftrace_caller() function. Because the two are hard to merge, it will most likely be a lot easier to have two separate definitions split up by #ifdefs. Same goes for the ftrace_stub() as that will now be inlined in ftrace_caller().

Before we get confused anymore, let's check out some pseudo code so you can implement your own stuff in assembly:

```

void mcount(void)
{
    return;
}

void ftrace_caller(void)
{
    /* save all state needed by the ABI (see paragraph above) */

    unsigned long frompc = ...;
    unsigned long selfpc = <return address> - MCOUNT_INSN_SIZE;

ftrace_call:
    ftrace_stub(frompc, selfpc);

    /* restore all state needed by the ABI */

ftrace_stub:
    return;
}

```

This might look a little odd at first, but keep in mind that we will be runtime patching multiple things. First, only functions that we actually want to trace will be patched to call `ftrace_caller()`. Second, since we only have one tracer active at a time, we will patch the `ftrace_caller()` function itself to call the specific tracer in question. That is the point of the `ftrace_call` label.

With that in mind, let's move on to the C code that will actually be doing the runtime patching. You'll need a little knowledge of your arch's opcodes in order to make it through the next section.

Every arch has an `init` callback function. If you need to do something early on to initialize some state, this is the time to do that. Otherwise, this simple function below should be sufficient for most people:

```

int __init ftrace_dyn_arch_init(void)
{
    return 0;
}

```

There are two functions that are used to do runtime patching of arbitrary functions. The first is used to turn the `mcount` call site into a `nop` (which is what helps us retain runtime performance when not tracing). The second is used to turn the `mcount` call site into a call to an arbitrary location (but typically that is `ftracer_caller()`). See the general function definition in `linux/ftrace.h` for the functions:

```

ftrace_make_nop()
ftrace_make_call()

```

The `rec->ip` value is the address of the `mcount` call site that was collected by the `scripts/recordmcount.pl` during build time.

The last function is used to do runtime patching of the active tracer. This will be modifying the assembly code at the location of the `ftrace_call` symbol inside of the `ftrace_caller()` function. So you should have sufficient padding at that location to support the new function calls you'll

be inserting. Some people will be using a “call” type instruction while others will be using a “branch” type instruction. Specifically, the function is:

`ftrace_update_ftrace_func()`

1.10 HAVE_DYNAMIC_FTRACE + HAVE_FUNCTION_GRAPH_TRACER

The function grapher needs a few tweaks in order to work with dynamic ftrace. Basically, you will need to:

- **update:**
 - `ftrace_caller()`
 - `ftrace_graph_call()`
 - `ftrace_graph_caller()`
- **implement:**
 - `ftrace_enable_ftrace_graph_caller()`
 - `ftrace_disable_ftrace_graph_caller()`

<details to be filled>

Quick notes:

- add a nop stub after the `ftrace_call` location named `ftrace_graph_call`; stub needs to be large enough to support a call to `ftrace_graph_caller()`
- update `ftrace_graph_caller()` to work with being called by the new `ftrace_caller()` since some semantics may have changed
- `ftrace_enable_ftrace_graph_caller()` will runtime patch the `ftrace_graph_call` location with a call to `ftrace_graph_caller()`
- `ftrace_disable_ftrace_graph_caller()` will runtime patch the `ftrace_graph_call` location with nops

NOTES ON ANALYSING BEHAVIOUR USING EVENTS AND TRACEPOINTS

Author Mel Gorman (PCL information heavily based on email from Ingo Molnar)

2.1 1. Introduction

Tracepoints (see *Using the Linux Kernel Tracepoints*) can be used without creating custom kernel modules to register probe functions using the event tracing infrastructure.

Simplistically, tracepoints represent important events that can be taken in conjunction with other tracepoints to build a “Big Picture” of what is going on within the system. There are a large number of methods for gathering and interpreting these events. Lacking any current Best Practises, this document describes some of the methods that can be used.

This document assumes that debugfs is mounted on /sys/kernel/debug and that the appropriate tracing options have been configured into the kernel. It is assumed that the PCL tool tools/perf has been installed and is in your path.

2.2 2. Listing Available Events

2.2.1 2.1 Standard Utilities

All possible events are visible from /sys/kernel/debug/tracing/events. Simply calling:

```
$ find /sys/kernel/debug/tracing/events -type d
```

will give a fair indication of the number of events available.

2.2.2 2.2 PCL (Performance Counters for Linux)

Discovery and enumeration of all counters and events, including tracepoints, are available with the perf tool. Getting a list of available events is a simple case of:

```
$ perf list 2>&1 | grep Tracepoint
ext4:ext4_free_inode           [Tracepoint event]
ext4:ext4_request_inode       [Tracepoint event]
ext4:ext4_allocate_inode      [Tracepoint event]
ext4:ext4_write_begin         [Tracepoint event]
```

```
ext4:ext4_ordered_write_end          [Tracepoint event]
[ .... remaining output snipped .... ]
```

2.3 3. Enabling Events

2.3.1 3.1 System-Wide Event Enabling

See [Event Tracing](#) for a proper description on how events can be enabled system-wide. A short example of enabling all events related to page allocation would look something like:

```
$ for i in `find /sys/kernel/debug/tracing/events -name "enable" | grep mm_`;  
do echo 1 > $i; done
```

2.3.2 3.2 System-Wide Event Enabling with SystemTap

In SystemTap, tracepoints are accessible using the `kernel.trace()` function call. The following is an example that reports every 5 seconds what processes were allocating the pages.

```
global page_allocs

probe kernel.trace("mm_page_alloc") {
    page_allocs[execname()]++
}

function print_count() {
    printf ("% -25s %-s\n", "#Pages Allocated", "Process Name")
    foreach (proc in page_allocs-)
        printf ("% -25d %-s\n", page_allocs[proc], proc)
    printf ("\n")
    delete page_allocs
}

probe timer.s(5) {
    print_count()
}
```

2.3.3 3.3 System-Wide Event Enabling with PCL

By specifying the `-a` switch and analysing sleep, the system-wide events for a duration of time can be examined.

```
$ perf stat -a \  
-e kmem:mm_page_alloc -e kmem:mm_page_free \  
-e kmem:mm_page_free_batched \  
sleep 10  
Performance counter stats for 'sleep 10':
```

```

9630  kmem:mm_page_alloc
2143  kmem:mm_page_free
7424  kmem:mm_page_free_batched

10.002577764  seconds time elapsed

```

Similarly, one could execute a shell and exit it as desired to get a report at that point.

2.3.4 3.4 Local Event Enabling

ftrace - Function Tracer describes how to enable events on a per-thread basis using `set_ftrace_pid`.

2.3.5 3.5 Local Event Enablement with PCL

Events can be activated and tracked for the duration of a process on a local basis using PCL such as follows.

```

$ perf stat -e kmem:mm_page_alloc -e kmem:mm_page_free \
              -e kmem:mm_page_free_batched ./hackbench 10
Time: 0.909

Performance counter stats for './hackbench 10':

    17803  kmem:mm_page_alloc
    12398  kmem:mm_page_free
     4827  kmem:mm_page_free_batched

0.973913387  seconds time elapsed

```

2.4 4. Event Filtering

ftrace - Function Tracer covers in-depth how to filter events in `ftrace`. Obviously using `grep` and `awk` of `trace_pipe` is an option as well as any script reading `trace_pipe`.

2.5 5. Analysing Event Variances with PCL

Any workload can exhibit variances between runs and it can be important to know what the standard deviation is. By and large, this is left to the performance analyst to do it by hand. In the event that the discrete event occurrences are useful to the performance analyst, then `perf` can be used.

```

$ perf stat --repeat 5 -e kmem:mm_page_alloc -e kmem:mm_page_free
                        -e kmem:mm_page_free_batched ./hackbench 10
Time: 0.890
Time: 0.895

```

```
Time: 0.915
Time: 1.001
Time: 0.899
```

Performance counter stats for './hackbench 10' (5 runs):

```
16630 kmem:mm_page_alloc      ( +-  3.542% )
11486 kmem:mm_page_free       ( +-  4.771% )
 4730 kmem:mm_page_free_batched ( +-  2.325% )
```

```
0.982653002 seconds time elapsed ( +-  1.448% )
```

In the event that some higher-level event is required that depends on some aggregation of discrete events, then a script would need to be developed.

Using `-repeat`, it is also possible to view how events are fluctuating over time on a system-wide basis using `-a` and `sleep`.

```
$ perf stat -e kmem:mm_page_alloc -e kmem:mm_page_free \
            -e kmem:mm_page_free_batched \
            -a --repeat 10 \
            sleep 1
```

Performance counter stats for 'sleep 1' (10 runs):

```
1066 kmem:mm_page_alloc      ( +- 26.148% )
 182 kmem:mm_page_free       ( +-  5.464% )
 890 kmem:mm_page_free_batched ( +- 30.079% )
```

```
1.002251757 seconds time elapsed ( +-  0.005% )
```

2.6 6. Higher-Level Analysis with Helper Scripts

When events are enabled the events that are triggering can be read from `/sys/kernel/debug/tracing/trace_pipe` in human-readable format although binary options exist as well. By post-processing the output, further information can be gathered on-line as appropriate. Examples of post-processing might include

- Reading information from `/proc` for the PID that triggered the event
- Deriving a higher-level event from a series of lower-level events.
- Calculating latencies between two events

`Documentation/trace/postprocess/trace-pagealloc-postprocess.pl` is an example script that can read `trace_pipe` from STDIN or a copy of a trace. When used on-line, it can be interrupted once to generate a report without exiting and twice to exit.

Simplistically, the script just reads STDIN and counts up events but it also can do more such as

- Derive high-level events from many low-level events. If a number of pages are freed to the main allocator from the per-CPU lists, it recognises that as one per-CPU drain even though there is no specific tracepoint for that event

- It can aggregate based on PID or individual process number
- In the event memory is getting externally fragmented, it reports on whether the fragmentation event was severe or moderate.
- When receiving an event about a PID, it can record who the parent was so that if large numbers of events are coming from very short-lived processes, the parent process responsible for creating all the helpers can be identified

2.7 7. Lower-Level Analysis with PCL

There may also be a requirement to identify what functions within a program were generating events within the kernel. To begin this sort of analysis, the data must be recorded. At the time of writing, this required root:

```
$ perf record -c 1 \
    -e kmem:mm_page_alloc -e kmem:mm_page_free \
    -e kmem:mm_page_free_batched \
    ./hackbench 10
Time: 0.894
[ perf record: Captured and wrote 0.733 MB perf.data (~32010 samples) ]
```

Note the use of ‘-c 1’ to set the event period to sample. The default sample period is quite high to minimise overhead but the information collected can be very coarse as a result.

This record outputted a file called perf.data which can be analysed using perf report.

```
$ perf report
# Samples: 30922
#
# Overhead      Command          Shared Object
# .....
#
# 87.27% hackbench [vdso]
# 6.85% hackbench /lib/i686/cmov/libc-2.9.so
# 2.62% hackbench /lib/ld-2.9.so
# 1.52% perf [vdso]
# 1.22% hackbench ./hackbench
# 0.48% hackbench [kernel]
# 0.02% perf /lib/i686/cmov/libc-2.9.so
# 0.01% perf /usr/bin/perf
# 0.01% perf /lib/ld-2.9.so
# 0.00% hackbench /lib/i686/cmov/libpthread-2.9.so
#
# (For more details, try: perf report --sort comm,dso,symbol)
#
```

According to this, the vast majority of events triggered on events within the VDSO. With simple binaries, this will often be the case so let’s take a slightly different example. In the course of writing this, it was noticed that X was generating an insane amount of page allocations so let’s look at it:

```
$ perf record -c 1 -f \
    -e kmem:mm_page_alloc -e kmem:mm_page_free \
    -e kmem:mm_page_free_batched \
    -p `pidof X`
```

This was interrupted after a few seconds and

```
$ perf report
# Samples: 27666
#
# Overhead   Command          Shared Object
# .....
#
# 51.95%     Xorg             [vdso]
# 47.95%     Xorg             /opt/gfx-test/lib/libpixman-1.so.0.13.1
# 0.09%      Xorg             /lib/i686/cmov/libc-2.9.so
# 0.01%      Xorg             [kernel]
#
# (For more details, try: perf report --sort comm,dso,symbol)
#
```

So, almost half of the events are occurring in a library. To get an idea which symbol:

```
$ perf report --sort comm,dso,symbol
# Samples: 27666
#
# Overhead   Command          Shared Object   Symbol
# .....
#
# 51.95%     Xorg             [vdso]          [.]
→0x000000fffffe424
# 47.93%     Xorg             /opt/gfx-test/lib/libpixman-1.so.0.13.1 [.]
→pixmanFillse2
# 0.09%      Xorg             /lib/i686/cmov/libc-2.9.so          [.] _int_malloc
# 0.01%      Xorg             /opt/gfx-test/lib/libpixman-1.so.0.13.1 [.] pixman_
→region32_copy_f
# 0.01%      Xorg             [kernel]        [k] read_hpet
# 0.01%      Xorg             /opt/gfx-test/lib/libpixman-1.so.0.13.1 [.] get_fast_path
# 0.00%      Xorg             [kernel]        [k] ftrace_trace_
→userstack
```

To see where within the function `pixmanFillse2` things are going wrong:

```
$ perf annotate pixmanFillse2
[ ... ]
0.00 :      34eeb:      0f 18 08      prefetcht0 (%eax)
      :      }
      :
      :      extern __inline void __attribute__((__gnu_inline__, __always_
→inline__, __
      :      __mm_store_si128 (__m128i *__P, __m128i __B) :      {
```

	:	*__P = __B;		
12.40	:	34eee:	66 0f 7f 80 40 ff ff	movdqa %xmm0, -0xc0(%eax)
0.00	:	34ef5:	ff	
12.40	:	34ef6:	66 0f 7f 80 50 ff ff	movdqa %xmm0, -0xb0(%eax)
0.00	:	34efd:	ff	
12.39	:	34efe:	66 0f 7f 80 60 ff ff	movdqa %xmm0, -0xa0(%eax)
0.00	:	34f05:	ff	
12.67	:	34f06:	66 0f 7f 80 70 ff ff	movdqa %xmm0, -0x90(%eax)
0.00	:	34f0d:	ff	
12.58	:	34f0e:	66 0f 7f 40 80	movdqa %xmm0, -0x80(%eax)
12.31	:	34f13:	66 0f 7f 40 90	movdqa %xmm0, -0x70(%eax)
12.40	:	34f18:	66 0f 7f 40 a0	movdqa %xmm0, -0x60(%eax)
12.31	:	34f1d:	66 0f 7f 40 b0	movdqa %xmm0, -0x50(%eax)

At a glance, it looks like the time is being spent copying pixmaps to the card. Further investigation would be needed to determine why pixmaps are being copied around so much but a starting point would be to take an ancient build of libpixmap out of the library path where it was totally forgotten about from months ago!

FTRACE - FUNCTION TRACER

Copyright 2008 Red Hat Inc.

Author Steven Rostedt <srostedt@redhat.com>

License The GNU Free Documentation License, Version 1.2 (dual licensed under the GPL v2)

Original Reviewers Elias Oltmanns, Randy Dunlap, Andrew Morton, John Kacur, and David Teigland.

- Written for: 2.6.28-rc2
- Updated for: 3.10
- Updated for: 4.13 - Copyright 2017 VMware Inc. Steven Rostedt
- Converted to rst format - Changbin Du <changbin.du@intel.com>

3.1 Introduction

Ftrace is an internal tracer designed to help out developers and designers of systems to find what is going on inside the kernel. It can be used for debugging or analyzing latencies and performance issues that take place outside of user-space.

Although ftrace is typically considered the function tracer, it is really a framework of several assorted tracing utilities. There's latency tracing to examine what occurs between interrupts disabled and enabled, as well as for preemption and from a time a task is woken to the task is actually scheduled in.

One of the most common uses of ftrace is the event tracing. Throughout the kernel is hundreds of static event points that can be enabled via the tracefs file system to see what is going on in certain parts of the kernel.

See *Event Tracing* for more information.

3.2 Implementation Details

See *Function Tracer Design* for details for arch porters and such.

3.3 The File System

Ftrace uses the tracefs file system to hold the control files as well as the files to display output. When tracefs is configured into the kernel (which selecting any ftrace option will do) the directory `/sys/kernel/tracing` will be created. To mount this directory, you can add to your `/etc/fstab` file:

tracefs	/sys/kernel/tracing	tracefs defaults	0	0
---------	---------------------	------------------	---	---

Or you can mount it at run time with:

```
mount -t tracefs nodev /sys/kernel/tracing
```

For quicker access to that directory you may want to make a soft link to it:

```
ln -s /sys/kernel/tracing /tracing
```

Attention: Before 4.1, all ftrace tracing control files were within the debugfs file system, which is typically located at `/sys/kernel/debug/tracing`. For backward compatibility, when mounting the debugfs file system, the tracefs file system will be automatically mounted at:

`/sys/kernel/debug/tracing`

All files located in the tracefs file system will be located in that debugfs file system directory as well.

Attention: Any selected ftrace option will also create the tracefs file system. The rest of the document will assume that you are in the ftrace directory (`cd /sys/kernel/tracing`) and will only concentrate on the files within that directory and not distract from the content with the extended `"/sys/kernel/tracing"` path name.

That's it! (assuming that you have ftrace configured into your kernel)

After mounting tracefs you will have access to the control and output files of ftrace. Here is a list of some of the key files:

Note: all time values are in microseconds.

`current_tracer:`

This is used to set or display the current tracer that is configured. Changing the current tracer clears the ring buffer content as well as the "snapshot" buffer.

`available_tracers:`

This holds the different types of tracers that have been compiled into the kernel. The tracers listed here can be configured by echoing their name into `current_tracer`.

`tracing_on`:

This sets or displays whether writing to the trace ring buffer is enabled. Echo 0 into this file to disable the tracer or 1 to enable it. Note, this only disables writing to the ring buffer, the tracing overhead may still be occurring.

The kernel function `tracing_off()` can be used within the kernel to disable writing to the ring buffer, which will set this file to "0". User space can re-enable tracing by echoing "1" into the file.

Note, the function and event trigger "traceoff" will also set this file to zero and stop tracing. Which can also be re-enabled by user space using this file.

`trace`:

This file holds the output of the trace in a human readable format (described below). Opening this file for writing with the `O_TRUNC` flag clears the ring buffer content. Note, this file is not a consumer. If tracing is off (no tracer running, or `tracing_on` is zero), it will produce the same output each time it is read. When tracing is on, it may produce inconsistent results as it tries to read the entire buffer without consuming it.

`trace_pipe`:

The output is the same as the "trace" file but this file is meant to be streamed with live tracing. Reads from this file will block until new data is retrieved. Unlike the "trace" file, this file is a consumer. This means reading from this file causes sequential reads to display more current data. Once data is read from this file, it is consumed, and will not be read again with a sequential read. The "trace" file is static, and if the tracer is not adding more data, it will display the same information every time it is read.

`trace_options`:

This file lets the user control the amount of data that is displayed in one of the above output files. Options also exist to modify how a tracer or events work (stack traces, timestamps, etc).

`options`:

This is a directory that has a file for every available trace option (also in `trace_options`). Options may also be set or cleared by writing a "1" or "0" respectively into the corresponding file with the option name.

`tracing_max_latency`:

Some of the tracers record the max latency. For example, the maximum time that interrupts are disabled. The maximum time is saved in this file. The max trace will also be stored, and displayed by "trace".

A new max trace will only be recorded if the latency is greater than the value in this file (in microseconds).

By echoing in a time into this file, no latency will be recorded unless it is greater than the time in this file.

`tracing_thresh:`

Some latency tracers will record a trace whenever the latency is greater than the number in this file. Only active when the file contains a number greater than 0. (in microseconds)

`buffer_size_kb:`

This sets or displays the number of kilobytes each CPU buffer holds. By default, the trace buffers are the same size for each CPU. The displayed number is the size of the CPU buffer and not total size of all buffers. The trace buffers are allocated in pages (blocks of memory that the kernel uses for allocation, usually 4 KB in size). A few extra pages may be allocated to accommodate buffer management meta-data. If the last page allocated has room for more bytes than requested, the rest of the page will be used, making the actual allocation bigger than requested or shown. (Note, the size may not be a multiple of the page size due to buffer management meta-data.)

Buffer sizes for individual CPUs may vary (see “`per_cpu/cpu0/buffer_size_kb`” below), and if they do this file will show “X”.

`buffer_total_size_kb:`

This displays the total combined size of all the trace buffers.

`free_buffer:`

If a process is performing tracing, and the ring buffer should be shrunk “freed” when the process is finished, even if it were to be killed by a signal, this file can be used for that purpose. On close of this file, the ring buffer will be resized to its minimum size. Having a process that is tracing also open this file, when the process exits its file descriptor for this file will be closed, and in doing so, the ring buffer will be “freed”.

It may also stop tracing if `disable_on_free` option is set.

`tracing_cpumask:`

This is a mask that lets the user only trace on specified CPUs. The format is a hex string representing the CPUs.

`set_fttrace_filter:`

When dynamic ftrace is configured in (see the section below “dynamic ftrace”), the code is dynamically modified (code text rewrite) to disable calling of the function profiler (mcount). This lets tracing be configured in with practically no overhead in performance. This also has a side effect of enabling or disabling specific functions to be traced. Echoing names of functions into this file will limit

the trace to only those functions. This influences the tracers “function” and “function_graph” and thus also function profiling (see “function_profile_enabled”).

The functions listed in “available_filter_functions” are what can be written into this file.

This interface also allows for commands to be used. See the “Filter commands” section for more details.

As a speed up, since processing strings can be quite expensive and requires a check of all functions registered to tracing, instead an index can be written into this file. A number (starting with “1”) written will instead select the same corresponding at the line position of the “available_filter_functions” file.

`set_fttrace_notrace:`

This has an effect opposite to that of `set_fttrace_filter`. Any function that is added here will not be traced. If a function exists in both `set_fttrace_filter` and `set_fttrace_notrace`, the function will `_not_` be traced.

`set_fttrace_pid:`

Have the function tracer only trace the threads whose PID are listed in this file.

If the “function-fork” option is set, then when a task whose PID is listed in this file forks, the child’s PID will automatically be added to this file, and the child will be traced by the function tracer as well. This option will also cause PIDs of tasks that exit to be removed from the file.

`set_fttrace_notrace_pid:`

Have the function tracer ignore threads whose PID are listed in this file.

If the “function-fork” option is set, then when a task whose PID is listed in this file forks, the child’s PID will automatically be added to this file, and the child will not be traced by the function tracer as well. This option will also cause PIDs of tasks that exit to be removed from the file.

If a PID is in both this file and “`set_fttrace_pid`”, then this file takes precedence, and the thread will not be traced.

`set_event_pid:`

Have the events only trace a task with a PID listed in this file. Note, `sched_switch` and `sched_wake_up` will also trace events listed in this file.

To have the PIDs of children of tasks with their PID in this file added on fork, enable the “event-fork” option. That option will also cause the PIDs of tasks to be removed from this file when the task exits.

`set_event_notrace_pid:`

Have the events not trace a task with a PID listed in this file. Note, `sched_switch` and `sched_wakeup` will trace threads not listed in this file, even if a thread's PID is in the file if the `sched_switch` or `sched_wakeup` events also trace a thread that should be traced.

To have the PIDs of children of tasks with their PID in this file added on fork, enable the "event-fork" option. That option will also cause the PIDs of tasks to be removed from this file when the task exits.

`set_graph_function`:

Functions listed in this file will cause the function graph tracer to only trace these functions and the functions that they call. (See the section "dynamic ftrace" for more details). Note, `set_ftrace_filter` and `set_ftrace_notrace` still affects what functions are being traced.

`set_graph_notrace`:

Similar to `set_graph_function`, but will disable function graph tracing when the function is hit until it exits the function. This makes it possible to ignore tracing functions that are called by a specific function.

`available_filter_functions`:

This lists the functions that ftrace has processed and can trace. These are the function names that you can pass to "`set_ftrace_filter`", "`set_ftrace_notrace`", "`set_graph_function`", or "`set_graph_notrace`". (See the section "dynamic ftrace" below for more details.)

`dyn_ftrace_total_info`:

This file is for debugging purposes. The number of functions that have been converted to nops and are available to be traced.

`enabled_functions`:

This file is more for debugging ftrace, but can also be useful in seeing if any function has a callback attached to it. Not only does the trace infrastructure use ftrace function trace utility, but other subsystems might too. This file displays all functions that have a callback attached to them as well as the number of callbacks that have been attached. Note, a callback may also call multiple functions which will not be listed in this count.

If the callback registered to be traced by a function with the "save regs" attribute (thus even more overhead), a 'R' will be displayed on the same line as the function that is returning registers.

If the callback registered to be traced by a function with the "ip modify" attribute (thus the `regs->ip` can be changed), an 'I' will be displayed on the same line as the function that can be overridden.

If the architecture supports it, it will also show what callback is being directly called by the function. If the count is greater than 1 it most likely will be `ftrace_ops_list_func()`.

If the callback of a function jumps to a trampoline that is specific to the callback and which is not the standard trampoline, its address will be printed as well as the function that the trampoline calls.

`function_profile_enabled`:

When set it will enable all functions with either the function tracer, or if configured, the function graph tracer. It will keep a histogram of the number of functions that were called and if the function graph tracer was configured, it will also keep track of the time spent in those functions. The histogram content can be displayed in the files:

`trace_stat/function<cpu> (function0, function1, etc).`

`trace_stat`:

A directory that holds different tracing stats.

`kprobe_events`:

Enable dynamic trace points. See *Kprobe-based Event Tracing*.

`kprobe_profile`:

Dynamic trace points stats. See *Kprobe-based Event Tracing*.

`max_graph_depth`:

Used with the function graph tracer. This is the max depth it will trace into a function. Setting this to a value of one will show only the first kernel function that is called from user space.

`printk_formats`:

This is for tools that read the raw format files. If an event in the ring buffer references a string, only a pointer to the string is recorded into the buffer and not the string itself. This prevents tools from knowing what that string was. This file displays the string and address for the string allowing tools to map the pointers to what the strings were.

`saved_cmdlines`:

Only the pid of the task is recorded in a trace event unless the event specifically saves the task comm as well. Ftrace makes a cache of pid mappings to comms to try to display comms for events. If a pid for a comm is not listed, then "<...>" is displayed in the output.

If the option "record-cmd" is set to "0", then comms of tasks will not be saved during recording. By default, it is enabled.

`saved_cmdlines_size`:

By default, 128 comms are saved (see "saved_cmdlines" above). To increase or decrease the amount of comms that are cached, echo the number of comms to cache into this file.

`saved_tgids`:

If the option "record-tgid" is set, on each scheduling context switch the Task Group ID of a task is saved in a table mapping the PID of the thread to its TGID. By default, the "record-tgid" option is disabled.

`snapshot`:

This displays the “snapshot” buffer and also lets the user take a snapshot of the current running trace. See the “Snapshot” section below for more details.

`stack_max_size`:

When the stack tracer is activated, this will display the maximum stack size it has encountered. See the “Stack Trace” section below.

`stack_trace`:

This displays the stack back trace of the largest stack that was encountered when the stack tracer is activated. See the “Stack Trace” section below.

`stack_trace_filter`:

This is similar to “`set_ftrace_filter`” but it limits what functions the stack tracer will check.

`trace_clock`:

Whenever an event is recorded into the ring buffer, a “timestamp” is added. This stamp comes from a specified clock. By default, ftrace uses the “local” clock. This clock is very fast and strictly per cpu, but on some systems it may not be monotonic with respect to other CPUs. In other words, the local clocks may not be in sync with local clocks on other CPUs.

Usual clocks for tracing:

```
# cat trace_clock
[local] global counter x86-tsc
```

The clock with the square brackets around it is the one in effect.

local: Default clock, but may not be in sync across CPUs

global: This clock is in sync with all CPUs but may be a bit slower than the local clock.

counter: This is not a clock at all, but literally an atomic counter. It counts up one by one, but is in sync with all CPUs. This is useful when you need to know exactly the order events occurred with respect to each other on different CPUs.

uptime: This uses the jiffies counter and the time stamp is relative to the time since boot up.

perf: This makes ftrace use the same clock that perf uses. Eventually perf will be able to read ftrace buffers and this will help out in interleaving the data.

x86-tsc: Architectures may define their own clocks. For example, x86 uses its own TSC cycle clock here.

ppc-tb: This uses the powerpc timebase register value. This is in sync across CPUs and can also be used to correlate events across hypervisor/guest if `tb_offset` is known.

mono: This uses the fast monotonic clock (CLOCK_MONOTONIC) which is monotonic and is subject to NTP rate adjustments.

mono_raw: This is the raw monotonic clock (CLOCK_MONOTONIC_RAW) which is monotonic but is not subject to any rate adjustments and ticks at the same rate as the hardware clocksource.

boot: This is the boot clock (CLOCK_BOOTTIME) and is based on the fast monotonic clock, but also accounts for time spent in suspend. Since the clock access is designed for use in tracing in the suspend path, some side effects are possible if clock is accessed after the suspend time is accounted before the fast mono clock is updated. In this case, the clock update appears to happen slightly sooner than it normally would have. Also on 32-bit systems, it's possible that the 64-bit boot offset sees a partial update. These effects are rare and post processing should be able to handle them. See comments in the `ktime_get_boot_fast_ns()` function for more information.

tai: This is the tai clock (CLOCK_TAI) and is derived from the wall-clock time. However, this clock does not experience discontinuities and backwards jumps caused by NTP inserting leap seconds. Since the clock access is designed for use in tracing, side effects are possible. The clock access may yield wrong readouts in case the internal TAI offset is updated e.g., caused by setting the system time or using `adjtimex()` with an offset. These effects are rare and post processing should be able to handle them. See comments in the `ktime_get_tai_fast_ns()` function for more information.

To set a clock, simply echo the clock name into this file:

```
# echo global > trace_clock
```

Setting a clock clears the ring buffer content as well as the “snapshot” buffer.

`trace_marker:`

This is a very useful file for synchronizing user space with events happening in the kernel. Writing strings into this file will be written into the ftrace buffer.

It is useful in applications to open this file at the start of the application and just reference the file descriptor for the file:

```
void trace_write(const char *fmt, ...)
{
    va_list ap;
    char buf[256];
    int n;

    if (trace_fd < 0)
        return;
```

```
    va_start(ap, fmt);
    n = vsnprintf(buf, 256, fmt, ap);
    va_end(ap);

    write(trace_fd, buf, n);
}
```

start:

```
trace_fd = open("trace_marker", WR_ONLY);
```

Note: Writing into the trace_marker file can also initiate triggers that are written into `/sys/kernel/tracing/events/ftrace/print/trigger`. See “Event triggers” in [Event Tracing](#) and an example in [Event Histograms](#) (Section 3.)

trace_marker_raw:

This is similar to trace_marker above, but is meant for binary data to be written to it, where a tool can be used to parse the data from trace_pipe_raw.

uprobe_events:

Add dynamic tracepoints in programs. See [Uprobe-tracer: Uprobe-based Event Tracing](#)

uprobe_profile:

Uprobe statistics. See `uprobetrace.txt`

instances:

This is a way to make multiple trace buffers where different events can be recorded in different buffers. See “Instances” section below.

events:

This is the trace event directory. It holds event tracepoints (also known as static tracepoints) that have been compiled into the kernel. It shows what event tracepoints exist and how they are grouped by system. There are “enable” files at various levels that can enable the tracepoints when a “1” is written to them.

See [Event Tracing](#) for more information.

set_event:

By echoing in the event into this file, will enable that event.

See [Event Tracing](#) for more information.

available_events:

A list of events that can be enabled in tracing.

See [Event Tracing](#) for more information.

timestamp_mode:

Certain tracers may change the timestamp mode used when logging trace events into the event buffer. Events with different modes can coexist within a buffer but the mode in effect when an event is logged determines which timestamp mode is used for that event. The default timestamp mode is 'delta'.

Usual timestamp modes for tracing:

```
# cat timestamp_mode [delta] absolute
```

The timestamp mode with the square brackets around it is the one in effect.

delta: Default timestamp mode - timestamp is a delta against a per-buffer timestamp.

absolute: The timestamp is a full timestamp, not a delta against some other value. As such it takes up more space and is less efficient.

hwlat_detector:

Directory for the Hardware Latency Detector. See "Hardware Latency Detector" section below.

per_cpu:

This is a directory that contains the trace per_cpu information.

per_cpu/cpu0/buffer_size_kb:

The ftrace buffer is defined per_cpu. That is, there's a separate buffer for each CPU to allow writes to be done atomically, and free from cache bouncing. These buffers may have different size buffers. This file is similar to the buffer_size_kb file, but it only displays or sets the buffer size for the specific CPU. (here cpu0).

per_cpu/cpu0/trace:

This is similar to the "trace" file, but it will only display the data specific for the CPU. If written to, it only clears the specific CPU buffer.

per_cpu/cpu0/trace_pipe

This is similar to the "trace_pipe" file, and is a consuming read, but it will only display (and consume) the data specific for the CPU.

per_cpu/cpu0/trace_pipe_raw

For tools that can parse the ftrace ring buffer binary format, the trace_pipe_raw file can be used to extract the data from the ring buffer directly. With the use of the splice() system call, the buffer data can be quickly transferred to a file or to the network where a server is collecting the data.

Like trace_pipe, this is a consuming reader, where multiple reads will always produce different data.

per_cpu/cpu0/snapshot:

This is similar to the main “snapshot” file, but will only snapshot the current CPU (if supported). It only displays the content of the snapshot for a given CPU, and if written to, only clears this CPU buffer.

`per_cpu/cpu0/snapshot_raw`:

Similar to the `trace_pipe_raw`, but will read the binary format from the snapshot buffer for the given CPU.

`per_cpu/cpu0/stats`:

This displays certain stats about the ring buffer:

entries: The number of events that are still in the buffer.

overflow: The number of lost events due to overwriting when the buffer was full.

commit overflow: Should always be zero. This gets set if so many events happened within a nested event (ring buffer is re-entrant), that it fills the buffer and starts dropping events.

bytes: Bytes actually read (not overwritten).

oldest event ts: The oldest timestamp in the buffer

now ts: The current timestamp

dropped events: Events lost due to overwrite option being off.

read events: The number of events read.

3.4 The Tracers

Here is the list of current tracers that may be configured.

“function”

Function call tracer to trace all kernel functions.

“function_graph”

Similar to the function tracer except that the function tracer probes the functions on their entry whereas the function graph tracer traces on both entry and exit of the functions. It then provides the ability to draw a graph of function calls similar to C code source.

“blk”

The block tracer. The tracer used by the blktrace user application.

“hwlat”

The Hardware Latency tracer is used to detect if the hardware produces any latency. See “Hardware Latency Detector” section below.

“irqsoff”

Traces the areas that disable interrupts and saves the trace with the longest max latency. See `tracing_max_latency`. When a new max is recorded, it replaces the old trace. It is best to view this trace with the `latency-format` option enabled, which happens automatically when the tracer is selected.

`"preemptoff"`

Similar to `irqsoff` but traces and records the amount of time for which preemption is disabled.

`"preemptirqsoff"`

Similar to `irqsoff` and `preemptoff`, but traces and records the largest time for which irqs and/or preemption is disabled.

`"wakeup"`

Traces and records the max latency that it takes for the highest priority task to get scheduled after it has been woken up. Traces all tasks as an average developer would expect.

`"wakeup_rt"`

Traces and records the max latency that it takes for just RT tasks (as the current `"wakeup"` does). This is useful for those interested in wake up timings of RT tasks.

`"wakeup_dl"`

Traces and records the max latency that it takes for a `SCHED_DEADLINE` task to be woken (as the `"wakeup"` and `"wakeup_rt"` does).

`"mmiotrace"`

A special tracer that is used to trace binary module. It will trace all the calls that a module makes to the hardware. Everything it writes and reads from the I/O as well.

`"branch"`

This tracer can be configured when tracing likely/unlikely calls within the kernel. It will trace when a likely and unlikely branch is hit and if it was correct in its prediction of being correct.

`"nop"`

This is the `"trace nothing"` tracer. To remove all tracers from tracing simply echo `"nop"` into `current_tracer`.

3.5 Error conditions

For most `ftrace` commands, failure modes are obvious and communicated using standard return codes.

For other more involved commands, extended error information may be available via the `tracing/error_log` file. For the commands that support it, reading the `tracing/error_log` file after an error will display more detailed information about what went wrong, if information is available. The `tracing/error_log` file is a circular error

log displaying a small number (currently, 8) of ftrace errors for the last (8) failed commands.

The extended error information and usage takes the form shown in this example:

```
# echo xxx > /sys/kernel/debug/tracing/events/sched/sched_wakeup/
↪ trigger
echo: write error: Invalid argument

# cat /sys/kernel/debug/tracing/error_log
[ 5348.887237] location: error: Couldn't yyy: zzz
  Command: xxx
           ^
[ 7517.023364] location: error: Bad rrr: sss
  Command: ppp qq
           ^
```

To clear the error log, echo the empty string into it:

```
# echo > /sys/kernel/debug/tracing/error_log
```

3.6 Examples of using the tracer

Here are typical examples of using the tracers when controlling them only with the tracefs interface (without using any user-land utilities).

3.7 Output format:

Here is an example of the output format of the file “trace”:

```
# tracer: function
#
# entries-in-buffer/entries-written: 140080/250280   #P:4
#
#          _-----=> irqsoff
#          /_-----=> need-resched
#          | /_-----=> hardirq/softirq
#          || /_--=> preempt-depth
#          ||| /      delay
#          ||| /
#          TASK-PID   CPU#   |      |      |      |      |      |      |      |      |      |
#          | |         |     |      |      |      |      |      |      |      |      |
#          bash-1977   [000] .... 17284.993652: sys_close <-system_call_
↪ fastpath
          bash-1977   [000] .... 17284.993653: __close_fd <-sys_close
          bash-1977   [000] .... 17284.993653: _raw_spin_lock <-__close_fd
          sshd-1974    [003] .... 17284.993653: __srcu_read_unlock <-fsnotify
          bash-1977   [000] .... 17284.993654: add_preempt_count <-_raw_spin_
↪ lock
          bash-1977   [000] ...1 17284.993655: _raw_spin_unlock <-__close_fd
```

```

bash-1977 [000] ...1 17284.993656: sub_preempt_count <-_raw_spin_
→unlock
bash-1977 [000] .... 17284.993657: filp_close <-__close_fd
bash-1977 [000] .... 17284.993657: dnotify_flush <-filp_close
sshd-1974 [003] .... 17284.993658: sys_select <-system_call_
→fastpath
....

```

A header is printed with the tracer name that is represented by the trace. In this case the tracer is “function”. Then it shows the number of events in the buffer as well as the total number of entries that were written. The difference is the number of entries that were lost due to the buffer filling up ($250280 - 140080 = 110200$ events lost).

The header explains the content of the events. Task name “bash”, the task PID “1977”, the CPU that it was running on “000”, the latency format (explained below), the timestamp in <secs>.<usecs> format, the function name that was traced “sys_close” and the parent function that called this function “system_call_fastpath”. The timestamp is the time at which the function was entered.

3.8 Latency trace format

When the latency-format option is enabled or when one of the latency tracers is set, the trace file gives somewhat more information to see why a latency happened. Here is a typical trace:

```

# tracer: irqsoff
#
# irqsoff latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 259 us, #4/4, CPU#2 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
#
# | task: ps-6143 (uid:0 nice:0 policy:0 rt_prio:0)
# -----
# => started at: __lock_task_sighand
# => ended at:  _raw_spin_unlock_irqrestore
#
#
#          _-----=> CPU#
#          /_-----=> irqs-off
#          | /_-----=> need-resched
#          || /_-----=> hardirq/softirq
#          ||| /_-----=> preempt-depth
#          |||| /_-----=> delay
# cmd      pid  | time | caller
# \      /      | \    | /
#   ps-6143  2d... 0us!: trace_hardirqs_off <-__lock_task_sighand
#   ps-6143  2d..1 259us+: trace_hardirqs_on <-_raw_spin_unlock_irqrestore
#   ps-6143  2d..1 263us+: time_hardirqs_on <-_raw_spin_unlock_irqrestore
#   ps-6143  2d..1 306us : <stack trace>
=> trace_hardirqs_on_caller
=> trace_hardirqs_on

```

```
=> _raw_spin_unlock_irqrestore
=> do_task_stat
=> proc_tgid_stat
=> proc_single_show
=> seq_read
=> vfs_read
=> sys_read
=> system_call_fastpath
```

This shows that the current tracer is “irqsoff” tracing the time for which interrupts were disabled. It gives the trace version (which never changes) and the version of the kernel upon which this was executed on (3.8). Then it displays the max latency in microseconds (259 us). The number of trace entries displayed and the total number (both are four: #4/4). VP, KP, SP, and HP are always zero and are reserved for later use. #P is the number of online CPUs (#P:4).

The task is the process that was running when the latency occurred. (ps pid: 6143).

The start and stop (the functions in which the interrupts were disabled and enabled respectively) that caused the latencies:

- `__lock_task_sighand` is where the interrupts were disabled.
- `_raw_spin_unlock_irqrestore` is where they were enabled again.

The next lines after the header are the trace itself. The header explains which is which.

cmd: The name of the process in the trace.

pid: The PID of that process.

CPU#: The CPU which the process was running on.

irqs-off: ‘d’ interrupts are disabled. ‘.’ otherwise.

Caution: If the architecture does not support a way to read the irq flags variable, an ‘X’ will always be printed here.

need-resched:

- ‘N’ both `TIF_NEED_RESCHED` and `PREEMPT_NEED_RESCHED` is set,
- ‘n’ only `TIF_NEED_RESCHED` is set,
- ‘p’ only `PREEMPT_NEED_RESCHED` is set,
- ‘.’ otherwise.

hardirq/softirq:

- ‘Z’ - NMI occurred inside a hardirq
- ‘z’ - NMI is running
- ‘H’ - hard irq occurred inside a softirq.
- ‘h’ - hard irq is running
- ‘s’ - soft irq is running

- `'.'` - normal context.

`preempt-depth`: The level of `preempt_disabled`

The above is mostly meaningful for kernel developers.

time: When the `latency-format` option is enabled, the trace file output includes a timestamp relative to the start of the trace. This differs from the output when `latency-format` is disabled, which includes an absolute timestamp.

delay: This is just to help catch your eye a bit better. And needs to be fixed to be only relative to the same CPU. The marks are determined by the difference between this current trace and the next trace.

- `'$'` - greater than 1 second
- `'@'` - greater than 100 millisecond
- `'*'` - greater than 10 millisecond
- `'#'` - greater than 1000 microsecond
- `'!'` - greater than 100 microsecond
- `'+'` - greater than 10 microsecond
- `''` - less than or equal to 10 microsecond.

The rest is the same as the `'trace'` file.

Note, the latency tracers will usually end with a back trace to easily find where the latency occurred.

3.9 trace_options

The `trace_options` file (or the options directory) is used to control what gets printed in the trace output, or manipulate the tracers. To see what is available, simply cat the file:

```
cat trace_options
  print-parent
  nosym-offset
  nosym-addr
  noverbose
  noraw
  nohex
  nobin
  noblock
  trace_printk
  annotate
  nouserstacktrace
  nosym-userobj
  noprintk-msg-only
  context-info
  nolatency-format
  record-cmd
  norecord-tgid
```

```
overwrite
nodisable_on_free
irq-info
markers
noevent-fork
function-trace
nofunction-fork
nodisplay-graph
nostacktrace
nobranch
```

To disable one of the options, echo in the option prepended with “no”:

```
echo noprint-parent > trace_options
```

To enable an option, leave off the “no”:

```
echo sym-offset > trace_options
```

Here are the available options:

print-parent On function traces, display the calling (parent) function as well as the function being traced.

```
print-parent:
bash-4000 [01] 1477.606694: simple_strtoul <-kstrtoul

noprint-parent:
bash-4000 [01] 1477.606694: simple_strtoul
```

sym-offset Display not only the function name, but also the offset in the function. For example, instead of seeing just “ktime_get”, you will see “ktime_get+0xb/0x20”.

```
sym-offset:
bash-4000 [01] 1477.606694: simple_strtoul+0x6/0xa0
```

sym-addr This will also display the function address as well as the function name.

```
sym-addr:
bash-4000 [01] 1477.606694: simple_strtoul <c0339346>
```

verbose This deals with the trace file when the latency-format option is enabled.

```
bash 4000 1 0 00000000 00010a95 [58127d26] 1720.415ms \
(+0.000ms): simple_strtoul (kstrtoul)
```

raw This will display raw numbers. This option is best for use with user applications that can translate the raw numbers better than having it done in the kernel.

hex Similar to raw, but the numbers will be in a hexadecimal format.

bin This will print out the formats in raw binary.

block When set, reading trace_pipe will not block when polled.

trace_printk Can disable `trace_printk()` from writing into the buffer.

annotate It is sometimes confusing when the CPU buffers are full and one CPU buffer had a lot of events recently, thus a shorter time frame, were another CPU may have only had a few events, which lets it have older events. When the trace is reported, it shows the oldest events first, and it may look like only one CPU ran (the one with the oldest events). When the `annotate` option is set, it will display when a new CPU buffer started:

```

        <idle>-0      [001] dNs4 21169.031481: wake_up_idle_cpu <-
↪add_timer_on
        <idle>-0      [001] dNs4 21169.031482: _raw_spin_unlock_
↪irqrestore <-add_timer_on
        <idle>-0      [001] .Ns4 21169.031484: sub_preempt_count
↪<_raw_spin_unlock_irqrestore
##### CPU 2 buffer started #####
        <idle>-0      [002] .N.1 21169.031484: rcu_idle_exit <-
↪cpu_idle
        <idle>-0      [001] .Ns3 21169.031484: _raw_spin_unlock <-
↪clocksource_watchdog
        <idle>-0      [001] .Ns3 21169.031485: sub_preempt_count
↪<_raw_spin_unlock

```

userstacktrace This option changes the trace. It records a stacktrace of the current user space thread after each trace event.

sym-userobj when user stacktrace are enabled, look up which object the address belongs to, and print a relative address. This is especially useful when ASLR is on, otherwise you don't get a chance to resolve the address to object/file/line after the app is no longer running

The lookup is performed when you read `trace,trace_pipe`. Example:

```

a.out-1623  [000] 40874.465068: /root/a.out[+0x480] <- /root/a.
↪out[+0
x494] <- /root/a.out[+0x4a8] <- /lib/libc-2.7.so[+0x1e1a6]

```

printk-msg-only When set, `trace_printk()`s will only show the format and not their parameters (if `trace_bprintk()` or `trace_bputs()` was used to save the `trace_printk()`).

context-info Show only the event data. Hides the comm, PID, timestamp, CPU, and other useful data.

latency-format This option changes the trace output. When it is enabled, the trace displays additional information about the latency, as described in "Latency trace format".

pause-on-trace When set, opening the trace file for read, will pause writing to the ring buffer (as if `tracing_on` was set to zero). This simulates the original behavior of the trace file. When the file is closed, tracing will be enabled again.

hash-ptr When set, "%p" in the event `printk` format displays the hashed pointer value instead of real address. This will be useful if you want to find out which hashed value is corresponding to the real value in trace log.

record-cmd When any event or tracer is enabled, a hook is enabled in the sched_switch trace point to fill comm cache with mapped pids and comms. But this may cause some overhead, and if you only care about pids, and not the name of the task, disabling this option can lower the impact of tracing. See “saved_cmdlines”.

record-tgid When any event or tracer is enabled, a hook is enabled in the sched_switch trace point to fill the cache of mapped Thread Group IDs (TGID) mapping to pids. See “saved_tgids”.

overwrite This controls what happens when the trace buffer is full. If “1” (default), the oldest events are discarded and overwritten. If “0”, then the newest events are discarded. (see per_cpu/cpu0/stats for overrun and dropped)

disable_on_free When the free_buffer is closed, tracing will stop (tracing_on set to 0).

irq-info Shows the interrupt, preempt count, need resched data. When disabled, the trace looks like:

```
# tracer: function
#
# entries-in-buffer/entries-written: 144405/9452052   #P:4
#
#          TASK-PID    CPU#    TIMESTAMP    FUNCTION
#          | |         |         |         |
#          <idle>-0     [002]    23636.756054: ttwu_do_activate.
↪constprop.89 <-try_to_wake_up
#          <idle>-0     [002]    23636.756054: activate_task <-ttwu_
↪do_activate.constprop.89
#          <idle>-0     [002]    23636.756055: enqueue_task <-
↪activate_task
```

markers When set, the trace_marker is writable (only by root). When disabled, the trace_marker will error with EINVAL on write.

event-fork When set, tasks with PIDs listed in set_event_pid will have the PIDs of their children added to set_event_pid when those tasks fork. Also, when tasks with PIDs in set_event_pid exit, their PIDs will be removed from the file.

This affects PIDs listed in set_event_notrace_pid as well.

function-trace The latency tracers will enable function tracing if this option is enabled (default it is). When it is disabled, the latency tracers do not trace functions. This keeps the overhead of the tracer down when performing latency tests.

function-fork When set, tasks with PIDs listed in set_fttrace_pid will have the PIDs of their children added to set_fttrace_pid when those tasks fork. Also, when tasks with PIDs in set_fttrace_pid exit, their PIDs will be removed from the file.

This affects PIDs in set_fttrace_notrace_pid as well.

display-graph When set, the latency tracers (irqsoff, wakeup, etc) will use function graph tracing instead of function tracing.

stacktrace When set, a stack trace is recorded after any trace event is recorded.

branch Enable branch tracing with the tracer. This enables branch tracer along with the currently set tracer. Enabling this with the “nop” tracer is the same as just enabling the “branch” tracer.

Tip: Some tracers have their own options. They only appear in this file when the tracer is active. They always appear in the options directory.

Here are the per tracer options:

Options for function tracer:

func_stack_trace When set, a stack trace is recorded after every function that is recorded. NOTE! Limit the functions that are recorded before enabling this, with “set_fttrace_filter” otherwise the system performance will be critically degraded. Remember to disable this option before clearing the function filter.

Options for function_graph tracer:

Since the function_graph tracer has a slightly different output it has its own options to control what is displayed.

funcgraph-overflow When set, the “overflow” of the graph stack is displayed after each function traced. The overflow, is when the stack depth of the calls is greater than what is reserved for each task. Each task has a fixed array of functions to trace in the call graph. If the depth of the calls exceeds that, the function is not traced. The overflow is the number of functions missed due to exceeding this array.

funcgraph-cpu When set, the CPU number of the CPU where the trace occurred is displayed.

funcgraph-overhead When set, if the function takes longer than A certain amount, then a delay marker is displayed. See “delay” above, under the header description.

funcgraph-proc Unlike other tracers, the process’ command line is not displayed by default, but instead only when a task is traced in and out during a context switch. Enabling this options has the command of each process displayed at every line.

funcgraph-duration At the end of each function (the return) the duration of the amount of time in the function is displayed in microseconds.

funcgraph-abstime When set, the timestamp is displayed at each line.

funcgraph-irqs When disabled, functions that happen inside an interrupt will not be traced.

funcgraph-tail When set, the return event will include the function that it represents. By default this is off, and only a closing curly bracket “}” is displayed for the return of a function.

sleep-time When running function graph tracer, to include the time a task schedules out in its function. When enabled, it will account time the task has been scheduled out as part of the function call.

graph-time When running function profiler with function graph tracer, to include the time to call nested functions. When this is not set, the time reported for the function will only include the time the function itself executed for, not the time for functions that it called.

Options for blk tracer:

blk_classic Shows a more minimalistic output.

3.10 irqsoff

When interrupts are disabled, the CPU can not react to any other external event (besides NMIs and SMIs). This prevents the timer interrupt from triggering or the mouse interrupt from letting the kernel know of a new mouse event. The result is a latency with the reaction time.

The irqsoff tracer tracks the time for which interrupts are disabled. When a new maximum latency is hit, the tracer saves the trace leading up to that latency point so that every time a new maximum is reached, the old saved trace is discarded and the new trace is saved.

To reset the maximum, echo 0 into tracing_max_latency. Here is an example:

```
# echo 0 > options/function-trace
# echo irqsoff > current_tracer
# echo 1 > tracing_on
# echo 0 > tracing_max_latency
# ls -ltr
[...]
# echo 0 > tracing_on
# cat trace
# tracer: irqsoff
#
# irqsoff latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 16 us, #4/4, CPU#0 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
# -----
#   | task: swapper/0-0 (uid:0 nice:0 policy:0 rt_prio:0)
#   | -----
# => started at: run_timer_softirq
# => ended at:   run_timer_softirq
#
#
#           _-----=> CPU#
#           /_-----=> irqs-off
#           | /_-----=> need-resched
#           || /_-----=> hardirq/softirq
#           ||| /_-----=> preempt-depth
#           |||| /_-----=> delay
# cmd      pid  ||||| time | caller
#   \      /  ||||| \    | /
<idle>-0    0d.s2   0us+: _raw_spin_lock_irq <-run_timer_softirq
<idle>-0    0dNs3   17us : _raw_spin_unlock_irq <-run_timer_softirq
<idle>-0    0dNs3   17us+: trace_hardirqs_on <-run_timer_softirq
```

```

<idle>-0      0dNs3   25us : <stack trace>
=> _raw_spin_unlock_irq
=> run_timer_softirq
=> __do_softirq
=> call_softirq
=> do_softirq
=> irq_exit
=> smp_apic_timer_interrupt
=> apic_timer_interrupt
=> rcu_idle_exit
=> cpu_idle
=> rest_init
=> start_kernel
=> x86_64_start_reservations
=> x86_64_start_kernel

```

Here we see that we had a latency of 16 microseconds (which is very good). The `_raw_spin_lock_irq` in `run_timer_softirq` disabled interrupts. The difference between the 16 and the displayed timestamp 25us occurred because the clock was incremented between the time of recording the max latency and the time of recording the function that had that latency.

Note the above example had function-trace not set. If we set function-trace, we get a much larger output:

with `echo 1 > options/function-trace`

```

# tracer: irqsoff
#
# irqsoff latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 71 us, #168/168, CPU#3 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
#
#   | task: bash-2042 (uid:0 nice:0 policy:0 rt_prio:0)
#   | -----
# => started at: ata_scsi_queuecmd
# => ended at:   ata_scsi_queuecmd
#
#
#           _-----=> CPU#
#           /_-----=> irqs-off
#           | /_-----=> need-resched
#           || /_---=> hardirq/softirq
#           ||| /_--=> preempt-depth
#           |||| /_   delay
# cmd      pid    ||||| time | caller
# \      /      ||||| \   | /
bash-2042  3d...  0us : _raw_spin_lock_irqsave <-ata_scsi_queuecmd
bash-2042  3d...  0us : add_preempt_count <-_raw_spin_lock_irqsave
bash-2042  3d..1  1us : ata_scsi_find_dev <-ata_scsi_queuecmd
bash-2042  3d..1  1us : __ata_scsi_find_dev <-ata_scsi_find_dev
bash-2042  3d..1  2us : ata_find_dev.part.14 <-__ata_scsi_find_dev

```

```
bash-2042    3d..1    2us : ata_qc_new_init <-__ata_scsi_queuecmd
bash-2042    3d..1    3us : ata_sg_init <-__ata_scsi_queuecmd
bash-2042    3d..1    4us : ata_scsi_rw_xlat <-__ata_scsi_queuecmd
bash-2042    3d..1    4us : ata_build_rw_tf <-ata_scsi_rw_xlat
[...]
```

bash-2042 3d..1 67us : delay_tsc <-__delay

bash-2042 3d..1 67us : add_preempt_count <-delay_tsc

bash-2042 3d..2 67us : sub_preempt_count <-delay_tsc

bash-2042 3d..1 67us : add_preempt_count <-delay_tsc

bash-2042 3d..2 68us : sub_preempt_count <-delay_tsc

bash-2042 3d..1 68us+: ata_bmdma_start <-ata_bmdma_qc_issue

bash-2042 3d..1 71us : _raw_spin_unlock_irqrestore <-ata_scsi_

→queuecmd

bash-2042 3d..1 71us : _raw_spin_unlock_irqrestore <-ata_scsi_

→queuecmd

bash-2042 3d..1 72us+: trace_hardirqs_on <-ata_scsi_queuecmd

bash-2042 3d..1 120us : <stack trace>

=> _raw_spin_unlock_irqrestore

=> ata_scsi_queuecmd

=> scsi_dispatch_cmd

=> scsi_request_fn

=> __blk_run_queue_uncond

=> __blk_run_queue

=> blk_queue_bio

=> submit_bio_noacct

=> submit_bio

=> submit_bh

=> __ext3_get_inode_loc

=> ext3_iget

=> ext3_lookup

=> lookup_real

=> __lookup_hash

=> walk_component

=> lookup_last

=> path_lookupat

=> filename_lookup

=> user_path_at_empty

=> user_path_at

=> vfs_fstatat

=> vfs_stat

=> sys_newstat

=> system_call_fastpath

Here we traced a 71 microsecond latency. But we also see all the functions that were called during that time. Note that by enabling function tracing, we incur an added overhead. This overhead may extend the latency times. But nevertheless, this trace has provided some very helpful debugging information.

If we prefer function graph output instead of function, we can set display-graph option:

with echo 1 > options/display-graph

```
# tracer: irqsoff
#
# irqsoff latency trace v1.1.5 on 4.20.0-rc6+
# -----
# latency: 3751 us, #274/274, CPU#0 | (M:desktop VP:0, KP:0, SP:0 HP:0 #P:4)
# -----
# | task: bash-1507 (uid:0 nice:0 policy:0 rt_prio:0)
# -----
# => started at: free_debug_processing
# => ended at:   return_to_handler
#
#
#                                     -----=> irqs-off
#                                     /-----=> need-resched
#                                     | /-----=> hardirq/softirq
#                                     || /-----=> preempt-depth
#                                     ||| /
#
# REL TIME      CPU  TASK/PID      |||| DURATION
->FUNCTION CALLS
#      |      |      |      |      |      |
->|      |      |      |      |      |      |
      0 us | 0)  bash-1507 | d... | 0.000 us | _raw_spin_lock_
->irqsave();
      0 us | 0)  bash-1507 | d..1 | 0.378 us | do_raw_spin_
->trylock();
      1 us | 0)  bash-1507 | d..2 |      | set_track() {
      2 us | 0)  bash-1507 | d..2 |      |   save_stack_
->trace() {
      2 us | 0)  bash-1507 | d..2 |      |   __save_
->stack_trace() {
      3 us | 0)  bash-1507 | d..2 |      |   _
->unwind_start() {
      3 us | 0)  bash-1507 | d..2 |      |   get_
->stack_info() {
      3 us | 0)  bash-1507 | d..2 | 0.351 us |   in_
->task_stack();
      4 us | 0)  bash-1507 | d..2 | 1.107 us |   }
[...
      3750 us | 0)  bash-1507 | d..1 | 0.516 us | do_raw_spin_
->unlock();
      3750 us | 0)  bash-1507 | d..1 | 0.000 us | _raw_spin_
->unlock_irqrestore();
      3764 us | 0)  bash-1507 | d..1 | 0.000 us | tracer_hardirqs_
->on();
      bash-1507 0d..1 3792us : <stack trace>
=> free_debug_processing
=> __slab_free
=> kmem_cache_free
```

```
=> vm_area_free
=> remove_vma
=> exit_mmap
=> mmput
=> begin_new_exec
=> load_elf_binary
=> search_binary_handler
=> __do_execve_file.isra.32
=> __x64_sys_execve
=> do_syscall_64
=> entry_SYSCALL_64_after_hwframe
```

3.11 preemptoff

When preemption is disabled, we may be able to receive interrupts but the task cannot be preempted and a higher priority task must wait for preemption to be enabled again before it can preempt a lower priority task.

The preemptoff tracer traces the places that disable preemption. Like the irqsoff tracer, it records the maximum latency for which preemption was disabled. The control of preemptoff tracer is much like the irqsoff tracer.

```
# echo 0 > options/function-trace
# echo preemptoff > current_tracer
# echo 1 > tracing_on
# echo 0 > tracing_max_latency
# ls -ltr
[...]
# echo 0 > tracing_on
# cat trace
# tracer: preemptoff
#
# preemptoff latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 46 us, #4/4, CPU#1 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
# -----
# | task: sshd-1991 (uid:0 nice:0 policy:0 rt_prio:0)
# -----
# => started at: do_IRQ
# => ended at:   do_IRQ
#
#
#           _-----=> CPU#
#         /_-----=> irqs-off
#       | /_-----=> need-resched
#     || /_-----=> hardirq/softirq
#   ||| /_-----=> preempt-depth
#  |||| /_-----=> delay
# cmd  pid  ||||| time  | caller
```

```
#      \  /      |||||  \   |  /
ssh-d-1991  1d.h.    0us+: irq_enter <-do_IRQ
ssh-d-1991  1d..1   46us : irq_exit <-do_IRQ
ssh-d-1991  1d..1   47us+: trace_preempt_on <-do_IRQ
ssh-d-1991  1d..1   52us : <stack trace>
=> sub_preempt_count
=> irq_exit
=> do_IRQ
=> ret_from_intr
```

This has some more changes. Preemption was disabled when an interrupt came in (notice the 'h'), and was enabled on exit. But we also see that interrupts have been disabled when entering the preempt off section and leaving it (the 'd'). We do not know if interrupts were enabled in the mean time or shortly after this was over.

```
# tracer: preemptoff
#
# preemptoff latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 83 us, #241/241, CPU#1 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
#
# | task: bash-1994 (uid:0 nice:0 policy:0 rt_prio:0)
# -----
# => started at: wake_up_new_task
# => ended at:   task_rq_unlock
#
#
#          _-----=> CPU#
#          /_-----=> irqs-off
#          | /_-----=> need-resched
#          || /_-----=> hardirq/softirq
#          ||| /_-----=> preempt-depth
#          |||| /_-----=> delay
# cmd      pid  ||||| time | caller
#  \      /  ||||| \   | /
bash-1994  1d..1    0us : _raw_spin_lock_irqsave <-wake_up_new_task
bash-1994  1d..1    0us : select_task_rq_fair <-select_task_rq
bash-1994  1d..1    1us : __rcu_read_lock <-select_task_rq_fair
bash-1994  1d..1    1us : source_load <-select_task_rq_fair
bash-1994  1d..1    1us : source_load <-select_task_rq_fair
[... ]
bash-1994  1d..1   12us : irq_enter <-smp_apic_timer_interrupt
bash-1994  1d..1   12us : rcu_irq_enter <-irq_enter
bash-1994  1d..1   13us : add_preempt_count <-irq_enter
bash-1994  1d.h1   13us : exit_idle <-smp_apic_timer_interrupt
bash-1994  1d.h1   13us : hrtimer_interrupt <-smp_apic_timer_interrupt
bash-1994  1d.h1   13us : _raw_spin_lock <-hrtimer_interrupt
bash-1994  1d.h1   14us : add_preempt_count <-_raw_spin_lock
bash-1994  1d.h2   14us : ktime_get_update_offsets <-hrtimer_interrupt
[... ]
bash-1994  1d.h1   35us : lapic_next_event <-clockevents_program_event
```

```
bash-1994    ld.h1    35us : irq_exit <-smp_apic_timer_interrupt
bash-1994    ld.h1    36us : sub_preempt_count <-irq_exit
bash-1994    ld..2    36us : do_softirq <-irq_exit
bash-1994    ld..2    36us : __do_softirq <-call_softirq
bash-1994    ld..2    36us : __local_bh_disable <-__do_softirq
bash-1994    ld.s2    37us : add_preempt_count <-_raw_spin_lock_irq
bash-1994    ld.s3    38us : _raw_spin_unlock <-run_timer_softirq
bash-1994    ld.s3    39us : sub_preempt_count <-_raw_spin_unlock
bash-1994    ld.s2    39us : call_timer_fn <-run_timer_softirq
[...]
```

```
bash-1994    ldNs2    81us : cpu_needs_another_gp <-rcu_process_callbacks
bash-1994    ldNs2    82us : __local_bh_enable <-__do_softirq
bash-1994    ldNs2    82us : sub_preempt_count <-__local_bh_enable
bash-1994    ldN.2    82us : idle_cpu <-irq_exit
bash-1994    ldN.2    83us : rcu_irq_exit <-irq_exit
bash-1994    ldN.2    83us : sub_preempt_count <-irq_exit
bash-1994    1.N.1    84us : _raw_spin_unlock_irqrestore <-task_rq_unlock
bash-1994    1.N.1    84us+: trace_preempt_on <-task_rq_unlock
bash-1994    1.N.1    104us : <stack trace>
=> sub_preempt_count
=> _raw_spin_unlock_irqrestore
=> task_rq_unlock
=> wake_up_new_task
=> do_fork
=> sys_clone
=> stub_clone
```

The above is an example of the preemptoff trace with function-trace set. Here we see that interrupts were not disabled the entire time. The `irq_enter` code lets us know that we entered an interrupt 'h'. Before that, the functions being traced still show that it is not in an interrupt, but we can see from the functions themselves that this is not the case.

3.12 preemptirqsoff

Knowing the locations that have interrupts disabled or preemption disabled for the longest times is helpful. But sometimes we would like to know when either preemption and/or interrupts are disabled.

Consider the following code:

```
local_irq_disable();
call_function_with_irqs_off();
preempt_disable();
call_function_with_irqs_and_preemption_off();
local_irq_enable();
call_function_with_preemption_off();
preempt_enable();
```

The `irqsoff` tracer will record the total length of `call_function_with_irqs_off()` and `call_function_with_irqs_and_preemption_off()`.

The preemptoff tracer will record the total length of `call_function_with_irqs_and_preemption_off()` and `call_function_with_preemption_off()`.

But neither will trace the time that interrupts and/or preemption is disabled. This total time is the time that we can not schedule. To record this time, use the preemptirqsoff tracer.

Again, using this trace is much like the irqsoff and preemptoff tracers.

```
# echo 0 > options/function-trace
# echo preemptirqsoff > current_tracer
# echo 1 > tracing_on
# echo 0 > tracing_max_latency
# ls -ltr
[...]
# echo 0 > tracing_on
# cat trace
# tracer: preemptirqsoff
#
# preemptirqsoff latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 100 us, #4/4, CPU#3 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
# -----
#   | task: ls-2230 (uid:0 nice:0 policy:0 rt_prio:0)
#   | -----
# => started at: ata_scsi_queuecmd
# => ended at:   ata_scsi_queuecmd
#
#
#           _-----=> CPU#
#           /_-----=> irqsoff
#           | /_-----=> need-resched
#           || /_-----=> hardirq/softirq
#           ||| /_-----=> preempt-depth
#           |||| /_-----=> delay
# cmd      pid  ||||| time | caller
#  \      /    ||||| \    | /
#   ls-2230  3d...  0us+: _raw_spin_lock_irqsave <-ata_scsi_queuecmd
#   ls-2230  3...1 100us : _raw_spin_unlock_irqrestore <-ata_scsi_queuecmd
#   ls-2230  3...1 101us+: trace_preempt_on <-ata_scsi_queuecmd
#   ls-2230  3...1 111us : <stack trace>
=> sub_preempt_count
=> _raw_spin_unlock_irqrestore
=> ata_scsi_queuecmd
=> scsi_dispatch_cmd
=> scsi_request_fn
=> __blk_run_queue_uncond
=> __blk_run_queue
=> blk_queue_bio
=> submit_bio_noacct
=> submit_bio
=> submit_bh
=> ext3_bread
```

```

=> ext3_dir_bread
=> htree_dirblock_to_tree
=> ext3_htree_fill_tree
=> ext3_readdir
=> vfs_readdir
=> sys_getdents
=> system_call_fastpath

```

The `trace_hardirqs_off_thunk` is called from assembly on x86 when interrupts are disabled in the assembly code. Without the function tracing, we do not know if interrupts were enabled within the preemption points. We do see that it started with preemption enabled.

Here is a trace with function-trace set:

```

# tracer: preemptirqsoff
#
# preemptirqsoff latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 161 us, #339/339, CPU#3 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
# -----
# | task: ls-2269 (uid:0 nice:0 policy:0 rt_prio:0)
# -----
# => started at: schedule
# => ended at:   mutex_unlock
#
#
#           _-----=> CPU#
#           /_-----=> irqsoff
#           | /_-----=> need-resched
#           || /_-----=> hardirq/softirq
#           ||| /_-----=> preempt-depth
#           |||| /_-----=> delay
# cmd      pid  ||||| time | caller
# \      /  ||||| \    | /
kworker/-59 3...1 0us : __schedule <-schedule
kworker/-59 3d..1 0us : rcu_preempt_qs <-rcu_note_context_switch
kworker/-59 3d..1 1us : add_preempt_count <-_raw_spin_lock_irq
kworker/-59 3d..2 1us : deactivate_task <-__schedule
kworker/-59 3d..2 1us : dequeue_task <-deactivate_task
kworker/-59 3d..2 2us : update_rq_clock <-dequeue_task
kworker/-59 3d..2 2us : dequeue_task_fair <-dequeue_task
kworker/-59 3d..2 2us : update_curr <-dequeue_task_fair
kworker/-59 3d..2 2us : update_min_vruntime <-update_curr
kworker/-59 3d..2 3us : cpuacct_charge <-update_curr
kworker/-59 3d..2 3us : __rcu_read_lock <-cpuacct_charge
kworker/-59 3d..2 3us : __rcu_read_unlock <-cpuacct_charge
kworker/-59 3d..2 3us : update_cfs_rq_blocked_load <-dequeue_task_fair
kworker/-59 3d..2 4us : clear_buddies <-dequeue_task_fair
kworker/-59 3d..2 4us : account_entity_dequeue <-dequeue_task_fair
kworker/-59 3d..2 4us : update_min_vruntime <-dequeue_task_fair
kworker/-59 3d..2 4us : update_cfs_shares <-dequeue_task_fair

```

```

kworker/-59      3d..2      5us : hrtick_update <-dequeue_task_fair
kworker/-59      3d..2      5us : wq_worker_sleeping <-__schedule
kworker/-59      3d..2      5us : kthread_data <-wq_worker_sleeping
kworker/-59      3d..2      5us : put_prev_task_fair <-__schedule
kworker/-59      3d..2      6us : pick_next_task_fair <-pick_next_task
kworker/-59      3d..2      6us : clear_buddies <-pick_next_task_fair
kworker/-59      3d..2      6us : set_next_entity <-pick_next_task_fair
kworker/-59      3d..2      6us : update_stats_wait_end <-set_next_entity
  ls-2269        3d..2      7us : finish_task_switch <-__schedule
  ls-2269        3d..2      7us : _raw_spin_unlock_irq <-finish_task_switch
  ls-2269        3d..2      8us : do_IRQ <-ret_from_intr
  ls-2269        3d..2      8us : irq_enter <-do_IRQ
  ls-2269        3d..2      8us : rcu_irq_enter <-irq_enter
  ls-2269        3d..2      9us : add_preempt_count <-irq_enter
  ls-2269        3d.h2      9us : exit_idle <-do_IRQ
[... ]
  ls-2269        3d.h3     20us : sub_preempt_count <-_raw_spin_unlock
  ls-2269        3d.h2     20us : irq_exit <-do_IRQ
  ls-2269        3d.h2     21us : sub_preempt_count <-irq_exit
  ls-2269        3d..3     21us : do_softirq <-irq_exit
  ls-2269        3d..3     21us : __do_softirq <-call_softirq
  ls-2269        3d..3     21us+: __local_bh_disable <-__do_softirq
  ls-2269        3d.s4     29us : sub_preempt_count <-_local_bh_enable_ip
  ls-2269        3d.s5     29us : sub_preempt_count <-_local_bh_enable_ip
  ls-2269        3d.s5     31us : do_IRQ <-ret_from_intr
  ls-2269        3d.s5     31us : irq_enter <-do_IRQ
  ls-2269        3d.s5     31us : rcu_irq_enter <-irq_enter
[... ]
  ls-2269        3d.s5     31us : rcu_irq_enter <-irq_enter
  ls-2269        3d.s5     32us : add_preempt_count <-irq_enter
  ls-2269        3d.H5     32us : exit_idle <-do_IRQ
  ls-2269        3d.H5     32us : handle_irq <-do_IRQ
  ls-2269        3d.H5     32us : irq_to_desc <-handle_irq
  ls-2269        3d.H5     33us : handle_fasteoi_irq <-handle_irq
[... ]
  ls-2269        3d.s5    158us : _raw_spin_unlock_irqrestore <-rtl8139_poll
  ls-2269        3d.s3    158us : net_rps_action_and_irq_enable.isra.65 <-net_rx_
→action
  ls-2269        3d.s3    159us : __local_bh_enable <-__do_softirq
  ls-2269        3d.s3    159us : sub_preempt_count <-__local_bh_enable
  ls-2269        3d..3    159us : idle_cpu <-irq_exit
  ls-2269        3d..3    159us : rcu_irq_exit <-irq_exit
  ls-2269        3d..3    160us : sub_preempt_count <-irq_exit
  ls-2269        3d...    161us : __mutex_unlock_slowpath <-mutex_unlock
  ls-2269        3d...    162us+: trace_hardirqs_on <-mutex_unlock
  ls-2269        3d...    186us : <stack trace>
=> __mutex_unlock_slowpath
=> mutex_unlock
=> process_output
=> n_tty_write

```

```
=> tty_write
=> vfs_write
=> sys_write
=> system_call_fastpath
```

This is an interesting trace. It started with kworker running and scheduling out and ls taking over. But as soon as ls released the rq lock and enabled interrupts (but not preemption) an interrupt triggered. When the interrupt finished, it started running softirqs. But while the softirq was running, another interrupt triggered. When an interrupt is running inside a softirq, the annotation is 'H'.

3.13 wakeup

One common case that people are interested in tracing is the time it takes for a task that is woken to actually wake up. Now for non Real-Time tasks, this can be arbitrary. But tracing it none the less can be interesting.

Without function tracing:

```
# echo 0 > options/function-trace
# echo wakeup > current_tracer
# echo 1 > tracing_on
# echo 0 > tracing_max_latency
# chrt -f 5 sleep 1
# echo 0 > tracing_on
# cat trace
# tracer: wakeup
#
# wakeup latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 15 us, #4/4, CPU#3 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
# -----
# | task: kworker/3:1H-312 (uid:0 nice:-20 policy:0 rt_prio:0)
# -----
#
#          -----=> CPU#
#          /-----=> irqs-off
#          | /-----=> need-resched
#          || /-----=> hardirq/softirq
#          ||| /-----=> preempt-depth
#          |||| /-----=> delay
# cmd      pid    | time | caller
# \      /      | \    | /
<idle>-0    3dNs7   0us :      0:120:R  + [003]   312:100:R kworker/3:1H
<idle>-0    3dNs7   1us+: ttwu_do_activate.constprop.87 <-try_to_wake_up
<idle>-0    3d..3   15us : __schedule <-schedule
<idle>-0    3d..3   15us :      0:120:R ==> [003]   312:100:R kworker/3:1H
```

The tracer only traces the highest priority task in the system to avoid tracing the normal circumstances. Here we see that the kworker with a nice priority of -20 (not very nice), took just

15 microseconds from the time it woke up, to the time it ran.

Non Real-Time tasks are not that interesting. A more interesting trace is to concentrate only on Real-Time tasks.

3.14 wakeup_rt

In a Real-Time environment it is very important to know the wakeup time it takes for the highest priority task that is woken up to the time that it executes. This is also known as “schedule latency”. I stress the point that this is about RT tasks. It is also important to know the scheduling latency of non-RT tasks, but the average schedule latency is better for non-RT tasks. Tools like LatencyTop are more appropriate for such measurements.

Real-Time environments are interested in the worst case latency. That is the longest latency it takes for something to happen, and not the average. We can have a very fast scheduler that may only have a large latency once in a while, but that would not work well with Real-Time tasks. The wakeup_rt tracer was designed to record the worst case wakeups of RT tasks. Non-RT tasks are not recorded because the tracer only records one worst case and tracing non-RT tasks that are unpredictable will overwrite the worst case latency of RT tasks (just run the normal wakeup tracer for a while to see that effect).

Since this tracer only deals with RT tasks, we will run this slightly differently than we did with the previous tracers. Instead of performing an ‘ls’, we will run ‘sleep 1’ under ‘chrt’ which changes the priority of the task.

```
# echo 0 > options/function-trace
# echo wakeup_rt > current_tracer
# echo 1 > tracing_on
# echo 0 > tracing_max_latency
# chrt -f 5 sleep 1
# echo 0 > tracing_on
# cat trace
# tracer: wakeup
#
# tracer: wakeup_rt
#
# wakeup_rt latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 5 us, #4/4, CPU#3 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
#
# | task: sleep-2389 (uid:0 nice:0 policy:1 rt_prio:5)
#
#
#          _-----=> CPU#
#          /_-----=> irqs-off
#          | /_-----=> need-resched
#          || /_-----=> hardirq/softirq
#          ||| /_-----=> preempt-depth
#          |||| /_-----=> delay
# cmd      pid      ||||| time | caller
#  \      /      ||||| \    | /
```

```

<idle>-0      3d.h4      0us :      0:120:R   + [003]  2389: 94:R sleep
<idle>-0      3d.h4      1us+: ttwu_do_activate.constprop.87 <-try_to_wake_up
<idle>-0      3d..3      5us :  __schedule <-schedule
<idle>-0      3d..3      5us :      0:120:R ==> [003]  2389: 94:R sleep

```

Running this on an idle system, we see that it only took 5 microseconds to perform the task switch. Note, since the trace point in the schedule is before the actual “switch”, we stop the tracing when the recorded task is about to schedule in. This may change if we add a new marker at the end of the scheduler.

Notice that the recorded task is ‘sleep’ with the PID of 2389 and it has an `rt_prio` of 5. This priority is user-space priority and not the internal kernel priority. The policy is 1 for `SCHED_FIFO` and 2 for `SCHED_RR`.

Note, that the trace data shows the internal priority (99 - `rtprio`).

```

<idle>-0      3d..3      5us :      0:120:R ==> [003]  2389: 94:R sleep

```

The 0:120:R means idle was running with a nice priority of 0 (120 - 120) and in the running state ‘R’. The sleep task was scheduled in with 2389: 94:R. That is the priority is the kernel `rtprio` (99 - 5 = 94) and it too is in the running state.

Doing the same with `chrt -r 5` and `function-trace` set.

```

echo 1 > options/function-trace

# tracer: wakeup_rt
#
# wakeup_rt latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 29 us, #85/85, CPU#3 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
# -----
# | task: sleep-2448 (uid:0 nice:0 policy:1 rt_prio:5)
# -----
#
#
#          _-----=> CPU#
#          /_-----=> irqsoft
#          | /_-----=> need-resched
#          || /_-----=> hardirq/softirq
#          ||| /_-----=> preempt-depth
#          |||| /_-----=> delay
# cmd      pid      | time | caller
# \      /          | \    | /
<idle>-0      3d.h4      1us+:      0:120:R   + [003]  2448: 94:R sleep
<idle>-0      3d.h4      2us : ttwu_do_activate.constprop.87 <-try_to_wake_up
<idle>-0      3d.h3      3us : check_preempt_curr <-ttwu_do_wakeup
<idle>-0      3d.h3      3us : resched_curr <-check_preempt_curr
<idle>-0      3dNh3      4us : task_woken_rt <-ttwu_do_wakeup
<idle>-0      3dNh3      4us : _raw_spin_unlock <-try_to_wake_up
<idle>-0      3dNh3      4us : sub_preempt_count <-_raw_spin_unlock
<idle>-0      3dNh2      5us : ttwu_stat <-try_to_wake_up
<idle>-0      3dNh2      5us : _raw_spin_unlock_irqrestore <-try_to_wake_up

```

```

<idle>-0      3dNh2      6us : sub_preempt_count <-_raw_spin_unlock_irqrestore
<idle>-0      3dNh1      6us : _raw_spin_lock <-__run_hrtimer
<idle>-0      3dNh1      6us : add_preempt_count <-_raw_spin_lock
<idle>-0      3dNh2      7us : _raw_spin_unlock <-hrtimer_interrupt
<idle>-0      3dNh2      7us : sub_preempt_count <-_raw_spin_unlock
<idle>-0      3dNh1      7us : tick_program_event <-hrtimer_interrupt
<idle>-0      3dNh1      7us : clockevents_program_event <-tick_program_event
<idle>-0      3dNh1      8us : ktime_get <-clockevents_program_event
<idle>-0      3dNh1      8us : lapic_next_event <-clockevents_program_event
<idle>-0      3dNh1      8us : irq_exit <-smp_apic_timer_interrupt
<idle>-0      3dNh1      9us : sub_preempt_count <-irq_exit
<idle>-0      3dN.2      9us : idle_cpu <-irq_exit
<idle>-0      3dN.2      9us : rcu_irq_exit <-irq_exit
<idle>-0      3dN.2     10us : rcu_eqs_enter_common.isra.45 <-rcu_irq_exit
<idle>-0      3dN.2     10us : sub_preempt_count <-irq_exit
<idle>-0      3.N.1     11us : rcu_idle_exit <-cpu_idle
<idle>-0      3dN.1     11us : rcu_eqs_exit_common.isra.43 <-rcu_idle_exit
<idle>-0      3.N.1     11us : tick_nohz_idle_exit <-cpu_idle
<idle>-0      3dN.1     12us : menu_hrtimer_cancel <-tick_nohz_idle_exit
<idle>-0      3dN.1     12us : ktime_get <-tick_nohz_idle_exit
<idle>-0      3dN.1     12us : tick_do_update_jiffies64 <-tick_nohz_idle_exit
<idle>-0      3dN.1     13us : cpu_load_update_nohz <-tick_nohz_idle_exit
<idle>-0      3dN.1     13us : _raw_spin_lock <-cpu_load_update_nohz
<idle>-0      3dN.1     13us : add_preempt_count <-_raw_spin_lock
<idle>-0      3dN.2     13us : __cpu_load_update <-cpu_load_update_nohz
<idle>-0      3dN.2     14us : sched_avg_update <-__cpu_load_update
<idle>-0      3dN.2     14us : _raw_spin_unlock <-cpu_load_update_nohz
<idle>-0      3dN.2     14us : sub_preempt_count <-_raw_spin_unlock
<idle>-0      3dN.1     15us : calc_load_nohz_stop <-tick_nohz_idle_exit
<idle>-0      3dN.1     15us : touch_softlockup_watchdog <-tick_nohz_idle_exit
<idle>-0      3dN.1     15us : hrtimer_cancel <-tick_nohz_idle_exit
<idle>-0      3dN.1     15us : hrtimer_try_to_cancel <-hrtimer_cancel
<idle>-0      3dN.1     16us : lock_hrtimer_base.isra.18 <-hrtimer_try_to_
→cancel
<idle>-0      3dN.1     16us : _raw_spin_lock_irqsave <-lock_hrtimer_base.
→isra.18
<idle>-0      3dN.1     16us : add_preempt_count <-_raw_spin_lock_irqsave
<idle>-0      3dN.2     17us : __remove_hrtimer <-remove_hrtimer.part.16
<idle>-0      3dN.2     17us : hrtimer_force_reprogram <-__remove_hrtimer
<idle>-0      3dN.2     17us : tick_program_event <-hrtimer_force_reprogram
<idle>-0      3dN.2     18us : clockevents_program_event <-tick_program_event
<idle>-0      3dN.2     18us : ktime_get <-clockevents_program_event
<idle>-0      3dN.2     18us : lapic_next_event <-clockevents_program_event
<idle>-0      3dN.2     19us : _raw_spin_unlock_irqrestore <-hrtimer_try_to_
→cancel
<idle>-0      3dN.2     19us : sub_preempt_count <-_raw_spin_unlock_irqrestore
<idle>-0      3dN.1     19us : hrtimer_forward <-tick_nohz_idle_exit
<idle>-0      3dN.1     20us : ktime_add_safe <-hrtimer_forward
<idle>-0      3dN.1     20us : ktime_add_safe <-hrtimer_forward
<idle>-0      3dN.1     20us : hrtimer_start_range_ns <-hrtimer_start_expires.
→constprop.11

```

```

<idle>-0      3dN.1   20us : __hrtimer_start_range_ns <-hrtimer_start_range_
→ns
<idle>-0      3dN.1   21us : lock_hrtimer_base.isra.18 <-__hrtimer_start_
→range_ns
<idle>-0      3dN.1   21us : _raw_spin_lock_irqsave <-lock_hrtimer_base.
→isra.18
<idle>-0      3dN.1   21us : add_preempt_count <-_raw_spin_lock_irqsave
<idle>-0      3dN.2   22us : ktime_add_safe <-__hrtimer_start_range_ns
<idle>-0      3dN.2   22us : enqueue_hrtimer <-__hrtimer_start_range_ns
<idle>-0      3dN.2   22us : tick_program_event <-__hrtimer_start_range_ns
<idle>-0      3dN.2   23us : clockevents_program_event <-tick_program_event
<idle>-0      3dN.2   23us : ktime_get <-clockevents_program_event
<idle>-0      3dN.2   23us : lapic_next_event <-clockevents_program_event
<idle>-0      3dN.2   24us : _raw_spin_unlock_irqrestore <-__hrtimer_start_
→range_ns
<idle>-0      3dN.2   24us : sub_preempt_count <-_raw_spin_unlock_irqrestore
<idle>-0      3dN.1   24us : account_idle_ticks <-tick_nohz_idle_exit
<idle>-0      3dN.1   24us : account_idle_time <-account_idle_ticks
<idle>-0      3dN.1   25us : sub_preempt_count <-cpu_idle
<idle>-0      3dN..   25us : schedule <-cpu_idle
<idle>-0      3dN..   25us : __schedule <-preempt_schedule
<idle>-0      3dN..   26us : add_preempt_count <-__schedule
<idle>-0      3dN.1   26us : rcu_note_context_switch <-__schedule
<idle>-0      3dN.1   26us : rcu_sched_qs <-rcu_note_context_switch
<idle>-0      3dN.1   27us : rcu_preempt_qs <-rcu_note_context_switch
<idle>-0      3dN.1   27us : _raw_spin_lock_irq <-__schedule
<idle>-0      3dN.1   27us : add_preempt_count <-_raw_spin_lock_irq
<idle>-0      3dN.2   28us : put_prev_task_idle <-__schedule
<idle>-0      3dN.2   28us : pick_next_task_stop <-pick_next_task
<idle>-0      3dN.2   28us : pick_next_task_rt <-pick_next_task
<idle>-0      3dN.2   29us : dequeue_pushable_task <-pick_next_task_rt
<idle>-0      3d..3   29us : __schedule <-preempt_schedule
<idle>-0      3d..3   30us :      0:120:R ==> [003] 2448: 94:R sleep

```

This isn't that big of a trace, even with function tracing enabled, so I included the entire trace.

The interrupt went off while when the system was idle. Somewhere before `task_woken_rt()` was called, the `NEED_RESCHED` flag was set, this is indicated by the first occurrence of the 'N' flag.

3.15 Latency tracing and events

As function tracing can induce a much larger latency, but without seeing what happens within the latency it is hard to know what caused it. There is a middle ground, and that is with enabling events.

```

# echo 0 > options/function-trace
# echo wakeup_rt > current_tracer
# echo 1 > events/enable
# echo 1 > tracing_on
# echo 0 > tracing_max_latency

```

```

# chrt -f 5 sleep 1
# echo 0 > tracing_on
# cat trace
# tracer: wakeup_rt
#
# wakeup_rt latency trace v1.1.5 on 3.8.0-test+
# -----
# latency: 6 us, #12/12, CPU#2 | (M:preempt VP:0, KP:0, SP:0 HP:0 #P:4)
# -----
# | task: sleep-5882 (uid:0 nice:0 policy:1 rt_prio:5)
# -----
#
#          _-----=> CPU#
#          /_-----=> irqs-off
#          | /_-----=> need-resched
#          || /_-----=> hardirq/softirq
#          ||| /_-----=> preempt-depth
#          |||| /_-----=> delay
# cmd      pid      | time | caller
# \      /      | \      /
<idle>-0      2d.h4      0us :      0:120:R  + [002]  5882: 94:R sleep
<idle>-0      2d.h4      0us : ttwu_do_activate.constprop.87 <-try_to_wake_up
<idle>-0      2d.h4      1us : sched_wakeup: comm=sleep pid=5882 prio=94
→ success=1 target_cpu=002
<idle>-0      2dNh2      1us : hrtimer_expire_exit: hrtimer=ffff88007796feb8
<idle>-0      2.N.2      2us : power_end: cpu_id=2
<idle>-0      2.N.2      3us : cpu_idle: state=4294967295 cpu_id=2
<idle>-0      2dN.3      4us : hrtimer_cancel: hrtimer=ffff88007d50d5e0
<idle>-0      2dN.3      4us : hrtimer_start: hrtimer=ffff88007d50d5e0
→ function=tick_sched_timer expires=34311211000000 softexpires=34311211000000
<idle>-0      2.N.2      5us : rcu_utilization: Start context switch
<idle>-0      2.N.2      5us : rcu_utilization: End context switch
<idle>-0      2d..3      6us : __schedule <-schedule
<idle>-0      2d..3      6us :      0:120:R ==> [002]  5882: 94:R sleep

```

3.16 Hardware Latency Detector

The hardware latency detector is executed by enabling the “hwlat” tracer.

NOTE, this tracer will affect the performance of the system as it will periodically make a CPU constantly busy with interrupts disabled.

```

# echo hwlat > current_tracer
# sleep 100
# cat trace
# tracer: hwlat
#
# entries-in-buffer/entries-written: 13/13   #P:8
#

```

```

#          -----=> irqs-off
#          /-----=> need-resched
#          | /-----=> hardirq/softirq
#          || /-----=> preempt-depth
#          ||| /-----=> delay
#          TASK-PID  CPU#  ||||  TIMESTAMP  FUNCTION
#          | |      | |   | |   |          |
<...>-1729 [001] d... 678.473449: #1    inner/outer(us): 11/
→12  ts:1581527483.343962693 count:6
<...>-1729 [004] d... 689.556542: #2    inner/outer(us): 16/9
→  ts:1581527494.889008092 count:1
<...>-1729 [005] d... 714.756290: #3    inner/outer(us): 16/
→16  ts:1581527519.678961629 count:5
<...>-1729 [001] d... 718.788247: #4    inner/outer(us): 9/
→17  ts:1581527523.889012713 count:1
<...>-1729 [002] d... 719.796341: #5    inner/outer(us): 13/9
→  ts:1581527524.912872606 count:1
<...>-1729 [006] d... 844.787091: #6    inner/outer(us): 9/
→12  ts:1581527649.889048502 count:2
<...>-1729 [003] d... 849.827033: #7    inner/outer(us): 18/9
→  ts:1581527654.889013793 count:1
<...>-1729 [007] d... 853.859002: #8    inner/outer(us): 9/
→12  ts:1581527658.889065736 count:1
<...>-1729 [001] d... 855.874978: #9    inner/outer(us): 9/
→11  ts:1581527660.861991877 count:1
<...>-1729 [001] d... 863.938932: #10   inner/outer(us): 9/
→11  ts:1581527668.970010500 count:1 nmi-total:7 nmi-count:1
<...>-1729 [007] d... 878.050780: #11   inner/outer(us): 9/
→12  ts:1581527683.385002600 count:1 nmi-total:5 nmi-count:1
<...>-1729 [007] d... 886.114702: #12   inner/outer(us): 9/
→12  ts:1581527691.385001600 count:1

```

The above output is somewhat the same in the header. All events will have interrupts disabled 'd'. Under the FUNCTION title there is:

#1 This is the count of events recorded that were greater than the tracing_threshold (See below).

inner/outer(us): 11/11

This shows two numbers as "inner latency" and "outer latency". The test runs in a loop checking a timestamp twice. The latency detected within the two timestamps is the "inner latency" and the latency detected after the previous timestamp and the next timestamp in the loop is the "outer latency".

ts:1581527483.343962693

The absolute timestamp that the first latency was recorded in the window.

count:6

The number of times a latency was detected during the window.

nmi-total:7 nmi-count:1

On architectures that support it, if an NMI comes in during the test, the time spent in NMI is reported in “nmi-total” (in microseconds).

All architectures that have NMIs will show the “nmi-count” if an NMI comes in during the test.

hwlat files:

tracing_threshold This gets automatically set to “10” to represent 10 microseconds. This is the threshold of latency that needs to be detected before the trace will be recorded.

Note, when hwlat tracer is finished (another tracer is written into “current_tracer”), the original value for tracing_threshold is placed back into this file.

hwlat_detector/width The length of time the test runs with interrupts disabled.

hwlat_detector/window The length of time of the window which the test runs. That is, the test will run for “width” microseconds per “window” microseconds

tracing_cpumask When the test is started. A kernel thread is created that runs the test. This thread will alternate between CPUs listed in the tracing_cpumask between each period (one “window”). To limit the test to specific CPUs set the mask in this file to only the CPUs that the test should run on.

3.17 function

This tracer is the function tracer. Enabling the function tracer can be done from the debug file system. Make sure the ftrace_enabled is set; otherwise this tracer is a nop. See the “ftrace_enabled” section below.

```
# sysctl kernel.ftrace_enabled=1
# echo function > current_tracer
# echo 1 > tracing_on
# usleep 1
# echo 0 > tracing_on
# cat trace
# tracer: function
#
# entries-in-buffer/entries-written: 24799/24799   #P:4
#
#          _-----=> irqsoff
#          /_-----=> need-resched
#          |/_-----=> hardirq/softirq
#          ||/_-----=> preempt-depth
#          |||/_-----=> delay
#
# TASK-PID   CPU#  | TIMESTAMP | FUNCTION
#   | |       |   |         |   |
bash-1994   [002]  ....    3082.063030: mutex_unlock <-rb_simple_write
bash-1994   [002]  ....    3082.063031: __mutex_unlock_slowpath <-
↪mutex_unlock
bash-1994   [002]  ....    3082.063031: __fsnotify_parent <-fsnotify_
↪modify
```

```

bash-1994 [002] .... 3082.063032: fsnotify <-fsnotify_modify
bash-1994 [002] .... 3082.063032: __srcu_read_lock <-fsnotify
bash-1994 [002] .... 3082.063032: add_preempt_count <-__srcu_
→read_lock
bash-1994 [002] ...1 3082.063032: sub_preempt_count <-__srcu_
→read_lock
bash-1994 [002] .... 3082.063033: __srcu_read_unlock <-fsnotify
[...]
```

Note: function tracer uses ring buffers to store the above entries. The newest data may overwrite the oldest data. Sometimes using echo to stop the trace is not sufficient because the tracing could have overwritten the data that you wanted to record. For this reason, it is sometimes better to disable tracing directly from a program. This allows you to stop the tracing at the point that you hit the part that you are interested in. To disable the tracing directly from a C program, something like following code snippet can be used:

```

int trace_fd;
[...]
int main(int argc, char *argv[]) {
    [...]
    trace_fd = open(tracing_file("tracing_on"), O_WRONLY);
    [...]
    if (condition_hit()) {
        write(trace_fd, "0", 1);
    }
    [...]
}
```

3.18 Single thread tracing

By writing into set_ftrace_pid you can trace a single thread. For example:

```

# cat set_ftrace_pid
no pid
# echo 3111 > set_ftrace_pid
# cat set_ftrace_pid
3111
# echo function > current_tracer
# cat trace | head
# tracer: function
#
#          TASK-PID    CPU#    TIMESTAMP    FUNCTION
#          | |         |         |         |
yum-updatesd-3111 [003]  1637.254676: finish_task_switch <-thread_return
yum-updatesd-3111 [003]  1637.254681: hrtimer_cancel <-schedule_hrtimeout_
→range
yum-updatesd-3111 [003]  1637.254682: hrtimer_try_to_cancel <-hrtimer_
→cancel
yum-updatesd-3111 [003]  1637.254683: lock_hrtimer_base <-hrtimer_try_to_
→cancel
```



```

yum-updatesd-3111 [003] 1637.254685: fget_light <-do_sys_poll
yum-updatesd-3111 [003] 1637.254686: pipe_poll <-do_sys_poll
# echo > set_fttrace_pid
# cat trace |head
# tracer: function
#
#           TASK-PID    CPU#    TIMESTAMP    FUNCTION
#           | |         |         |         |
##### CPU 3 buffer started #####
yum-updatesd-3111 [003] 1701.957688: free_poll_entry <-poll_freewait
yum-updatesd-3111 [003] 1701.957689: remove_wait_queue <-free_poll_entry
yum-updatesd-3111 [003] 1701.957691: fput <-free_poll_entry
yum-updatesd-3111 [003] 1701.957692: audit_syscall_exit <-sysret_audit
yum-updatesd-3111 [003] 1701.957693: path_put <-audit_syscall_exit

```

If you want to trace a function when executing, you could use something like this simple program.

```

#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
#include <string.h>

#define _STR(x) #x
#define STR(x) _STR(x)
#define MAX_PATH 256

const char *find_tracefs(void)
{
    static char tracefs[MAX_PATH+1];
    static int tracefs_found;
    char type[100];
    FILE *fp;

    if (tracefs_found)
        return tracefs;

    if ((fp = fopen("/proc/mounts", "r")) == NULL) {
        perror("/proc/mounts");
        return NULL;
    }

    while (fscanf(fp, "%*s %"
                  STR(MAX_PATH)
                  "s %99s %*s %*d %*d\n",
                  tracefs, type) == 2) {
        if (strcmp(type, "tracefs") == 0)
            break;
    }
}

```

```
    }
    fclose(fp);

    if (strcmp(type, "tracefs") != 0) {
        fprintf(stderr, "tracefs not mounted");
        return NULL;
    }

    strcat(tracefs, "/tracing/");
    tracefs_found = 1;

    return tracefs;
}

const char *tracing_file(const char *file_name)
{
    static char trace_file[MAX_PATH+1];
    snprintf(trace_file, MAX_PATH, "%s/%s", find_tracefs(), file_name);
    return trace_file;
}

int main (int argc, char **argv)
{
    if (argc < 1)
        exit(-1);

    if (fork() > 0) {
        int fd, ffd;
        char line[64];
        int s;

        ffd = open(tracing_file("current_tracer"), O_WRONLY);
        if (ffd < 0)
            exit(-1);
        write(ffd, "nop", 3);

        fd = open(tracing_file("set_ftrace_pid"), O_WRONLY);
        s = sprintf(line, "%d\n", getpid());
        write(fd, line, s);

        write(ffd, "function", 8);

        close(fd);
        close(ffd);

        execvp(argv[1], argv+1);
    }

    return 0;
}
```

Or this simple script!

```
#!/bin/bash

tracefs=`sed -ne 's/^tracefs \(.*\) tracefs.*\/\1/p' /proc/mounts`
echo 0 > $tracefs/tracing_on
echo $$ > $tracefs/set_ftrace_pid
echo function > $tracefs/current_tracer
echo 1 > $tracefs/tracing_on
exec "$@"
```

3.19 function graph tracer

This tracer is similar to the function tracer except that it probes a function on its entry and its exit. This is done by using a dynamically allocated stack of return addresses in each `task_struct`. On function entry the tracer overwrites the return address of each function traced to set a custom probe. Thus the original return address is stored on the stack of return address in the `task_struct`.

Probing on both ends of a function leads to special features such as:

- measure of a function's time execution
- having a reliable call stack to draw function calls graph

This tracer is useful in several situations:

- you want to find the reason of a strange kernel behavior and need to see what happens in detail on any areas (or specific ones).
- you are experiencing weird latencies but it's difficult to find its origin.
- you want to find quickly which path is taken by a specific function
- you just want to peek inside a working kernel and want to see what happens there.

```
# tracer: function_graph
#
# CPU    DURATION                FUNCTION CALLS
# |      |      |              | | | |
0)      |      |      sys_open() {
0)      |      |      do_sys_open() {
0)      |      |      getname() {
0)      |      |      kmem_cache_alloc() {
0)  1.382 us |      |      __might_sleep();
0)  2.478 us |      |      }
0)      |      |      strncpy_from_user() {
0)      |      |      might_fault() {
0)  1.389 us |      |      __might_sleep();
0)  2.553 us |      |      }
0)  3.807 us |      |      }
0)  7.876 us |      |      }
0)      |      |      alloc_fd() {
```

0)	0.668 us		_spin_lock();
0)	0.570 us		expand_files();
0)	0.586 us		_spin_unlock();

There are several columns that can be dynamically enabled/disabled. You can use every combination of options you want, depending on your needs.

- The cpu number on which the function executed is default enabled. It is sometimes better to only trace one cpu (see `tracing_cpu_mask` file) or you might sometimes see unordered function calls while cpu tracing switch.
 - hide: `echo nofuncgraph-cpu > trace_options`
 - show: `echo funcgraph-cpu > trace_options`
- The duration (function's time of execution) is displayed on the closing bracket line of a function or on the same line than the current function in case of a leaf one. It is default enabled.
 - hide: `echo nofuncgraph-duration > trace_options`
 - show: `echo funcgraph-duration > trace_options`
- The overhead field precedes the duration field in case of reached duration thresholds.
 - hide: `echo nofuncgraph-overhead > trace_options`
 - show: `echo funcgraph-overhead > trace_options`
 - depends on: `funcgraph-duration`

ie:

3)	#	1837.709 us		}	/* __switch_to */
3)				finish_task_switch()	{
3)	0.313 us			_raw_spin_unlock_irq();	
3)	3.177 us			}	
3)	#	1889.063 us		}	/* __schedule */
3)	!	140.417 us		}	/* __schedule */
3)	#	2034.948 us		}	/* schedule */
3)	*	33998.59 us		}	/* schedule_preempt_disabled */
[...]					
1)	0.260 us			msecs_to_jiffies();	
1)	0.313 us			__rcu_read_unlock();	
1)	+	61.770 us		}	
1)	+	64.479 us		}	
1)	0.313 us			rcu_bh_qs();	
1)	0.313 us			__local_bh_enable();	
1)	!	217.240 us		}	
1)	0.365 us			idle_cpu();	
1)				rcu_irq_exit()	{
1)	0.417 us			rcu_eqs_enter_common.isra.47();	
1)	3.125 us			}	
1)	!	227.812 us		}	

```

1) ! 457.395 us |    }
1) @ 119760.2 us |    }

[...]

2)                |    handle_IPI() {
1)    6.979 us    |    }
2)    0.417 us    |    scheduler_ipi();
1)    9.791 us    |    }
1) + 12.917 us    |    }
2)    3.490 us    |    }
1) + 15.729 us    |    }
1) + 18.542 us    |    }
2) $ 3594274 us   |    }

```

Flags:

```

+ means that the function exceeded 10 usecs.
! means that the function exceeded 100 usecs.
# means that the function exceeded 1000 usecs.
* means that the function exceeded 10 msecs.
@ means that the function exceeded 100 msecs.
$ means that the function exceeded 1 sec.

```

- The task/pid field displays the thread cmdline and pid which executed the function. It is default disabled.
 - hide: `echo nofuncgraph-proc > trace_options`
 - show: `echo funcgraph-proc > trace_options`

ie:

```

# tracer: function_graph
#
# CPU    TASK/PID          DURATION          FUNCTION CALLS
# |      |      |          |      |          |      |      |      |
0)      sh-4802          |          |          d_free() {
0)      sh-4802          |          |          call_rcu() {
0)      sh-4802          |          |          __call_rcu() {
0)      sh-4802          |  0.616 us  |          rcu_process_gp_
↪end();
0)      sh-4802          |  0.586 us  |          check_for_new_
↪grace_period();
0)      sh-4802          |  2.899 us  |          }
0)      sh-4802          |  4.040 us  |          }
0)      sh-4802          |  5.151 us  |          }
0)      sh-4802          | + 49.370 us |          }

```

- The absolute time field is an absolute timestamp given by the system clock since it started. A snapshot of this time is given on each entry/exit of functions
 - hide: `echo nofuncgraph-abstime > trace_options`

- show: echo funcgraph-abstime > trace_options

ie:

```
#
#      TIME      CPU  DURATION      FUNCTION CALLS
#      |         |   |         |         |         |
360.774522 | 1)  0.541 us      |
↪}
360.774522 | 1)  4.663 us      |
360.774523 | 1)  0.541 us      |
↪wake_up_bit();
360.774524 | 1)  6.796 us      |
360.774524 | 1)  7.952 us      |
360.774525 | 1)  9.063 us      |
360.774525 | 1)  0.615 us      |
↪mark_dirty();
360.774527 | 1)  0.578 us      |
↪brelse();
360.774528 | 1)
↪reiserfs_prepare_for_journal() {
360.774528 | 1)
↪unlock_buffer() {
360.774529 | 1)
↪wake_up_bit() {
360.774529 | 1)
↪bit_waitqueue() {
360.774530 | 1)  0.594 us      |
↪__phys_addr();
```

The function name is always displayed after the closing bracket for a function if the start of that function is not in the trace buffer.

Display of the function name after the closing bracket may be enabled for functions whose start is in the trace buffer, allowing easier searching with `grep` for function durations. It is default disabled.

- hide: echo nofuncgraph-tail > trace_options
- show: echo funcgraph-tail > trace_options

Example with nofuncgraph-tail (default):

```
0)      |      putname() {
0)      |      kmem_cache_free() {
0)  0.518 us |      __phys_addr();
0)  1.757 us |      }
0)  2.861 us |      }
```

Example with funcgraph-tail:

```
0)      |      putname() {
0)      |      kmem_cache_free() {
0)  0.518 us |      __phys_addr();
```

```
0)    1.757 us    |          } /* kmem_cache_free() */
0)    2.861 us    |          } /* putname() */
```

You can put some comments on specific functions by using `trace_printk()`. For example, if you want to put a comment inside the `__might_sleep()` function, you just have to include `<linux/ftrace.h>` and call `trace_printk()` inside `__might_sleep()`:

```
trace_printk("I'm a comment!\n")
```

will produce:

```
1)          |          __might_sleep() {
1)          |          /* I'm a comment! */
1)    1.449 us |          }
```

You might find other useful features for this tracer in the following “dynamic ftrace” section such as tracing only specific functions or tasks.

3.20 dynamic ftrace

If `CONFIG_DYNAMIC_FTRACE` is set, the system will run with virtually no overhead when function tracing is disabled. The way this works is the `mcount` function call (placed at the start of every kernel function, produced by the `-pg` switch in `gcc`), starts off pointing to a simple return. (Enabling `FTRACE` will include the `-pg` switch in the compiling of the kernel.)

At compile time every C file object is run through the `recordmcount` program (located in the `scripts` directory). This program will parse the ELF headers in the C object to find all the locations in the `.text` section that call `mcount`. Starting with `gcc` version 4.6, the `-mfentry` has been added for x86, which calls “`__fentry__`” instead of “`mcount`”. Which is called before the creation of the stack frame.

Note, not all sections are traced. They may be prevented by either a `notrace`, or blocked another way and all inline functions are not traced. Check the “`available_filter_functions`” file to see what functions can be traced.

A section called “`__mcount_loc`” is created that holds references to all the `mcount/fentry` call sites in the `.text` section. The `recordmcount` program re-links this section back into the original object. The final linking stage of the kernel will add all these references into a single table.

On boot up, before `SMP` is initialized, the dynamic ftrace code scans this table and updates all the locations into nops. It also records the locations, which are added to the `available_filter_functions` list. Modules are processed as they are loaded and before they are executed. When a module is unloaded, it also removes its functions from the ftrace function list. This is automatic in the module unload code, and the module author does not need to worry about it.

When tracing is enabled, the process of modifying the function tracepoints is dependent on architecture. The old method is to use `kstop_machine` to prevent races with the CPUs executing code being modified (which can cause the CPU to do undesirable things, especially if the modified code crosses cache (or page) boundaries), and the nops are patched back to calls. But this time, they do not call `mcount` (which is just a function stub). They now call into the ftrace infrastructure.

The new method of modifying the function tracepoints is to place a breakpoint at the location to be modified, sync all CPUs, modify the rest of the instruction not covered by the breakpoint. Sync all CPUs again, and then remove the breakpoint with the finished version to the ftrace call site.

Some archs do not even need to monkey around with the synchronization, and can just slap the new code on top of the old without any problems with other CPUs executing it at the same time.

One special side-effect to the recording of the functions being traced is that we can now selectively choose which functions we wish to trace and which ones we want the mcount calls to remain as nops.

Two files are used, one for enabling and one for disabling the tracing of specified functions. They are:

set_ftrace_filter

and

set_ftrace_notrace

A list of available functions that you can add to these files is listed in:

available_filter_functions

```
# cat available_filter_functions
put_prev_task_idle
kmem_cache_create
pick_next_task_rt
cpus_read_lock
pick_next_task_fair
mutex_lock
[...]
```

If I am only interested in sys_nanosleep and hrtimer_interrupt:

```
# echo sys_nanosleep hrtimer_interrupt > set_ftrace_filter
# echo function > current_tracer
# echo 1 > tracing_on
# usleep 1
# echo 0 > tracing_on
# cat trace
# tracer: function
#
# entries-in-buffer/entries-written: 5/5   #P:4
#
#          _-----=> irqsoff
#          /_-----=> need-resched
#          |/_-----=> hardirq/softirq
#          ||/_-----=> preempt-depth
#          |||/_-----=> delay
#          TASK-PID   CPU#  | TIMESTAMP | FUNCTION
#          | |        |   |   |   |         |
#          usleep-2665 [001] ....  4186.475355: sys_nanosleep <-system_call_
->fastpath
#          <idle>-0    [001] d.h1  4186.475409: hrtimer_interrupt <-smp_apic_
->timer_interrupt
```



```

        usleep-2665 [001] d.h1 4186.475426: hrtimer_interrupt <-smp_apic_
→timer_interrupt
        <idle>-0    [003] d.h1 4186.475426: hrtimer_interrupt <-smp_apic_
→timer_interrupt
        <idle>-0    [002] d.h1 4186.475427: hrtimer_interrupt <-smp_apic_
→timer_interrupt

```

To see which functions are being traced, you can cat the file:

```

# cat set_ftrace_filter
hrtimer_interrupt
sys_nanosleep

```

Perhaps this is not enough. The filters also allow glob(7) matching.

<match>* will match functions that begin with <match>

***<match>** will match functions that end with <match>

<match> will match functions that have <match> in it

<match1>*<match2> will match functions that begin with <match1> and end with <match2>

Note: It is better to use quotes to enclose the wild cards, otherwise the shell may expand the parameters into names of files in the local directory.

```

# echo 'hrtimer_*' > set_ftrace_filter

```

Produces:

```

# tracer: function
#
# entries-in-buffer/entries-written: 897/897   #P:4
#
#          _-----=> irqs-off
#          /_-----=> need-resched
#          | /_----=> hardirq/softirq
#          || /_--=> preempt-depth
#          ||| /_   delay
#          TASK-PID  CPU#  ||||   TIMESTAMP  FUNCTION
#          |   |   |   |   |   |   |
        <idle>-0    [003] dN.1  4228.547803: hrtimer_cancel <-tick_nohz_
→idle_exit
        <idle>-0    [003] dN.1  4228.547804: hrtimer_try_to_cancel <-
→hrtimer_cancel
        <idle>-0    [003] dN.2  4228.547805: hrtimer_force_reprogram <-__
→remove_hrtimer
        <idle>-0    [003] dN.1  4228.547805: hrtimer_forward <-tick_nohz_
→idle_exit
        <idle>-0    [003] dN.1  4228.547805: hrtimer_start_range_ns <-
→hrtimer_start_expires.constprop.11

```

```
        <idle>-0      [003] d..1  4228.547858: hrtimer_get_next_event <-get_
→next_timer_interrupt
        <idle>-0      [003] d..1  4228.547859: hrtimer_start <-__tick_nohz_
→idle_enter
        <idle>-0      [003] d..2  4228.547860: hrtimer_force_reprogram <-__rem
```

Notice that we lost the `sys_nanosleep`.

```
# cat set_ftrace_filter
hrtimer_run_queues
hrtimer_run_pending
hrtimer_init
hrtimer_cancel
hrtimer_try_to_cancel
hrtimer_forward
hrtimer_start
hrtimer_reprogram
hrtimer_force_reprogram
hrtimer_get_next_event
hrtimer_interrupt
hrtimer_nanosleep
hrtimer_wakeup
hrtimer_get_remaining
hrtimer_get_res
hrtimer_init_sleeper
```

This is because the `>` and `>>` act just like they do in bash. To rewrite the filters, use `>` To append to the filters, use `>>`

To clear out a filter so that all functions will be recorded again:

```
# echo > set_ftrace_filter
# cat set_ftrace_filter
#
```

Again, now we want to append.

```
# echo sys_nanosleep > set_ftrace_filter
# cat set_ftrace_filter
sys_nanosleep
# echo 'hrtimer_*' >> set_ftrace_filter
# cat set_ftrace_filter
hrtimer_run_queues
hrtimer_run_pending
hrtimer_init
hrtimer_cancel
hrtimer_try_to_cancel
hrtimer_forward
hrtimer_start
hrtimer_reprogram
hrtimer_force_reprogram
hrtimer_get_next_event
```

```
hrtimer_interrupt
sys_nanosleep
hrtimer_nanosleep
hrtimer_wakeup
hrtimer_get_remaining
hrtimer_get_res
hrtimer_init_sleeper
```

The `set_ftrace_notrace` prevents those functions from being traced.

```
# echo '*preempt*' '*lock*' > set_ftrace_notrace
```

Produces:

```
# tracer: function
#
# entries-in-buffer/entries-written: 39608/39608   #P:4
#
#          _-----=> irqs-off
#          / _-----=> need-resched
#          | / _---=> hardirq/softirq
#          || / _--=> preempt-depth
#          ||| /      delay
#          |||| /
#          TASK-PID  CPU#  ||||  TIMESTAMP  FUNCTION
#          | |       |   ||||  |           |
bash-1994  [000]  ....   4342.324896: file_ra_state_init <-do_dentry_
↪open
bash-1994  [000]  ....   4342.324897: open_check_o_direct <-do_last
bash-1994  [000]  ....   4342.324897: ima_file_check <-do_last
bash-1994  [000]  ....   4342.324898: process_measurement <-ima_file_
↪check
bash-1994  [000]  ....   4342.324898: ima_get_action <-process_
↪measurement
bash-1994  [000]  ....   4342.324898: ima_match_policy <-ima_get_
↪action
bash-1994  [000]  ....   4342.324899: do_truncate <-do_last
bash-1994  [000]  ....   4342.324899: should_remove_suid <-do_
↪truncate
bash-1994  [000]  ....   4342.324899: notify_change <-do_truncate
bash-1994  [000]  ....   4342.324900: current_fs_time <-notify_change
bash-1994  [000]  ....   4342.324900: current_kernel_time <-current_
↪fs_time
bash-1994  [000]  ....   4342.324900: timespec_trunc <-current_fs_
↪time
```

We can see that there's no more lock or preempt tracing.

3.21 Selecting function filters via index

Because processing of strings is expensive (the address of the function needs to be looked up before comparing to the string being passed in), an index can be used as well to enable functions. This is useful in the case of setting thousands of specific functions at a time. By passing in a list of numbers, no string processing will occur. Instead, the function at the specific location in the internal array (which corresponds to the functions in the “available_filter_functions” file), is selected.

```
# echo 1 > set_ftrace_filter
```

Will select the first function listed in “available_filter_functions”

```
# head -1 available_filter_functions
trace_initcall_finish_cb

# cat set_ftrace_filter
trace_initcall_finish_cb

# head -50 available_filter_functions | tail -1
x86_pmu_commit_txn

# echo 1 50 > set_ftrace_filter
# cat set_ftrace_filter
trace_initcall_finish_cb
x86_pmu_commit_txn
```

3.22 Dynamic ftrace with the function graph tracer

Although what has been explained above concerns both the function tracer and the function-graph-tracer, there are some special features only available in the function-graph tracer.

If you want to trace only one function and all of its children, you just have to echo its name into `set_graph_function`:

```
echo __do_fault > set_graph_function
```

will produce the following “expanded” trace of the `__do_fault()` function:

```
0)          | __do_fault() {
0)          |   filemap_fault() {
0)          |     find_lock_page() {
0)  0.804 us |       find_get_page();
0)          |       __might_sleep() {
0)  1.329 us |         }
0)  3.904 us |       }
0)  4.979 us |     }
0)  0.653 us |   _spin_lock();
0)  0.578 us |   page_add_file_rmap();
0)  0.525 us |   native_set_pte_at();
```

```

0) 0.585 us | _spin_unlock();
0)          | unlock_page() {
0) 0.541 us |     page_waitqueue();
0) 0.639 us |     __wake_up_bit();
0) 2.786 us | }
0) + 14.237 us | }
0)          | __do_fault() {
0)          |     filemap_fault() {
0)          |         find_lock_page() {
0) 0.698 us |             find_get_page();
0)          |             __might_sleep() {
0) 1.412 us |                 }
0) 3.950 us |             }
0) 5.098 us |         }
0) 0.631 us |     _spin_lock();
0) 0.571 us |     page_add_file_rmap();
0) 0.526 us |     native_set_pte_at();
0) 0.586 us |     _spin_unlock();
0)          |     unlock_page() {
0) 0.533 us |         page_waitqueue();
0) 0.638 us |         __wake_up_bit();
0) 2.793 us |     }
0) + 14.012 us | }

```

You can also expand several functions at once:

```

echo sys_open > set_graph_function
echo sys_close >> set_graph_function

```

Now if you want to go back to trace all functions you can clear this special filter via:

```

echo > set_graph_function

```

3.23 ftrace_enabled

Note, the `proc sysctl ftrace_enable` is a big on/off switch for the function tracer. By default it is enabled (when function tracing is enabled in the kernel). If it is disabled, all function tracing is disabled. This includes not only the function tracers for `ftrace`, but also for any other uses (perf, kprobes, stack tracing, profiling, etc). It cannot be disabled if there is a callback with `FTRACE_OPS_FL_PERMANENT` set registered.

Please disable this with care.

This can be disabled (and enabled) with:

```

sysctl kernel.ftrace_enabled=0
sysctl kernel.ftrace_enabled=1

```

or

```
echo 0 > /proc/sys/kernel/ftrace_enabled
echo 1 > /proc/sys/kernel/ftrace_enabled
```

3.24 Filter commands

A few commands are supported by the `set_ftrace_filter` interface. Trace commands have the following format:

```
<function>:<command>:<parameter>
```

The following commands are supported:

- `mod`: This command enables function filtering per module. The parameter defines the module. For example, if only the `write*` functions in the `ext3` module are desired, run:

```
echo 'write*:mod:ext3' > set_ftrace_filter
```

This command interacts with the filter in the same way as filtering based on function names. Thus, adding more functions in a different module is accomplished by appending (`>>`) to the filter file. Remove specific module functions by prepending `!`:

```
echo '!writeback*:mod:ext3' >> set_ftrace_filter
```

`Mod` command supports module globbing. Disable tracing for all functions except a specific module:

```
echo '!*:mod:!ext3' >> set_ftrace_filter
```

Disable tracing for all modules, but still trace kernel:

```
echo '!*:mod:*' >> set_ftrace_filter
```

Enable filter only for kernel:

```
echo '*write*:mod:!*' >> set_ftrace_filter
```

Enable filter for module globbing:

```
echo '*write*:mod:*snd*' >> set_ftrace_filter
```

- `tracemon/tracemonoff`: These commands turn tracing on and off when the specified functions are hit. The parameter determines how many times the tracing system is turned on and off. If unspecified, there is no limit. For example, to disable tracing when a `schedule bug` is hit the first 5 times, run:

```
echo '__schedule_bug:tracemonoff:5' > set_ftrace_filter
```

To always disable tracing when `__schedule_bug` is hit:

```
echo '__schedule_bug:tracemonoff' > set_ftrace_filter
```

These commands are cumulative whether or not they are appended to `set_fttrace_filter`. To remove a command, prepend it by `!` and drop the parameter:

```
echo '!__schedule_bug:traceoff:0' > set_fttrace_filter
```

The above removes the `traceoff` command for `__schedule_bug` that have a counter. To remove commands without counters:

```
echo '!__schedule_bug:traceoff' > set_fttrace_filter
```

- `snapshot`: Will cause a snapshot to be triggered when the function is hit.

```
echo 'native_flush_tlb_others:snapshot' > set_fttrace_filter
```

To only snapshot once:

```
echo 'native_flush_tlb_others:snapshot:1' > set_fttrace_filter
```

To remove the above commands:

```
echo '!native_flush_tlb_others:snapshot' > set_fttrace_filter
echo '!native_flush_tlb_others:snapshot:0' > set_fttrace_filter
```

- `enable_event/disable_event`: These commands can enable or disable a trace event. Note, because function tracing callbacks are very sensitive, when these commands are registered, the trace point is activated, but disabled in a “soft” mode. That is, the tracepoint will be called, but just will not be traced. The event tracepoint stays in this mode as long as there’s a command that triggers it.

```
echo 'try_to_wake_up:enable_event:sched:sched_switch:2' > \
    set_fttrace_filter
```

The format is:

```
<function>:enable_event:<system>:<event>[:count]
<function>:disable_event:<system>:<event>[:count]
```

To remove the events commands:

```
echo '!try_to_wake_up:enable_event:sched:sched_switch:0' > \
    set_fttrace_filter
echo '!schedule:disable_event:sched:sched_switch' > \
    set_fttrace_filter
```

- `dump`: When the function is hit, it will dump the contents of the `ftrace` ring buffer to the console. This is useful if you need to debug something, and want to dump the trace when a certain function is hit. Perhaps it’s a function that is called before a triple fault happens and does not allow you to get a regular dump.
- `cpudump`: When the function is hit, it will dump the contents of the `ftrace` ring buffer for the current CPU to the console. Unlike the “dump” command, it only prints out the contents of the ring buffer for the CPU that executed the function that triggered the dump.
- `stacktrace`: When the function is hit, a stack trace is recorded.

3.25 trace_pipe

The `trace_pipe` outputs the same content as the trace file, but the effect on the tracing is different. Every read from `trace_pipe` is consumed. This means that subsequent reads will be different. The trace is live.

```
# echo function > current_tracer
# cat trace_pipe > /tmp/trace.out &
[1] 4153
# echo 1 > tracing_on
# usleep 1
# echo 0 > tracing_on
# cat trace
# tracer: function
#
# entries-in-buffer/entries-written: 0/0   #P:4
#
#          _-----=> irqs-off
#         /_-----=> need-resched
#        |/_-----=> hardirq/softirq
#       ||/_-----=> preempt-depth
#      |||/_-----=> delay
#     TASK-PID   CPU#  | |||   TIMESTAMP   FUNCTION
#     | |       |   |   | |||   |         |
#
#
# cat /tmp/trace.out
bash-1994 [000] .... 5281.568961: mutex_unlock <-rb_simple_write
bash-1994 [000] .... 5281.568963: __mutex_unlock_slowpath <-mutex_
→unlock
bash-1994 [000] .... 5281.568963: __fsnotify_parent <-fsnotify_
→modify
bash-1994 [000] .... 5281.568964: fsnotify <-fsnotify_modify
bash-1994 [000] .... 5281.568964: __srcu_read_lock <-fsnotify
bash-1994 [000] .... 5281.568964: add_preempt_count <-__srcu_read_
→lock
bash-1994 [000] ...1 5281.568965: sub_preempt_count <-__srcu_read_
→lock
bash-1994 [000] .... 5281.568965: __srcu_read_unlock <-fsnotify
bash-1994 [000] .... 5281.568967: sys_dup2 <-system_call_fastpath
```

Note, reading the `trace_pipe` file will block until more input is added. This is contrary to the trace file. If any process opened the trace file for reading, it will actually disable tracing and prevent new entries from being added. The `trace_pipe` file does not have this limitation.

3.26 trace entries

Having too much or not enough data can be troublesome in diagnosing an issue in the kernel. The file `buffer_size_kb` is used to modify the size of the internal trace buffers. The number listed is the number of entries that can be recorded per CPU. To know the full size, multiply the number of possible CPUs with the number of entries.

```
# cat buffer_size_kb
1408 (units kilobytes)
```

Or simply read `buffer_total_size_kb`

```
# cat buffer_total_size_kb
5632
```

To modify the buffer, simple echo in a number (in 1024 byte segments).

```
# echo 10000 > buffer_size_kb
# cat buffer_size_kb
10000 (units kilobytes)
```

It will try to allocate as much as possible. If you allocate too much, it can cause Out-Of-Memory to trigger.

```
# echo 10000000000000 > buffer_size_kb
-bash: echo: write error: Cannot allocate memory
# cat buffer_size_kb
85
```

The `per_cpu` buffers can be changed individually as well:

```
# echo 10000 > per_cpu/cpu0/buffer_size_kb
# echo 100 > per_cpu/cpu1/buffer_size_kb
```

When the `per_cpu` buffers are not the same, the `buffer_size_kb` at the top level will just show an X

```
# cat buffer_size_kb
X
```

This is where the `buffer_total_size_kb` is useful:

```
# cat buffer_total_size_kb
12916
```

Writing to the top level `buffer_size_kb` will reset all the buffers to be the same again.

3.27 Snapshot

CONFIG_TRACER_SNAPSHOT makes a generic snapshot feature available to all non latency tracers. (Latency tracers which record max latency, such as “irqsoff” or “wakeup”, can’t use this feature, since those are already using the snapshot mechanism internally.)

Snapshot preserves a current trace buffer at a particular point in time without stopping tracing. Ftrace swaps the current buffer with a spare buffer, and tracing continues in the new current (=previous spare) buffer.

The following tracefs files in “tracing” are related to this feature:

snapshot:

This is used to take a snapshot and to read the output of the snapshot. Echo 1 into this file to allocate a spare buffer and to take a snapshot (swap), then read the snapshot from this file in the same format as “trace” (described above in the section “The File System”). Both reads snapshot and tracing are executable in parallel. When the spare buffer is allocated, echoing 0 frees it, and echoing else (positive) values clear the snapshot contents. More details are shown in the table below.

status\input	0	1	else
not allocated	(do nothing)	alloc+swap	(do nothing)
allocated	free	swap	clear

Here is an example of using the snapshot feature.

```
# echo 1 > events/sched/enabled
# echo 1 > snapshot
# cat snapshot
# tracer: nop
#
# entries-in-buffer/entries-written: 71/71   #P:8
#
#          _-----=> irqsoff
#          /_-----=> need-resched
#          | /_-----=> hardirq/softirq
#          || /_-----=> preempt-depth
#          ||| /_-----=> delay
#          TASK-PID   CPU#  | TIMESTAMP | FUNCTION
#          | |       | |   | |          | |
<idle>-0    [005] d...  2440.603828: sched_switch: prev_
->comm=swapper/5 prev_pid=0 prev_prio=120 prev_state=R ==> next_
->comm=snapshot-test-2 next_pid=2242 next_prio=120
      sleep-2242 [005] d...  2440.603846: sched_switch: prev_
->comm=snapshot-test-2 prev_pid=2242 prev_prio=120 prev_state=R ==> next_
->comm=kworker/5:1 next_pid=60 next_prio=120
[...]
```

```

<idle>-0    [002] d...  2440.707230: sched_switch: prev_comm=swapper/
->2 prev_pid=0 prev_prio=120 prev_state=R ==> next_comm=snapshot-test-2 next_
->pid=2229 next_prio=120
```

```
# cat trace
# tracer: nop
#
# entries-in-buffer/entries-written: 77/77   #P:8
#
#
#          _-----=> irqsoft
#         /_-----=> need-resched
#        |/_-----=> hardirq/softirq
#       ||/_-----=> preempt-depth
#      |||/_-----=> delay
#     TASK-PID  CPU#  |         |   TIMESTAMP  FUNCTION
#     |   |   |   |   |   |   |   |   |
# <idle>-0     [007] d... 2440.707395: sched_switch: prev_
->comm=swapper/7 prev_pid=0 prev_prio=120 prev_state=R ==> next_comm=snapshot-
->test-2 next_pid=2243 next_prio=120
snapshot-test-2-2229 [002] d... 2440.707438: sched_switch: prev_
->comm=snapshot-test-2 prev_pid=2229 prev_prio=120 prev_state=S ==> next_
->comm=swapper/2 next_pid=0 next_prio=120
[...]
```

If you try to use this snapshot feature when current tracer is one of the latency tracers, you will get the following results.

```
# echo wakeup > current_tracer
# echo 1 > snapshot
bash: echo: write error: Device or resource busy
# cat snapshot
cat: snapshot: Device or resource busy
```

3.28 Instances

In the tracefs tracing directory, there is a directory called “instances”. This directory can have new directories created inside of it using `mkdir`, and removing directories with `rmdir`. The directory created with `mkdir` in this directory will already contain files and other directories after it is created.

```
# mkdir instances/foo
# ls instances/foo
buffer_size_kb  buffer_total_size_kb  events  free_buffer  per_cpu
set_event      snapshot  trace   trace_clock  trace_marker  trace_options
trace_pipe     tracing_on
```

As you can see, the new directory looks similar to the tracing directory itself. In fact, it is very similar, except that the buffer and events are agnostic from the main directory, or from any other instances that are created.

The files in the new directory work just like the files with the same name in the tracing directory except the buffer that is used is a separate and new buffer. The files affect that buffer but do not affect the main buffer with the exception of `trace_options`. Currently, the `trace_options` affect

all instances and the top level buffer the same, but this may change in future releases. That is, options may become specific to the instance they reside in.

Notice that none of the function tracer files are there, nor is `current_tracer` and `available_tracers`. This is because the buffers can currently only have events enabled for them.

```
# mkdir instances/foo
# mkdir instances/bar
# mkdir instances/zoot
# echo 100000 > buffer_size_kb
# echo 1000 > instances/foo/buffer_size_kb
# echo 5000 > instances/bar/per_cpu/cpu1/buffer_size_kb
# echo function > current_trace
# echo 1 > instances/foo/events/sched/sched_wakeup/enable
# echo 1 > instances/foo/events/sched/sched_wakeup_new/enable
# echo 1 > instances/foo/events/sched/sched_switch/enable
# echo 1 > instances/bar/events/irq/enable
# echo 1 > instances/zoot/events/syscalls/enable
# cat trace_pipe
CPU:2 [LOST 11745 EVENTS]
        bash-2044 [002] .... 10594.481032: _raw_spin_lock_irqsave <-get_
→page_from_freelist
        bash-2044 [002] d... 10594.481032: add_preempt_count <-_raw_spin_
→lock_irqsave
        bash-2044 [002] d..1 10594.481032: __rmqueue <-get_page_from_
→freelist
        bash-2044 [002] d..1 10594.481033: _raw_spin_unlock <-get_page_
→from_freelist
        bash-2044 [002] d..1 10594.481033: sub_preempt_count <-_raw_spin_
→unlock
        bash-2044 [002] d... 10594.481033: get_pageblock_flags_group <-
→get_pageblock_migratetype
        bash-2044 [002] d... 10594.481034: __mod_zone_page_state <-get_
→page_from_freelist
        bash-2044 [002] d... 10594.481034: zone_statistics <-get_page_
→from_freelist
        bash-2044 [002] d... 10594.481034: __inc_zone_state <-zone_
→statistics
        bash-2044 [002] d... 10594.481034: __inc_zone_state <-zone_
→statistics
        bash-2044 [002] .... 10594.481035: arch_dup_task_struct <-copy_
→process
[...]
```

```
# cat instances/foo/trace_pipe
        bash-1998 [000] d..4   136.676759: sched_wakeup: comm=kworker/0:1
→pid=59 prio=120 success=1 target_cpu=000
        bash-1998 [000] dN.4   136.676760: sched_wakeup: comm=bash
→pid=1998 prio=120 success=1 target_cpu=000
        <idle>-0    [003] d.h3    136.676906: sched_wakeup: comm=rcu_preempt
→pid=9 prio=120 success=1 target_cpu=003
        <idle>-0    [003] d..3    136.676909: sched_switch: prev_
→comm=swapper/3 prev_pid=0 prev_prio=120 prev_state=R ==> next_comm=rcu_
→preempt next_pid=9 next_prio=120
```

```

    rcu_preempt-9      [003] d..3   136.676916: sched_switch: prev_comm=rcu_
→preempt prev_pid=9 prev_prio=120 prev_state=S ==> next_comm=swapper/3 next_
→pid=0 next_prio=120
        bash-1998     [000] d..4   136.677014: sched_wakeup: comm=kworker/0:1_
→pid=59 prio=120 success=1 target_cpu=000
        bash-1998     [000] dN.4   136.677016: sched_wakeup: comm=bash_
→pid=1998 prio=120 success=1 target_cpu=000
        bash-1998     [000] d..3   136.677018: sched_switch: prev_comm=bash_
→prev_pid=1998 prev_prio=120 prev_state=R+ ==> next_comm=kworker/0:1 next_
→pid=59 next_prio=120
        kworker/0:1-59 [000] d..4   136.677022: sched_wakeup: comm=sshd_
→pid=1995 prio=120 success=1 target_cpu=001
        kworker/0:1-59 [000] d..3   136.677025: sched_switch: prev_
→comm=kworker/0:1 prev_pid=59 prev_prio=120 prev_state=S ==> next_comm=bash_
→next_pid=1998 next_prio=120
[...]
```

```

# cat instances/bar/trace_pipe
    migration/1-14     [001] d.h3   138.732674: softirq_raise: vec=3_
→[action=NET_RX]
        <idle>-0       [001] dNh3   138.732725: softirq_raise: vec=3_
→[action=NET_RX]
        bash-1998     [000] d.h1   138.733101: softirq_raise: vec=1_
→[action=TIMER]
        bash-1998     [000] d.h1   138.733102: softirq_raise: vec=9_
→[action=RCU]
        bash-1998     [000] ..s2   138.733105: softirq_entry: vec=1_
→[action=TIMER]
        bash-1998     [000] ..s2   138.733106: softirq_exit: vec=1_
→[action=TIMER]
        bash-1998     [000] ..s2   138.733106: softirq_entry: vec=9_
→[action=RCU]
        bash-1998     [000] ..s2   138.733109: softirq_exit: vec=9_
→[action=RCU]
        sshd-1995     [001] d.h1   138.733278: irq_handler_entry: irq=21_
→name=uhci_hcd:usb4
        sshd-1995     [001] d.h1   138.733280: irq_handler_exit: irq=21_
→ret=unhandled
        sshd-1995     [001] d.h1   138.733281: irq_handler_entry: irq=21_
→name=eth0
        sshd-1995     [001] d.h1   138.733283: irq_handler_exit: irq=21_
→ret=handled
[...]
```

```

# cat instances/zoot/trace
# tracer: nop
#
# entries-in-buffer/entries-written: 18996/18996   #P:4
#
#
# _-----=> irqs-off
```

```

#           / _----=> need-resched
#           | / _----=> hardirq/softirq
#           || / _----=> preempt-depth
#           ||| / _----=> delay
#           TASK-PID  CPU#  ||||  TIMESTAMP  FUNCTION
#           | |      |  ||||  |           |
bash-1998  [000] d...  140.733501: sys_write -> 0x2
bash-1998  [000] d...  140.733504: sys_dup2(oldfd: a, newfd: 1)
bash-1998  [000] d...  140.733506: sys_dup2 -> 0x1
bash-1998  [000] d...  140.733508: sys_fcntl(fd: a, cmd: 1, arg: 0)
bash-1998  [000] d...  140.733509: sys_fcntl -> 0x1
bash-1998  [000] d...  140.733510: sys_close(fd: a)
bash-1998  [000] d...  140.733510: sys_close -> 0x0
bash-1998  [000] d...  140.733514: sys_rt_sigprocmask(how: 0, nset: 0, oset: 6e2768, sigsetsize: 8)
bash-1998  [000] d...  140.733515: sys_rt_sigprocmask -> 0x0
bash-1998  [000] d...  140.733516: sys_rt_sigaction(sig: 2, act: 7fff718846f0, oact: 7fff71884650, sigsetsize: 8)
bash-1998  [000] d...  140.733516: sys_rt_sigaction -> 0x0

```

You can see that the trace of the top most trace buffer shows only the function tracing. The foo instance displays wakeups and task switches.

To remove the instances, simply delete their directories:

```

# rmdir instances/foo
# rmdir instances/bar
# rmdir instances/zoot

```

Note, if a process has a trace file open in one of the instance directories, the rmdir will fail with EBUSY.

3.29 Stack trace

Since the kernel has a fixed sized stack, it is important not to waste it in functions. A kernel developer must be conscience of what they allocate on the stack. If they add too much, the system can be in danger of a stack overflow, and corruption will occur, usually leading to a system panic.

There are some tools that check this, usually with interrupts periodically checking usage. But if you can perform a check at every function call that will become very useful. As ftrace provides a function tracer, it makes it convenient to check the stack size at every function call. This is enabled via the stack tracer.

CONFIG_STACK_TRACER enables the ftrace stack tracing functionality. To enable it, write a '1' into /proc/sys/kernel/stack_tracer_enabled.

```

# echo 1 > /proc/sys/kernel/stack_tracer_enabled

```

You can also enable it from the kernel command line to trace the stack size of the kernel during boot up, by adding “stacktrace” to the kernel command line parameter.

After running it for a few minutes, the output looks like:

```
# cat stack_max_size
2928

# cat stack_trace
      Depth    Size  Location      (18 entries)
      -----
0)      2928     224  update_sd_lb_stats+0xbc/0x4ac
1)      2704     160  find_busiest_group+0x31/0x1f1
2)      2544     256  load_balance+0xd9/0x662
3)      2288      80  idle_balance+0xbb/0x130
4)      2208     128  __schedule+0x26e/0x5b9
5)      2080      16  schedule+0x64/0x66
6)      2064     128  schedule_timeout+0x34/0xe0
7)      1936     112  wait_for_common+0x97/0xf1
8)      1824      16  wait_for_completion+0x1d/0x1f
9)      1808     128  flush_work+0xfe/0x119
10)     1680      16  tty_flush_to_ldisc+0x1e/0x20
11)     1664      48  input_available_p+0x1d/0x5c
12)     1616      48  n_tty_poll+0x6d/0x134
13)     1568      64  tty_poll+0x64/0x7f
14)     1504     880  do_select+0x31e/0x511
15)       624     400  core_sys_select+0x177/0x216
16)       224      96  sys_select+0x91/0xb9
17)       128     128  system_call_fastpath+0x16/0x1b
```

Note, if -mfentry is being used by gcc, functions get traced before they set up the stack frame. This means that leaf level functions are not tested by the stack tracer when -mfentry is used.

Currently, -mfentry is used by gcc 4.6.0 and above on x86 only.

3.30 More

More details can be found in the source code, in the *kernel/trace/*.c* files.

USING FTRACE TO HOOK TO FUNCTIONS

Written for: 4.14

4.1 Introduction

The ftrace infrastructure was originally created to attach callbacks to the beginning of functions in order to record and trace the flow of the kernel. But callbacks to the start of a function can have other use cases. Either for live kernel patching, or for security monitoring. This document describes how to use ftrace to implement your own function callbacks.

4.2 The ftrace context

Warning: The ability to add a callback to almost any function within the kernel comes with risks. A callback can be called from any context (normal, softirq, irq, and NMI). Callbacks can also be called just before going to idle, during CPU bring up and takedown, or going to user space. This requires extra care to what can be done inside a callback. A callback can be called outside the protective scope of RCU.

There are helper functions to help against recursion, and making sure RCU is watching. These are explained below.

4.3 The ftrace_ops structure

To register a function callback, a ftrace_ops is required. This structure is used to tell ftrace what function should be called as the callback as well as what protections the callback will perform and not require ftrace to handle.

There is only one field that is needed to be set when registering an ftrace_ops with ftrace:

```
struct ftrace_ops ops = {
    .func           = my_callback_func,
    .flags          = MY_FTRACE_FLAGS
    .private        = any_private_data_structure,
};
```

Both `.flags` and `.private` are optional. Only `.func` is required.

To enable tracing call:

```
register_ftrace_function(&ops);
```

To disable tracing call:

```
unregister_ftrace_function(&ops);
```

The above is defined by including the header:

```
#include <linux/ftrace.h>
```

The registered callback will start being called some time after the `register_ftrace_function()` is called and before it returns. The exact time that callbacks start being called is dependent upon architecture and scheduling of services. The callback itself will have to handle any synchronization if it must begin at an exact moment.

The `unregister_ftrace_function()` will guarantee that the callback is no longer being called by functions after the `unregister_ftrace_function()` returns. Note that to perform this guarantee, the `unregister_ftrace_function()` may take some time to finish.

4.4 The callback function

The prototype of the callback function is as follows (as of v4.14):

```
void callback_func(unsigned long ip, unsigned long parent_ip,  
                  struct ftrace_ops *op, struct pt_regs *regs);
```

@ip This is the instruction pointer of the function that is being traced. (where the fentry or mcount is within the function)

@parent_ip This is the instruction pointer of the function that called the the function being traced (where the call of the function occurred).

@op This is a pointer to `ftrace_ops` that was used to register the callback. This can be used to pass data to the callback via the private pointer.

@regs If the `FTRACE_OPS_FL_SAVE_REGS` or `FTRACE_OPS_FL_SAVE_REGS_IF_SUPPORTED` flags are set in the `ftrace_ops` structure, then this will be pointing to the `pt_regs` structure like it would be if an breakpoint was placed at the start of the function where ftrace was tracing. Otherwise it either contains garbage, or NULL.

4.5 Protect your callback

As functions can be called from anywhere, and it is possible that a function called by a callback may also be traced, and call that same callback, recursion protection must be used. There are two helper functions that can help in this regard. If you start your code with:

```
int bit;

bit = ftrace_test_recursion_trylock(ip, parent_ip);
if (bit < 0)
    return;
```

and end it with:

```
ftrace_test_recursion_unlock(bit);
```

The code in between will be safe to use, even if it ends up calling a function that the callback is tracing. Note, on success, `ftrace_test_recursion_trylock()` will disable preemption, and the `ftrace_test_recursion_unlock()` will enable it again (if it was previously enabled). The instruction pointer (`ip`) and its parent (`parent_ip`) is passed to `ftrace_test_recursion_trylock()` to record where the recursion happened (if `CONFIG_FTRACE_RECORD_RECURSION` is set).

Alternatively, if the `FTRACE_OPS_FL_RECURSION` flag is set on the `ftrace_ops` (as explained below), then a helper trampoline will be used to test for recursion for the callback and no recursion test needs to be done. But this is at the expense of a slightly more overhead from an extra function call.

If your callback accesses any data or critical section that requires RCU protection, it is best to make sure that RCU is “watching”, otherwise that data or critical section will not be protected as expected. In this case add:

```
if (!rcu_is_watching())
    return;
```

Alternatively, if the `FTRACE_OPS_FL_RCU` flag is set on the `ftrace_ops` (as explained below), then a helper trampoline will be used to test for `rcu_is_watching` for the callback and no other test needs to be done. But this is at the expense of a slightly more overhead from an extra function call.

4.6 The ftrace FLAGS

The `ftrace_ops` flags are all defined and documented in `include/linux/ftrace.h`. Some of the flags are used for internal infrastructure of ftrace, but the ones that users should be aware of are the following:

FTRACE_OPS_FL_SAVE_REGS If the callback requires reading or modifying the `pt_regs` passed to the callback, then it must set this flag. Registering a `ftrace_ops` with this flag set on an architecture that does not support passing of `pt_regs` to the callback will fail.

FTRACE_OPS_FL_SAVE_REGS_IF_SUPPORTED Similar to `SAVE_REGS` but the registering of a `ftrace_ops` on an architecture that does not support passing of `regs` will not fail with this

flag set. But the callback must check if `regs` is `NULL` or not to determine if the architecture supports it.

FTRACE_OPS_FL_RECURSION By default, it is expected that the callback can handle recursion. But if the callback is not that worried about overhead, then setting this bit will add the recursion protection around the callback by calling a helper function that will do the recursion protection and only call the callback if it did not recurse.

Note, if this flag is not set, and recursion does occur, it could cause the system to crash, and possibly reboot via a triple fault.

Not, if this flag is set, then the callback will always be called with preemption disabled. If it is not set, then it is possible (but not guaranteed) that the callback will be called in preemptable context.

FTRACE_OPS_FL_IPMODIFY Requires **FTRACE_OPS_FL_SAVE_REGS** set. If the callback is to “hijack” the traced function (have another function called instead of the traced function), it requires setting this flag. This is what live kernel patches uses. Without this flag the `pt_regs->ip` can not be modified.

Note, only one `ftrace_ops` with **FTRACE_OPS_FL_IPMODIFY** set may be registered to any given function at a time.

FTRACE_OPS_FL_RCU If this is set, then the callback will only be called by functions where RCU is “watching”. This is required if the callback function performs any `rcu_read_lock()` operation.

RCU stops watching when the system goes idle, the time when a CPU is taken down and comes back online, and when entering from kernel to user space and back to kernel space. During these transitions, a callback may be executed and RCU synchronization will not protect it.

FTRACE_OPS_FL_PERMANENT If this is set on any `ftrace_ops`, then the tracing cannot disabled by writing 0 to the `proc sysctl ftrace_enabled`. Equally, a callback with the flag set cannot be registered if `ftrace_enabled` is 0.

Livepatch uses it not to lose the function redirection, so the system stays protected.

4.7 Filtering which functions to trace

If a callback is only to be called from specific functions, a filter must be set up. The filters are added by name, or ip if it is known.

```
int ftrace_set_filter(struct ftrace_ops *ops, unsigned char *buf,
                    int len, int reset);
```

@ops The ops to set the filter with

@buf The string that holds the function filter text.

@len The length of the string.

@reset Non-zero to reset all filters before applying this filter.

Filters denote which functions should be enabled when tracing is enabled. If `@buf` is `NULL` and `reset` is set, all functions will be enabled for tracing.

The @buf can also be a glob expression to enable all functions that match a specific pattern. See Filter Commands in Documentation/trace/ftrace.rst.

To just trace the schedule function:

```
ret = ftrace_set_filter(&ops, "schedule", strlen("schedule"), 0);
```

To add more functions, call the ftrace_set_filter() more than once with the @reset parameter set to zero. To remove the current filter set and replace it with new functions defined by @buf, have @reset be non-zero.

To remove all the filtered functions and trace all functions:

```
ret = ftrace_set_filter(&ops, NULL, 0, 1);
```

Sometimes more than one function has the same name. To trace just a specific function in this case, ftrace_set_filter_ip() can be used.

```
ret = ftrace_set_filter_ip(&ops, ip, 0, 0);
```

Although the ip must be the address where the call to fentry or mcount is located in the function. This function is used by perf and kprobes that gets the ip address from the user (usually using debug info from the kernel).

If a glob is used to set the filter, functions can be added to a “notrace” list that will prevent those functions from calling the callback. The “notrace” list takes precedence over the “filter” list. If the two lists are non-empty and contain the same functions, the callback will not be called by any function.

An empty “notrace” list means to allow all functions defined by the filter to be traced.

```
int ftrace_set_notrace(struct ftrace_ops *ops, unsigned char *buf,
                      int len, int reset);
```

This takes the same parameters as ftrace_set_filter() but will add the functions it finds to not be traced. This is a separate list from the filter list, and this function does not modify the filter list.

A non-zero @reset will clear the “notrace” list before adding functions that match @buf to it.

Clearing the “notrace” list is the same as clearing the filter list

```
ret = ftrace_set_notrace(&ops, NULL, 0, 1);
```

The filter and notrace lists may be changed at any time. If only a set of functions should call the callback, it is best to set the filters before registering the callback. But the changes may also happen after the callback has been registered.

If a filter is in place, and the @reset is non-zero, and @buf contains a matching glob to functions, the switch will happen during the time of the ftrace_set_filter() call. At no time will all functions call the callback.

```
ftrace_set_filter(&ops, "schedule", strlen("schedule"), 1);

register_ftrace_function(&ops);
```

```
msleep(10);  
ftrace_set_filter(&ops, "try_to_wake_up", strlen("try_to_wake_up"), 1);
```

is not the same as:

```
ftrace_set_filter(&ops, "schedule", strlen("schedule"), 1);  
register_ftrace_function(&ops);  
msleep(10);  
ftrace_set_filter(&ops, NULL, 0, 1);  
ftrace_set_filter(&ops, "try_to_wake_up", strlen("try_to_wake_up"), 0);
```

As the latter will have a short time where all functions will call the callback, between the time of the reset, and the time of the new setting of the filter.

FPROBE - FUNCTION ENTRY/EXIT PROBE

5.1 Introduction

Fprobe is a function entry/exit probe mechanism based on ftrace. Instead of using ftrace full feature, if you only want to attach callbacks on function entry and exit, similar to the kprobes and kretprobes, you can use fprobe. Compared with kprobes and kretprobes, fprobe gives faster instrumentation for multiple functions with single handler. This document describes how to use fprobe.

5.2 The usage of fprobe

The fprobe is a wrapper of ftrace (+ kretprobe-like return callback) to attach callbacks to multiple function entry and exit. User needs to set up the `struct fprobe` and pass it to `register_fprobe()`.

Typically, *fprobe* data structure is initialized with the *entry_handler* and/or *exit_handler* as below.

```
struct fprobe fp = {  
    .entry_handler = my_entry_callback,  
    .exit_handler  = my_exit_callback,  
};
```

To enable the fprobe, call one of `register_fprobe()`, `register_fprobe_ips()`, and `register_fprobe_syms()`. These functions register the fprobe with different types of parameters.

The `register_fprobe()` enables a fprobe by function-name filters. E.g. this enables @fp on "func*" function except "func2"::

```
register_fprobe(&fp, "func*", "func2");
```

The `register_fprobe_ips()` enables a fprobe by ftrace-location addresses. E.g.

```
unsigned long ips[] = { 0x.... };  
  
register_fprobe_ips(&fp, ips, ARRAY_SIZE(ips));
```

And the `register_fprobe_syms()` enables a fprobe by symbol names. E.g.

```
char syms[] = {"func1", "func2", "func3"};

register_fprobe_syms(&fp, syms, ARRAY_SIZE(syms));
```

To disable (remove from functions) this fprobe, call:

```
unregister_fprobe(&fp);
```

You can temporally (soft) disable the fprobe by:

```
disable_fprobe(&fp);
```

and resume by:

```
enable_fprobe(&fp);
```

The above is defined by including the header:

```
#include <linux/fprobe.h>
```

Same as ftrace, the registered callbacks will start being called some time after the `register_fprobe()` is called and before it returns. See Documentation/trace/ftrace.rst.

Also, the `unregister_fprobe()` will guarantee that the both enter and exit handlers are no longer being called by functions after `unregister_fprobe()` returns as same as `unregister_ftrace_function()`.

5.3 The fprobe entry/exit handler

The prototype of the entry/exit callback function is as follows:

```
void callback_func(struct fprobe *fp, unsigned long entry_ip, struct pt_regs_
↳ *regs);
```

Note that both entry and exit callbacks have same ptototype. The @entry_ip is saved at function entry and passed to exit handler.

@fp This is the address of *fprobe* data structure related to this handler. You can embed the *fprobe* to your data structure and get it by `container_of()` macro from @fp. The @fp must not be NULL.

@entry_ip This is the ftrace address of the traced function (both entry and exit). Note that this may not be the actual entry address of the function but the address where the ftrace is instrumented.

@regs This is the *pt_regs* data structure at the entry and exit. Note that the instruction pointer of @regs may be different from the @entry_ip in the entry_handler. If you need traced instruction pointer, you need to use @entry_ip. On the other hand, in the exit_handler, the instruction pointer of @regs is set to the current return address.

5.4 Share the callbacks with kprobes

Since the recursion safeness of the fprobe (and ftrace) is a bit different from the kprobes, this may cause an issue if user wants to run the same code from the fprobe and the kprobes.

Kprobes has per-cpu 'current_kprobe' variable which protects the kprobe handler from recursion in all cases. On the other hand, fprobe uses only ftrace_test_recursion_trylock(). This allows interrupt context to call another (or same) fprobe while the fprobe user handler is running.

This is not a matter if the common callback code has its own recursion detection, or it can handle the recursion in the different contexts (normal/interrupt/NMI.) But if it relies on the 'current_kprobe' recursion lock, it has to check kprobe_running() and use kprobe_busy_*() APIs.

Fprobe has FPROBE_FL_KPROBE_SHARED flag to do this. If your common callback code will be shared with kprobes, please set FPROBE_FL_KPROBE_SHARED *before* registering the fprobe, like:

```
fprobe.flags = FPROBE_FL_KPROBE_SHARED;
register_fprobe(&fprobe, "func*", NULL);
```

This will protect your common callback from the nested call.

5.5 The missed counter

The *fprobe* data structure has *fprobe::nmissed* counter field as same as kprobes. This counter counts up when;

- fprobe fails to take ftrace_recursion lock. This usually means that a function which is traced by other ftrace users is called from the entry_handler.
- fprobe fails to setup the function exit because of the shortage of rethook (the shadow stack for hooking the function return.)

The *fprobe::nmissed* field counts up in both cases. Therefore, the former skips both of entry and exit callback and the latter skips the exit callback, but in both case the counter will increase by 1.

Note that if you set the FTRACE_OPS_FL_RECURSION and/or FTRACE_OPS_FL_RCU to *fprobe::ops::flags* (ftrace_ops::flags) when registering the fprobe, this counter may not work correctly, because ftrace skips the fprobe function which increase the counter.

5.6 Functions and structures

struct **fprobe**
ftrace based probe.

Definition

```
struct fprobe {
#ifdef CONFIG_FUNCTION_TRACER;
```

```
    struct ftrace_ops      ops;
#endif;
    unsigned long          nmissd;
    unsigned int           flags;
    struct rethook         *rethook;
    void (*entry_handler)(struct fprobe *fp, unsigned long entry_ip, struct pt_
↪regs *regs);
    void (*exit_handler)(struct fprobe *fp, unsigned long entry_ip, struct pt_
↪regs *regs);
};
```

Members

ops The ftrace_ops.

nmissd The counter for missing events.

flags The status flag.

rethook The rethook data structure. (internal data)

entry_handler The callback function for function entry.

exit_handler The callback function for function exit.

void **disable_fprobe**(struct *fprobe* *fp)
Disable fprobe

Parameters

struct fprobe *fp The fprobe to be disabled.

Description

This will soft-disable **fp**. Note that this doesn't remove the ftrace hooks from the function entry.

void **enable_fprobe**(struct *fprobe* *fp)
Enable fprobe

Parameters

struct fprobe *fp The fprobe to be enabled.

Description

This will soft-enable **fp**.

int **register_fprobe**(struct *fprobe* *fp, const char *filter, const char *notfilter)
Register fprobe to ftrace by pattern.

Parameters

struct fprobe *fp A fprobe data structure to be registered.

const char *filter A wildcard pattern of probed symbols.

const char *notfilter A wildcard pattern of NOT probed symbols.

Description

Register **fp** to ftrace for enabling the probe on the symbols matched to **filter**. If **notfilter** is not NULL, the symbols matched the **notfilter** are not probed.

Return 0 if **fp** is registered successfully, -errno if not.

int **register_fprobe_ips**(struct *fprobe* *fp, unsigned long *addrs, int num)
Register fprobe to ftrace by address.

Parameters

struct fprobe *fp A fprobe data structure to be registered.

unsigned long *addrs An array of target ftrace location addresses.

int num The number of entries of **addrs**.

Description

Register **fp** to ftrace for enabling the probe on the address given by **addrs**. The **addrs** must be the addresses of ftrace location address, which may be the symbol address + arch-dependent offset. If you unsure what this mean, please use other registration functions.

Return 0 if **fp** is registered successfully, -errno if not.

int **register_fprobe_syms**(struct *fprobe* *fp, const char **syms, int num)
Register fprobe to ftrace by symbols.

Parameters

struct fprobe *fp A fprobe data structure to be registered.

const char **syms An array of target symbols.

int num The number of entries of **syms**.

Description

Register **fp** to the symbols given by **syms** array. This will be useful if you are sure the symbols exist in the kernel.

Return 0 if **fp** is registered successfully, -errno if not.

int **unregister_fprobe**(struct *fprobe* *fp)
Unregister fprobe from ftrace

Parameters

struct fprobe *fp A fprobe data structure to be unregistered.

Description

Unregister fprobe (and remove ftrace hooks from the function entries).

Return 0 if **fp** is unregistered successfully, -errno if not.

KERNEL PROBES (KPROBES)

Author Jim Keniston <jkenisto@us.ibm.com>

Author Prasanna S Panchamukhi <prasanna.panchamukhi@gmail.com>

Author Masami Hiramatsu <mhiramat@redhat.com>

6.1 Concepts: Kprobes and Return Probes

Kprobes enables you to dynamically break into any kernel routine and collect debugging and performance information non-disruptively. You can trap at almost any kernel code address¹, specifying a handler routine to be invoked when the breakpoint is hit.

There are currently two types of probes: kprobes, and kretprobes (also called return probes). A kprobe can be inserted on virtually any instruction in the kernel. A return probe fires when a specified function returns.

In the typical case, Kprobes-based instrumentation is packaged as a kernel module. The module's init function installs ("registers") one or more probes, and the exit function unregisters them. A registration function such as `register_kprobe()` specifies where the probe is to be inserted and what handler is to be called when the probe is hit.

There are also `register_/unregister_*probes()` functions for batch registration/unregistration of a group of *probes. These functions can speed up unregistration process when you have to unregister a lot of probes at once.

The next four subsections explain how the different types of probes work and how jump optimization works. They explain certain things that you'll need to know in order to make the best use of Kprobes - e.g., the difference between a `pre_handler` and a `post_handler`, and how to use the `maxactive` and `nmissed` fields of a kretprobe. But if you're in a hurry to start using Kprobes, you can skip ahead to [Architectures Supported](#).

¹ some parts of the kernel code can not be trapped, see [Blacklist](#))

6.1.1 How Does a Kprobe Work?

When a kprobe is registered, Kprobes makes a copy of the probed instruction and replaces the first byte(s) of the probed instruction with a breakpoint instruction (e.g., `int3` on i386 and `x86_64`).

When a CPU hits the breakpoint instruction, a trap occurs, the CPU's registers are saved, and control passes to Kprobes via the `notifier_call_chain` mechanism. Kprobes executes the "pre_handler" associated with the kprobe, passing the handler the addresses of the kprobe struct and the saved registers.

Next, Kprobes single-steps its copy of the probed instruction. (It would be simpler to single-step the actual instruction in place, but then Kprobes would have to temporarily remove the breakpoint instruction. This would open a small time window when another CPU could sail right past the probepoint.)

After the instruction is single-stepped, Kprobes executes the "post_handler," if any, that is associated with the kprobe. Execution then continues with the instruction following the probepoint.

6.1.2 Changing Execution Path

Since kprobes can probe into a running kernel code, it can change the register set, including instruction pointer. This operation requires maximum care, such as keeping the stack frame, recovering the execution path etc. Since it operates on a running kernel and needs deep knowledge of computer architecture and concurrent computing, you can easily shoot your foot.

If you change the instruction pointer (and set up other related registers) in `pre_handler`, you must return `!0` so that kprobes stops single stepping and just returns to the given address. This also means `post_handler` should not be called anymore.

Note that this operation may be harder on some architectures which use TOC (Table of Contents) for function call, since you have to setup a new TOC for your function in your module, and recover the old one after returning from it.

6.1.3 Return Probes

How Does a Return Probe Work?

When you call `register_kretprobe()`, Kprobes establishes a kprobe at the entry to the function. When the probed function is called and this probe is hit, Kprobes saves a copy of the return address, and replaces the return address with the address of a "trampoline." The trampoline is an arbitrary piece of code – typically just a `nop` instruction. At boot time, Kprobes registers a kprobe at the trampoline.

When the probed function executes its return instruction, control passes to the trampoline and that probe is hit. Kprobes' trampoline handler calls the user-specified return handler associated with the `kretprobe`, then sets the saved instruction pointer to the saved return address, and that's where execution resumes upon return from the trap.

While the probed function is executing, its return address is stored in an object of type `kretprobe_instance`. Before calling `register_kretprobe()`, the user sets the `maxactive` field of the

kretprobe struct to specify how many instances of the specified function can be probed simultaneously. `register_kretprobe()` pre-allocates the indicated number of `kretprobe_instance` objects.

For example, if the function is non-recursive and is called with a spinlock held, `maxactive = 1` should be enough. If the function is non-recursive and can never relinquish the CPU (e.g., via a semaphore or preemption), `NR_CPUS` should be enough. If `maxactive <= 0`, it is set to a default value. If `CONFIG_PREEMPT` is enabled, the default is `max(10, 2*NR_CPUS)`. Otherwise, the default is `NR_CPUS`.

It's not a disaster if you set `maxactive` too low; you'll just miss some probes. In the `kretprobe` struct, the `nmissed` field is set to zero when the return probe is registered, and is incremented every time the probed function is entered but there is no `kretprobe_instance` object available for establishing the return probe.

Kretprobe entry-handler

Kretprobes also provides an optional user-specified handler which runs on function entry. This handler is specified by setting the `entry_handler` field of the `kretprobe` struct. Whenever the kprobe placed by `kretprobe` at the function entry is hit, the user-defined `entry_handler`, if any, is invoked. If the `entry_handler` returns 0 (success) then a corresponding return handler is guaranteed to be called upon function return. If the `entry_handler` returns a non-zero error then Kprobes leaves the return address as is, and the `kretprobe` has no further effect for that particular function instance.

Multiple entry and return handler invocations are matched using the unique `kretprobe_instance` object associated with them. Additionally, a user may also specify per return-instance private data to be part of each `kretprobe_instance` object. This is especially useful when sharing private data between corresponding user entry and return handlers. The size of each private data object can be specified at `kretprobe` registration time by setting the `data_size` field of the `kretprobe` struct. This data can be accessed through the `data` field of each `kretprobe_instance` object.

In case probed function is entered but there is no `kretprobe_instance` object available, then in addition to incrementing the `nmissed` count, the user `entry_handler` invocation is also skipped.

6.1.4 How Does Jump Optimization Work?

If your kernel is built with `CONFIG_OTPROBES=y` (currently this flag is automatically set 'y' on x86/x86-64, non-preemptive kernel) and the “`debug.kprobes_optimization`” kernel parameter is set to 1 (see `sysctl(8)`), Kprobes tries to reduce probe-hit overhead by using a jump instruction instead of a breakpoint instruction at each probepoint.

Init a Kprobe

When a probe is registered, before attempting this optimization, Kprobes inserts an ordinary, breakpoint-based kprobe at the specified address. So, even if it's not possible to optimize this particular probepoint, there'll be a probe there.

Safety Check

Before optimizing a probe, Kprobes performs the following safety checks:

- Kprobes verifies that the region that will be replaced by the jump instruction (the “optimized region”) lies entirely within one function. (A jump instruction is multiple bytes, and so may overlay multiple instructions.)
- Kprobes analyzes the entire function and verifies that there is no jump into the optimized region. Specifically:
 - the function contains no indirect jump;
 - the function contains no instruction that causes an exception (since the fixup code triggered by the exception could jump back into the optimized region – Kprobes checks the exception tables to verify this);
 - there is no near jump to the optimized region (other than to the first byte).
- For each instruction in the optimized region, Kprobes verifies that the instruction can be executed out of line.

Preparing Detour Buffer

Next, Kprobes prepares a “detour” buffer, which contains the following instruction sequence:

- code to push the CPU's registers (emulating a breakpoint trap)
- a call to the trampoline code which calls user's probe handlers.
- code to restore registers
- the instructions from the optimized region
- a jump back to the original execution path.

Pre-optimization

After preparing the detour buffer, Kprobes verifies that none of the following situations exist:

- The probe has a `post_handler`.
- Other instructions in the optimized region are probed.
- The probe is disabled.

In any of the above cases, Kprobes won't start optimizing the probe. Since these are temporary situations, Kprobes tries to start optimizing it again if the situation is changed.

If the kprobe can be optimized, Kprobes enqueues the kprobe to an optimizing list, and kicks the kprobe-optimizer workqueue to optimize it. If the to-be-optimized probepoint is hit before being optimized, Kprobes returns control to the original instruction path by setting the CPU's instruction pointer to the copied code in the detour buffer – thus at least avoiding the single-step.

Optimization

The Kprobe-optimizer doesn't insert the jump instruction immediately; rather, it calls `synchronize_rcu()` for safety first, because it's possible for a CPU to be interrupted in the middle of executing the optimized region³. As you know, `synchronize_rcu()` can ensure that all interruptions that were active when `synchronize_rcu()` was called are done, but only if `CONFIG_PREEMPT=n`. So, this version of kprobe optimization supports only kernels with `CONFIG_PREEMPT=n`⁴.

After that, the Kprobe-optimizer calls `stop_machine()` to replace the optimized region with a jump instruction to the detour buffer, using `text_poke_smp()`.

Unoptimization

When an optimized kprobe is unregistered, disabled, or blocked by another kprobe, it will be unoptimized. If this happens before the optimization is complete, the kprobe is just dequeued from the optimized list. If the optimization has been done, the jump is replaced with the original code (except for an `int3` breakpoint in the first byte) by using `text_poke_smp()`.

NOTE for geeks: The jump optimization changes the kprobe's `pre_handler` behavior. Without optimization, the `pre_handler` can change the kernel's execution path by changing `regs->ip` and returning 1. However, when the probe is optimized, that modification is ignored. Thus, if you want to tweak the kernel's execution path, you need to suppress optimization, using one of the following techniques:

- Specify an empty function for the kprobe's `post_handler`.

or

- Execute `'sysctl -w debug.kprobes_optimization=n'`

³ Please imagine that the 2nd instruction is interrupted and then the optimizer replaces the 2nd instruction with the jump *address* while the interrupt handler is running. When the interrupt returns to original address, there is no valid instruction, and it causes an unexpected result.

⁴ This optimization-safety checking may be replaced with the stop-machine method that `ksplce` uses for supporting a `CONFIG_PREEMPT=y` kernel.

6.1.5 Blacklist

Kprobes can probe most of the kernel except itself. This means that there are some functions where kprobes cannot probe. Probing (trapping) such functions can cause a recursive trap (e.g. double fault) or the nested probe handler may never be called. Kprobes manages such functions as a blacklist. If you want to add a function into the blacklist, you just need to (1) include `linux/kprobes.h` and (2) use `NOKPROBE_SYMBOL()` macro to specify a blacklisted function. Kprobes checks the given probe address against the blacklist and rejects registering it, if the given address is in the blacklist.

6.2 Architectures Supported

Kprobes and return probes are implemented on the following architectures:

- i386 (Supports jump optimization)
- x86_64 (AMD-64, EM64T) (Supports jump optimization)
- ppc64
- ia64 (Does not support probes on instruction slot1.)
- sparc64 (Return probes not yet implemented.)
- arm
- ppc
- mips
- s390
- parisc

6.3 Configuring Kprobes

When configuring the kernel using `make menuconfig/xconfig/oldconfig`, ensure that `CONFIG_KPROBES` is set to “y”. Under “General setup”, look for “Kprobes”.

So that you can load and unload Kprobes-based instrumentation modules, make sure “Loadable module support” (`CONFIG_MODULES`) and “Module unloading” (`CONFIG_MODULE_UNLOAD`) are set to “y”.

Also make sure that `CONFIG_KALLSYMS` and perhaps even `CONFIG_KALLSYMS_ALL` are set to “y”, since `kallsyms_lookup_name()` is used by the in-kernel kprobe address resolution code.

If you need to insert a probe in the middle of a function, you may find it useful to “Compile the kernel with debug info” (`CONFIG_DEBUG_INFO`), so you can use “`objdump -d -l vmlinux`” to see the source-to-object code mapping.

6.4 API Reference

The Kprobes API includes a “register” function and an “unregister” function for each type of probe. The API also includes “register_*probes” and “unregister_*probes” functions for (un)registering arrays of probes. Here are terse, mini-man-page specifications for these functions and the associated probe handlers that you’ll write. See the files in the samples/kprobes/ sub-directory for examples.

6.4.1 register_kprobe

```
#include <linux/kprobes.h>
int register_kprobe(struct kprobe *kp);
```

Sets a breakpoint at the address `kp->addr`. When the breakpoint is hit, Kprobes calls `kp->pre_handler`. After the probed instruction is single-stepped, Kprobe calls `kp->post_handler`. Any or all handlers can be NULL. If `kp->flags` is set `KPROBE_FLAG_DISABLED`, that `kp` will be registered but disabled, so, its handlers aren’t hit until calling `enable_kprobe(kp)`.

Note:

1. With the introduction of the “symbol_name” field to struct kprobe, the probepoint address resolution will now be taken care of by the kernel. The following will now work:

```
kp.symbol_name = "symbol_name";
```

(64-bit powerpc intricacies such as function descriptors are handled transparently)

2. Use the “offset” field of struct kprobe if the offset into the symbol to install a probepoint is known. This field is used to calculate the probepoint.
3. Specify either the kprobe “symbol_name” OR the “addr”. If both are specified, kprobe registration will fail with -EINVAL.
4. With CISC architectures (such as i386 and x86_64), the kprobes code does not validate if the `kprobe.addr` is at an instruction boundary. Use “offset” with caution.

`register_kprobe()` returns 0 on success, or a negative `errno` otherwise.

User’s pre-handler (`kp->pre_handler`):

```
#include <linux/kprobes.h>
#include <linux/ptrace.h>
int pre_handler(struct kprobe *p, struct pt_regs *regs);
```

Called with `p` pointing to the kprobe associated with the breakpoint, and `regs` pointing to the struct containing the registers saved when the breakpoint was hit. Return 0 here unless you’re a Kprobes geek.

User’s post-handler (`kp->post_handler`):

```
#include <linux/kprobes.h>
#include <linux/ptrace.h>
```

```
void post_handler(struct kprobe *p, struct pt_regs *regs,
                  unsigned long flags);
```

p and regs are as described for the pre_handler. flags always seems to be zero.

6.4.2 register_kretprobe

```
#include <linux/kprobes.h>
int register_kretprobe(struct kretprobe *rp);
```

Establishes a return probe for the function whose address is rp->kp.addr. When that function returns, Kprobes calls rp->handler. You must set rp->maxactive appropriately before you call register_kretprobe(); see “How Does a Return Probe Work?” for details.

register_kretprobe() returns 0 on success, or a negative errno otherwise.

User’s return-probe handler (rp->handler):

```
#include <linux/kprobes.h>
#include <linux/ptrace.h>
int kretprobe_handler(struct kretprobe_instance *ri,
                     struct pt_regs *regs);
```

regs is as described for kprobe.pre_handler. ri points to the kretprobe_instance object, of which the following fields may be of interest:

- ret_addr: the return address
- rp: points to the corresponding kretprobe object
- task: points to the corresponding task struct
- **data: points to per return-instance private data; see “Kretprobe entry-handler” for details.**

The regs_return_value(regs) macro provides a simple abstraction to extract the return value from the appropriate register as defined by the architecture’s ABI.

The handler’s return value is currently ignored.

6.4.3 unregister_*probe

```
#include <linux/kprobes.h>
void unregister_kprobe(struct kprobe *kp);
void unregister_kretprobe(struct kretprobe *rp);
```

Removes the specified probe. The unregister function can be called at any time after the probe has been registered.

Note: If the functions find an incorrect probe (ex. an unregistered probe), they clear the addr field of the probe.

6.4.4 register_*probes

```
#include <linux/kprobes.h>
int register_kprobes(struct kprobe **kps, int num);
int register_kretprobes(struct kretprobe **rps, int num);
```

Registers each of the num probes in the specified array. If any error occurs during registration, all probes in the array, up to the bad probe, are safely unregistered before the register_*probes function returns.

- kps/rps: an array of pointers to *probe data structures
- num: the number of the array entries.

Note: You have to allocate(or define) an array of pointers and set all of the array entries before using these functions.

6.4.5 unregister_*probes

```
#include <linux/kprobes.h>
void unregister_kprobes(struct kprobe **kps, int num);
void unregister_kretprobes(struct kretprobe **rps, int num);
```

Removes each of the num probes in the specified array at once.

Note: If the functions find some incorrect probes (ex. unregistered probes) in the specified array, they clear the addr field of those incorrect probes. However, other probes in the array are unregistered correctly.

6.4.6 disable_*probe

```
#include <linux/kprobes.h>
int disable_kprobe(struct kprobe *kp);
int disable_kretprobe(struct kretprobe *rp);
```

Temporarily disables the specified *probe. You can enable it again by using enable_*probe(). You must specify the probe which has been registered.

6.4.7 enable_*probe

```
#include <linux/kprobes.h>
int enable_kprobe(struct kprobe *kp);
int enable_kretprobe(struct kretprobe *rp);
```

Enables *probe which has been disabled by disable_*probe(). You must specify the probe which has been registered.

6.5 Kprobes Features and Limitations

Kprobes allows multiple probes at the same address. Also, a probepoint for which there is a post_handler cannot be optimized. So if you install a kprobe with a post_handler, at an optimized probepoint, the probepoint will be unoptimized automatically.

In general, you can install a probe anywhere in the kernel. In particular, you can probe interrupt handlers. Known exceptions are discussed in this section.

The register_*probe functions will return -EINVAL if you attempt to install a probe in the code that implements Kprobes (mostly kernel/kprobes.c and arch/*/kernel/kprobes.c, but also functions such as do_page_fault and notifier_call_chain).

If you install a probe in an inline-able function, Kprobes makes no attempt to chase down all inline instances of the function and install probes there. gcc may inline a function without being asked, so keep this in mind if you're not seeing the probe hits you expect.

A probe handler can modify the environment of the probed function – e.g., by modifying kernel data structures, or by modifying the contents of the pt_regs struct (which are restored to the registers upon return from the breakpoint). So Kprobes can be used, for example, to install a bug fix or to inject faults for testing. Kprobes, of course, has no way to distinguish the deliberately injected faults from the accidental ones. Don't drink and probe.

Kprobes makes no attempt to prevent probe handlers from stepping on each other – e.g., probing printk() and then calling printk() from a probe handler. If a probe handler hits a probe, that second probe's handlers won't be run in that instance, and the kprobe.nmissed member of the second probe will be incremented.

As of Linux v2.6.15-rc1, multiple handlers (or multiple instances of the same handler) may run concurrently on different CPUs.

Kprobes does not use mutexes or allocate memory except during registration and unregistration.

Probe handlers are run with preemption disabled or interrupt disabled, which depends on the architecture and optimization state. (e.g., kretprobe handlers and optimized kprobe handlers run without interrupt disabled on x86/x86-64). In any case, your handler should not yield the CPU (e.g., by attempting to acquire a semaphore, or waiting I/O).

Since a return probe is implemented by replacing the return address with the trampoline's address, stack backtraces and calls to __builtin_return_address() will typically yield the trampoline's address instead of the real return address for kretprobed functions. (As far as we can tell, __builtin_return_address() is used only for instrumentation and error reporting.)

If the number of times a function is called does not match the number of times it returns, registering a return probe on that function may produce undesirable results. In such a case,

a line: kretprobe BUG!: Processing kretprobe d000000000041aa8 @ c00000000004f48c gets printed. With this information, one will be able to correlate the exact instance of the kretprobe that caused the problem. We have the `do_exit()` case covered. `do_execve()` and `do_fork()` are not an issue. We're unaware of other specific cases where this could be a problem.

If, upon entry to or exit from a function, the CPU is running on a stack other than that of the current task, registering a return probe on that function may produce undesirable results. For this reason, Kprobes doesn't support return probes (or kprobes) on the x86_64 version of `__switch_to()`; the registration functions return `-EINVAL`.

On x86/x86-64, since the Jump Optimization of Kprobes modifies instructions widely, there are some limitations to optimization. To explain it, we introduce some terminology. Imagine a 3-instruction sequence consisting of a two 2-byte instructions and one 3-byte instruction.

```

      IA
      |
[-2][-1][0][1][2][3][4][5][6][7]
      [ins1][ins2][ ins3 ]
      [<-   DCR   ->]
      [<- JTPR ->]

ins1: 1st Instruction
ins2: 2nd Instruction
ins3: 3rd Instruction
IA:   Insertion Address
JTPR: Jump Target Prohibition Region
DCR:  Detoured Code Region

```

The instructions in DCR are copied to the out-of-line buffer of the kprobe, because the bytes in DCR are replaced by a 5-byte jump instruction. So there are several limitations.

- a) The instructions in DCR must be relocatable.
- b) The instructions in DCR must not include a call instruction.
- c) JTPR must not be targeted by any jump or call instruction.
- d) DCR must not straddle the border between functions.

Anyway, these limitations are checked by the in-kernel instruction decoder, so you don't need to worry about that.

6.6 Probe Overhead

On a typical CPU in use in 2005, a kprobe hit takes 0.5 to 1.0 microseconds to process. Specifically, a benchmark that hits the same probepoint repeatedly, firing a simple handler each time, reports 1-2 million hits per second, depending on the architecture. A return-probe hit typically takes 50-75% longer than a kprobe hit. When you have a return probe set on a function, adding a kprobe at the entry to that function adds essentially no overhead.

Here are sample overhead figures (in usec) for different architectures:

```

k = kprobe; r = return probe; kr = kprobe + return probe
on same function

```

```
i386: Intel Pentium M, 1495 MHz, 2957.31 bogomips
k = 0.57 usec; r = 0.92; kr = 0.99

x86_64: AMD Opteron 246, 1994 MHz, 3971.48 bogomips
k = 0.49 usec; r = 0.80; kr = 0.82

ppc64: POWER5 (gr), 1656 MHz (SMT disabled, 1 virtual CPU per physical CPU)
k = 0.77 usec; r = 1.26; kr = 1.45
```

6.6.1 Optimized Probe Overhead

Typically, an optimized kprobe hit takes 0.07 to 0.1 microseconds to process. Here are sample overhead figures (in usec) for x86 architectures:

```
k = unoptimized kprobe, b = boosted (single-step skipped), o = optimized ↵
↵kprobe,
r = unoptimized kretprobe, rb = boosted kretprobe, ro = optimized kretprobe.

i386: Intel(R) Xeon(R) E5410, 2.33GHz, 4656.90 bogomips
k = 0.80 usec; b = 0.33; o = 0.05; r = 1.10; rb = 0.61; ro = 0.33

x86-64: Intel(R) Xeon(R) E5410, 2.33GHz, 4656.90 bogomips
k = 0.99 usec; b = 0.43; o = 0.06; r = 1.24; rb = 0.68; ro = 0.30
```

6.7 TODO

- SystemTap (<http://sourceware.org/systemtap>): Provides a simplified programming interface for probe-based instrumentation. Try it out.
- Kernel return probes for sparc64.
- Support for other architectures.
- User-space probes.
- Watchpoint probes (which fire on data references).

6.8 Kprobes Example

See `samples/kprobes/kprobe_example.c`

6.9 Kretprobes Example

See `samples/kprobes/kretprobe_example.c`

6.10 Deprecated Features

Jprobes is now a deprecated feature. People who are depending on it should migrate to other tracing features or use older kernels. Please consider to migrate your tool to one of the following options:

- Use trace-event to trace target function with arguments.

trace-event is a low-overhead (and almost no visible overhead if it is off) statically defined event interface. You can define new events and trace it via ftrace or any other tracing tools.

See the following urls:

- <https://lwn.net/Articles/379903/>
- <https://lwn.net/Articles/381064/>
- <https://lwn.net/Articles/383362/>

- Use ftrace dynamic events (kprobe event) with perf-probe.

If you build your kernel with debug info (`CONFIG_DEBUG_INFO=y`), you can find which register/stack is assigned to which local variable or arguments by using perf-probe and set up new event to trace it.

See following documents:

- *Kprobe-based Event Tracing*
- *Event Tracing*
- `tools/perf/Documentation/perf-probe.txt`

6.11 The kprobes debugfs interface

With recent kernels (> 2.6.20) the list of registered kprobes is visible under the `/sys/kernel/debug/kprobes/` directory (assuming debugfs is mounted at `//sys/kernel/debug`).

`/sys/kernel/debug/kprobes/list`: Lists all registered probes on the system:

```
c015d71a k  vfs_read+0x0
c03dedc5 r  tcp_v4_rcv+0x0
```

The first column provides the kernel address where the probe is inserted. The second column identifies the type of probe (k - kprobe and r - kretprobe) while the third column specifies the symbol+offset of the probe. If the probed function belongs to a module, the module name is also specified. Following columns show probe status. If the probe is on a virtual address that is no longer valid (module init sections, module virtual addresses that correspond to modules that've been unloaded), such probes are marked with [GONE]. If the probe is temporarily disabled, such

probes are marked with [DISABLED]. If the probe is optimized, it is marked with [OPTIMIZED]. If the probe is ftrace-based, it is marked with [FTRACE].

`/sys/kernel/debug/kprobes/enabled`: Turn kprobes ON/OFF forcibly.

Provides a knob to globally and forcibly turn registered kprobes ON or OFF. By default, all kprobes are enabled. By echoing “0” to this file, all registered probes will be disarmed, till such time a “1” is echoed to this file. Note that this knob just disarms and arms all kprobes and doesn’t change each probe’s disabling state. This means that disabled kprobes (marked [DISABLED]) will be not enabled if you turn ON all kprobes by this knob.

6.12 The kprobes sysctl interface

`/proc/sys/debug/kprobes-optimization`: Turn kprobes optimization ON/OFF.

When `CONFIG_OTPROBES=y`, this sysctl interface appears and it provides a knob to globally and forcibly turn jump optimization (see section *How Does Jump Optimization Work?*) ON or OFF. By default, jump optimization is allowed (ON). If you echo “0” to this file or set “`debug.kprobes_optimization`” to 0 via sysctl, all optimized probes will be unoptimized, and any new probes registered after that will not be optimized.

Note that this knob *changes* the optimized state. This means that optimized probes (marked [OPTIMIZED]) will be unoptimized ([OPTIMIZED] tag will be removed). If the knob is turned on, they will be optimized again.

6.13 References

For additional information on Kprobes, refer to the following URLs:

- <https://lwn.net/Articles/132196/>
- <https://www.kernel.org/doc/ols/2006/ols2006v2-pages-109-124.pdf>

KPROBE-BASED EVENT TRACING

Author Masami Hiramatsu

7.1 Overview

These events are similar to tracepoint based events. Instead of Tracepoint, this is based on kprobes (kprobe and kretprobe). So it can probe wherever kprobes can probe (this means, all functions except those with `__kprobes/nokprobe_inline` annotation and those marked `NOKPROBE_SYMBOL`). Unlike the Tracepoint based event, this can be added and removed dynamically, on the fly.

To enable this feature, build your kernel with `CONFIG_KPROBE_EVENTS=y`.

Similar to the events tracer, this doesn't need to be activated via `current_tracer`. Instead of that, add probe points via `/sys/kernel/debug/tracing/kprobe_events`, and enable it via `/sys/kernel/debug/tracing/events/kprobes/<EVENT>/enable`.

You can also use `/sys/kernel/debug/tracing/dynamic_events` instead of `kprobe_events`. That interface will provide unified access to other dynamic events too.

7.2 Synopsis of kprobe_events

```
p[:[GRP/]EVENT] [MOD:]SYM[+offs]|MEMADDR [FETCHARGS] : Set a probe
r[MAXACTIVE][:[GRP/]EVENT] [MOD:]SYM[+0] [FETCHARGS] : Set a return probe
p[:[GRP/]EVENT] [MOD:]SYM[+0]%return [FETCHARGS]      : Set a return probe
-:[GRP/]EVENT                                          : Clear a probe

GRP           : Group name. If omitted, use "kprobes" for it.
EVENT         : Event name. If omitted, the event name is generated
                based on SYM+offs or MEMADDR.
MOD           : Module name which has given SYM.
SYM[+offs]    : Symbol+offset where the probe is inserted.
SYM%return    : Return address of the symbol
MEMADDR       : Address where the probe is inserted.
MAXACTIVE     : Maximum number of instances of the specified function that
                can be probed simultaneously, or 0 for the default value
                as defined in Documentation/trace/kprobes.rst section 1.3.1.

FETCHARGS     : Arguments. Each probe can have up to 128 args.
```

```

%REG          : Fetch register REG
@ADDR         : Fetch memory at ADDR (ADDR should be in kernel)
@SYM[+|-offs] : Fetch memory at SYM +|- offs (SYM should be a data symbol)
$stackN       : Fetch Nth entry of stack (N >= 0)
$stack        : Fetch stack address.
$argN         : Fetch the Nth function argument. (N >= 1) (\*1)
$retval       : Fetch return value.(\*2)
$comm        : Fetch current task comm.
+|--[u]OFFS(FETCHARG) : Fetch memory at FETCHARG +|- OFFS address.(\*3)(\*4)
\IMM          : Store an immediate value to the argument.
NAME=FETCHARG : Set NAME as the argument name of FETCHARG.
FETCHARG:TYPE : Set TYPE as the type of FETCHARG. Currently, basic types
                (u8/u16/u32/u64/s8/s16/s32/s64), hexadecimal types
                (x8/x16/x32/x64), "string", "ustring" and bitfield
                are supported.

```

(*1) only for the probe on function entry (offs == 0).

(*2) only for return probe.

(*3) this is useful for fetching a field of data structures.

(*4) "u" means user-space dereference. See :ref:`user_mem_access`.

7.3 Types

Several types are supported for fetch-args. Kprobe tracer will access memory by given type. Prefix 's' and 'u' means those types are signed and unsigned respectively. 'x' prefix implies it is unsigned. Traced arguments are shown in decimal ('s' and 'u') or hexadecimal ('x'). Without type casting, 'x32' or 'x64' is used depends on the architecture (e.g. x86-32 uses x32, and x86-64 uses x64). These value types can be an array. To record array data, you can add '[N]' (where N is a fixed number, less than 64) to the base type. E.g. 'x16[4]' means an array of x16 (2bytes hex) with 4 elements. Note that the array can be applied to memory type fetchargs, you can not apply it to registers/stack-entries etc. (for example, '\$stack1:x8[8]' is wrong, but '+8(\$stack):x8[8]' is OK.) String type is a special type, which fetches a "null-terminated" string from kernel space. This means it will fail and store NULL if the string container has been paged out. "ustring" type is an alternative of string for user-space. See [User Memory Access](#) for more info.. The string array type is a bit different from other types. For other base types, <base-type>[1] is equal to <base-type> (e.g. +0(%di):x32[1] is same as +0(%di):x32.) But string[1] is not equal to string. The string type itself represents "char array", but string array type represents "char * array". So, for example, +0(%di):string[1] is equal to +0(+0(%di)):string. Bitfield is another special type, which takes 3 parameters, bit-width, bit- offset, and container-size (usually 32). The syntax is:

```
b<bit-width>@<bit-offset>/<container-size>
```

Symbol type('symbol') is an alias of u32 or u64 type (depends on BITS_PER_LONG) which shows given pointer in "symbol+offset" style. For \$comm, the default type is "string"; any other type is invalid.

7.4 User Memory Access

Kprobe events supports user-space memory access. For that purpose, you can use either user-space dereference syntax or ‘ustring’ type.

The user-space dereference syntax allows you to access a field of a data structure in user-space. This is done by adding the “u” prefix to the dereference syntax. For example, `+u4(%si)` means it will read memory from the address in the register `%si` offset by 4, and the memory is expected to be in user-space. You can use this for strings too, e.g. `+u0(%si):string` will read a string from the address in the register `%si` that is expected to be in user-space. ‘ustring’ is a shortcut way of performing the same task. That is, `+0(%si):ustring` is equivalent to `+u0(%si):string`.

Note that `kprobe-event` provides the user-memory access syntax but it doesn’t use it transparently. This means if you use normal dereference or string type for user memory, it might fail, and may always fail on some archs. The user has to carefully check if the target data is in kernel or user space.

7.5 Per-Probe Event Filtering

Per-probe event filtering feature allows you to set different filter on each probe and gives you what arguments will be shown in trace buffer. If an event name is specified right after ‘p:’ or ‘r:’ in `kprobe_events`, it adds an event under `tracing/events/kprobes/<EVENT>`, at the directory you can see ‘id’, ‘enable’, ‘format’, ‘filter’ and ‘trigger’.

enable: You can enable/disable the probe by writing 1 or 0 on it.

format: This shows the format of this probe event.

filter: You can write filtering rules of this event.

id: This shows the id of this probe event.

trigger: This allows to install trigger commands which are executed when the event is hit (for details, see [Event Tracing](#), section 6).

7.6 Event Profiling

You can check the total number of probe hits and probe miss-hits via `/sys/kernel/debug/tracing/kprobe_profile`. The first column is event name, the second is the number of probe hits, the third is the number of probe miss-hits.

7.7 Kernel Boot Parameter

You can add and enable new kprobe events when booting up the kernel by “`kprobe_event=`” parameter. The parameter accepts a semicolon-delimited kprobe events, which format is similar to the `kprobe_events`. The difference is that the probe definition parameters are comma-delimited instead of space. For example, adding myprobe event on `do_sys_open` like below

```
p:myprobe do_sys_open dfd=%ax filename=%dx flags=%cx mode=+4($stack)
```

should be below for kernel boot parameter (just replace spaces with comma)

```
p:myprobe,do_sys_open,dfd=%ax,filename=%dx,flags=%cx,mode=+4($stack)
```

7.8 Usage examples

To add a probe as a new event, write a new definition to `kprobe_events` as below:

```
echo 'p:myprobe do_sys_open dfd=%ax filename=%dx flags=%cx mode=+4($stack)' > /  
↪sys/kernel/debug/tracing/kprobe_events
```

This sets a kprobe on the top of `do_sys_open()` function with recording 1st to 4th arguments as “myprobe” event. Note, which register/stack entry is assigned to each function argument depends on arch-specific ABI. If you are unsure the ABI, please try to use probe subcommand of `perf-tools` (you can find it under `tools/perf/`). As this example shows, users can choose more familiar names for each arguments.

```
echo 'r:myretprobe do_sys_open $retval' >> /sys/kernel/debug/tracing/kprobe_  
↪events
```

This sets a kretprobe on the return point of `do_sys_open()` function with recording return value as “myretprobe” event. You can see the format of these events via `/sys/kernel/debug/tracing/events/kprobes/<EVENT>/format`.

```
cat /sys/kernel/debug/tracing/events/kprobes/myprobe/format  
name: myprobe  
ID: 780  
format:  
    field:unsigned short common_type;          offset:0;          size:2; ↵  
↪signed:0;  
    field:unsigned char common_flags;          offset:2;          size:1; ↵  
↪signed:0;  
    field:unsigned char common_preempt_count;    offset:3; size:1;  
↪signed:0;  
    field:int common_pid;    offset:4;          size:4; signed:1;  
  
    field:unsigned long __probe_ip; offset:12;          size:4; signed:0;  
    field:int __probe_nargs;    offset:16;          size:4; signed:1;  
    field:unsigned long dfd;    offset:20;          size:4; signed:0;  
    field:unsigned long filename; offset:24;          size:4; signed:0;  
    field:unsigned long flags;    offset:28;          size:4; signed:0;  
    field:unsigned long mode;    offset:32;          size:4; signed:0;  
  
print fmt: "(%lx) dfd=%lx filename=%lx flags=%lx mode=%lx", REC->__probe_ip,  
REC->dfd, REC->filename, REC->flags, REC->mode
```

You can see that the event has 4 arguments as in the expressions you specified.

```
echo > /sys/kernel/debug/tracing/kprobe_events
```

This clears all probe points.

Or,

```
echo - :myprobe >> kprobe_events
```

This clears probe points selectively.

Right after definition, each event is disabled by default. For tracing these events, you need to enable it.

```
echo 1 > /sys/kernel/debug/tracing/events/kprobes/myprobe/enable
echo 1 > /sys/kernel/debug/tracing/events/kprobes/myretprobe/enable
```

Use the following command to start tracing in an interval.

```
# echo 1 > tracing_on
Open something...
# echo 0 > tracing_on
```

And you can see the traced information via `/sys/kernel/debug/tracing/trace`.

```
cat /sys/kernel/debug/tracing/trace
# tracer: nop
#
#           TASK-PID    CPU#    TIMESTAMP    FUNCTION
#           | |         |         |         |
<...>-1447  [001] 1038282.286875: myprobe: (do_sys_open+0x0/0xd6)
↪dfd=3 filename=7fffd1ec4440 flags=8000 mode=0
<...>-1447  [001] 1038282.286878: myretprobe: (sys_openat+0xc/0xe <-
↪ do_sys_open) $retval=fffffffffffffffe
<...>-1447  [001] 1038282.286885: myprobe: (do_sys_open+0x0/0xd6)
↪dfd=ffffff9c filename=40413c flags=8000 mode=1b6
<...>-1447  [001] 1038282.286915: myretprobe: (sys_open+0x1b/0x1d <-
↪ do_sys_open) $retval=3
<...>-1447  [001] 1038282.286969: myprobe: (do_sys_open+0x0/0xd6)
↪dfd=ffffff9c filename=4041c6 flags=98800 mode=10
<...>-1447  [001] 1038282.286976: myretprobe: (sys_open+0x1b/0x1d <-
↪ do_sys_open) $retval=3
```

Each line shows when the kernel hits an event, and `<- SYMBOL` means kernel returns from `SYMBOL`(e.g. “`sys_open+0x1b/0x1d <- do_sys_open`” means kernel returns from `do_sys_open` to `sys_open+0x1b`).

UPROBE-TRACER: UPROBE-BASED EVENT TRACING

Author Srikar Dronamraju

8.1 Overview

Uprobe based trace events are similar to kprobe based trace events. To enable this feature, build your kernel with `CONFIG_UPROBE_EVENTS=y`.

Similar to the kprobe-event tracer, this doesn't need to be activated via `current_tracer`. Instead of that, add probe points via `/sys/kernel/debug/tracing/uprobe_events`, and enable it via `/sys/kernel/debug/tracing/events/uprobes/<EVENT>/enable`.

However unlike kprobe-event tracer, the uprobe event interface expects the user to calculate the offset of the probepoint in the object.

You can also use `/sys/kernel/debug/tracing/dynamic_events` instead of `uprobe_events`. That interface will provide unified access to other dynamic events too.

8.2 Synopsis of uprobe_tracer

```
p[:[GRP/]EVENT] PATH:OFFSET [FETCHARGS] : Set a uprobe
r[:[GRP/]EVENT] PATH:OFFSET [FETCHARGS] : Set a return uprobe (uretprobe)
p[:[GRP/]EVENT] PATH:OFFSET%return [FETCHARGS] : Set a return uprobe.
→(uretprobe)
-:[GRP/]EVENT : Clear uprobe or uretprobe event
```

```
GRP : Group name. If omitted, "uprobes" is the default value.
EVENT : Event name. If omitted, the event name is generated based
on PATH+OFFSET.
PATH : Path to an executable or a library.
OFFSET : Offset where the probe is inserted.
OFFSET%return : Offset where the return probe is inserted.
```

```
FETCHARGS : Arguments. Each probe can have up to 128 args.
%REG : Fetch register REG
@ADDR : Fetch memory at ADDR (ADDR should be in userspace)
@+OFFSET : Fetch memory at OFFSET (OFFSET from same file as PATH)
$stackN : Fetch Nth entry of stack (N >= 0)
$stack : Fetch stack address.
```

```
$retval      : Fetch return value.(\\*1)
$comm        : Fetch current task comm.
+|--[u]OFFS(FETCHARG) : Fetch memory at FETCHARG +|- OFFS address.(\\*2)(\\*3)
\\IMM         : Store an immediate value to the argument.
NAME=FETCHARG : Set NAME as the argument name of FETCHARG.
FETCHARG:TYPE : Set TYPE as the type of FETCHARG. Currently, basic types
                (u8/u16/u32/u64/s8/s16/s32/s64), hexadecimal types
                (x8/x16/x32/x64), "string" and bitfield are supported.
```

(*1) only for return probe.

(*2) this is useful for fetching a field of data structures.

(*3) Unlike kprobe event, "u" prefix will just be ignored, because uprobe events can access only user-space memory.

8.3 Types

Several types are supported for fetch-args. Uprobe tracer will access memory by given type. Prefix 's' and 'u' means those types are signed and unsigned respectively. 'x' prefix implies it is unsigned. Traced arguments are shown in decimal ('s' and 'u') or hexadecimal ('x'). Without type casting, 'x32' or 'x64' is used depends on the architecture (e.g. x86-32 uses x32, and x86-64 uses x64). String type is a special type, which fetches a "null-terminated" string from user space. Bitfield is another special type, which takes 3 parameters, bit-width, bit- offset, and container-size (usually 32). The syntax is:

```
b<bit-width>@<bit-offset>/<container-size>
```

For \$comm, the default type is "string"; any other type is invalid.

8.4 Event Profiling

You can check the total number of probe hits per event via /sys/kernel/debug/tracing/uprobe_profile. The first column is the filename, the second is the event name, the third is the number of probe hits.

8.5 Usage examples

- Add a probe as a new uprobe event, write a new definition to uprobe_events as below (sets a uprobe at an offset of 0x4245c0 in the executable /bin/bash):

```
echo 'p /bin/bash:0x4245c0' > /sys/kernel/debug/tracing/uprobe_events
```

- Add a probe as a new uretprobe event:

```
echo 'r /bin/bash:0x4245c0' > /sys/kernel/debug/tracing/uprobe_events
```

- Unset registered event:

```
echo '-:p_bash_0x4245c0' >> /sys/kernel/debug/tracing/uprobe_events
```

- Print out the events that are registered:

```
cat /sys/kernel/debug/tracing/uprobe_events
```

- Clear all events:

```
echo > /sys/kernel/debug/tracing/uprobe_events
```

Following example shows how to dump the instruction pointer and %ax register at the probed text address. Probe zfree function in /bin/zsh:

```
# cd /sys/kernel/debug/tracing/
# cat /proc`pgrep zsh`/maps | grep /bin/zsh | grep r-xp
00400000-0048a000 r-xp 00000000 08:03 130904 /bin/zsh
# objdump -T /bin/zsh | grep -w zfree
0000000000446420 g      DF .text 0000000000000012 Base      zfree
```

0x46420 is the offset of zfree in object /bin/zsh that is loaded at 0x00400000. Hence the command to uprobe would be:

```
# echo 'p:zfree_entry /bin/zsh:0x46420 %ip %ax' > uprobe_events
```

And the same for the uretprobe would be:

```
# echo 'r:zfree_exit /bin/zsh:0x46420 %ip %ax' >> uprobe_events
```

Note: User has to explicitly calculate the offset of the probe-point in the object.

We can see the events that are registered by looking at the uprobe_events file.

```
# cat uprobe_events
p:uprobes/zfree_entry /bin/zsh:0x00046420 arg1=%ip arg2=%ax
r:uprobes/zfree_exit /bin/zsh:0x00046420 arg1=%ip arg2=%ax
```

Format of events can be seen by viewing the file events/uprobes/zfree_entry/format.

```
# cat events/uprobes/zfree_entry/format
name: zfree_entry
ID: 922
format:
    field:unsigned short common_type;          offset:0;  size:2; signed:0;
    field:unsigned char common_flags;          offset:2;  size:1; signed:0;
    field:unsigned char common_preempt_count;  offset:3;  size:1; signed:0;
    field:int common_pid;                      offset:4;  size:4; signed:1;
    field:int common_padding;                  offset:8;  size:4; signed:1;

    field:unsigned long __probe_ip;            offset:12; size:4; signed:0;
    field:u32 arg1;                            offset:16; size:4; signed:0;
    field:u32 arg2;                            offset:20; size:4; signed:0;
```

```
print fmt: "(%lx) arg1=%lx arg2=%lx", REC->__probe_ip, REC->arg1, REC->arg2
```

Right after definition, each event is disabled by default. For tracing these events, you need to enable it by:

```
# echo 1 > events/uprobes/enable
```

Lets start tracing, sleep for some time and stop tracing.

```
# echo 1 > tracing_on
# sleep 20
# echo 0 > tracing_on
```

Also, you can disable the event by:

```
# echo 0 > events/uprobes/enable
```

And you can see the traced information via `/sys/kernel/debug/tracing/trace`.

```
# cat trace
# tracer: nop
#
#          TASK-PID    CPU#    TIMESTAMP    FUNCTION
#          | |         |         |         |
zsh-24842 [006] 258544.995456: zfree_entry: (0x446420)
↪arg1=446420 arg2=79
zsh-24842 [007] 258545.000270: zfree_exit: (0x446540 <-
↪0x446420) arg1=446540 arg2=0
zsh-24842 [002] 258545.043929: zfree_entry: (0x446420)
↪arg1=446420 arg2=79
zsh-24842 [004] 258547.046129: zfree_exit: (0x446540 <-
↪0x446420) arg1=446540 arg2=0
```

Output shows us uprobe was triggered for a pid 24842 with ip being 0x446420 and contents of ax register being 79. And uretprobe was triggered with ip at 0x446540 with counterpart function entry at 0x446420.

USING THE LINUX KERNEL TRACEPOINTS

Author Mathieu Desnoyers

This document introduces Linux Kernel Tracepoints and their use. It provides examples of how to insert tracepoints in the kernel and connect probe functions to them and provides some examples of probe functions.

9.1 Purpose of tracepoints

A tracepoint placed in code provides a hook to call a function (probe) that you can provide at runtime. A tracepoint can be “on” (a probe is connected to it) or “off” (no probe is attached). When a tracepoint is “off” it has no effect, except for adding a tiny time penalty (checking a condition for a branch) and space penalty (adding a few bytes for the function call at the end of the instrumented function and adds a data structure in a separate section). When a tracepoint is “on”, the function you provide is called each time the tracepoint is executed, in the execution context of the caller. When the function provided ends its execution, it returns to the caller (continuing from the tracepoint site).

You can put tracepoints at important locations in the code. They are lightweight hooks that can pass an arbitrary number of parameters, which prototypes are described in a tracepoint declaration placed in a header file.

They can be used for tracing and performance accounting.

9.2 Usage

Two elements are required for tracepoints :

- A tracepoint definition, placed in a header file.
- The tracepoint statement, in C code.

In order to use tracepoints, you should include `linux/tracepoint.h`.

In `include/trace/events/subsys.h`:

```
#undef TRACE_SYSTEM
#define TRACE_SYSTEM subsys

#if !defined(_TRACE_SUBSYS_H) || defined(TRACE_HEADER_MULTI_READ)
#define _TRACE_SUBSYS_H
```

```
#include <linux/tracepoint.h>

DECLARE_TRACE(subsys_eventname,
              TP_PROTO(int firstarg, struct task_struct *p),
              TP_ARGS(firstarg, p));

#endif /* _TRACE_SUBSYS_H */

/* This part must be outside protection */
#include <trace/define_trace.h>
```

In `subsys/file.c` (where the tracing statement must be added):

```
#include <trace/events/subsys.h>

#define CREATE_TRACE_POINTS
DEFINE_TRACE(subsys_eventname);

void somefct(void)
{
    ...
    trace_subsys_eventname(arg, task);
    ...
}
```

Where :

- `subsys_eventname` is an identifier unique to your event
 - `subsys` is the name of your subsystem.
 - `eventname` is the name of the event to trace.
- `TP_PROTO(int firstarg, struct task_struct *p)` is the prototype of the function called by this tracepoint.
- `TP_ARGS(firstarg, p)` are the parameters names, same as found in the prototype.
- if you use the header in multiple source files, `#define CREATE_TRACE_POINTS` should appear only in one source file.

Connecting a function (probe) to a tracepoint is done by providing a probe (function to call) for the specific tracepoint through `register_trace_subsys_eventname()`. Removing a probe is done through `unregister_trace_subsys_eventname()`; it will remove the probe.

`tracepoint_synchronize_unregister()` must be called before the end of the module exit function to make sure there is no caller left using the probe. This, and the fact that preemption is disabled around the probe call, make sure that probe removal and module unload are safe.

The tracepoint mechanism supports inserting multiple instances of the same tracepoint, but a single definition must be made of a given tracepoint name over all the kernel to make sure no type conflict will occur. Name mangling of the tracepoints is done using the prototypes to make sure typing is correct. Verification of probe type correctness is done at the registration site by

the compiler. Tracepoints can be put in inline functions, inlined static functions, and unrolled loops as well as regular functions.

The naming scheme “`subsys_event`” is suggested here as a convention intended to limit collisions. Tracepoint names are global to the kernel: they are considered as being the same whether they are in the core kernel image or in modules.

If the tracepoint has to be used in kernel modules, an `EXPORT_TRACEPOINT_SYMBOL_GPL()` or `EXPORT_TRACEPOINT_SYMBOL()` can be used to export the defined tracepoints.

If you need to do a bit of work for a tracepoint parameter, and that work is only used for the tracepoint, that work can be encapsulated within an if statement with the following:

```
if (trace_foo_bar_enabled()) {
    int i;
    int tot = 0;

    for (i = 0; i < count; i++)
        tot += calculate_nuggets();

    trace_foo_bar(tot);
}
```

All `trace_<tracepoint>()` calls have a matching `trace_<tracepoint>_enabled()` function defined that returns true if the tracepoint is enabled and false otherwise. The `trace_<tracepoint>()` should always be within the block of the `if (trace_<tracepoint>_enabled())` to prevent races between the tracepoint being enabled and the check being seen.

The advantage of using the `trace_<tracepoint>_enabled()` is that it uses the `static_key` of the tracepoint to allow the if statement to be implemented with jump labels and avoid conditional branches.

Note: The convenience macro `TRACE_EVENT` provides an alternative way to define tracepoints. Check <http://lwn.net/Articles/379903>, <http://lwn.net/Articles/381064> and <http://lwn.net/Articles/383362> for a series of articles with more details.

If you require calling a tracepoint from a header file, it is not recommended to call one directly or to use the `trace_<tracepoint>_enabled()` function call, as tracepoints in header files can have side effects if a header is included from a file that has `CREATE_TRACE_POINTS` set, as well as the `trace_<tracepoint>()` is not that small of an inline and can bloat the kernel if used by other inlined functions. Instead, include `tracepoint-defs.h` and use `tracepoint_enabled()`.

In a C file:

```
void do_trace_foo_bar_wrapper(args)
{
    trace_foo_bar(args);
}
```

In the header file:

```
DECLARE_TRACEPOINT(foo_bar);
```

```
static inline void some_inline_function()
{
    [...]
    if (tracepoint_enabled(foo_bar))
        do_trace_foo_bar_wrapper(args);
    [...]
}
```


EVENT TRACING

Author Theodore Ts'o

Updated Li Zefan and Tom Zanussi

10.1 1. Introduction

Tracepoints (see *Using the Linux Kernel Tracepoints*) can be used without creating custom kernel modules to register probe functions using the event tracing infrastructure.

Not all tracepoints can be traced using the event tracing system; the kernel developer must provide code snippets which define how the tracing information is saved into the tracing buffer, and how the tracing information should be printed.

10.2 2. Using Event Tracing

10.2.1 2.1 Via the 'set_event' interface

The events which are available for tracing can be found in the file `/sys/kernel/debug/tracing/available_events`.

To enable a particular event, such as 'sched_wakeup', simply echo it to `/sys/kernel/debug/tracing/set_event`. For example:

```
# echo sched_wakeup >> /sys/kernel/debug/tracing/set_event
```

Note: '>>' is necessary, otherwise it will firstly disable all the events.

To disable an event, echo the event name to the `set_event` file prefixed with an exclamation point:

```
# echo '!sched_wakeup' >> /sys/kernel/debug/tracing/set_event
```

To disable all events, echo an empty line to the `set_event` file:

```
# echo > /sys/kernel/debug/tracing/set_event
```

To enable all events, echo `*:*` or `*:` to the `set_event` file:

```
# echo *: * > /sys/kernel/debug/tracing/set_event
```

The events are organized into subsystems, such as ext4, irq, sched, etc., and a full event name looks like this: <subsystem>:<event>. The subsystem name is optional, but it is displayed in the available_events file. All of the events in a subsystem can be specified via the syntax <subsystem>:*; for example, to enable all irq events, you can use the command:

```
# echo 'irq:*' > /sys/kernel/debug/tracing/set_event
```

10.2.2 2.2 Via the ‘enable’ toggle

The events available are also listed in /sys/kernel/debug/tracing/events/ hierarchy of directories. To enable event ‘sched_wakeup’:

```
# echo 1 > /sys/kernel/debug/tracing/events/sched/sched_wakeup/enable
```

To disable it:

```
# echo 0 > /sys/kernel/debug/tracing/events/sched/sched_wakeup/enable
```

To enable all events in sched subsystem:

```
# echo 1 > /sys/kernel/debug/tracing/events/sched/enable
```

To enable all events:

```
# echo 1 > /sys/kernel/debug/tracing/events/enable
```

When reading one of these enable files, there are four results:

- 0 - all events this file affects are disabled
- 1 - all events this file affects are enabled
- X - there is a mixture of events enabled and disabled
- ? - this file does not affect any event

10.2.3 2.3 Boot option

In order to facilitate early boot debugging, use boot option:

```
trace_event=[event-list]
```

event-list is a comma separated list of events. See section 2.1 for event format.

10.3 3. Defining an event-enabled tracepoint

See The example provided in `samples/trace_events`

10.4 4. Event formats

Each trace event has a ‘format’ file associated with it that contains a description of each field in a logged event. This information can be used to parse the binary trace stream, and is also the place to find the field names that can be used in event filters (see section 5).

It also displays the format string that will be used to print the event in text mode, along with the event name and ID used for profiling.

Every event has a set of common fields associated with it; these are the fields prefixed with `common_`. The other fields vary between events and correspond to the fields defined in the `TRACE_EVENT` definition for that event.

Each field in the format has the form:

```
field:field-type field-name; offset:N; size:N;
```

where `offset` is the offset of the field in the trace record and `size` is the size of the data item, in bytes.

For example, here’s the information displayed for the ‘`sched_wakeup`’ event:

```
# cat /sys/kernel/debug/tracing/events/sched/sched_wakeup/format

name: sched_wakeup
ID: 60
format:
    field:unsigned short common_type;          offset:0;          size:2;
    field:unsigned char common_flags;          offset:2;          size:1;
    field:unsigned char common_preempt_count;    offset:3;          size:1;
    field:int common_pid;      offset:4;          size:4;
    field:int common_tgid;     offset:8;          size:4;

    field:char comm[TASK_COMM_LEN]; offset:12;          size:16;
    field:pid_t pid;            offset:28;          size:4;
    field:int prio; offset:32;          size:4;
    field:int success;          offset:36;          size:4;
    field:int cpu;  offset:40;          size:4;

print fmt: "task %s:%d [%d] success=%d [%03d]", REC->comm, REC->pid,
           REC->prio, REC->success, REC->cpu
```

This event contains 10 fields, the first 5 common and the remaining 5 event-specific. All the fields for this event are numeric, except for ‘`comm`’ which is a string, a distinction important for event filtering.

10.5 5. Event filtering

Trace events can be filtered in the kernel by associating boolean ‘filter expressions’ with them. As soon as an event is logged into the trace buffer, its fields are checked against the filter expression associated with that event type. An event with field values that ‘match’ the filter will appear in the trace output, and an event whose values don’t match will be discarded. An event with no filter associated with it matches everything, and is the default when no filter has been set for an event.

10.5.1 5.1 Expression syntax

A filter expression consists of one or more ‘predicates’ that can be combined using the logical operators ‘&&’ and ‘||’. A predicate is simply a clause that compares the value of a field contained within a logged event with a constant value and returns either 0 or 1 depending on whether the field value matched (1) or didn’t match (0):

```
field-name relational-operator value
```

Parentheses can be used to provide arbitrary logical groupings and double-quotes can be used to prevent the shell from interpreting operators as shell metacharacters.

The field-names available for use in filters can be found in the ‘format’ files for trace events (see section 4).

The relational-operators depend on the type of the field being tested:

The operators available for numeric fields are:

`==, !=, <, <=, >, >=, &`

And for string fields they are:

`==, !=, ~`

The glob (`~`) accepts a wild card character (`*`,`?`) and character classes (`[]`). For example:

```
prev_comm ~ "*sh"  
prev_comm ~ "sh*"  
prev_comm ~ "*sh*"  
prev_comm ~ "ba*sh"
```

If the field is a pointer that points into user space (for example “filename” from `sys_enter_openat`), then you have to append “.ustring” to the field name:

```
filename.ustring ~ "password"
```

As the kernel will have to know how to retrieve the memory that the pointer is at from user space.

10.5.2 5.2 Setting filters

A filter for an individual event is set by writing a filter expression to the ‘filter’ file for the given event.

For example:

```
# cd /sys/kernel/debug/tracing/events/sched/sched_wakeup
# echo "common_preempt_count > 4" > filter
```

A slightly more involved example:

```
# cd /sys/kernel/debug/tracing/events/signal/signal_generate
# echo "((sig >= 10 && sig < 15) || sig == 17) && comm != bash" > filter
```

If there is an error in the expression, you’ll get an ‘Invalid argument’ error when setting it, and the erroneous string along with an error message can be seen by looking at the filter e.g.:

```
# cd /sys/kernel/debug/tracing/events/signal/signal_generate
# echo "((sig >= 10 && sig < 15) || dsig == 17) && comm != bash" > filter
-bash: echo: write error: Invalid argument
# cat filter
((sig >= 10 && sig < 15) || dsig == 17) && comm != bash
^
parse_error: Field not found
```

Currently the caret (‘^’) for an error always appears at the beginning of the filter string; the error message should still be useful though even without more accurate position info.

10.5.3 5.2.1 Filter limitations

If a filter is placed on a string pointer (char *) that does not point to a string on the ring buffer, but instead points to kernel or user space memory, then, for safety reasons, at most 1024 bytes of the content is copied onto a temporary buffer to do the compare. If the copy of the memory faults (the pointer points to memory that should not be accessed), then the string compare will be treated as not matching.

10.5.4 5.3 Clearing filters

To clear the filter for an event, write a ‘0’ to the event’s filter file.

To clear the filters for all events in a subsystem, write a ‘0’ to the subsystem’s filter file.

10.5.5 5.3 Subsystem filters

For convenience, filters for every event in a subsystem can be set or cleared as a group by writing a filter expression into the filter file at the root of the subsystem. Note however, that if a filter for any event within the subsystem lacks a field specified in the subsystem filter, or if the filter can't be applied for any other reason, the filter for that event will retain its previous setting. This can result in an unintended mixture of filters which could lead to confusing (to the user who might think different filters are in effect) trace output. Only filters that reference just the common fields can be guaranteed to propagate successfully to all events.

Here are a few subsystem filter examples that also illustrate the above points:

Clear the filters on all events in the sched subsystem:

```
# cd /sys/kernel/debug/tracing/events/sched
# echo 0 > filter
# cat sched_switch/filter
none
# cat sched_wakeup/filter
none
```

Set a filter using only common fields for all events in the sched subsystem (all events end up with the same filter):

```
# cd /sys/kernel/debug/tracing/events/sched
# echo common_pid == 0 > filter
# cat sched_switch/filter
common_pid == 0
# cat sched_wakeup/filter
common_pid == 0
```

Attempt to set a filter using a non-common field for all events in the sched subsystem (all events but those that have a `prev_pid` field retain their old filters):

```
# cd /sys/kernel/debug/tracing/events/sched
# echo prev_pid == 0 > filter
# cat sched_switch/filter
prev_pid == 0
# cat sched_wakeup/filter
common_pid == 0
```

10.5.6 5.4 PID filtering

The `set_event_pid` file in the same directory as the top events directory exists, will filter all events from tracing any task that does not have the PID listed in the `set_event_pid` file.

```
# cd /sys/kernel/debug/tracing
# echo $$ > set_event_pid
# echo 1 > events/enable
```

Will only trace events for the current task.

To add more PIDs without losing the PIDs already included, use '>>'.

```
# echo 123 244 1 >> set_event_pid
```

10.6 6. Event triggers

Trace events can be made to conditionally invoke trigger ‘commands’ which can take various forms and are described in detail below; examples would be enabling or disabling other trace events or invoking a stack trace whenever the trace event is hit. Whenever a trace event with attached triggers is invoked, the set of trigger commands associated with that event is invoked. Any given trigger can additionally have an event filter of the same form as described in section 5 (Event filtering) associated with it - the command will only be invoked if the event being invoked passes the associated filter. If no filter is associated with the trigger, it always passes.

Triggers are added to and removed from a particular event by writing trigger expressions to the ‘trigger’ file for the given event.

A given event can have any number of triggers associated with it, subject to any restrictions that individual commands may have in that regard.

Event triggers are implemented on top of “soft” mode, which means that whenever a trace event has one or more triggers associated with it, the event is activated even if it isn’t actually enabled, but is disabled in a “soft” mode. That is, the tracepoint will be called, but just will not be traced, unless of course it’s actually enabled. This scheme allows triggers to be invoked even for events that aren’t enabled, and also allows the current event filter implementation to be used for conditionally invoking triggers.

The syntax for event triggers is roughly based on the syntax for `set_fttrace_filter` ‘fttrace filter commands’ (see the ‘Filter commands’ section of *fttrace - Function Tracer*), but there are major differences and the implementation isn’t currently tied to it in any way, so beware about making generalizations between the two.

Note: Writing into `trace_marker` (See *fttrace - Function Tracer*) can also enable triggers that are written into `/sys/kernel/tracing/events/fttrace/print/trigger`

10.6.1 6.1 Expression syntax

Triggers are added by echoing the command to the ‘trigger’ file:

```
# echo 'command[:count] [if filter]' > trigger
```

Triggers are removed by echoing the same command but starting with ‘!’ to the ‘trigger’ file:

```
# echo '!command[:count] [if filter]' > trigger
```

The `[if filter]` part isn’t used in matching commands when removing, so leaving that off in a ‘!’ command will accomplish the same thing as having it in.

The filter syntax is the same as that described in the ‘Event filtering’ section above.

For ease of use, writing to the trigger file using ‘>’ currently just adds or removes a single trigger and there’s no explicit ‘>>’ support (‘>’ actually behaves like ‘>>’) or truncation support to remove all triggers (you have to use ‘!’ for each one added.)

10.6.2 6.2 Supported trigger commands

The following commands are supported:

- enable_event/disable_event

These commands can enable or disable another trace event whenever the triggering event is hit. When these commands are registered, the other trace event is activated, but disabled in a “soft” mode. That is, the tracepoint will be called, but just will not be traced. The event tracepoint stays in this mode as long as there’s a trigger in effect that can trigger it.

For example, the following trigger causes kmem events to be traced when a read system call is entered, and the :1 at the end specifies that this enablement happens only once:

```
# echo 'enable_event:kmem:kmalloc:1' > \
    /sys/kernel/debug/tracing/events/syscalls/sys_enter_read/trigger
```

The following trigger causes kmem events to stop being traced when a read system call exits. This disablement happens on every read system call exit:

```
# echo 'disable_event:kmem:kmalloc' > \
    /sys/kernel/debug/tracing/events/syscalls/sys_exit_read/trigger
```

The format is:

```
enable_event:<system>:<event>[:count]
disable_event:<system>:<event>[:count]
```

To remove the above commands:

```
# echo '!enable_event:kmem:kmalloc:1' > \
    /sys/kernel/debug/tracing/events/syscalls/sys_enter_read/trigger

# echo '!disable_event:kmem:kmalloc' > \
    /sys/kernel/debug/tracing/events/syscalls/sys_exit_read/trigger
```

Note that there can be any number of enable/disable_event triggers per triggering event, but there can only be one trigger per triggered event. e.g. sys_enter_read can have triggers enabling both kmem:kmalloc and sched:sched_switch, but can’t have two kmem:kmalloc versions such as kmem:kmalloc and kmem:kmalloc:1 or ‘kmem:kmalloc if bytes_req == 256’ and ‘kmem:kmalloc if bytes_alloc == 256’ (they could be combined into a single filter on kmem:kmalloc though).

- stacktrace

This command dumps a stacktrace in the trace buffer whenever the triggering event occurs.

For example, the following trigger dumps a stacktrace every time the kmem tracepoint is hit:


```
# echo 'stacktrace' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger
```

The following trigger dumps a stacktrace the first 5 times a kmalloc request happens with a size $\geq 64K$:

```
# echo 'stacktrace:5 if bytes_req >= 65536' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger
```

The format is:

```
stacktrace[:count]
```

To remove the above commands:

```
# echo '!stacktrace' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger

# echo '!stacktrace:5 if bytes_req >= 65536' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger
```

The latter can also be removed more simply by the following (without the filter):

```
# echo '!stacktrace:5' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger
```

Note that there can be only one stacktrace trigger per triggering event.

- snapshot

This command causes a snapshot to be triggered whenever the triggering event occurs.

The following command creates a snapshot every time a block request queue is unplugged with a depth > 1 . If you were tracing a set of events or functions at the time, the snapshot trace buffer would capture those events when the trigger event occurred:

```
# echo 'snapshot if nr_rq > 1' > \
    /sys/kernel/debug/tracing/events/block/block_unplug/trigger
```

To only snapshot once:

```
# echo 'snapshot:1 if nr_rq > 1' > \
    /sys/kernel/debug/tracing/events/block/block_unplug/trigger
```

To remove the above commands:

```
# echo '!snapshot if nr_rq > 1' > \
    /sys/kernel/debug/tracing/events/block/block_unplug/trigger

# echo '!snapshot:1 if nr_rq > 1' > \
    /sys/kernel/debug/tracing/events/block/block_unplug/trigger
```

Note that there can be only one snapshot trigger per triggering event.

- `traceon/traceoff`

These commands turn tracing on and off when the specified events are hit. The parameter determines how many times the tracing system is turned on and off. If unspecified, there is no limit.

The following command turns tracing off the first time a block request queue is unplugged with a depth > 1. If you were tracing a set of events or functions at the time, you could then examine the trace buffer to see the sequence of events that led up to the trigger event:

```
# echo 'traceoff:1 if nr_rq > 1' > \
    /sys/kernel/debug/tracing/events/block/block_unplug/trigger
```

To always disable tracing when `nr_rq > 1`:

```
# echo 'traceoff if nr_rq > 1' > \
    /sys/kernel/debug/tracing/events/block/block_unplug/trigger
```

To remove the above commands:

```
# echo '!traceoff:1 if nr_rq > 1' > \
    /sys/kernel/debug/tracing/events/block/block_unplug/trigger

# echo '!traceoff if nr_rq > 1' > \
    /sys/kernel/debug/tracing/events/block/block_unplug/trigger
```

Note that there can be only one `traceon` or `traceoff` trigger per triggering event.

- `hist`

This command aggregates event hits into a hash table keyed on one or more trace event format fields (or stacktrace) and a set of running totals derived from one or more trace event format fields and/or event counts (hitcount).

See [Event Histograms](#) for details and examples.

10.7 7. In-kernel trace event API

In most cases, the command-line interface to trace events is more than sufficient. Sometimes, however, applications might find the need for more complex relationships than can be expressed through a simple series of linked command-line expressions, or putting together sets of commands may be simply too cumbersome. An example might be an application that needs to ‘listen’ to the trace stream in order to maintain an in-kernel state machine detecting, for instance, when an illegal kernel state occurs in the scheduler.

The trace event subsystem provides an in-kernel API allowing modules or other kernel code to generate user-defined ‘synthetic’ events at will, which can be used to either augment the existing trace stream and/or signal that a particular important state has occurred.

A similar in-kernel API is also available for creating `kprobe` and `kretprobe` events.

Both the synthetic event and `kretprobe` event APIs are built on top of a lower-level “`dyn-event_cmd`” event command API, which is also available for more specialized applications, or as the basis of other higher-level trace event APIs.

The API provided for these purposes is describe below and allows the following:

- dynamically creating synthetic event definitions
- dynamically creating kprobe and kretprobe event definitions
- tracing synthetic events from in-kernel code
- the low-level “dynevent_cmd” API

10.7.1 7.1 Dyamically creating synthetic event definitions

There are a couple ways to create a new synthetic event from a kernel module or other kernel code.

The first creates the event in one step, using `synth_event_create()`. In this method, the name of the event to create and an array defining the fields is supplied to `synth_event_create()`. If successful, a synthetic event with that name and fields will exist following that call. For example, to create a new “schedtest” synthetic event:

```
ret = synth_event_create("schedtest", sched_fields,
                        ARRAY_SIZE(sched_fields), THIS_MODULE);
```

The `sched_fields` param in this example points to an array of struct `synth_field_desc`, each of which describes an event field by type and name:

```
static struct synth_field_desc sched_fields[] = {
    { .type = "pid_t",           .name = "next_pid_field" },
    { .type = "char[16]",       .name = "next_comm_field" },
    { .type = "u64",            .name = "ts_ns" },
    { .type = "u64",            .name = "ts_ms" },
    { .type = "unsigned int",    .name = "cpu" },
    { .type = "char[64]",       .name = "my_string_field" },
    { .type = "int",            .name = "my_int_field" },
};
```

See `synth_field_size()` for available types.

If `field_name` contains `[n]`, the field is considered to be a static array.

If `field_names` contains`[]` (no subscript), the field is considered to be a dynamic array, which will only take as much space in the event as is required to hold the array.

Because space for an event is reserved before assigning field values to the event, using dynamic arrays implies that the piecewise in-kernel API described below can’t be used with dynamic arrays. The other non-piecewise in-kernel APIs can, however, be used with dynamic arrays.

If the event is created from within a module, a pointer to the module must be passed to `synth_event_create()`. This will ensure that the trace buffer won’t contain unreadable events when the module is removed.

At this point, the event object is ready to be used for generating new events.

In the second method, the event is created in several steps. This allows events to be created dynamically and without the need to create and populate an array of fields beforehand.

To use this method, an empty or partially empty synthetic event should first be created using `synth_event_gen_cmd_start()` or `synth_event_gen_cmd_array_start()`. For `synth_event_gen_cmd_start()`, the name of the event along with one or more pairs of args each pair representing a ‘type field_name;’ field specification should be supplied. For `synth_event_gen_cmd_array_start()`, the name of the event along with an array of struct `synth_field_desc` should be supplied. Before calling `synth_event_gen_cmd_start()` or `synth_event_gen_cmd_array_start()`, the user should create and initialize a `dynevent_cmd` object using `synth_event_cmd_init()`.

For example, to create a new “schedtest” synthetic event with two fields:

```
struct dynevent_cmd cmd;
char *buf;

/* Create a buffer to hold the generated command */
buf = kzalloc(MAX_DYNEVENT_CMD_LEN, GFP_KERNEL);

/* Before generating the command, initialize the cmd object */
synth_event_cmd_init(&cmd, buf, MAX_DYNEVENT_CMD_LEN);

ret = synth_event_gen_cmd_start(&cmd, "schedtest", THIS_MODULE,
                               "pid_t", "next_pid_field",
                               "u64", "ts_ns");
```

Alternatively, using an array of struct `synth_field_desc` fields containing the same information:

```
ret = synth_event_gen_cmd_array_start(&cmd, "schedtest", THIS_MODULE,
                                     fields, n_fields);
```

Once the synthetic event object has been created, it can then be populated with more fields. Fields are added one by one using `synth_event_add_field()`, supplying the `dynevent_cmd` object, a field type, and a field name. For example, to add a new int field named “intfield”, the following call should be made:

```
ret = synth_event_add_field(&cmd, "int", "intfield");
```

See `synth_field_size()` for available types. If `field_name` contains [n] the field is considered to be an array.

A group of fields can also be added all at once using an array of `synth_field_desc` with `add_synth_fields()`. For example, this would add just the first four `sched_fields`:

```
ret = synth_event_add_fields(&cmd, sched_fields, 4);
```

If you already have a string of the form ‘type field_name’, `synth_event_add_field_str()` can be used to add it as-is; it will also automatically append a ‘;’ to the string.

Once all the fields have been added, the event should be finalized and registered by calling the `synth_event_gen_cmd_end()` function:

```
ret = synth_event_gen_cmd_end(&cmd);
```

At this point, the event object is ready to be used for tracing new events.

10.7.2 7.2 Tracing synthetic events from in-kernel code

To trace a synthetic event, there are several options. The first option is to trace the event in one call, using `synth_event_trace()` with a variable number of values, or `synth_event_trace_array()` with an array of values to be set. A second option can be used to avoid the need for a pre-formed array of values or list of arguments, via `synth_event_trace_start()` and `synth_event_trace_end()` along with `synth_event_add_next_val()` or `synth_event_add_val()` to add the values piecewise.

10.7.3 7.2.1 Tracing a synthetic event all at once

To trace a synthetic event all at once, the `synth_event_trace()` or `synth_event_trace_array()` functions can be used.

The `synth_event_trace()` function is passed the `trace_event_file` representing the synthetic event (which can be retrieved using `trace_get_event_file()` using the synthetic event name, “synthetic” as the system name, and the trace instance name (NULL if using the global trace array)), along with an variable number of u64 args, one for each synthetic event field, and the number of values being passed.

So, to trace an event corresponding to the synthetic event definition above, code like the following could be used:

```
ret = synth_event_trace(create_synth_test, 7, /* number of values */
                        444, /* next_pid_field */
                        (u64)"clackers", /* next_comm_field */
                        1000000, /* ts_ns */
                        1000, /* ts_ms */
                        smp_processor_id(), /* cpu */
                        (u64)"Thneed", /* my_string_field */
                        999); /* my_int_field */
```

All vals should be cast to u64, and string vals are just pointers to strings, cast to u64. Strings will be copied into space reserved in the event for the string, using these pointers.

Alternatively, the `synth_event_trace_array()` function can be used to accomplish the same thing. It is passed the `trace_event_file` representing the synthetic event (which can be retrieved using `trace_get_event_file()` using the synthetic event name, “synthetic” as the system name, and the trace instance name (NULL if using the global trace array)), along with an array of u64, one for each synthetic event field.

To trace an event corresponding to the synthetic event definition above, code like the following could be used:

```
u64 vals[7];

vals[0] = 777; /* next_pid_field */
vals[1] = (u64)"tiddlywinks"; /* next_comm_field */
vals[2] = 1000000; /* ts_ns */
vals[3] = 1000; /* ts_ms */
vals[4] = smp_processor_id(); /* cpu */
vals[5] = (u64)"thneed"; /* my_string_field */
vals[6] = 398; /* my_int_field */
```

The 'vals' array is just an array of u64, the number of which must match the number of field in the synthetic event, and which must be in the same order as the synthetic event fields.

All vals should be cast to u64, and string vals are just pointers to strings, cast to u64. Strings will be copied into space reserved in the event for the string, using these pointers.

In order to trace a synthetic event, a pointer to the trace event file is needed. The `trace_get_event_file()` function can be used to get it - it will find the file in the given trace instance (in this case NULL since the top trace array is being used) while at the same time preventing the instance containing it from going away:

```
schedtest_event_file = trace_get_event_file(NULL, "synthetic",
                                             "schedtest");
```

Before tracing the event, it should be enabled in some way, otherwise the synthetic event won't actually show up in the trace buffer.

To enable a synthetic event from the kernel, `trace_array_set_clr_event()` can be used (which is not specific to synthetic events, so does need the "synthetic" system name to be specified explicitly).

To enable the event, pass 'true' to it:

```
trace_array_set_clr_event(schedtest_event_file->tr,
                          "synthetic", "schedtest", true);
```

To disable it pass false:

```
trace_array_set_clr_event(schedtest_event_file->tr,
                          "synthetic", "schedtest", false);
```

Finally, `synth_event_trace_array()` can be used to actually trace the event, which should be visible in the trace buffer afterwards:

```
ret = synth_event_trace_array(schedtest_event_file, vals,
                              ARRAY_SIZE(vals));
```

To remove the synthetic event, the event should be disabled, and the trace instance should be 'put' back using `trace_put_event_file()`:

```
trace_array_set_clr_event(schedtest_event_file->tr,
                          "synthetic", "schedtest", false);
trace_put_event_file(schedtest_event_file);
```

If those have been successful, `synth_event_delete()` can be called to remove the event:

```
ret = synth_event_delete("schedtest");
```

10.7.4 7.2.2 Tracing a synthetic event piecewise

To trace a synthetic using the piecewise method described above, the `synth_event_trace_start()` function is used to 'open' the synthetic event trace:

```
struct synth_event_trace_state trace_state;

ret = synth_event_trace_start(schedtest_event_file, &trace_state);
```

It's passed the `trace_event_file` representing the synthetic event using the same methods as described above, along with a pointer to a struct `synth_event_trace_state` object, which will be zeroed before use and used to maintain state between this and following calls.

Once the event has been opened, which means space for it has been reserved in the trace buffer, the individual fields can be set. There are two ways to do that, either one after another for each field in the event, which requires no lookups, or by name, which does. The tradeoff is flexibility in doing the assignments vs the cost of a lookup per field.

To assign the values one after the other without lookups, `synth_event_add_next_val()` should be used. Each call is passed the same `synth_event_trace_state` object used in the `synth_event_trace_start()`, along with the value to set the next field in the event. After each field is set, the 'cursor' points to the next field, which will be set by the subsequent call, continuing until all the fields have been set in order. The same sequence of calls as in the above examples using this method would be (without error-handling code):

```
/* next_pid_field */
ret = synth_event_add_next_val(777, &trace_state);

/* next_comm_field */
ret = synth_event_add_next_val((u64)"slinky", &trace_state);

/* ts_ns */
ret = synth_event_add_next_val(1000000, &trace_state);

/* ts_ms */
ret = synth_event_add_next_val(1000, &trace_state);

/* cpu */
ret = synth_event_add_next_val(smp_processor_id(), &trace_state);

/* my_string_field */
ret = synth_event_add_next_val((u64)"thneed_2.01", &trace_state);

/* my_int_field */
ret = synth_event_add_next_val(395, &trace_state);
```

To assign the values in any order, `synth_event_add_val()` should be used. Each call is passed the same `synth_event_trace_state` object used in the `synth_event_trace_start()`, along with the field name of the field to set and the value to set it to. The same sequence of calls as in the above examples using this method would be (without error-handling code):

```
ret = synth_event_add_val("next_pid_field", 777, &trace_state);
ret = synth_event_add_val("next_comm_field", (u64)"silly putty",
```

```
                                &trace_state);
ret = synth_event_add_val("ts_ns", 1000000, &trace_state);
ret = synth_event_add_val("ts_ms", 1000, &trace_state);
ret = synth_event_add_val("cpu", smp_processor_id(), &trace_state);
ret = synth_event_add_val("my_string_field", (u64)"thneed_9",
                                &trace_state);
ret = synth_event_add_val("my_int_field", 3999, &trace_state);
```

Note that `synth_event_add_next_val()` and `synth_event_add_val()` are incompatible if used within the same trace of an event - either one can be used but not both at the same time.

Finally, the event won't be actually traced until it's 'closed', which is done using `synth_event_trace_end()`, which takes only the struct `synth_event_trace_state` object used in the previous calls:

```
ret = synth_event_trace_end(&trace_state);
```

Note that `synth_event_trace_end()` must be called at the end regardless of whether any of the add calls failed (say due to a bad field name being passed in).

10.7.5 7.3 Dynamically creating kprobe and kretprobe event definitions

To create a kprobe or kretprobe trace event from kernel code, the `kprobe_event_gen_cmd_start()` or `kretprobe_event_gen_cmd_start()` functions can be used.

To create a kprobe event, an empty or partially empty kprobe event should first be created using `kprobe_event_gen_cmd_start()`. The name of the event and the probe location should be specified along with one or args each representing a probe field should be supplied to this function. Before calling `kprobe_event_gen_cmd_start()`, the user should create and initialize a `dynevent_cmd` object using `kprobe_event_cmd_init()`.

For example, to create a new "schedtest" kprobe event with two fields:

```
struct dynevent_cmd cmd;
char *buf;

/* Create a buffer to hold the generated command */
buf = kzalloc(MAX_DYNEVENT_CMD_LEN, GFP_KERNEL);

/* Before generating the command, initialize the cmd object */
kprobe_event_cmd_init(&cmd, buf, MAX_DYNEVENT_CMD_LEN);

/*
 * Define the gen_kprobe_test event with the first 2 kprobe
 * fields.
 */
ret = kprobe_event_gen_cmd_start(&cmd, "gen_kprobe_test", "do_sys_open",
                                "dfd=%ax", "filename=%dx");
```

Once the kprobe event object has been created, it can then be populated with more fields. Fields can be added using `kprobe_event_add_fields()`, supplying the `dynevent_cmd` object along with

a variable arg list of probe fields. For example, to add a couple additional fields, the following call could be made:

```
ret = kprobe_event_add_fields(&cmd, "flags=%cx", "mode=+4($stack)");
```

Once all the fields have been added, the event should be finalized and registered by calling the `kprobe_event_gen_cmd_end()` or `kretprobe_event_gen_cmd_end()` functions, depending on whether a `kprobe` or `kretprobe` command was started:

```
ret = kprobe_event_gen_cmd_end(&cmd);
```

or:

```
ret = kretprobe_event_gen_cmd_end(&cmd);
```

At this point, the event object is ready to be used for tracing new events.

Similarly, a `kretprobe` event can be created using `kretprobe_event_gen_cmd_start()` with a probe name and location and additional params such as `$retval`:

```
ret = kretprobe_event_gen_cmd_start(&cmd, "gen_kretprobe_test",
                                   "do_sys_open", "$retval");
```

Similar to the synthetic event case, code like the following can be used to enable the newly created `kprobe` event:

```
gen_kprobe_test = trace_get_event_file(NULL, "kprobes", "gen_kprobe_test");
ret = trace_array_set_clr_event(gen_kprobe_test->tr,
                               "kprobes", "gen_kprobe_test", true);
```

Finally, also similar to synthetic events, the following code can be used to give the `kprobe` event file back and delete the event:

```
trace_put_event_file(gen_kprobe_test);
ret = kprobe_event_delete("gen_kprobe_test");
```

10.7.6 7.4 The “dynevent_cmd” low-level API

Both the in-kernel synthetic event and `kprobe` interfaces are built on top of a lower-level “dynevent_cmd” interface. This interface is meant to provide the basis for higher-level interfaces such as the synthetic and `kprobe` interfaces, which can be used as examples.

The basic idea is simple and amounts to providing a general-purpose layer that can be used to generate trace event commands. The generated command strings can then be passed to the command-parsing and event creation code that already exists in the trace event subsystem for creating the corresponding trace events.

In a nutshell, the way it works is that the higher-level interface code creates a struct `dynevent_cmd` object, then uses a couple functions, `dynevent_arg_add()` and `dynevent_arg_pair_add()` to build up a command string, which finally causes the command to be

executed using the `dynevent_create()` function. The details of the interface are described below.

The first step in building a new command string is to create and initialize an instance of a `dynevent_cmd`. Here, for instance, we create a `dynevent_cmd` on the stack and initialize it:

```
struct dynevent_cmd cmd;
char *buf;
int ret;

buf = kzalloc(MAX_DYNEVENT_CMD_LEN, GFP_KERNEL);

dynevent_cmd_init(cmd, buf, maxlen, DYNEVENT_TYPE_F00,
                  foo_event_run_command);
```

The `dynevent_cmd` initialization needs to be given a user-specified buffer and the length of the buffer (`MAX_DYNEVENT_CMD_LEN` can be used for this purpose - at 2k it's generally too big to be comfortably put on the stack, so is dynamically allocated), a `dynevent` type id, which is meant to be used to check that further API calls are for the correct command type, and a pointer to an event-specific `run_command()` callback that will be called to actually execute the event-specific command function.

Once that's done, the command string can be built up by successive calls to argument-adding functions.

To add a single argument, define and initialize a `struct dynevent_arg` or `struct dynevent_arg_pair` object. Here's an example of the simplest possible arg addition, which is simply to append the given string as a whitespace-separated argument to the command:

```
struct dynevent_arg arg;

dynevent_arg_init(&arg, NULL, 0);

arg.str = name;

ret = dynevent_arg_add(cmd, &arg);
```

The `arg` object is first initialized using `dynevent_arg_init()` and in this case the parameters are `NULL` or `0`, which means there's no optional sanity-checking function or separator appended to the end of the `arg`.

Here's another more complicated example using an 'arg pair', which is used to create an argument that consists of a couple components added together as a unit, for example, a 'type field_name;' `arg` or a simple expression `arg` e.g. 'flags=%cx':

```
struct dynevent_arg_pair arg_pair;

dynevent_arg_pair_init(&arg_pair, dynevent_foo_check_arg_fn, 0, ';');

arg_pair.lhs = type;
arg_pair.rhs = name;

ret = dynevent_arg_pair_add(cmd, &arg_pair);
```

Again, the `arg_pair` is first initialized, in this case with a callback function used to check the sanity of the args (for example, that neither part of the pair is NULL), along with a character to be used to add an operator between the pair (here none) and a separator to be appended onto the end of the arg pair (here `';`').

There's also a `dynevent_str_add()` function that can be used to simply add a string as-is, with no spaces, delimiters, or arg check.

Any number of `dynevent_*_add()` calls can be made to build up the string (until its length surpasses `cmd->maxlen`). When all the arguments have been added and the command string is complete, the only thing left to do is run the command, which happens by simply calling `dynevent_create()`:

```
ret = dynevent_create(&cmd);
```

At that point, if the return value is 0, the dynamic event has been created and is ready to use.

See the `dynevent_cmd` function definitions themselves for the details of the API.

SUBSYSTEM TRACE POINTS: KMEM

The kmem tracing system captures events related to object and page allocation within the kernel. Broadly speaking there are five major subheadings.

- Slab allocation of small objects of unknown type (kmalloc)
- Slab allocation of small objects of known type
- Page allocation
- Per-CPU Allocator Activity
- External Fragmentation

This document describes what each of the tracepoints is and why they might be useful.

11.1 1. Slab allocation of small objects of unknown type

```
kmalloc                call_site=%lx ptr=%p bytes_req=%zu bytes_alloc=%zu gfp_  
→ flags=%s  
kmalloc_node call_site=%lx ptr=%p bytes_req=%zu bytes_alloc=%zu gfp_flags=%s  
→ node=%d  
kfree                call_site=%lx ptr=%p
```

Heavy activity for these events may indicate that a specific cache is justified, particularly if kmalloc slab pages are getting significantly internal fragmented as a result of the allocation pattern. By correlating kmalloc with kfree, it may be possible to identify memory leaks and where the allocation sites were.

11.2 2. Slab allocation of small objects of known type

```
kmem_cache_alloc      call_site=%lx ptr=%p bytes_req=%zu bytes_alloc=%zu gfp_  
→ flags=%s  
kmem_cache_alloc_node call_site=%lx ptr=%p bytes_req=%zu bytes_alloc=%zu gfp_  
→ flags=%s node=%d  
kmem_cache_free       call_site=%lx ptr=%p
```

These events are similar in usage to the kmalloc-related events except that it is likely easier to pin the event down to a specific cache. At the time of writing, no information is available

on what slab is being allocated from, but the `call_site` can usually be used to extrapolate that information.

11.3 3. Page allocation

```
mm_page_alloc          page=%p pfn=%lu order=%d migratetype=%d gfp_flags=%s
mm_page_alloc_zone_locked page=%p pfn=%lu order=%u migratetype=%d cpu=%d
↳ percpu_refill=%d
mm_page_free           page=%p pfn=%lu order=%d
mm_page_free_batched   page=%p pfn=%lu order=%d cold=%d
```

These four events deal with page allocation and freeing. `mm_page_alloc` is a simple indicator of page allocator activity. Pages may be allocated from the per-CPU allocator (high performance) or the buddy allocator.

If pages are allocated directly from the buddy allocator, the `mm_page_alloc_zone_locked` event is triggered. This event is important as high amounts of activity imply high activity on the `zone->lock`. Taking this lock impairs performance by disabling interrupts, dirtying cache lines between CPUs and serialising many CPUs.

When a page is freed directly by the caller, the only `mm_page_free` event is triggered. Significant amounts of activity here could indicate that the callers should be batching their activities.

When pages are freed in batch, the also `mm_page_free_batched` is triggered. Broadly speaking, pages are taken off the LRU lock in bulk and freed in batch with a page list. Significant amounts of activity here could indicate that the system is under memory pressure and can also indicate contention on the `lruvec->lru_lock`.

11.4 4. Per-CPU Allocator Activity

```
mm_page_alloc_zone_locked page=%p pfn=%lu order=%u migratetype=%d cpu=%d
↳ percpu_refill=%d
mm_page_pcpu_drain        page=%p pfn=%lu order=%d cpu=%d migratetype=%d
```

In front of the page allocator is a per-cpu page allocator. It exists only for order-0 pages, reduces contention on the `zone->lock` and reduces the amount of writing on struct page.

When a per-CPU list is empty or pages of the wrong type are allocated, the `zone->lock` will be taken once and the per-CPU list refilled. The event triggered is `mm_page_alloc_zone_locked` for each page allocated with the event indicating whether it is for a `percpu_refill` or not.

When the per-CPU list is too full, a number of pages are freed, each one which triggers a `mm_page_pcpu_drain` event.

The individual nature of the events is so that pages can be tracked between allocation and freeing. A number of drain or refill pages that occur consecutively imply the `zone->lock` being taken once. Large amounts of per-CPU refills and drains could imply an imbalance between CPUs where too much work is being concentrated in one place. It could also indicate that the per-CPU lists should be a larger size. Finally, large amounts of refills on one CPU and drains on another could be a factor in causing large amounts of cache line bounces due to writes

between CPUs and worth investigating if pages can be allocated and freed on the same CPU through some algorithm change.

11.5 5. External Fragmentation

```
mm_page_alloc_extfrag      page=%p pfn=%lu alloc_order=%d fallback_order=%d
↪ pageblock_order=%d alloc_migratetype=%d fallback_migratetype=%d fragmenting=
↪ %d change_ownership=%d
```

External fragmentation affects whether a high-order allocation will be successful or not. For some types of hardware, this is important although it is avoided where possible. If the system is using huge pages and needs to be able to resize the pool over the lifetime of the system, this value is important.

Large numbers of this event implies that memory is fragmenting and high-order allocations will start failing at some time in the future. One means of reducing the occurrence of this event is to increase the size of `min_free_kbytes` in increments of `3*pageblock_size*nr_online_nodes` where `pageblock_size` is usually the size of the default hugepage size.

SUBSYSTEM TRACE POINTS: POWER

The power tracing system captures events related to power transitions within the kernel. Broadly speaking there are three major subheadings:

- Power state switch which reports events related to suspend (S-states), cpuidle (C-states) and cpufreq (P-states)
- System clock related changes
- Power domains related changes and transitions

This document describes what each of the tracepoints is and why they might be useful.

Cf. `include/trace/events/power.h` for the events definitions.

12.1 1. Power state switch events

12.1.1 1.1 Trace API

A 'cpu' event class gathers the CPU-related events: cpuidle and cpufreq.

cpu_idle	"state=%lu cpu_id=%lu"
cpu_frequency	"state=%lu cpu_id=%lu"
cpu_frequency_limits	"min=%lu max=%lu cpu_id=%lu"

A suspend event is used to indicate the system going in and out of the suspend mode:

machine_suspend	"state=%lu"
-----------------	-------------

Note: the value of '-1' or '4294967295' for state means an exit from the current state, i.e. `trace_cpu_idle(4, smp_processor_id())` means that the system enters the idle state 4, while `trace_cpu_idle(PWR_EVENT_EXIT, smp_processor_id())` means that the system exits the previous idle state.

The event which has 'state=4294967295' in the trace is very important to the user space tools which are using it to detect the end of the current state, and so to correctly draw the states diagrams and to calculate accurate statistics etc.

12.2 2. Clocks events

The clock events are used for clock enable/disable and for clock rate change.

clock_enable	"%s state=%lu cpu_id=%lu"
clock_disable	"%s state=%lu cpu_id=%lu"
clock_set_rate	"%s state=%lu cpu_id=%lu"

The first parameter gives the clock name (e.g. "gpio1_iclk"). The second parameter is '1' for enable, '0' for disable, the target clock rate for set_rate.

12.3 3. Power domains events

The power domain events are used for power domains transitions

power_domain_target	"%s state=%lu cpu_id=%lu"
---------------------	---------------------------

The first parameter gives the power domain name (e.g. "mpu_pwrdom"). The second parameter is the power domain target state.

12.4 4. PM QoS events

The PM QoS events are used for QoS add/update/remove request and for target/flags update.

pm_qos_update_target	"action=%s prev_value=%d curr_value=%d"
pm_qos_update_flags	"action=%s prev_value=0x%x curr_value=0x%x"

The first parameter gives the QoS action name (e.g. "ADD_REQ"). The second parameter is the previous QoS value. The third parameter is the current QoS value to update.

There are also events used for device PM QoS add/update/remove request.

dev_pm_qos_add_request	"device=%s type=%s new_value=%d"
dev_pm_qos_update_request	"device=%s type=%s new_value=%d"
dev_pm_qos_remove_request	"device=%s type=%s new_value=%d"

The first parameter gives the device name which tries to add/update/remove QoS requests. The second parameter gives the request type (e.g. "DEV_PM_QOS_RESUME_LATENCY"). The third parameter is value to be added/updated/removed.

And, there are events used for CPU latency QoS add/update/remove request.

pm_qos_add_request	"value=%d"
pm_qos_update_request	"value=%d"
pm_qos_remove_request	"value=%d"

The parameter is the value to be added/updated/removed.

NMI TRACE EVENTS

These events normally show up here:

```
/sys/kernel/debug/tracing/events/nmi
```

13.1 nmi_handler

You might want to use this tracepoint if you suspect that your NMI handlers are hogging large amounts of CPU time. The kernel will warn if it sees long-running handlers:

```
INFO: NMI handler took too long to run: 9.207 msecs
```

and this tracepoint will allow you to drill down and get some more details.

Let's say you suspect that `perf_event_nmi_handler()` is causing you some problems and you only want to trace that handler specifically. You need to find its address:

```
$ grep perf_event_nmi_handler /proc/kallsyms
ffffffff81625600 t perf_event_nmi_handler
```

Let's also say you are only interested in when that function is really hogging a lot of CPU time, like a millisecond at a time. Note that the kernel's output is in milliseconds, but the input to the filter is in nanoseconds! You can filter on 'delta_ns':

```
cd /sys/kernel/debug/tracing/events/nmi/nmi_handler
echo 'handler==0xffffffff81625600 && delta_ns>1000000' > filter
echo 1 > enable
```

Your output would then look like:

```
$ cat /sys/kernel/debug/tracing/trace_pipe
<idle>-0      [000] d.h3    505.397558: nmi_handler: perf_event_nmi_handler()
↳delta_ns: 3236765 handled: 1
<idle>-0      [000] d.h3    505.805893: nmi_handler: perf_event_nmi_handler()
↳delta_ns: 3174234 handled: 1
<idle>-0      [000] d.h3    506.158206: nmi_handler: perf_event_nmi_handler()
↳delta_ns: 3084642 handled: 1
<idle>-0      [000] d.h3    506.334346: nmi_handler: perf_event_nmi_handler()
↳delta_ns: 3080351 handled: 1
```


MSR TRACE EVENTS

The x86 kernel supports tracing most MSR (Model Specific Register) accesses. To see the definition of the MSRs on Intel systems please see the SDM at <https://www.intel.com/sdm> (Volume 3)

Available trace points:

`/sys/kernel/debug/tracing/events/msr/`

Trace MSR reads:

`read_msr`

- `msr`: MSR number
- `val`: Value written
- `failed`: 1 if the access failed, otherwise 0

Trace MSR writes:

`write_msr`

- `msr`: MSR number
- `val`: Value written
- `failed`: 1 if the access failed, otherwise 0

Trace RDPMC in kernel:

`rdpmc`

The trace data can be post processed with the `postprocess/decode_msr.py` script:

```
cat /sys/kernel/debug/tracing/trace | decode_msr.py /usr/src/linux/include/asm/  
↳msr-index.h
```

to add symbolic MSR names.

IN-KERNEL MEMORY-MAPPED I/O TRACING

Home page and links to optional user space tools:

<https://nouveau.freedesktop.org/wiki/MmioTrace>

MMIO tracing was originally developed by Intel around 2003 for their Fault Injection Test Harness. In Dec 2006 - Jan 2007, using the code from Intel, Jeff Muizelaar created a tool for tracing MMIO accesses with the Nouveau project in mind. Since then many people have contributed.

Mmiotrace was built for reverse engineering any memory-mapped IO device with the Nouveau project as the first real user. Only x86 and x86_64 architectures are supported.

Out-of-tree mmiotrace was originally modified for mainline inclusion and ftrace framework by Pekka Paalanen <pq@iki.fi>.

15.1 Preparation

Mmiotrace feature is compiled in by the CONFIG_MMIOTRACE option. Tracing is disabled by default, so it is safe to have this set to yes. SMP systems are supported, but tracing is unreliable and may miss events if more than one CPU is on-line, therefore mmiotrace takes all but one CPU off-line during run-time activation. You can re-enable CPUs by hand, but you have been warned, there is no way to automatically detect if you are losing events due to CPUs racing.

15.2 Usage Quick Reference

```
$ mount -t debugfs debugfs /sys/kernel/debug
$ echo mmiotrace > /sys/kernel/debug/tracing/current_tracer
$ cat /sys/kernel/debug/tracing/trace_pipe > mydump.txt &
Start X or whatever.
$ echo "X is up" > /sys/kernel/debug/tracing/trace_marker
$ echo nop > /sys/kernel/debug/tracing/current_tracer
Check for lost events.
```

15.3 Usage

Make sure debugfs is mounted to /sys/kernel/debug. If not (requires root privileges):

```
$ mount -t debugfs debugfs /sys/kernel/debug
```

Check that the driver you are about to trace is not loaded.

Activate mmiotrace (requires root privileges):

```
$ echo mmiotrace > /sys/kernel/debug/tracing/current_tracer
```

Start storing the trace:

```
$ cat /sys/kernel/debug/tracing/trace_pipe > mydump.txt &
```

The 'cat' process should stay running (sleeping) in the background.

Load the driver you want to trace and use it. Mmiotrace will only catch MMIO accesses to areas that are ioremapped while mmiotrace is active.

During tracing you can place comments (markers) into the trace by `$ echo "X is up" > /sys/kernel/debug/tracing/trace_marker` This makes it easier to see which part of the (huge) trace corresponds to which action. It is recommended to place descriptive markers about what you do.

Shut down mmiotrace (requires root privileges):

```
$ echo nop > /sys/kernel/debug/tracing/current_tracer
```

The 'cat' process exits. If it does not, kill it by issuing 'fg' command and pressing ctrl+c.

Check that mmiotrace did not lose events due to a buffer filling up. Either:

```
$ grep -i lost mydump.txt
```

which tells you exactly how many events were lost, or use:

```
$ dmesg
```

to view your kernel log and look for "mmiotrace has lost events" warning. If events were lost, the trace is incomplete. You should enlarge the buffers and try again. Buffers are enlarged by first seeing how large the current buffers are:

```
$ cat /sys/kernel/debug/tracing/buffer_size_kb
```

gives you a number. Approximately double this number and write it back, for instance:

```
$ echo 128000 > /sys/kernel/debug/tracing/buffer_size_kb
```

Then start again from the top.

If you are doing a trace for a driver project, e.g. Nouveau, you should also do the following before sending your results:


```
$ lspci -vvv > lspci.txt
$ dmesg > dmesg.txt
$ tar zcf pciid-nick-mmioTRACE.tar.gz mydump.txt lspci.txt dmesg.txt
```

and then send the .tar.gz file. The trace compresses considerably. Replace “pciid” and “nick” with the PCI ID or model name of your piece of hardware under investigation and your nick-name.

15.4 How MmioTRACE Works

Access to hardware IO-memory is gained by mapping addresses from PCI bus by calling one of the `ioremap_*()` functions. MmioTRACE is hooked into the `__ioremap()` function and gets called whenever a mapping is created. Mapping is an event that is recorded into the trace log. Note that ISA range mappings are not caught, since the mapping always exists and is returned directly.

MMIO accesses are recorded via page faults. Just before `__ioremap()` returns, the mapped pages are marked as not present. Any access to the pages causes a fault. The page fault handler calls mmioTRACE to handle the fault. MmioTRACE marks the page present, sets TF flag to achieve single stepping and exits the fault handler. The instruction that faulted is executed and debug trap is entered. Here mmioTRACE again marks the page as not present. The instruction is decoded to get the type of operation (read/write), data width and the value read or written. These are stored to the trace log.

Setting the page present in the page fault handler has a race condition on SMP machines. During the single stepping other CPUs may run freely on that page and events can be missed without a notice. Re-enabling other CPUs during tracing is discouraged.

15.5 Trace Log Format

The raw log is text and easily filtered with e.g. `grep` and `awk`. One record is one line in the log. A record starts with a keyword, followed by keyword- dependent arguments. Arguments are separated by a space, or continue until the end of line. The format for version 20070824 is as follows:

15.6 Explanation Keyword Space-separated arguments

read event R width, timestamp, map id, physical, value, PC, PID write event W width, timestamp, map id, physical, value, PC, PID ioremap event MAP timestamp, map id, physical, virtual, length, PC, PID iounmap event UNMAP timestamp, map id, PC, PID marker MARK timestamp, text version VERSION the string “20070824” info for reader LSPCI one line from `lspci -v` PCI address map PCIDEV space-separated /proc/bus/pci/devices data unk. opcode UNKNOWN timestamp, map id, physical, data, PC, PID

Timestamp is in seconds with decimals. Physical is a PCI bus address, virtual is a kernel virtual address. Width is the data width in bytes and value is the data value. Map id is an arbitrary id number identifying the mapping that was used in an operation. PC is the program counter and

PID is process id. PC is zero if it is not recorded. PID is always zero as tracing MMIO accesses originating in user space memory is not yet supported.

For instance, the following awk filter will pass all 32-bit writes that target physical addresses in the range [0xfb73ce40, 0xfb800000]

```
$ awk '/W 4 / { adr=strtonum($5); if (adr >= 0xfb73ce40 &&
adr < 0xfb800000) print; }'
```

15.7 Tools for Developers

The user space tools include utilities for:

- replacing numeric addresses and values with hardware register names
- replaying MMIO logs, i.e., re-executing the recorded writes

EVENT HISTOGRAMS

Documentation written by Tom Zanussi

16.1 1. Introduction

Histogram triggers are special event triggers that can be used to aggregate trace event data into histograms. For information on trace events and event triggers, see *Event Tracing*.

16.2 2. Histogram Trigger Command

A histogram trigger command is an event trigger command that aggregates event hits into a hash table keyed on one or more trace event format fields (or stacktrace) and a set of running totals derived from one or more trace event format fields and/or event counts (hitcount).

The format of a hist trigger is as follows:

```
hist:keys=<field1[,field2,...]>[:values=<field1[,field2,...]>]  
[:sort=<field1[,field2,...]>][:size=#entries][:pause][:continue]  
[:clear][:name=histname1][:<handler>.<action>] [if <filter>]
```

When a matching event is hit, an entry is added to a hash table using the key(s) and value(s) named. Keys and values correspond to fields in the event's format description. Values must correspond to numeric fields - on an event hit, the value(s) will be added to a sum kept for that field. The special string 'hitcount' can be used in place of an explicit value field - this is simply a count of event hits. If 'values' isn't specified, an implicit 'hitcount' value will be automatically created and used as the only value. Keys can be any field, or the special string 'stacktrace', which will use the event's kernel stacktrace as the key. The keywords 'keys' or 'key' can be used to specify keys, and the keywords 'values', 'vals', or 'val' can be used to specify values. Compound keys consisting of up to two fields can be specified by the 'keys' keyword. Hashing a compound key produces a unique entry in the table for each unique combination of component keys, and can be useful for providing more fine-grained summaries of event data. Additionally, sort keys consisting of up to two fields can be specified by the 'sort' keyword. If more than one field is specified, the result will be a 'sort within a sort': the first key is taken to be the primary sort key and the second the secondary key. If a hist trigger is given a name using the 'name' parameter, its histogram data

will be shared with other triggers of the same name, and trigger hits will update this common data. Only triggers with 'compatible' fields can be combined in this way; triggers are 'compatible' if the fields named in the trigger share the same number and type of fields and those fields also have the same names. Note that any two events always share the compatible 'hitcount' and 'stacktrace' fields and can therefore be combined using those fields, however pointless that may be.

'hist' triggers add a 'hist' file to each event's subdirectory. Reading the 'hist' file for the event will dump the hash table in its entirety to stdout. If there are multiple hist triggers attached to an event, there will be a table for each trigger in the output. The table displayed for a named trigger will be the same as any other instance having the same name. Each printed hash table entry is a simple list of the keys and values comprising the entry; keys are printed first and are delineated by curly braces, and are followed by the set of value fields for the entry. By default, numeric fields are displayed as base-10 integers. This can be modified by appending any of the following modifiers to the field name:

.hex	display a number as a hex value
.sym	display an address as a symbol
.sym-offset	display an address as a symbol and offset
.syscall	display a syscall id as a system call name
.execname	display a common_pid as a program name
.log2	display log2 value rather than raw number
.buckets=size	display grouping of values rather than raw number
.usecs	display a common_timestamp in microseconds

Note that in general the semantics of a given field aren't interpreted when applying a modifier to it, but there are some restrictions to be aware of in this regard:

- only the 'hex' modifier can be used for values (because values are essentially sums, and the other modifiers don't make sense in that context).
- the 'execname' modifier can only be used on a 'common_pid'. The reason for this is that the execname is simply the 'comm' value saved for the 'current' process when an event was triggered, which is the same as the common_pid value saved by the event tracing code. Trying to apply that comm value to other pid values wouldn't be correct, and typically events that care save pid-specific comm fields in the event itself.

A typical usage scenario would be the following to enable a hist trigger, read its current contents, and then turn it off:

```
# echo 'hist:keys=skbaddr.hex:vals=len' > \
  /sys/kernel/debug/tracing/events/net/netif_rx/trigger

# cat /sys/kernel/debug/tracing/events/net/netif_rx/hist

# echo '!hist:keys=skbaddr.hex:vals=len' > \
  /sys/kernel/debug/tracing/events/net/netif_rx/trigger
```

The trigger file itself can be read to show the details of the currently attached hist trigger. This information is also displayed at the top of the 'hist' file when read.

By default, the size of the hash table is 2048 entries. The 'size' parameter can be used to specify more or fewer than that. The units are in terms of hashtable entries - if a run uses more entries than specified, the results will show the number of 'drops', the number of hits that were ignored. The size should be a power of 2 between 128 and 131072 (any non- power-of-2 number specified will be rounded up).

The 'sort' parameter can be used to specify a value field to sort on. The default if unspecified is 'hitcount' and the default sort order is 'ascending'. To sort in the opposite direction, append '.descending' to the sort key.

The 'pause' parameter can be used to pause an existing hist trigger or to start a hist trigger but not log any events until told to do so. 'continue' or 'cont' can be used to start or restart a paused hist trigger.

The 'clear' parameter will clear the contents of a running hist trigger and leave its current paused/active state.

Note that the 'pause', 'cont', and 'clear' parameters should be applied using 'append' shell operator ('>>') if applied to an existing trigger, rather than via the '>' operator, which will cause the trigger to be removed through truncation.

- enable_hist/disable_hist

The enable_hist and disable_hist triggers can be used to have one event conditionally start and stop another event's already-attached hist trigger. Any number of enable_hist and disable_hist triggers can be attached to a given event, allowing that event to kick off and stop aggregations on a host of other events.

The format is very similar to the enable/disable_event triggers:

```
enable_hist:<system>:<event>[:count]
disable_hist:<system>:<event>[:count]
```

Instead of enabling or disabling the tracing of the target event into the trace buffer as the enable/disable_event triggers do, the enable/disable_hist triggers enable or disable the aggregation of the target event into a hash table.

A typical usage scenario for the enable_hist/disable_hist triggers would be to first set up a paused hist trigger on some event, followed by an enable_hist/disable_hist pair that turns the hist aggregation on and off when conditions of interest are hit:

```
# echo 'hist:keys=skbaddr.hex:vals=len:pause' > \
  /sys/kernel/debug/tracing/events/net/netif_receive_skb/trigger

# echo 'enable_hist:net:netif_receive_skb if filename==/usr/bin/wget' > \
  /sys/kernel/debug/tracing/events/sched/sched_process_exec/trigger

# echo 'disable_hist:net:netif_receive_skb if comm==wget' > \
  /sys/kernel/debug/tracing/events/sched/sched_process_exit/trigger
```

The above sets up an initially paused hist trigger which is unpaused and starts aggregating events when a given program is executed, and which stops aggregating when the process exits and the hist trigger is paused again.

The examples below provide a more concrete illustration of the concepts and typical usage patterns discussed above.

16.2.1 ‘special’ event fields

There are a number of ‘special event fields’ available for use as keys or values in a hist trigger. These look like and behave as if they were actual event fields, but aren’t really part of the event’s field definition or format file. They are however available for any event, and can be used anywhere an actual event field could be. They are:

common_timestamp	u64	timestamp (from ring buffer) associated with the event, in nanoseconds. May be modified by .usecs to have timestamps interpreted as microseconds.
common_cpu	int	the cpu on which the event occurred.

16.2.2 Extended error information

For some error conditions encountered when invoking a hist trigger command, extended error information is available via the tracing/error_log file. See Error Conditions in Documentation/trace/ftrace.rst for details.

16.2.3 6.2 ‘hist’ trigger examples

The first set of examples creates aggregations using the kmalloc event. The fields that can be used for the hist trigger are listed in the kmalloc event’s format file:

```
# cat /sys/kernel/debug/tracing/events/kmem/kmalloc/format
name: kmalloc
ID: 374
format:
    field:unsigned short common_type;          offset:0;          size:2;
    ↪ signed:0;
    field:unsigned char common_flags;          offset:2;          size:1;
    ↪ signed:0;
    field:unsigned char common_preempt_count;    offset:3;
    ↪ size:1; signed:0;
    field:int common_pid;                      offset:4;
    ↪ size:4; signed:1;

    field:unsigned long call_site;              offset:8;
    ↪ size:8; signed:0;
    field:const void * ptr;                    offset:16;
    ↪ size:8; signed:0;
    field:size_t bytes_req;                     offset:24;
    ↪ size:8; signed:0;
    field:size_t bytes_alloc;                   offset:32;
    ↪ size:8; signed:0;
    field:gfp_t gfp_flags;                      offset:40;
    ↪ size:4; signed:0;
```

We'll start by creating a hist trigger that generates a simple table that lists the total number of bytes requested for each function in the kernel that made one or more calls to `kmalloc`:

```
# echo 'hist:key=call_site:val=bytes_req.buckets=32' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger
```

This tells the tracing system to create a 'hist' trigger using the `call_site` field of the `kmalloc` event as the key for the table, which just means that each unique `call_site` address will have an entry created for it in the table. The '`val=bytes_req`' parameter tells the hist trigger that for each unique entry (`call_site`) in the table, it should keep a running total of the number of bytes requested by that `call_site`.

We'll let it run for awhile and then dump the contents of the 'hist' file in the `kmalloc` event's subdirectory (for readability, a number of entries have been omitted):

```
# cat /sys/kernel/debug/tracing/events/kmem/kmalloc/hist
# trigger info: hist:keys=call_site:vals=bytes_
↪ req:sort=hitcount:size=2048 [active]

{ call_site: 18446744072106379007 } hitcount:      1 bytes_req: ↪
↪ 176
{ call_site: 18446744071579557049 } hitcount:      1 bytes_req: ↪
↪ 1024
{ call_site: 18446744071580608289 } hitcount:      1 bytes_req: ↪
↪ 16384
{ call_site: 18446744071581827654 } hitcount:      1 bytes_req: ↪
↪ 24
{ call_site: 18446744071580700980 } hitcount:      1 bytes_req: ↪
↪ 8
{ call_site: 18446744071579359876 } hitcount:      1 bytes_req: ↪
↪ 152
{ call_site: 18446744071580795365 } hitcount:      3 bytes_req: ↪
↪ 144
{ call_site: 18446744071581303129 } hitcount:      3 bytes_req: ↪
↪ 144
{ call_site: 18446744071580713234 } hitcount:      4 bytes_req: ↪
↪ 2560
{ call_site: 18446744071580933750 } hitcount:      4 bytes_req: ↪
↪ 736
.
.
.
{ call_site: 18446744072106047046 } hitcount:     69 bytes_req: ↪
↪ 5576
{ call_site: 18446744071582116407 } hitcount:     73 bytes_req: ↪
↪ 2336
{ call_site: 18446744072106054684 } hitcount:    136 bytes_req: ↪
↪ 140504
{ call_site: 18446744072106224230 } hitcount:    136 bytes_req: ↪
↪ 19584
{ call_site: 18446744072106078074 } hitcount:    153 bytes_req: ↪
↪ 2448
```

```
{ call_site: 18446744072106062406 } hitcount:      153  bytes_req:  153
↪ 36720
{ call_site: 18446744071582507929 } hitcount:      153  bytes_req:  153
↪ 37088
{ call_site: 18446744072102520590 } hitcount:      273  bytes_req:  273
↪ 10920
{ call_site: 18446744071582143559 } hitcount:      358  bytes_req:  358
↪ 716
{ call_site: 18446744072106465852 } hitcount:      417  bytes_req:  417
↪ 56712
{ call_site: 18446744072102523378 } hitcount:      485  bytes_req:  485
↪ 27160
{ call_site: 18446744072099568646 } hitcount:     1676  bytes_req:  1676
↪ 33520

Totals:
  Hits: 4610
  Entries: 45
  Dropped: 0
```

The output displays a line for each entry, beginning with the key specified in the trigger, followed by the value(s) also specified in the trigger. At the beginning of the output is a line that displays the trigger info, which can also be displayed by reading the ‘trigger’ file:

```
# cat /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger
hist:keys=call_site:vals=bytes_req:sort=hitcount:size=2048 [active]
```

At the end of the output are a few lines that display the overall totals for the run. The ‘Hits’ field shows the total number of times the event trigger was hit, the ‘Entries’ field shows the total number of used entries in the hash table, and the ‘Dropped’ field shows the number of hits that were dropped because the number of used entries for the run exceeded the maximum number of entries allowed for the table (normally 0, but if not a hint that you may want to increase the size of the table using the ‘size’ parameter).

Notice in the above output that there’s an extra field, ‘hitcount’, which wasn’t specified in the trigger. Also notice that in the trigger info output, there’s a parameter, ‘sort=hitcount’, which wasn’t specified in the trigger either. The reason for that is that every trigger implicitly keeps a count of the total number of hits attributed to a given entry, called the ‘hitcount’. That hitcount information is explicitly displayed in the output, and in the absence of a user-specified sort parameter, is used as the default sort field.

The value ‘hitcount’ can be used in place of an explicit value in the ‘values’ parameter if you don’t really need to have any particular field summed and are mainly interested in hit frequencies.

To turn the hist trigger off, simply call up the trigger in the command history and re-execute it with a ‘!’ prepended:

```
# echo '!hist:key=call_site:val=bytes_req' > \
```



```
/sys/kernel/debug/tracing/events/kmem/kmalloc/trigger
```

Finally, notice that the `call_site` as displayed in the output above isn't really very useful. It's an address, but normally addresses are displayed in hex. To have a numeric field displayed as a hex value, simply append `'.hex'` to the field name in the trigger:

```
# echo 'hist:key=call_site.hex:val=bytes_req' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger

# cat /sys/kernel/debug/tracing/events/kmem/kmalloc/hist
# trigger info: hist:keys=call_site.hex:vals=bytes_
    req:sort=hitcount:size=2048 [active]

{ call_site: ffffffff8026b291 } hitcount:      1  bytes_req:      4
    ↪ 433
{ call_site: ffffffff807186ff } hitcount:      1  bytes_req:      4
    ↪ 176
{ call_site: ffffffff811ae721 } hitcount:      1  bytes_req:      4
    ↪ 16384
{ call_site: ffffffff811c5134 } hitcount:      1  bytes_req:      4
    ↪ 8
{ call_site: ffffffff804a9ebb } hitcount:      1  bytes_req:      4
    ↪ 511
{ call_site: ffffffff8122e0a6 } hitcount:      1  bytes_req:      4
    ↪ 12
{ call_site: ffffffff8107da84 } hitcount:      1  bytes_req:      4
    ↪ 152
{ call_site: ffffffff812d8246 } hitcount:      1  bytes_req:      4
    ↪ 24
{ call_site: ffffffff811dc1e5 } hitcount:      3  bytes_req:      4
    ↪ 144
{ call_site: ffffffff802515e8 } hitcount:      3  bytes_req:      4
    ↪ 648
{ call_site: ffffffff81258159 } hitcount:      3  bytes_req:      4
    ↪ 144
{ call_site: ffffffff811c80f4 } hitcount:      4  bytes_req:      4
    ↪ 544
.
.
.
{ call_site: ffffffff806c7646 } hitcount:     106  bytes_req:      4
    ↪ 8024
{ call_site: ffffffff806cb246 } hitcount:     132  bytes_req:      4
    ↪ 31680
{ call_site: ffffffff806cef7a } hitcount:     132  bytes_req:      4
    ↪ 2112
{ call_site: ffffffff8137e399 } hitcount:     132  bytes_req:      4
    ↪ 23232
{ call_site: ffffffff806c941c } hitcount:     185  bytes_req:      4
    ↪ 171360
{ call_site: ffffffff806f2a66 } hitcount:     185  bytes_req:      4
    ↪ 26640
```

```

{ call_site: ffffffff8036a70e } hitcount:      265  bytes_req:      10600
↳
{ call_site: ffffffff81325447 } hitcount:      292  bytes_req:      584
↳
{ call_site: ffffffff8072da3c } hitcount:      446  bytes_req:      60656
↳
{ call_site: ffffffff8036b1f2 } hitcount:      526  bytes_req:      29456
↳
{ call_site: ffffffff80099c06 } hitcount:     1780  bytes_req:     35600
↳

Totals:
  Hits: 4775
  Entries: 46
  Dropped: 0

```

Even that's only marginally more useful - while hex values do look more like addresses, what users are typically more interested in when looking at text addresses are the corresponding symbols instead. To have an address displayed as symbolic value instead, simply append `.sym` or `.sym-offset` to the field name in the trigger:

```

# echo 'hist:key=call_site.sym:val=bytes_req' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger

# cat /sys/kernel/debug/tracing/events/kmem/kmalloc/hist
# trigger info: hist:keys=call_site.sym:vals=bytes_
↳ req:sort=hitcount:size=2048 [active]

{ call_site: [fffffff810adcb9] syslog_print_all
↳      } hitcount:      1  bytes_req:      1024
{ call_site: [fffffff8154bc62] usb_control_msg
↳      } hitcount:      1  bytes_req:      8
{ call_site: [fffffff800bf6fe] hidraw_send_report [hid]
↳      } hitcount:      1  bytes_req:      7
{ call_site: [fffffff8154acbe] usb_alloc_urb
↳      } hitcount:      1  bytes_req:      192
{ call_site: [fffffff800bf1ca] hidraw_report_event [hid]
↳      } hitcount:      1  bytes_req:      7
{ call_site: [fffffff811e3a25] __seq_open_private
↳      } hitcount:      1  bytes_req:      40
{ call_site: [fffffff8109524a] alloc_fair_sched_group
↳      } hitcount:      2  bytes_req:      128
{ call_site: [fffffff811febd5] fsnotify_alloc_group
↳      } hitcount:      2  bytes_req:      528
{ call_site: [fffffff81440f58] __tty_buffer_request_room
↳      } hitcount:      2  bytes_req:      2624
{ call_site: [fffffff81200ba6] inotify_new_group
↳      } hitcount:      2  bytes_req:      96
{ call_site: [fffffff805e19af] ieee80211_start_tx_ba_session
↳ [mac80211]      } hitcount:      2  bytes_req:      464
{ call_site: [fffffff81672406] tcp_get_metrics
↳      } hitcount:      2  bytes_req:      304

```

```

{ call_site: [ffffffff81097ec2] alloc_rt_sched_group          ↵
↵      } hitcount:          2  bytes_req:          128
{ call_site: [ffffffff81089b05] sched_create_group            ↵
↵      } hitcount:          2  bytes_req:         1424
.
.
.
{ call_site: [ffffffffffa04a580c] intel_crtc_page_flip [i915]  ↵
↵      } hitcount:        1185  bytes_req:       123240
{ call_site: [ffffffffffa0287592] drm_mode_page_flip_ioctl [drm] ↵
↵      } hitcount:        1185  bytes_req:       104280
{ call_site: [ffffffffffa04c4a3c] intel_plane_duplicate_state [i915] ↵
↵      } hitcount:        1402  bytes_req:       190672
{ call_site: [ffffffff812891ca] ext4_find_extent              ↵
↵      } hitcount:        1518  bytes_req:       146208
{ call_site: [ffffffffffa029070e] drm_vma_node_allow [drm]    ↵
↵      } hitcount:        1746  bytes_req:        69840
{ call_site: [ffffffffffa045e7c4] i915_gem_do_execbuffer.isra.23 [i915] ↵
↵      } hitcount:        2021  bytes_req:       792312
{ call_site: [ffffffffffa02911f2] drm_modeset_lock_crtc [drm]  ↵
↵      } hitcount:        2592  bytes_req:       145152
{ call_site: [ffffffffffa0489a66] intel_ring_begin [i915]     ↵
↵      } hitcount:        2629  bytes_req:       378576
{ call_site: [ffffffffffa046041c] i915_gem_execbuffer2 [i915]  ↵
↵      } hitcount:        2629  bytes_req:       3783248
{ call_site: [ffffffff81325607] apparmor_file_alloc_security  ↵
↵      } hitcount:        5192  bytes_req:       10384
{ call_site: [ffffffffffa00b7c06] hid_report_raw_event [hid]   ↵
↵      } hitcount:        5529  bytes_req:       110584
{ call_site: [ffffffff8131ebf7] aa_alloc_task_context         ↵
↵      } hitcount:       21943  bytes_req:       702176
{ call_site: [ffffffff8125847d] ext4_htree_store_dirent       ↵
↵      } hitcount:       55759  bytes_req:      5074265

Totals:
  Hits: 109928
  Entries: 71
  Dropped: 0

```

Because the default sort key above is 'hitcount', the above shows a the list of call_sites by increasing hitcount, so that at the bottom we see the functions that made the most kmalloc calls during the run. If instead we we wanted to see the top kmalloc callers in terms of the number of bytes requested rather than the number of calls, and we wanted the top caller to appear at the top, we can use the 'sort' parameter, along with the 'descending' modifier:

```

# echo 'hist:key=call_site.sym:val=bytes_req:sort=bytes_req.descending
↵ ' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger

# cat /sys/kernel/debug/tracing/events/kmem/kmalloc/hist

```

```
# trigger info: hist:keys=call_site.sym:vals=bytes_req:sort=bytes_req.
↳descending:size=2048 [active]

{ call_site: [fffffffa046041c] i915_gem_execbuffer2 [i915]
↳      } hitcount:      2186 bytes_req:      3397464
{ call_site: [fffffffa045e7c4] i915_gem_do_execbuffer.isra.23 [i915]
↳      } hitcount:      1790 bytes_req:      712176
{ call_site: [ffffff8125847d] ext4_htree_store_dirent
↳      } hitcount:      8132 bytes_req:      513135
{ call_site: [ffffff811e2a1b] seq_buf_alloc
↳      } hitcount:       106 bytes_req:      440128
{ call_site: [fffffffa0489a66] intel_ring_begin [i915]
↳      } hitcount:      2186 bytes_req:      314784
{ call_site: [ffffff812891ca] ext4_find_extent
↳      } hitcount:      2174 bytes_req:      208992
{ call_site: [ffffff811ae8e1] __kmallocc
↳      } hitcount:        8 bytes_req:      131072
{ call_site: [fffffffa04c4a3c] intel_plane_duplicate_state [i915]
↳      } hitcount:       859 bytes_req:      116824
{ call_site: [fffffffa02911f2] drm_modeset_lock_crtc [drm]
↳      } hitcount:      1834 bytes_req:      102704
{ call_site: [fffffffa04a580c] intel_crtc_page_flip [i915]
↳      } hitcount:       972 bytes_req:      101088
{ call_site: [fffffffa0287592] drm_mode_page_flip_ioctl [drm]
↳      } hitcount:       972 bytes_req:      85536
{ call_site: [fffffffa00b7c06] hid_report_raw_event [hid]
↳      } hitcount:      3333 bytes_req:      66664
{ call_site: [ffffff8137e559] sg_kmallocc
↳      } hitcount:       209 bytes_req:      61632
.
.
.
{ call_site: [ffffff81095225] alloc_fair_sched_group
↳      } hitcount:        2 bytes_req:       128
{ call_site: [ffffff81097ec2] alloc_rt_sched_group
↳      } hitcount:        2 bytes_req:       128
{ call_site: [ffffff812d8406] copy_semundo
↳      } hitcount:        2 bytes_req:        48
{ call_site: [ffffff81200ba6] inotify_new_group
↳      } hitcount:        1 bytes_req:        48
{ call_site: [fffffffa027121a] drm_getmagic [drm]
↳      } hitcount:        1 bytes_req:        48
{ call_site: [ffffff811e3a25] __seq_open_private
↳      } hitcount:        1 bytes_req:       40
{ call_site: [ffffff811c52f4] bprm_change_interp
↳      } hitcount:        2 bytes_req:       16
{ call_site: [ffffff8154bc62] usb_control_msg
↳      } hitcount:        1 bytes_req:        8
{ call_site: [fffffffa00bf1ca] hidraw_report_event [hid]
↳      } hitcount:        1 bytes_req:        7
```

```

{ call_site: [fffffffa00bf6fe] hidraw_send_report [hid]
  ↪      } hitcount:          1  bytes_req:          7

Totals:
  Hits: 32133
  Entries: 81
  Dropped: 0

```

To display the offset and size information in addition to the symbol name, just use 'sym-offset' instead:

```

# echo 'hist:key=call_site.sym-offset:val=bytes_req:sort=bytes_req.
  ↪descending' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger

# cat /sys/kernel/debug/tracing/events/kmem/kmalloc/hist
# trigger info: hist:keys=call_site.sym-offset:vals=bytes_
  ↪req:sort=bytes_req.descending:size=2048 [active]

{ call_site: [fffffffa046041c] i915_gem_execbuffer2+0x6c/0x2c0 [i915]
  ↪      } hitcount:          4569  bytes_req:      3163720
{ call_site: [fffffffa0489a66] intel_ring_begin+0xc6/0x1f0 [i915]
  ↪      } hitcount:          4569  bytes_req:       657936
{ call_site: [fffffffa045e7c4] i915_gem_do_execbuffer.isra.23+0x694/
  ↪0x1020 [i915]      } hitcount:          1519  bytes_req:       472936
{ call_site: [fffffffa045e646] i915_gem_do_execbuffer.isra.23+0x516/
  ↪0x1020 [i915]      } hitcount:          3050  bytes_req:       211832
{ call_site: [ffffff811e2a1b] seq_buf_alloc+0x1b/0x50
  ↪      } hitcount:           34  bytes_req:       148384
{ call_site: [fffffffa04a580c] intel_crtc_page_flip+0xbc/0x870 [i915]
  ↪      } hitcount:          1385  bytes_req:       144040
{ call_site: [ffffff811ae8e1] __kmalloc+0x191/0x1b0
  ↪      } hitcount:           8  bytes_req:       131072
{ call_site: [fffffffa0287592] drm_mode_page_flip_ioctl+0x282/0x360
  ↪[drm]      } hitcount:          1385  bytes_req:       121880
{ call_site: [fffffffa02911f2] drm_modeset_lock_crtc+0x32/0x100 [drm]
  ↪      } hitcount:          1848  bytes_req:       103488
{ call_site: [fffffffa04c4a3c] intel_plane_duplicate_state+0x2c/0xa0
  ↪[i915]      } hitcount:          461  bytes_req:        62696
{ call_site: [fffffffa029070e] drm_vma_node_allow+0x2e/0xd0 [drm]
  ↪      } hitcount:          1541  bytes_req:        61640
{ call_site: [ffffff815f8d7b] sk_prot_alloc+0xcb/0x1b0
  ↪      } hitcount:           57  bytes_req:       57456
.
.
.
{ call_site: [ffffff8109524a] alloc_fair_sched_group+0x5a/0x1a0
  ↪      } hitcount:           2  bytes_req:         128
{ call_site: [fffffffa027b921] drm_vm_open_locked+0x31/0xa0 [drm]
  ↪      } hitcount:           3  bytes_req:          96
{ call_site: [ffffff8122e266] proc_self_follow_link+0x76/0xb0
  ↪      } hitcount:           8  bytes_req:          96

```

```

{ call_site: [ffffffff81213e80] load_elf_binary+0x240/0x1650
  ↳      } hitcount:          3  bytes_req:          84
{ call_site: [ffffffff8154bc62] usb_control_msg+0x42/0x110
  ↳      } hitcount:          1  bytes_req:           8
{ call_site: [ffffffffffa00bf6fe] hidraw_send_report+0x7e/0x1a0 [hid]
  ↳      } hitcount:          1  bytes_req:           7
{ call_site: [ffffffffffa00bf1ca] hidraw_report_event+0x8a/0x120 [hid]
  ↳      } hitcount:          1  bytes_req:           7

Totals:
  Hits: 26098
  Entries: 64
  Dropped: 0

```

We can also add multiple fields to the ‘values’ parameter. For example, we might want to see the total number of bytes allocated alongside bytes requested, and display the result sorted by bytes allocated in a descending order:

```

# echo 'hist:keys=call_site.sym:values=bytes_req,bytes_
  ↳alloc:sort=bytes_alloc.descending' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger

# cat /sys/kernel/debug/tracing/events/kmem/kmalloc/hist
# trigger info: hist:keys=call_site.sym:vals=bytes_req,bytes_
  ↳alloc:sort=bytes_alloc.descending:size=2048 [active]

{ call_site: [ffffffffffa046041c] i915_gem_execbuffer2 [i915]
  ↳      } hitcount:          7403  bytes_req:    4084360  bytes_alloc:
  ↳ 5958016
{ call_site: [ffffffff811e2a1b] seq_buf_alloc
  ↳      } hitcount:           541  bytes_req:    2213968  bytes_alloc:
  ↳ 2228224
{ call_site: [ffffffffffa0489a66] intel_ring_begin [i915]
  ↳      } hitcount:          7404  bytes_req:    1066176  bytes_alloc:
  ↳ 1421568
{ call_site: [ffffffffffa045e7c4] i915_gem_do_execbuffer.isra.23 [i915]
  ↳      } hitcount:          1565  bytes_req:     557368  bytes_alloc:
  ↳ 1037760
{ call_site: [ffffffff8125847d] ext4_htree_store_dirent
  ↳      } hitcount:          9557  bytes_req:     595778  bytes_alloc:
  ↳ 695744
{ call_site: [ffffffffffa045e646] i915_gem_do_execbuffer.isra.23 [i915]
  ↳      } hitcount:          5839  bytes_req:     430680  bytes_alloc:
  ↳ 470400
{ call_site: [ffffffffffa04c4a3c] intel_plane_duplicate_state [i915]
  ↳      } hitcount:          2388  bytes_req:     324768  bytes_alloc:
  ↳ 458496
{ call_site: [ffffffffffa02911f2] drm_modeset_lock_crtc [drm]
  ↳      } hitcount:          3911  bytes_req:     219016  bytes_alloc:
  ↳ 250304
{ call_site: [ffffffff815f8d7b] sk_prot_alloc
  ↳      } hitcount:           235  bytes_req:     236880  bytes_alloc:
  ↳ 240640

```

```

{ call_site: [ffffffff8137e559] sg_kmalloc
  ↳      } hitcount:      557 bytes_req:      169024 bytes_alloc:
  ↳      221760
{ call_site: [fffffffffa00b7c06] hid_report_raw_event [hid]
  ↳      } hitcount:      9378 bytes_req:      187548 bytes_alloc:
  ↳      206312
{ call_site: [fffffffffa04a580c] intel_crtc_page_flip [i915]
  ↳      } hitcount:      1519 bytes_req:      157976 bytes_alloc:
  ↳      194432
.
.
.
{ call_site: [ffffffff8109bd3b] sched_autogroup_create_attach
  ↳      } hitcount:        2 bytes_req:        144 bytes_alloc:
  ↳      192
{ call_site: [ffffffff81097ee8] alloc_rt_sched_group
  ↳      } hitcount:        2 bytes_req:        128 bytes_alloc:
  ↳      128
{ call_site: [ffffffff8109524a] alloc_fair_sched_group
  ↳      } hitcount:        2 bytes_req:        128 bytes_alloc:
  ↳      128
{ call_site: [ffffffff81095225] alloc_fair_sched_group
  ↳      } hitcount:        2 bytes_req:        128 bytes_alloc:
  ↳      128
{ call_site: [ffffffff81097ec2] alloc_rt_sched_group
  ↳      } hitcount:        2 bytes_req:        128 bytes_alloc:
  ↳      128
{ call_site: [ffffffff81213e80] load_elf_binary
  ↳      } hitcount:        3 bytes_req:         84 bytes_alloc:
  ↳      96
{ call_site: [ffffffff81079a2e] kthread_create_on_node
  ↳      } hitcount:        1 bytes_req:         56 bytes_alloc:
  ↳      64
{ call_site: [fffffffffa00bf6fe] hidraw_send_report [hid]
  ↳      } hitcount:        1 bytes_req:         7 bytes_alloc:
  ↳      8
{ call_site: [ffffffff8154bc62] usb_control_msg
  ↳      } hitcount:        1 bytes_req:         8 bytes_alloc:
  ↳      8
{ call_site: [fffffffffa00bf1ca] hidraw_report_event [hid]
  ↳      } hitcount:        1 bytes_req:         7 bytes_alloc:
  ↳      8

Totals:
  Hits: 66598
  Entries: 65
  Dropped: 0

```

Finally, to finish off our kmalloc example, instead of simply having the hist trigger display symbolic call_sites, we can have the hist trigger additionally display the complete set of kernel stack traces that led to each call_site. To do that, we simply use

the special value 'stacktrace' for the key parameter:

```
# echo 'hist:keys=stacktrace:values=bytes_req,bytes_alloc:sort=bytes_
↪alloc' > \
    /sys/kernel/debug/tracing/events/kmem/kmalloc/trigger
```

The above trigger will use the kernel stack trace in effect when an event is triggered as the key for the hash table. This allows the enumeration of every kernel callpath that led up to a particular event, along with a running total of any of the event fields for that event. Here we tally bytes requested and bytes allocated for every callpath in the system that led up to a kmalloc (in this case every callpath to a kmalloc for a kernel compile):

```
# cat /sys/kernel/debug/tracing/events/kmem/kmalloc/hist
# trigger info: hist:keys=stacktrace:vals=bytes_req,bytes_
↪alloc:sort=bytes_alloc:size=2048 [active]

{ stacktrace:
    __kmalloc_track_caller+0x10b/0x1a0
    kmemdup+0x20/0x50
    hidraw_report_event+0x8a/0x120 [hid]
    hid_report_raw_event+0x3ea/0x440 [hid]
    hid_input_report+0x112/0x190 [hid]
    hid_irq_in+0xc2/0x260 [usbhid]
    __usb_hcd_giveback_urb+0x72/0x120
    usb_giveback_urb_bh+0x9e/0xe0
    tasklet_hi_action+0xf8/0x100
    __do_softirq+0x114/0x2c0
    irq_exit+0xa5/0xb0
    do_IRQ+0x5a/0xf0
    ret_from_intr+0x0/0x30
    cpuidle_enter+0x17/0x20
    cpu_startup_entry+0x315/0x3e0
    rest_init+0x7c/0x80
} hitcount:          3 bytes_req:          21 bytes_alloc:          24
{ stacktrace:
    __kmalloc_track_caller+0x10b/0x1a0
    kmemdup+0x20/0x50
    hidraw_report_event+0x8a/0x120 [hid]
    hid_report_raw_event+0x3ea/0x440 [hid]
    hid_input_report+0x112/0x190 [hid]
    hid_irq_in+0xc2/0x260 [usbhid]
    __usb_hcd_giveback_urb+0x72/0x120
    usb_giveback_urb_bh+0x9e/0xe0
    tasklet_hi_action+0xf8/0x100
    __do_softirq+0x114/0x2c0
    irq_exit+0xa5/0xb0
    do_IRQ+0x5a/0xf0
    ret_from_intr+0x0/0x30
} hitcount:          3 bytes_req:          21 bytes_alloc:          24
{ stacktrace:
    kmem_cache_alloc_trace+0xeb/0x150
```



```

aa_alloc_task_context+0x27/0x40
apparmor_cred_prepare+0x1f/0x50
security_prepare_creds+0x16/0x20
prepare_creds+0xdf/0x1a0
SyS_capset+0xb5/0x200
system_call_fastpath+0x12/0x6a
} hitcount:          1 bytes_req:          32 bytes_alloc:          32
.
.
.
{ stacktrace:
    __kmalloc+0x11b/0x1b0
    i915_gem_execbuffer2+0x6c/0x2c0 [i915]
    drm_ioctl+0x349/0x670 [drm]
    do_vfs_ioctl+0x2f0/0x4f0
    SyS_ioctl+0x81/0xa0
    system_call_fastpath+0x12/0x6a
} hitcount:        17726 bytes_req:   13944120 bytes_alloc:   19593808
{ stacktrace:
    __kmalloc+0x11b/0x1b0
    load_elf_phdrs+0x76/0xa0
    load_elf_binary+0x102/0x1650
    search_binary_handler+0x97/0x1d0
    do_execveat_common.isra.34+0x551/0x6e0
    SyS_execve+0x3a/0x50
    return_from_execve+0x0/0x23
} hitcount:        33348 bytes_req:   17152128 bytes_alloc:   20226048
{ stacktrace:
    kmem_cache_alloc_trace+0xeb/0x150
    apparmor_file_alloc_security+0x27/0x40
    security_file_alloc+0x16/0x20
    get_empty_filp+0x93/0x1c0
    path_openat+0x31/0x5f0
    do_filp_open+0x3a/0x90
    do_sys_open+0x128/0x220
    SyS_open+0x1e/0x20
    system_call_fastpath+0x12/0x6a
} hitcount:       4766422 bytes_req:    9532844 bytes_alloc:   38131376
{ stacktrace:
    __kmalloc+0x11b/0x1b0
    seq_buf_alloc+0x1b/0x50
    seq_read+0x2cc/0x370
    proc_reg_read+0x3d/0x80
    __vfs_read+0x28/0xe0
    vfs_read+0x86/0x140
    SyS_read+0x46/0xb0
    system_call_fastpath+0x12/0x6a
} hitcount:        19133 bytes_req:   78368768 bytes_alloc:   78368768

Totals:

```

```
Hits: 6085872
Entries: 253
Dropped: 0
```

If you key a hist trigger on `common_pid`, in order for example to gather and display sorted totals for each process, you can use the special `.execname` modifier to display the executable names for the processes in the table rather than raw pids. The example below keeps a per-process sum of total bytes read:

```
# echo 'hist:key=common_pid.execname:val=count:sort=count.descending' >
↪ \
    /sys/kernel/debug/tracing/events/syscalls/sys_enter_read/trigger

# cat /sys/kernel/debug/tracing/events/syscalls/sys_enter_read/hist
# trigger info: hist:keys=common_pid.execname:vals=count:sort=count.
↪descending:size=2048 [active]

{ common_pid: gnome-terminal [    3196] } hitcount:      280 ↪
↪count:      1093512
{ common_pid: Xorg           [    1309] } hitcount:      525 ↪
↪count:      256640
{ common_pid: compiz         [    2889] } hitcount:       59 ↪
↪count:      254400
{ common_pid: bash           [    8710] } hitcount:        3 ↪
↪count:      66369
{ common_pid: dbus-daemon-lau [    8703] } hitcount:       49 ↪
↪count:      47739
{ common_pid: irqbalance     [    1252] } hitcount:       27 ↪
↪count:      27648
{ common_pid: 0lifupdown     [    8705] } hitcount:        3 ↪
↪count:      17216
{ common_pid: dbus-daemon    [     772] } hitcount:       10 ↪
↪count:      12396
{ common_pid: Socket Thread  [    8342] } hitcount:       11 ↪
↪count:      11264
{ common_pid: nm-dhcp-client. [    8701] } hitcount:        6 ↪
↪count:      7424
{ common_pid: gmain          [    1315] } hitcount:       18 ↪
↪count:      6336
.
.
.
{ common_pid: postgres       [    1892] } hitcount:        2 ↪
↪count:        32
{ common_pid: postgres       [    1891] } hitcount:        2 ↪
↪count:        32
{ common_pid: gmain          [    8704] } hitcount:        2 ↪
↪count:        32
{ common_pid: upstart-dbus-br [    2740] } hitcount:       21 ↪
↪count:        21
{ common_pid: nm-dispatcher.a [    8696] } hitcount:        1 ↪
↪count:        16
```

```

{ common_pid: indicator-datet [ 2904] } hitcount: 1
↪count: 16
{ common_pid: gdbus [ 2998] } hitcount: 1
↪count: 16
{ common_pid: rtkit-daemon [ 2052] } hitcount: 1
↪count: 8
{ common_pid: init [ 1] } hitcount: 2
↪count: 2

Totals:
  Hits: 2116
  Entries: 51
  Dropped: 0

```

Similarly, if you key a hist trigger on syscall id, for example to gather and display a list of systemwide syscall hits, you can use the special .syscall modifier to display the syscall names rather than raw ids. The example below keeps a running total of syscall counts for the system during the run:

```

# echo 'hist:key=id.syscall:val=hitcount' > \
    /sys/kernel/debug/tracing/events/raw_syscalls/sys_enter/trigger

# cat /sys/kernel/debug/tracing/events/raw_syscalls/sys_enter/hist
# trigger info: hist:keys=id.
↪syscall:vals=hitcount:sort=hitcount:size=2048 [active]

{ id: sys_fsync [ 74] } hitcount: 1
{ id: sys_newuname [ 63] } hitcount: 1
{ id: sys_prctl [157] } hitcount: 1
{ id: sys_statfs [137] } hitcount: 1
{ id: sys_symlink [ 88] } hitcount: 1
{ id: sys_sendmmsg [307] } hitcount: 1
{ id: sys_semctl [ 66] } hitcount: 1
{ id: sys_readlink [ 89] } hitcount: 3
{ id: sys_bind [ 49] } hitcount: 3
{ id: sys_getsockname [ 51] } hitcount: 3
{ id: sys_unlink [ 87] } hitcount: 3
{ id: sys_rename [ 82] } hitcount: 4
{ id: unknown_syscall [ 58] } hitcount: 4
{ id: sys_connect [ 42] } hitcount: 4
{ id: sys_getpid [ 39] } hitcount: 4
.
.
.
{ id: sys_rt_sigprocmask [ 14] } hitcount: 952
{ id: sys_futex [202] } hitcount: 1534
{ id: sys_write [ 1] } hitcount: 2689
{ id: sys_setitimer [ 38] } hitcount: 2797
{ id: sys_read [ 0] } hitcount: 3202
{ id: sys_select [ 23] } hitcount: 3773
{ id: sys_writev [ 20] } hitcount: 4531

```

```

{ id: sys_poll           [ 7] } hitcount:      8314
{ id: sys_recvmsg        [ 47] } hitcount:     13738
{ id: sys_ioctl          [ 16] } hitcount:     21843

```

Totals:

```

    Hits: 67612
    Entries: 72
    Dropped: 0

```

The syscall counts above provide a rough overall picture of system call activity on the system; we can see for example that the most popular system call on this system was the 'sys_ioctl' system call.

We can use 'compound' keys to refine that number and provide some further insight as to which processes exactly contribute to the overall ioctl count.

The command below keeps a hitcount for every unique combination of system call id and pid - the end result is essentially a table that keeps a per-pid sum of system call hits. The results are sorted using the system call id as the primary key, and the hitcount sum as the secondary key:

```

# echo 'hist:key=id.syscall,common_pid.execname:val=hitcount:sort=id,
↪hitcount' > \
    /sys/kernel/debug/tracing/events/raw_syscalls/sys_enter/trigger

```

```

# cat /sys/kernel/debug/tracing/events/raw_syscalls/sys_enter/hist
# trigger info: hist:keys=id.syscall,common_pid.
↪execname:vals=hitcount:sort=id.syscall,hitcount:size=2048 [active]

```

```

{ id: sys_read           [ 0], common_pid: rtkit-daemon  ↪
↪[ 1877] } hitcount: 1
{ id: sys_read           [ 0], common_pid: gdbus        ↪
↪[ 2976] } hitcount: 1
{ id: sys_read           [ 0], common_pid: console-kit-dae ↪
↪[ 3400] } hitcount: 1
{ id: sys_read           [ 0], common_pid: postgres      ↪
↪[ 1865] } hitcount: 1
{ id: sys_read           [ 0], common_pid: deja-dup-monito ↪
↪[ 3543] } hitcount: 2
{ id: sys_read           [ 0], common_pid: NetworkManager ↪
↪[ 890] } hitcount: 2
{ id: sys_read           [ 0], common_pid: evolution-calen ↪
↪[ 3048] } hitcount: 2
{ id: sys_read           [ 0], common_pid: postgres      ↪
↪[ 1864] } hitcount: 2
{ id: sys_read           [ 0], common_pid: nm-applet     ↪
↪[ 3022] } hitcount: 2
{ id: sys_read           [ 0], common_pid: whoopsie      ↪
↪[ 1212] } hitcount: 2
.
.
.

```

```

{ id: sys_ioctl          [ 16], common_pid: bash           1
↳[      8479] } hitcount:
{ id: sys_ioctl          [ 16], common_pid: bash           12
↳[      3472] } hitcount:
{ id: sys_ioctl          [ 16], common_pid: gnome-terminal 16
↳[      3199] } hitcount:
{ id: sys_ioctl          [ 16], common_pid: Xorg            1808
↳[      1267] } hitcount:
{ id: sys_ioctl          [ 16], common_pid: compiz          5580
↳[      2994] } hitcount:
.
.
.
{ id: sys_waitid         [247], common_pid: upstart-dbus-br 3
↳[      2690] } hitcount:
{ id: sys_waitid         [247], common_pid: upstart-dbus-br 16
↳[      2688] } hitcount:
{ id: sys_inotify_add_watch [254], common_pid: gmain        2
↳[      975] } hitcount:
{ id: sys_inotify_add_watch [254], common_pid: gmain        4
↳[     3204] } hitcount:
{ id: sys_inotify_add_watch [254], common_pid: gmain        4
↳[     2888] } hitcount:
{ id: sys_inotify_add_watch [254], common_pid: gmain        4
↳[     3003] } hitcount:
{ id: sys_inotify_add_watch [254], common_pid: gmain        4
↳[     2873] } hitcount:
{ id: sys_inotify_add_watch [254], common_pid: gmain        6
↳[     3196] } hitcount:
{ id: sys_openat         [257], common_pid: java            2
↳[     2623] } hitcount:
{ id: sys_eventfd2       [290], common_pid: ibus-ui-gtk3     4
↳[     2760] } hitcount:
{ id: sys_eventfd2       [290], common_pid: compiz          6
↳[     2994] } hitcount:

Totals:
  Hits: 31536
  Entries: 323
  Dropped: 0

```

The above list does give us a breakdown of the ioctl syscall by pid, but it also gives us quite a bit more than that, which we don't really care about at the moment. Since we know the syscall id for sys_ioctl (16, displayed next to the sys_ioctl name), we can use that to filter out all the other syscalls:

```

# echo 'hist:key=id.syscall,common_pid.execname:val=hitcount:sort=id,
↳hitcount if id == 16' > \
    /sys/kernel/debug/tracing/events/raw_syscalls/sys_enter/trigger

# cat /sys/kernel/debug/tracing/events/raw_syscalls/sys_enter/hist

```

```
# trigger info: hist:keys=id.syscall,common_pid.
↪execname:vals=hitcount:sort=id.syscall,hitcount:size=2048 if id ==
↪16 [active]

{ id: sys_ioctl          [ 16], common_pid: gmain          ↪
↪[      2769] } hitcount: 1
{ id: sys_ioctl          [ 16], common_pid: evolution-addre ↪
↪[      8571] } hitcount: 1
{ id: sys_ioctl          [ 16], common_pid: gmain          ↪
↪[      3003] } hitcount: 1
{ id: sys_ioctl          [ 16], common_pid: gmain          ↪
↪[      2781] } hitcount: 1
{ id: sys_ioctl          [ 16], common_pid: gmain          ↪
↪[      2829] } hitcount: 1
{ id: sys_ioctl          [ 16], common_pid: bash           ↪
↪[      8726] } hitcount: 1
{ id: sys_ioctl          [ 16], common_pid: bash           ↪
↪[      8508] } hitcount: 1
{ id: sys_ioctl          [ 16], common_pid: gmain          ↪
↪[      2970] } hitcount: 1
{ id: sys_ioctl          [ 16], common_pid: gmain          ↪
↪[      2768] } hitcount: 1
.
.
.
{ id: sys_ioctl          [ 16], common_pid: pool           ↪
↪[      8559] } hitcount: 45
{ id: sys_ioctl          [ 16], common_pid: pool           ↪
↪[      8555] } hitcount: 48
{ id: sys_ioctl          [ 16], common_pid: pool           ↪
↪[      8551] } hitcount: 48
{ id: sys_ioctl          [ 16], common_pid: avahi-daemon    ↪
↪[      896] } hitcount: 66
{ id: sys_ioctl          [ 16], common_pid: Xorg            ↪
↪[      1267] } hitcount: 26674
{ id: sys_ioctl          [ 16], common_pid: compiz          ↪
↪[      2994] } hitcount: 73443

Totals:
  Hits: 101162
  Entries: 103
  Dropped: 0
```

The above output shows that ‘compiz’ and ‘Xorg’ are far and away the heaviest ioctl callers (which might lead to questions about whether they really need to be making all those calls and to possible avenues for further investigation.)

The compound key examples used a key and a sum value (hitcount) to sort the output, but we can just as easily use two keys instead. Here’s an example where we use a compound key composed of the the common_pid and size event fields. Sorting with pid as the primary key and ‘size’ as the secondary key allows us to display an ordered

summary of the recvfrom sizes, with counts, received by each process:

```
# echo 'hist:key=common_pid.execname,size:val=hitcount:sort=common_pid,
↪size' > \
    /sys/kernel/debug/tracing/events/syscalls/sys_enter_recvfrom/
↪trigger

# cat /sys/kernel/debug/tracing/events/syscalls/sys_enter_recvfrom/hist
# trigger info: hist:keys=common_pid.execname,
↪size:vals=hitcount:sort=common_pid.execname,size:size=2048 [active]

{ common_pid: smbd          [      784], size:         4 }_
↪hitcount:          1
{ common_pid: dnsmasq       [     1412], size:       4096 }_
↪hitcount:         672
{ common_pid: postgres     [     1796], size:       1000 }_
↪hitcount:          6
{ common_pid: postgres     [     1867], size:       1000 }_
↪hitcount:         10
{ common_pid: bamfdaemon   [     2787], size:         28 }_
↪hitcount:          2
{ common_pid: bamfdaemon   [     2787], size:      14360 }_
↪hitcount:          1
{ common_pid: compiz       [     2994], size:         8 }_
↪hitcount:          1
{ common_pid: compiz       [     2994], size:        20 }_
↪hitcount:         11
{ common_pid: gnome-terminal [    3199], size:         4 }_
↪hitcount:          2
{ common_pid: firefox      [    8817], size:         4 }_
↪hitcount:          1
{ common_pid: firefox      [    8817], size:         8 }_
↪hitcount:          5
{ common_pid: firefox      [    8817], size:        588 }_
↪hitcount:          2
{ common_pid: firefox      [    8817], size:        628 }_
↪hitcount:          1
{ common_pid: firefox      [    8817], size:       6944 }_
↪hitcount:          1
{ common_pid: firefox      [    8817], size:    408880 }_
↪hitcount:          2
{ common_pid: firefox      [    8822], size:         8 }_
↪hitcount:          2
{ common_pid: firefox      [    8822], size:        160 }_
↪hitcount:          2
{ common_pid: firefox      [    8822], size:        320 }_
↪hitcount:          2
{ common_pid: firefox      [    8822], size:        352 }_
↪hitcount:          1
.
.
```

```

.
{ common_pid: pool          [      8923], size:      1960 }_
↪hitcount:                  10
{ common_pid: pool          [      8923], size:      2048 }_
↪hitcount:                  10
{ common_pid: pool          [      8924], size:      1960 }_
↪hitcount:                  10
{ common_pid: pool          [      8924], size:      2048 }_
↪hitcount:                  10
{ common_pid: pool          [      8928], size:      1964 }_
↪hitcount:                   4
{ common_pid: pool          [      8928], size:      1965 }_
↪hitcount:                   2
{ common_pid: pool          [      8928], size:      2048 }_
↪hitcount:                   6
{ common_pid: pool          [      8929], size:      1982 }_
↪hitcount:                   1
{ common_pid: pool          [      8929], size:      2048 }_
↪hitcount:                   1

Totals:
  Hits: 2016
  Entries: 224
  Dropped: 0

```

The above example also illustrates the fact that although a compound key is treated as a single entity for hashing purposes, the sub-keys it's composed of can be accessed independently.

The next example uses a string field as the hash key and demonstrates how you can manually pause and continue a hist trigger. In this example, we'll aggregate fork counts and don't expect a large number of entries in the hash table, so we'll drop it to a much smaller number, say 256:

```

# echo 'hist:key=child_comm:val=hitcount:size=256' > \
    /sys/kernel/debug/tracing/events/sched/sched_process_fork/
↪trigger

# cat /sys/kernel/debug/tracing/events/sched/sched_process_fork/hist
# trigger info: hist:keys=child_
↪comm:vals=hitcount:sort=hitcount:size=256 [active]

{ child_comm: dconf worker          } hitcount:          _
↪1
{ child_comm: ibus-daemon           } hitcount:          _
↪1
{ child_comm: whoopsie              } hitcount:          _
↪1
{ child_comm: smbd                  } hitcount:          _
↪1
{ child_comm: gdbus                 } hitcount:          _
↪1

```



```

{ child_comm: kthreadd                } hitcount:      1
↪1
{ child_comm: dconf worker            } hitcount:      1
↪1
{ child_comm: evolution-alarm         } hitcount:      2
↪2
{ child_comm: Socket Thread           } hitcount:      2
↪2
{ child_comm: postgres               } hitcount:      2
↪2
{ child_comm: bash                   } hitcount:      3
↪3
{ child_comm: compiz                  } hitcount:      3
↪3
{ child_comm: evolution-sourc        } hitcount:      4
↪4
{ child_comm: dhclient                } hitcount:      4
↪4
{ child_comm: pool                    } hitcount:      5
↪5
{ child_comm: nm-dispatcher.a         } hitcount:      8
↪8
{ child_comm: firefox                 } hitcount:      8
↪8
{ child_comm: dbus-daemon             } hitcount:      8
↪8
{ child_comm: glib-pacrunner          } hitcount:     10
↪10
{ child_comm: evolution               } hitcount:     23
↪23

Totals:
  Hits: 89
  Entries: 20
  Dropped: 0

```

If we want to pause the hist trigger, we can simply append `:pause` to the command that started the trigger. Notice that the trigger info displays as `[paused]`:

```

# echo 'hist:key=child_comm:val=hitcount:size=256:pause' >> \
  /sys/kernel/debug/tracing/events/sched/sched_process_fork/
↪trigger

# cat /sys/kernel/debug/tracing/events/sched/sched_process_fork/hist
# trigger info: hist:keys=child_
↪comm:vals=hitcount:sort=hitcount:size=256 [paused]

{ child_comm: dconf worker            } hitcount:      1
↪1
{ child_comm: kthreadd                } hitcount:      1
↪1

```

```

{ child_comm: dconf worker          } hitcount:      1
↪1
{ child_comm: gdbus                 } hitcount:      1
↪1
{ child_comm: ibus-daemon            } hitcount:      1
↪1
{ child_comm: Socket Thread         } hitcount:      2
↪2
{ child_comm: evolution-alarm        } hitcount:      2
↪2
{ child_comm: smbd                   } hitcount:      2
↪2
{ child_comm: bash                   } hitcount:      3
↪3
{ child_comm: whoopsie               } hitcount:      3
↪3
{ child_comm: compiz                 } hitcount:      3
↪3
{ child_comm: evolution-sourc        } hitcount:      4
↪4
{ child_comm: pool                   } hitcount:      5
↪5
{ child_comm: postgres               } hitcount:      6
↪6
{ child_comm: firefox                } hitcount:      8
↪8
{ child_comm: dhclient               } hitcount:     10
↪10
{ child_comm: emacs                  } hitcount:     12
↪12
{ child_comm: dbus-daemon            } hitcount:     20
↪20
{ child_comm: nm-dispatcher.a        } hitcount:     20
↪20
{ child_comm: evolution              } hitcount:     35
↪35
{ child_comm: glib-pacrunner         } hitcount:     59
↪59

Totals:
  Hits: 199
  Entries: 21
  Dropped: 0

```

To manually continue having the trigger aggregate events, append `:cont` instead. Notice that the trigger info displays as `[active]` again, and the data has changed:

```

# echo 'hist:key=child_comm:val=hitcount:size=256:cont' >> \
    /sys/kernel/debug/tracing/events/sched/sched_process_fork/
↪trigger

```

```

# cat /sys/kernel/debug/tracing/events/sched/sched_process_fork/hist
# trigger info: hist:keys=child_
  ↪ comm:vals=hitcount:sort=hitcount:size=256 [active]

{ child_comm: dconf worker                } hitcount: 1
  ↪ 1
{ child_comm: dconf worker                } hitcount: 1
  ↪ 1
{ child_comm: kthreadd                    } hitcount: 1
  ↪ 1
{ child_comm: gdbus                       } hitcount: 1
  ↪ 1
{ child_comm: ibus-daemon                  } hitcount: 1
  ↪ 1
{ child_comm: Socket Thread                } hitcount: 2
  ↪ 2
{ child_comm: evolution-alarm              } hitcount: 2
  ↪ 2
{ child_comm: smbd                         } hitcount: 2
  ↪ 2
{ child_comm: whoopsie                     } hitcount: 3
  ↪ 3
{ child_comm: compiz                       } hitcount: 3
  ↪ 3
{ child_comm: evolution-source             } hitcount: 4
  ↪ 4
{ child_comm: bash                         } hitcount: 5
  ↪ 5
{ child_comm: pool                         } hitcount: 5
  ↪ 5
{ child_comm: postgres                    } hitcount: 6
  ↪ 6
{ child_comm: firefox                      } hitcount: 8
  ↪ 8
{ child_comm: dhclient                     } hitcount: 11
  ↪ 11
{ child_comm: emacs                        } hitcount: 12
  ↪ 12
{ child_comm: dbus-daemon                  } hitcount: 22
  ↪ 22
{ child_comm: nm-dispatcher.a              } hitcount: 22
  ↪ 22
{ child_comm: evolution                    } hitcount: 35
  ↪ 35
{ child_comm: glib-pacrunner                } hitcount: 59
  ↪ 59

Totals:
  Hits: 206
  Entries: 21

```

```
Dropped: 0
```

The previous example showed how to start and stop a hist trigger by appending ‘pause’ and ‘continue’ to the hist trigger command. A hist trigger can also be started in a paused state by initially starting the trigger with ‘:pause’ appended. This allows you to start the trigger only when you’re ready to start collecting data and not before. For example, you could start the trigger in a paused state, then unpause it and do something you want to measure, then pause the trigger again when done.

Of course, doing this manually can be difficult and error-prone, but it is possible to automatically start and stop a hist trigger based on some condition, via the `enable_hist` and `disable_hist` triggers.

For example, suppose we wanted to take a look at the relative weights in terms of skb length for each callpath that leads to a `netif_receive_skb` event when downloading a decent-sized file using `wget`.

First we set up an initially paused stacktrace trigger on the `netif_receive_skb` event:

```
# echo 'hist:key=stacktrace:vals=len:pause' > \
    /sys/kernel/debug/tracing/events/net/netif_receive_skb/trigger
```

Next, we set up an ‘`enable_hist`’ trigger on the `sched_process_exec` event, with an ‘if filename==/usr/bin/wget’ filter. The effect of this new trigger is that it will ‘unpause’ the hist trigger we just set up on `netif_receive_skb` if and only if it sees a `sched_process_exec` event with a filename of ‘/usr/bin/wget’. When that happens, all `netif_receive_skb` events are aggregated into a hash table keyed on stacktrace:

```
# echo 'enable_hist:net:netif_receive_skb if filename==/usr/bin/wget' > \
    ↪ \
        /sys/kernel/debug/tracing/events/sched/sched_process_exec/
    ↪ trigger
```

The aggregation continues until the `netif_receive_skb` is paused again, which is what the following `disable_hist` event does by creating a similar setup on the `sched_process_exit` event, using the filter ‘`comm==wget`’:

```
# echo 'disable_hist:net:netif_receive_skb if comm==wget' > \
    /sys/kernel/debug/tracing/events/sched/sched_process_exit/
    ↪ trigger
```

Whenever a process exits and the `comm` field of the `disable_hist` trigger filter matches ‘`comm==wget`’, the `netif_receive_skb` hist trigger is disabled.

The overall effect is that `netif_receive_skb` events are aggregated into the hash table for only the duration of the `wget`. Executing a `wget` command and then listing the ‘hist’ file will display the output generated by the `wget` command:

```
$ wget https://www.kernel.org/pub/linux/kernel/v3.x/patch-3.19.xz

# cat /sys/kernel/debug/tracing/events/net/netif_receive_skb/hist
# trigger info: hist:keys=stacktrace:vals=len:sort=hitcount:size=2048
    ↪ [paused]
```

```

{ stacktrace:
    __netif_receive_skb_core+0x46d/0x990
    __netif_receive_skb+0x18/0x60
    netif_receive_skb_internal+0x23/0x90
    napi_gro_receive+0xc8/0x100
    ieee80211_deliver_skb+0xd6/0x270 [mac80211]
    ieee80211_rx_handlers+0xccf/0x22f0 [mac80211]
    ieee80211_prepare_and_rx_handle+0x4e7/0xc40 [mac80211]
    ieee80211_rx+0x31d/0x900 [mac80211]
    iwlnagn_rx_reply_rx+0x3db/0x6f0 [iwldvm]
    iwl_rx_dispatch+0x8e/0xf0 [iwldvm]
    iwl_pcie_irq_handler+0xe3c/0x12f0 [iwlwifi]
    irq_thread_fn+0x20/0x50
    irq_thread+0x11f/0x150
    kthread+0xd2/0xf0
    ret_from_fork+0x42/0x70
} hitcount:      85  len:      28884
{ stacktrace:
    __netif_receive_skb_core+0x46d/0x990
    __netif_receive_skb+0x18/0x60
    netif_receive_skb_internal+0x23/0x90
    napi_gro_complete+0xa4/0xe0
    dev_gro_receive+0x23a/0x360
    napi_gro_receive+0x30/0x100
    ieee80211_deliver_skb+0xd6/0x270 [mac80211]
    ieee80211_rx_handlers+0xccf/0x22f0 [mac80211]
    ieee80211_prepare_and_rx_handle+0x4e7/0xc40 [mac80211]
    ieee80211_rx+0x31d/0x900 [mac80211]
    iwlnagn_rx_reply_rx+0x3db/0x6f0 [iwldvm]
    iwl_rx_dispatch+0x8e/0xf0 [iwldvm]
    iwl_pcie_irq_handler+0xe3c/0x12f0 [iwlwifi]
    irq_thread_fn+0x20/0x50
    irq_thread+0x11f/0x150
    kthread+0xd2/0xf0
} hitcount:      98  len:      664329
{ stacktrace:
    __netif_receive_skb_core+0x46d/0x990
    __netif_receive_skb+0x18/0x60
    process_backlog+0xa8/0x150
    net_rx_action+0x15d/0x340
    __do_softirq+0x114/0x2c0
    do_softirq_own_stack+0x1c/0x30
    do_softirq+0x65/0x70
    __local_bh_enable_ip+0xb5/0xc0
    ip_finish_output+0x1f4/0x840
    ip_output+0x6b/0xc0
    ip_local_out_sk+0x31/0x40
    ip_send_skb+0x1a/0x50
    udp_send_skb+0x173/0x2a0
    udp_sendmsg+0x2bf/0x9f0

```

```

    inet_sendmsg+0x64/0xa0
    sock_sendmsg+0x3d/0x50
} hitcount:      115  len:      13030
{ stacktrace:
    __netif_receive_skb_core+0x46d/0x990
    __netif_receive_skb+0x18/0x60
    netif_receive_skb_internal+0x23/0x90
    napi_gro_complete+0xa4/0xe0
    napi_gro_flush+0x6d/0x90
    iwl_pcie_irq_handler+0x92a/0x12f0 [iwlwifi]
    irq_thread_fn+0x20/0x50
    irq_thread+0x11f/0x150
    kthread+0xd2/0xf0
    ret_from_fork+0x42/0x70
} hitcount:      934  len:      5512212

Totals:
  Hits: 1232
  Entries: 4
  Dropped: 0

```

The above shows all the `netif_receive_skb` callpaths and their total lengths for the duration of the `wget` command.

The ‘clear’ hist trigger param can be used to clear the hash table. Suppose we wanted to try another run of the previous example but this time also wanted to see the complete list of events that went into the histogram. In order to avoid having to set everything up again, we can just clear the histogram first:

```
# echo 'hist:key=stacktrace:vals=len:clear' >> \
    /sys/kernel/debug/tracing/events/net/netif_receive_skb/trigger
```

Just to verify that it is in fact cleared, here’s what we now see in the hist file:

```
# cat /sys/kernel/debug/tracing/events/net/netif_receive_skb/hist
# trigger info: hist:keys=stacktrace:vals=len:sort=hitcount:size=2048
↪[paused]

Totals:
  Hits: 0
  Entries: 0
  Dropped: 0

```

Since we want to see the detailed list of every `netif_receive_skb` event occurring during the new run, which are in fact the same events being aggregated into the hash table, we add some additional ‘enable_event’ events to the triggering `sched_process_exec` and `sched_process_exit` events as such:

```
# echo 'enable_event:net:netif_receive_skb if filename==/usr/bin/wget'
↪> \
    /sys/kernel/debug/tracing/events/sched/sched_process_exec/
↪trigger
```

```
# echo 'disable_event:net:netif_receive_skb if comm==wget' > \
    /sys/kernel/debug/tracing/events/sched/sched_process_exit/
↪ trigger
```

If you read the trigger files for the `sched_process_exec` and `sched_process_exit` triggers, you should see two triggers for each: one enabling/disabling the hist aggregation and the other enabling/disabling the logging of events:

```
# cat /sys/kernel/debug/tracing/events/sched/sched_process_exec/trigger
enable_event:net:netif_receive_skb:unlimited if filename==/usr/bin/wget
enable_hist:net:netif_receive_skb:unlimited if filename==/usr/bin/wget

# cat /sys/kernel/debug/tracing/events/sched/sched_process_exit/trigger
enable_event:net:netif_receive_skb:unlimited if comm==wget
disable_hist:net:netif_receive_skb:unlimited if comm==wget
```

In other words, whenever either of the `sched_process_exec` or `sched_process_exit` events is hit and matches ‘wget’, it enables or disables both the histogram and the event log, and what you end up with is a hash table and set of events just covering the specified duration. Run the `wget` command again:

```
$ wget https://www.kernel.org/pub/linux/kernel/v3.x/patch-3.19.xz
```

Displaying the ‘hist’ file should show something similar to what you saw in the last run, but this time you should also see the individual events in the trace file:

```
# cat /sys/kernel/debug/tracing/trace

# tracer: nop
#
# entries-in-buffer/entries-written: 183/1426   #P:4
#
#
#          _-----=> irqs-off
#          /_-----=> need-resched
#          | /_-----=> hardirq/softirq
#          || /_-----=> preempt-depth
#          ||| /_-----=> delay
#
# TASK-PID   CPU#  | |||   TIMESTAMP  FUNCTION
# | | | | |
  wget-15108 [000] ..s1 31769.606929: netif_receive_skb:↪
↪ dev=lo skbaddr=ffff88009c353100 len=60
  wget-15108 [000] ..s1 31769.606999: netif_receive_skb:↪
↪ dev=lo skbaddr=ffff88009c353200 len=60
  dnsmasq-1382 [000] ..s1 31769.677652: netif_receive_skb:↪
↪ dev=lo skbaddr=ffff88009c352b00 len=130
  dnsmasq-1382 [000] ..s1 31769.685917: netif_receive_skb:↪
↪ dev=lo skbaddr=ffff88009c352200 len=138
##### CPU 2 buffer started #####
  irq/29-iwlwifi-559 [002] ..s. 31772.031529: netif_receive_skb:↪
↪ dev=wlan0 skbaddr=ffff88009d433d00 len=2948
```

```

irq/29-iwlwifi-559 [002] ..s. 31772.031572: netif_receive_skb:
↳dev=wlan0 skbaddr=ffff88009d432200 len=1500
irq/29-iwlwifi-559 [002] ..s. 31772.032196: netif_receive_skb:
↳dev=wlan0 skbaddr=ffff88009d433100 len=2948
irq/29-iwlwifi-559 [002] ..s. 31772.032761: netif_receive_skb:
↳dev=wlan0 skbaddr=ffff88009d433000 len=2948
irq/29-iwlwifi-559 [002] ..s. 31772.033220: netif_receive_skb:
↳dev=wlan0 skbaddr=ffff88009d432e00 len=1500
.
.
.
```

The following example demonstrates how multiple hist triggers can be attached to a given event. This capability can be useful for creating a set of different summaries derived from the same set of events, or for comparing the effects of different filters, among other things:

```

# echo 'hist:keys=skbaddr.hex:vals=len if len < 0' >> \
  /sys/kernel/debug/tracing/events/net/netif_receive_skb/trigger
# echo 'hist:keys=skbaddr.hex:vals=len if len > 4096' >> \
  /sys/kernel/debug/tracing/events/net/netif_receive_skb/trigger
# echo 'hist:keys=skbaddr.hex:vals=len if len == 256' >> \
  /sys/kernel/debug/tracing/events/net/netif_receive_skb/trigger
# echo 'hist:keys=skbaddr.hex:vals=len' >> \
  /sys/kernel/debug/tracing/events/net/netif_receive_skb/trigger
# echo 'hist:keys=len:vals=common_preempt_count' >> \
  /sys/kernel/debug/tracing/events/net/netif_receive_skb/trigger
```

The above set of commands create four triggers differing only in their filters, along with a completely different though fairly nonsensical trigger. Note that in order to append multiple hist triggers to the same file, you should use the '>>' operator to append them ('>' will also add the new hist trigger, but will remove any existing hist triggers beforehand).

Displaying the contents of the 'hist' file for the event shows the contents of all five histograms:

```

# cat /sys/kernel/debug/tracing/events/net/netif_receive_skb/hist

# event histogram
#
# trigger info: hist:keys=len:vals=hitcount,common_preempt_
↳count:sort=hitcount:size=2048 [active]
#

{ len:          176 } hitcount:          1  common_preempt_count:
↳ 0
{ len:          223 } hitcount:          1  common_preempt_count:
↳ 0
{ len:          4854 } hitcount:          1  common_preempt_count:
↳ 0
```



```

{ len:          395 } hitcount:          1  common_preempt_count:
↳ 0
{ len:          177 } hitcount:          1  common_preempt_count:
↳ 0
{ len:          446 } hitcount:          1  common_preempt_count:
↳ 0
{ len:         1601 } hitcount:          1  common_preempt_count:
↳ 0
.
.
.
{ len:         1280 } hitcount:         66  common_preempt_count:
↳ 0
{ len:          116 } hitcount:         81  common_preempt_count:
↳ 40
{ len:          708 } hitcount:        112  common_preempt_count:
↳ 0
{ len:           46 } hitcount:        221  common_preempt_count:
↳ 0
{ len:         1264 } hitcount:        458  common_preempt_count:
↳ 0

Totals:
  Hits: 1428
  Entries: 147
  Dropped: 0

# event histogram
#
# trigger info: hist:keys=skbaddr.hex:vals=hitcount,
↳ len:sort=hitcount:size=2048 [active]
#

{ skbaddr: ffff8800baee5e00 } hitcount:          1  len:          130
{ skbaddr: ffff88005f3d5600 } hitcount:          1  len:         1280
{ skbaddr: ffff88005f3d4900 } hitcount:          1  len:         1280
{ skbaddr: ffff88009fed6300 } hitcount:          1  len:          115
{ skbaddr: ffff88009fe0ad00 } hitcount:          1  len:          115
{ skbaddr: ffff88008cdb1900 } hitcount:          1  len:           46
{ skbaddr: ffff880064b5ef00 } hitcount:          1  len:          118
{ skbaddr: ffff880044e3c700 } hitcount:          1  len:           60
{ skbaddr: ffff880100065900 } hitcount:          1  len:           46
{ skbaddr: ffff8800d46bd500 } hitcount:          1  len:          116
{ skbaddr: ffff88005f3d5f00 } hitcount:          1  len:         1280
{ skbaddr: ffff880100064700 } hitcount:          1  len:          365
{ skbaddr: ffff8800badb6f00 } hitcount:          1  len:           60
.
.
.
```

```

{ skbaddr: ffff88009fe0be00 } hitcount:      27  len:      24677
{ skbaddr: ffff88009fe0a400 } hitcount:      27  len:      23052
{ skbaddr: ffff88009fe0b700 } hitcount:      31  len:      25589
{ skbaddr: ffff88009fe0b600 } hitcount:      32  len:      27326
{ skbaddr: ffff88006a462800 } hitcount:      68  len:      71678
{ skbaddr: ffff88006a463700 } hitcount:      70  len:      72678
{ skbaddr: ffff88006a462b00 } hitcount:      71  len:      77589
{ skbaddr: ffff88006a463600 } hitcount:      73  len:      71307
{ skbaddr: ffff88006a462200 } hitcount:      81  len:      81032

```

Totals:

```

    Hits: 1451
    Entries: 318
    Dropped: 0

```

event histogram

#

```

# trigger info: hist:keys=skbaddr.hex:vals=hitcount,
↳ len:sort=hitcount:size=2048 if len == 256 [active]

```

#

Totals:

```

    Hits: 0
    Entries: 0
    Dropped: 0

```

event histogram

#

```

# trigger info: hist:keys=skbaddr.hex:vals=hitcount,
↳ len:sort=hitcount:size=2048 if len > 4096 [active]

```

#

```

{ skbaddr: ffff88009fd2c300 } hitcount:      1  len:      7212
{ skbaddr: ffff8800d2bcce00 } hitcount:      1  len:      7212
{ skbaddr: ffff8800d2bcd700 } hitcount:      1  len:      7212
{ skbaddr: ffff8800d2bcda00 } hitcount:      1  len:     21492
{ skbaddr: ffff8800ae2e2d00 } hitcount:      1  len:      7212
{ skbaddr: ffff8800d2bcdcb00 } hitcount:      1  len:      7212
{ skbaddr: ffff88006a4df500 } hitcount:      1  len:      4854
{ skbaddr: ffff88008ce47b00 } hitcount:      1  len:     18636
{ skbaddr: ffff8800ae2e2200 } hitcount:      1  len:     12924
{ skbaddr: ffff88005f3e1000 } hitcount:      1  len:      4356
{ skbaddr: ffff8800d2bcdcb00 } hitcount:      2  len:     24420
{ skbaddr: ffff8800d2bcc200 } hitcount:      2  len:     12996

```

Totals:

```

    Hits: 14

```

```

    Entries: 12
    Dropped: 0

# event histogram
#
# trigger info: hist:keys=skbaddr.hex:vals=hitcount,
→ len:sort=hitcount:size=2048 if len < 0 [active]
#

Totals:
    Hits: 0
    Entries: 0
    Dropped: 0

```

Named triggers can be used to have triggers share a common set of histogram data. This capability is mostly useful for combining the output of events generated by tracepoints contained inside inline functions, but names can be used in a hist trigger on any event. For example, these two triggers when hit will update the same 'len' field in the shared 'foo' histogram data:

```

# echo 'hist:name=foo:keys=skbaddr.hex:vals=len' > \
    /sys/kernel/debug/tracing/events/net/netif_receive_skb/trigger
# echo 'hist:name=foo:keys=skbaddr.hex:vals=len' > \
    /sys/kernel/debug/tracing/events/net/netif_rx/trigger

```

You can see that they're updating common histogram data by reading each event's hist files at the same time:

```

# cat /sys/kernel/debug/tracing/events/net/netif_receive_skb/hist;
  cat /sys/kernel/debug/tracing/events/net/netif_rx/hist

# event histogram
#
# trigger info: hist:name=foo:keys=skbaddr.hex:vals=hitcount,
→ len:sort=hitcount:size=2048 [active]
#

{ skbaddr: ffff88000ad53500 } hitcount:      1  len:      46
{ skbaddr: ffff8800af5a1500 } hitcount:      1  len:      76
{ skbaddr: ffff8800d62a1900 } hitcount:      1  len:      46
{ skbaddr: ffff8800d2bccb00 } hitcount:      1  len:     468
{ skbaddr: ffff8800d3c69900 } hitcount:      1  len:      46
{ skbaddr: ffff88009ff09100 } hitcount:      1  len:      52
{ skbaddr: ffff88010f13ab00 } hitcount:      1  len:     168
{ skbaddr: ffff88006a54f400 } hitcount:      1  len:      46
{ skbaddr: ffff8800d2bcc500 } hitcount:      1  len:     260
{ skbaddr: ffff880064505000 } hitcount:      1  len:      46
{ skbaddr: ffff8800baf24e00 } hitcount:      1  len:      32
{ skbaddr: ffff88009fe0ad00 } hitcount:      1  len:      46

```

```

{ skbaddr: ffff8800d3edff00 } hitcount:      1 len:      44
{ skbaddr: ffff88009fe0b400 } hitcount:      1 len:     168
{ skbaddr: ffff8800a1c55a00 } hitcount:      1 len:      40
{ skbaddr: ffff8800d2bcd100 } hitcount:      1 len:      40
{ skbaddr: ffff880064505f00 } hitcount:      1 len:     174
{ skbaddr: ffff8800a8bff200 } hitcount:      1 len:     160
{ skbaddr: ffff880044e3cc00 } hitcount:      1 len:      76
{ skbaddr: ffff8800a8bfe700 } hitcount:      1 len:      46
{ skbaddr: ffff8800d2bcd000 } hitcount:      1 len:      32
{ skbaddr: ffff8800a1f64800 } hitcount:      1 len:      46
{ skbaddr: ffff8800d2bcde00 } hitcount:      1 len:     988
{ skbaddr: ffff88006a5dea00 } hitcount:      1 len:      46
{ skbaddr: ffff88002e37a200 } hitcount:      1 len:      44
{ skbaddr: ffff8800a1f32c00 } hitcount:      2 len:     676
{ skbaddr: ffff88000ad52600 } hitcount:      2 len:     107
{ skbaddr: ffff8800a1f91e00 } hitcount:      2 len:      92
{ skbaddr: ffff8800af5a0200 } hitcount:      2 len:     142
{ skbaddr: ffff8800d2bcc600 } hitcount:      2 len:     220
{ skbaddr: ffff8800ba36f500 } hitcount:      2 len:      92
{ skbaddr: ffff8800d021f800 } hitcount:      2 len:      92
{ skbaddr: ffff8800a1f33600 } hitcount:      2 len:     675
{ skbaddr: ffff8800a8bfff00 } hitcount:      3 len:     138
{ skbaddr: ffff8800d62a1300 } hitcount:      3 len:     138
{ skbaddr: ffff88002e37a100 } hitcount:      4 len:     184
{ skbaddr: ffff880064504400 } hitcount:      4 len:     184
{ skbaddr: ffff8800a8bfec00 } hitcount:      4 len:     184
{ skbaddr: ffff88000ad53700 } hitcount:      5 len:     230
{ skbaddr: ffff8800d2bcd000 } hitcount:      5 len:     196
{ skbaddr: ffff8800a1f90000 } hitcount:      6 len:     276
{ skbaddr: ffff88006a54f900 } hitcount:      6 len:     276

```

Totals:

Hits: 81

Entries: 42

Dropped: 0

event histogram

#

trigger info: hist:name=foo:keys=skbaddr.hex:vals=hitcount,
 ↪ len:sort=hitcount:size=2048 [active]

#

```

{ skbaddr: ffff88000ad53500 } hitcount:      1 len:      46
{ skbaddr: ffff8800af5a1500 } hitcount:      1 len:      76
{ skbaddr: ffff8800d62a1900 } hitcount:      1 len:      46
{ skbaddr: ffff8800d2bccb00 } hitcount:      1 len:     468
{ skbaddr: ffff8800d3c69900 } hitcount:      1 len:      46
{ skbaddr: ffff88009ff09100 } hitcount:      1 len:      52
{ skbaddr: ffff88010f13ab00 } hitcount:      1 len:     168
{ skbaddr: ffff88006a54f400 } hitcount:      1 len:      46
{ skbaddr: ffff8800d2bcc500 } hitcount:      1 len:     260

```

```

{ skbaddr: ffff880064505000 } hitcount:      1 len:      46
{ skbaddr: ffff8800baf24e00 } hitcount:      1 len:      32
{ skbaddr: ffff88009fe0ad00 } hitcount:      1 len:      46
{ skbaddr: ffff8800d3edff00 } hitcount:      1 len:      44
{ skbaddr: ffff88009fe0b400 } hitcount:      1 len:     168
{ skbaddr: ffff8800a1c55a00 } hitcount:      1 len:      40
{ skbaddr: ffff8800d2bcd100 } hitcount:      1 len:      40
{ skbaddr: ffff880064505f00 } hitcount:      1 len:     174
{ skbaddr: ffff8800a8bff200 } hitcount:      1 len:     160
{ skbaddr: ffff880044e3cc00 } hitcount:      1 len:      76
{ skbaddr: ffff8800a8bfe700 } hitcount:      1 len:      46
{ skbaddr: ffff8800d2bcd000 } hitcount:      1 len:      32
{ skbaddr: ffff8800a1f64800 } hitcount:      1 len:      46
{ skbaddr: ffff8800d2bcde00 } hitcount:      1 len:     988
{ skbaddr: ffff88006a5dea00 } hitcount:      1 len:      46
{ skbaddr: ffff88002e37a200 } hitcount:      1 len:      44
{ skbaddr: ffff8800a1f32c00 } hitcount:      2 len:     676
{ skbaddr: ffff88000ad52600 } hitcount:      2 len:     107
{ skbaddr: ffff8800a1f91e00 } hitcount:      2 len:      92
{ skbaddr: ffff8800af5a0200 } hitcount:      2 len:     142
{ skbaddr: ffff8800d2bcc600 } hitcount:      2 len:     220
{ skbaddr: ffff8800ba36f500 } hitcount:      2 len:      92
{ skbaddr: ffff8800d021f800 } hitcount:      2 len:      92
{ skbaddr: ffff8800a1f33600 } hitcount:      2 len:     675
{ skbaddr: ffff8800a8bfff00 } hitcount:      3 len:     138
{ skbaddr: ffff8800d62a1300 } hitcount:      3 len:     138
{ skbaddr: ffff88002e37a100 } hitcount:      4 len:     184
{ skbaddr: ffff880064504400 } hitcount:      4 len:     184
{ skbaddr: ffff8800a8bfec00 } hitcount:      4 len:     184
{ skbaddr: ffff88000ad53700 } hitcount:      5 len:     230
{ skbaddr: ffff8800d2bcd000 } hitcount:      5 len:     196
{ skbaddr: ffff8800a1f90000 } hitcount:      6 len:     276
{ skbaddr: ffff88006a54f900 } hitcount:      6 len:     276

```

Totals:

```

    Hits: 81
    Entries: 42
    Dropped: 0

```

And here's an example that shows how to combine histogram data from any two events even if they don't share any 'compatible' fields other than 'hitcount' and 'stacktrace'. These commands create a couple of triggers named 'bar' using those fields:

```

# echo 'hist:name=bar:key=stacktrace:val=hitcount' > \
    /sys/kernel/debug/tracing/events/sched/sched_process_fork/
↪ trigger
# echo 'hist:name=bar:key=stacktrace:val=hitcount' > \
    /sys/kernel/debug/tracing/events/net/netif_rx/trigger

```

And displaying the output of either shows some interesting if somewhat confusing output:

```

# cat /sys/kernel/debug/tracing/events/sched/sched_process_fork/hist
# cat /sys/kernel/debug/tracing/events/net/netif_rx/hist

# event histogram
#
# trigger info:
↪ hist:name=bar:keys=stacktrace:vals=hitcount:sort=hitcount:size=2048
↪ [active]
#

{ stacktrace:
    kernel_clone+0x18e/0x330
    kernel_thread+0x29/0x30
    kthreadd+0x154/0x1b0
    ret_from_fork+0x3f/0x70
} hitcount:      1
{ stacktrace:
    netif_rx_internal+0xb2/0xd0
    netif_rx_ni+0x20/0x70
    dev_loopback_xmit+0xaa/0xd0
    ip_mc_output+0x126/0x240
    ip_local_out_sk+0x31/0x40
    igmp_send_report+0x1e9/0x230
    igmp_timer_expire+0xe9/0x120
    call_timer_fn+0x39/0xf0
    run_timer_softirq+0x1e1/0x290
    __do_softirq+0xfd/0x290
    irq_exit+0x98/0xb0
    smp_apic_timer_interrupt+0x4a/0x60
    apic_timer_interrupt+0x6d/0x80
    cpuidle_enter+0x17/0x20
    call_cpuidle+0x3b/0x60
    cpu_startup_entry+0x22d/0x310
} hitcount:      1
{ stacktrace:
    netif_rx_internal+0xb2/0xd0
    netif_rx_ni+0x20/0x70
    dev_loopback_xmit+0xaa/0xd0
    ip_mc_output+0x17f/0x240
    ip_local_out_sk+0x31/0x40
    ip_send_skb+0x1a/0x50
    udp_send_skb+0x13e/0x270
    udp_sendmsg+0x2bf/0x980
    inet_sendmsg+0x67/0xa0
    sock_sendmsg+0x38/0x50
    SYSC_sendto+0xef/0x170
    Sys_sendto+0xe/0x10
    entry_SYSCALL_64_fastpath+0x12/0x6a
} hitcount:      2
{ stacktrace:

```

```

    netif_rx_internal+0xb2/0xd0
    netif_rx+0x1c/0x60
    loopback_xmit+0x6c/0xb0
    dev_hard_start_xmit+0x219/0x3a0
    __dev_queue_xmit+0x415/0x4f0
    dev_queue_xmit_sk+0x13/0x20
    ip_finish_output2+0x237/0x340
    ip_finish_output+0x113/0x1d0
    ip_output+0x66/0xc0
    ip_local_out_sk+0x31/0x40
    ip_send_skb+0x1a/0x50
    udp_send_skb+0x16d/0x270
    udp_sendmsg+0x2bf/0x980
    inet_sendmsg+0x67/0xa0
    sock_sendmsg+0x38/0x50
    __sys_sendmsg+0x14e/0x270
} hitcount:      76
{ stacktrace:
    netif_rx_internal+0xb2/0xd0
    netif_rx+0x1c/0x60
    loopback_xmit+0x6c/0xb0
    dev_hard_start_xmit+0x219/0x3a0
    __dev_queue_xmit+0x415/0x4f0
    dev_queue_xmit_sk+0x13/0x20
    ip_finish_output2+0x237/0x340
    ip_finish_output+0x113/0x1d0
    ip_output+0x66/0xc0
    ip_local_out_sk+0x31/0x40
    ip_send_skb+0x1a/0x50
    udp_send_skb+0x16d/0x270
    udp_sendmsg+0x2bf/0x980
    inet_sendmsg+0x67/0xa0
    sock_sendmsg+0x38/0x50
    __sys_sendmsg+0x269/0x270
} hitcount:      77
{ stacktrace:
    netif_rx_internal+0xb2/0xd0
    netif_rx+0x1c/0x60
    loopback_xmit+0x6c/0xb0
    dev_hard_start_xmit+0x219/0x3a0
    __dev_queue_xmit+0x415/0x4f0
    dev_queue_xmit_sk+0x13/0x20
    ip_finish_output2+0x237/0x340
    ip_finish_output+0x113/0x1d0
    ip_output+0x66/0xc0
    ip_local_out_sk+0x31/0x40
    ip_send_skb+0x1a/0x50
    udp_send_skb+0x16d/0x270
    udp_sendmsg+0x2bf/0x980
    inet_sendmsg+0x67/0xa0

```

```
        sock_sendmsg+0x38/0x50
        SYSC_sendto+0xef/0x170
} hitcount:      88
{ stacktrace:
    kernel_clone+0x18e/0x330
    Sys_clone+0x19/0x20
    entry_SYSCALL_64_fastpath+0x12/0x6a
} hitcount:      244

Totals:
  Hits: 489
  Entries: 7
  Dropped: 0
```

16.2.4 2.2 Inter-event hist triggers

Inter-event hist triggers are hist triggers that combine values from one or more other events and create a histogram using that data. Data from an inter-event histogram can in turn become the source for further combined histograms, thus providing a chain of related histograms, which is important for some applications.

The most important example of an inter-event quantity that can be used in this manner is latency, which is simply a difference in timestamps between two events. Although latency is the most important inter-event quantity, note that because the support is completely general across the trace event subsystem, any event field can be used in an inter-event quantity.

An example of a histogram that combines data from other histograms into a useful chain would be a ‘wakeups witch latency’ histogram that combines a ‘wake up latency’ histogram and a ‘switch latency’ histogram.

Normally, a hist trigger specification consists of a (possibly compound) key along with one or more numeric values, which are continually updated sums associated with that key. A histogram specification in this case consists of individual key and value specifications that refer to trace event fields associated with a single event type.

The inter-event hist trigger extension allows fields from multiple events to be referenced and combined into a multi-event histogram specification. In support of this overall goal, a few enabling features have been added to the hist trigger support:

- In order to compute an inter-event quantity, a value from one event needs to be saved and then referenced from another event. This requires the introduction of support for histogram ‘variables’.
- The computation of inter-event quantities and their combination require some minimal amount of support for applying simple expressions to variables (+ and -).
- A histogram consisting of inter-event quantities isn’t logically a histogram on either event (so having the ‘hist’ file for either event host the histogram output doesn’t really make sense). To address the idea that the histogram is associated with a combination of events, support is added allowing the creation of ‘synthetic’ events that are events derived from other events. These synthetic events are full-fledged events just like any other and can be used as such, as for instance to create the ‘combination’ histograms mentioned previously.

- A set of ‘actions’ can be associated with histogram entries - these can be used to generate the previously mentioned synthetic events, but can also be used for other purposes, such as for example saving context when a ‘max’ latency has been hit.
- Trace events don’t have a ‘timestamp’ associated with them, but there is an implicit timestamp saved along with an event in the underlying ftrace ring buffer. This timestamp is now exposed as a synthetic field named ‘common_timestamp’ which can be used in histograms as if it were any other event field; it isn’t an actual field in the trace format but rather is a synthesized value that nonetheless can be used as if it were an actual field. By default it is in units of nanoseconds; appending ‘.usecs’ to a common_timestamp field changes the units to microseconds.

A note on inter-event timestamps: If common_timestamp is used in a histogram, the trace buffer is automatically switched over to using absolute timestamps and the “global” trace clock, in order to avoid bogus timestamp differences with other clocks that aren’t coherent across CPUs. This can be overridden by specifying one of the other trace clocks instead, using the “clock=XXX” hist trigger attribute, where XXX is any of the clocks listed in the tracing/trace_clock pseudo-file.

These features are described in more detail in the following sections.

16.2.5 2.2.1 Histogram Variables

Variables are simply named locations used for saving and retrieving values between matching events. A ‘matching’ event is defined as an event that has a matching key - if a variable is saved for a histogram entry corresponding to that key, any subsequent event with a matching key can access that variable.

A variable’s value is normally available to any subsequent event until it is set to something else by a subsequent event. The one exception to that rule is that any variable used in an expression is essentially ‘read-once’ - once it’s used by an expression in a subsequent event, it’s reset to its ‘unset’ state, which means it can’t be used again unless it’s set again. This ensures not only that an event doesn’t use an uninitialized variable in a calculation, but that that variable is used only once and not for any unrelated subsequent match.

The basic syntax for saving a variable is to simply prefix a unique variable name not corresponding to any keyword along with an ‘=’ sign to any event field.

Either keys or values can be saved and retrieved in this way. This creates a variable named ‘ts0’ for a histogram entry with the key ‘next_pid’:

```
# echo 'hist:keys=next_pid:vals=$ts0:ts0=common_timestamp ... >> \
    event/trigger'
```

The ts0 variable can be accessed by any subsequent event having the same pid as ‘next_pid’.

Variable references are formed by prepending the variable name with the ‘\$’ sign. Thus for example, the ts0 variable above would be referenced as ‘\$ts0’ in expressions.

Because ‘vals=’ is used, the common_timestamp variable value above will also be summed as a normal histogram value would (though for a timestamp it makes little sense).

The below shows that a key value can also be saved in the same way:

```
# echo 'hist:timer_pid=common_pid:key=timer_pid ...' >> event/trigger
```

If a variable isn't a key variable or prefixed with 'vals=', the associated event field will be saved in a variable but won't be summed as a value:

```
# echo 'hist:keys=next_pid:ts1=common_timestamp ...' >> event/trigger
```

Multiple variables can be assigned at the same time. The below would result in both ts0 and b being created as variables, with both common_timestamp and field1 additionally being summed as values:

```
# echo 'hist:keys=pid:vals=$ts0,$b:ts0=common_timestamp,b=field1 ...' >> \
    event/trigger
```

Note that variable assignments can appear either preceding or following their use. The command below behaves identically to the command above:

```
# echo 'hist:keys=pid:ts0=common_timestamp,b=field1:vals=$ts0,$b ...' >> \
    event/trigger
```

Any number of variables not bound to a 'vals=' prefix can also be assigned by simply separating them with colons. Below is the same thing but without the values being summed in the histogram:

```
# echo 'hist:keys=pid:ts0=common_timestamp:b=field1 ...' >> event/trigger
```

Variables set as above can be referenced and used in expressions on another event.

For example, here's how a latency can be calculated:

```
# echo 'hist:keys=pid,prio:ts0=common_timestamp ...' >> event1/trigger
# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp-$ts0 ...' >> event2/
↳ trigger
```

In the first line above, the event's timestamp is saved into the variable ts0. In the next line, ts0 is subtracted from the second event's timestamp to produce the latency, which is then assigned into yet another variable, 'wakeup_lat'. The hist trigger below in turn makes use of the wakeup_lat variable to compute a combined latency using the same key and variable from yet another event:

```
# echo 'hist:key=pid:wakeupswitch_lat=$wakeup_lat+$switchtime_lat ...' >>
↳ event3/trigger
```

Expressions support the use of addition, subtraction, multiplication and division operators (+, -, *, /).

Note if division by zero cannot be detected at parse time (i.e. the divisor is not a constant), the result will be -1.

Numeric constants can also be used directly in an expression:

```
# echo 'hist:keys=next_pid:timestamp_secs=common_timestamp/1000000 ...' >>
↳ event/trigger
```

or assigned to a variable and referenced in a subsequent expression:

```
# echo 'hist:keys=next_pid:us_per_sec=1000000 ...' >> event/trigger
# echo 'hist:keys=next_pid:timestamp_secs=common_timestamp/$us_per_sec ...' >>
↪event/trigger
```

16.2.6 2.2.2 Synthetic Events

Synthetic events are user-defined events generated from hist trigger variables or fields associated with one or more other events. Their purpose is to provide a mechanism for displaying data spanning multiple events consistent with the existing and already familiar usage for normal events.

To define a synthetic event, the user writes a simple specification consisting of the name of the new event along with one or more variables and their types, which can be any valid field type, separated by semicolons, to the tracing/synthetic_events file.

See `synth_field_size()` for available types.

If `field_name` contains `[n]`, the field is considered to be a static array.

If `field_names` contains `[]` (no subscript), the field is considered to be a dynamic array, which will only take as much space in the event as is required to hold the array.

A string field can be specified using either the static notation:

```
char name[32];
```

Or the dynamic:

```
char name[];
```

The size limit for either is 256.

For instance, the following creates a new event named 'wakeup_latency' with 3 fields: `lat`, `pid`, and `prio`. Each of those fields is simply a variable reference to a variable on another event:

```
# echo 'wakeup_latency \
      u64 lat; \
      pid_t pid; \
      int prio' >> \
      /sys/kernel/debug/tracing/synthetic_events
```

Reading the `tracing/synthetic_events` file lists all the currently defined synthetic events, in this case the event defined above:

```
# cat /sys/kernel/debug/tracing/synthetic_events
wakeup_latency u64 lat; pid_t pid; int prio
```

An existing synthetic event definition can be removed by prepending the command that defined it with a '!':

```
# echo '!wakeup_latency u64 lat pid_t pid int prio' >> \
      /sys/kernel/debug/tracing/synthetic_events
```

At this point, there isn't yet an actual 'wakeup_latency' event instantiated in the event subsystem - for this to happen, a 'hist trigger action' needs to be instantiated and bound to actual fields and variables defined on other events (see Section 2.2.3 below on how that is done using hist trigger 'onmatch' action). Once that is done, the 'wakeup_latency' synthetic event instance is created.

The new event is created under the tracing/events/synthetic/ directory and looks and behaves just like any other event:

```
# ls /sys/kernel/debug/tracing/events/synthetic/wakeup_latency
enable filter format hist id trigger
```

A histogram can now be defined for the new synthetic event:

```
# echo 'hist:keys=pid,prio,lat.log2:sort=lat' >> \
/sys/kernel/debug/tracing/events/synthetic/wakeup_latency/trigger
```

The above shows the latency "lat" in a power of 2 grouping.

Like any other event, once a histogram is enabled for the event, the output can be displayed by reading the event's 'hist' file.

```
# cat /sys/kernel/debug/tracing/events/synthetic/wakeup_latency/hist
```

```
# event histogram # # trigger info: hist:keys=pid,prio,lat.log2:vals=hitcount:sort=lat.log2:size=
[active] #
```

```
{ pid: 2035, prio: 9, lat: ~ 2^2 } hitcount: 43 { pid: 2034, prio: 9, lat: ~ 2^2 }
hitcount: 60 { pid: 2029, prio: 9, lat: ~ 2^2 } hitcount: 965 { pid: 2034, prio: 120,
lat: ~ 2^2 } hitcount: 9 { pid: 2033, prio: 120, lat: ~ 2^2 } hitcount: 5 { pid: 2030,
prio: 9, lat: ~ 2^2 } hitcount: 335 { pid: 2030, prio: 120, lat: ~ 2^2 } hitcount: 10
{ pid: 2032, prio: 120, lat: ~ 2^2 } hitcount: 1 { pid: 2035, prio: 120, lat: ~ 2^2 }
hitcount: 2 { pid: 2031, prio: 9, lat: ~ 2^2 } hitcount: 176 { pid: 2028, prio: 120,
lat: ~ 2^2 } hitcount: 15 { pid: 2033, prio: 9, lat: ~ 2^2 } hitcount: 91 { pid: 2032,
prio: 9, lat: ~ 2^2 } hitcount: 125 { pid: 2029, prio: 120, lat: ~ 2^2 } hitcount: 4
{ pid: 2031, prio: 120, lat: ~ 2^2 } hitcount: 3 { pid: 2029, prio: 120, lat: ~ 2^3 }
hitcount: 2 { pid: 2035, prio: 9, lat: ~ 2^3 } hitcount: 41 { pid: 2030, prio: 120, lat:
~ 2^3 } hitcount: 1 { pid: 2032, prio: 9, lat: ~ 2^3 } hitcount: 32 { pid: 2031, prio:
9, lat: ~ 2^3 } hitcount: 44 { pid: 2034, prio: 9, lat: ~ 2^3 } hitcount: 40 { pid:
2030, prio: 9, lat: ~ 2^3 } hitcount: 29 { pid: 2033, prio: 9, lat: ~ 2^3 } hitcount:
31 { pid: 2029, prio: 9, lat: ~ 2^3 } hitcount: 31 { pid: 2028, prio: 120, lat: ~ 2^3 }
hitcount: 18 { pid: 2031, prio: 120, lat: ~ 2^3 } hitcount: 2 { pid: 2028, prio: 120,
lat: ~ 2^4 } hitcount: 1 { pid: 2029, prio: 9, lat: ~ 2^4 } hitcount: 4 { pid: 2031,
prio: 120, lat: ~ 2^7 } hitcount: 1 { pid: 2032, prio: 120, lat: ~ 2^7 } hitcount: 1
```

```
Totals: Hits: 2122 Entries: 30 Dropped: 0
```

The latency values can also be grouped linearly by a given size with the ".buckets" modifier and specify a size (in this case groups of 10).

```
# echo 'hist:keys=pid,prio,lat.buckets=10:sort=lat' >>
```

```
/sys/kernel/debug/tracing/events/synthetic/wakeup_latency/trigger
```

```
# event histogram # # trigger info: hist:keys=pid,prio,lat.buckets=10:vals=hitcount:sort=lat.bu
[active] #
```

```
{ pid: 2067, prio: 9, lat: ~ 0-9 } hitcount: 220 { pid: 2068, prio: 9, lat: ~ 0-9 }
hitcount: 157 { pid: 2070, prio: 9, lat: ~ 0-9 } hitcount: 100 { pid: 2067, prio: 120,
lat: ~ 0-9 } hitcount: 6 { pid: 2065, prio: 120, lat: ~ 0-9 } hitcount: 2 { pid: 2066,
prio: 120, lat: ~ 0-9 } hitcount: 2 { pid: 2069, prio: 9, lat: ~ 0-9 } hitcount: 122
{ pid: 2069, prio: 120, lat: ~ 0-9 } hitcount: 8 { pid: 2070, prio: 120, lat: ~ 0-9 }
hitcount: 1 { pid: 2068, prio: 120, lat: ~ 0-9 } hitcount: 7 { pid: 2066, prio: 9, lat:
~ 0-9 } hitcount: 365 { pid: 2064, prio: 120, lat: ~ 0-9 } hitcount: 35 { pid: 2065,
prio: 9, lat: ~ 0-9 } hitcount: 998 { pid: 2071, prio: 9, lat: ~ 0-9 } hitcount: 85 {
pid: 2065, prio: 9, lat: ~ 10-19 } hitcount: 2 { pid: 2064, prio: 120, lat: ~ 10-19 }
hitcount: 2
```

Totals: Hits: 2112 Entries: 16 Dropped: 0

16.2.7 2.2.3 Hist trigger ‘handlers’ and ‘actions’

A hist trigger ‘action’ is a function that’s executed (in most cases conditionally) whenever a histogram entry is added or updated.

When a histogram entry is added or updated, a hist trigger ‘handler’ is what decides whether the corresponding action is actually invoked or not.

Hist trigger handlers and actions are paired together in the general form:

```
<handler>.<action>
```

To specify a handler.action pair for a given event, simply specify that handler.action pair between colons in the hist trigger specification.

In theory, any handler can be combined with any action, but in practice, not every handler.action combination is currently supported; if a given handler.action combination isn’t supported, the hist trigger will fail with -EINVAL;

The default ‘handler.action’ if none is explicitly specified is as it always has been, to simply update the set of values associated with an entry. Some applications, however, may want to perform additional actions at that point, such as generate another event, or compare and save a maximum.

The supported handlers and actions are listed below, and each is described in more detail in the following paragraphs, in the context of descriptions of some common and useful handler.action combinations.

The available handlers are:

- onmatch(matching.event) - invoke action on any addition or update
- onmax(var) - invoke action if var exceeds current max
- onchange(var) - invoke action if var changes

The available actions are:

- trace(<synthetic_event_name>,param list) - generate synthetic event
- save(field,...) - save current event fields
- snapshot() - snapshot the trace buffer

The following commonly-used handler.action pairs are available:

- `onmatch(matching.event).trace(<synthetic_event_name>,param list)`

The `'onmatch(matching.event).trace(<synthetic_event_name>,param list)'` hist trigger action is invoked whenever an event matches and the histogram entry would be added or updated. It causes the named synthetic event to be generated with the values given in the `'param list'`. The result is the generation of a synthetic event that consists of the values contained in those variables at the time the invoking event was hit. For example, if the synthetic event name is `'wakeup_latency'`, a `wakeup_latency` event is generated using `onmatch(event).trace(wakeup_latency,arg1,arg2)`.

There is also an equivalent alternative form available for generating synthetic events. In this form, the synthetic event name is used as if it were a function name. For example, using the `'wakeup_latency'` synthetic event name again, the `wakeup_latency` event would be generated by invoking it as if it were a function call, with the event field values passed in as arguments: `onmatch(event).wakeup_latency(arg1,arg2)`. The syntax for this form is:

`onmatch(matching.event).<synthetic_event_name>(param list)`

In either case, the `'param list'` consists of one or more parameters which may be either variables or fields defined on either the `'matching.event'` or the target event. The variables or fields specified in the param list may be either fully-qualified or unqualified. If a variable is specified as unqualified, it must be unique between the two events. A field name used as a param can be unqualified if it refers to the target event, but must be fully qualified if it refers to the matching event. A fully-qualified name is of the form `'system.event_name.$var_name'` or `'system.event_name.field'`.

The `'matching.event'` specification is simply the fully qualified event name of the event that matches the target event for the `onmatch()` functionality, in the form `'system.event_name'`. Histogram keys of both events are compared to find if events match. In case multiple histogram keys are used, they all must match in the specified order.

Finally, the number and type of variables/fields in the `'param list'` must match the number and types of the fields in the synthetic event being generated.

As an example the below defines a simple synthetic event and uses a variable defined on the `sched_wakeup_new` event as a parameter when invoking the synthetic event. Here we define the synthetic event:

```
# echo 'wakeup_new_test pid_t pid' >> \
    /sys/kernel/debug/tracing/synthetic_events

# cat /sys/kernel/debug/tracing/synthetic_events
wakeup_new_test pid_t pid
```

The following hist trigger both defines the missing `testpid` variable and specifies an `onmatch()` action that generates a `wakeup_new_test` synthetic event whenever a `sched_wakeup_new` event occurs, which because of the `'if comm == "cyclicttest"'` filter only happens when the executable is `cyclicttest`:

```
# echo 'hist:keys=$testpid:testpid=pid:onmatch(sched.sched_wakeup_new).\
wakeup_new_test($testpid) if comm=="cyclicttest"' >> \
    /sys/kernel/debug/tracing/events/sched/sched_wakeup_new/trigger
```

Or, equivalently, using the `'trace'` keyword syntax:

```
# echo 'hist:keys=$testpid:testpid=pid:onmatch(sched.sched_wakeup_new).
    trace(wakeup_new_test,$testpid)          if      comm=="cyclictest"      >>
    /sys/kernel/debug/tracing/events/sched/sched_wakeup_new/trigger
```

Creating and displaying a histogram based on those events is now just a matter of using the fields and new synthetic event in the tracing/events/synthetic directory, as usual:

```
# echo 'hist:keys=pid:sort=pid' >> \
    /sys/kernel/debug/tracing/events/synthetic/wakeup_new_test/trigger
```

Running ‘cyclictest’ should cause wakeup_new events to generate wakeup_new_test synthetic events which should result in histogram output in the wakeup_new_test event’s hist file:

```
# cat /sys/kernel/debug/tracing/events/synthetic/wakeup_new_test/hist
```

A more typical usage would be to use two events to calculate a latency. The following example uses a set of hist triggers to produce a ‘wakeup_latency’ histogram.

First, we define a ‘wakeup_latency’ synthetic event:

```
# echo 'wakeup_latency u64 lat; pid_t pid; int prio' >> \
    /sys/kernel/debug/tracing/synthetic_events
```

Next, we specify that whenever we see a sched_waking event for a cyclictest thread, save the timestamp in a ‘ts0’ variable:

```
# echo 'hist:keys=$saved_pid:saved_pid=pid:ts0=common_timestamp.usecs \
    if comm=="cyclictest" >> \
    /sys/kernel/debug/tracing/events/sched/sched_waking/trigger
```

Then, when the corresponding thread is actually scheduled onto the CPU by a sched_switch event (saved_pid matches next_pid), calculate the latency and use that along with another variable and an event field to generate a wakeup_latency synthetic event:

```
# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-$ts0:\
    onmatch(sched.sched_waking).wakeup_latency($wakeup_lat,\
        $saved_pid,next_prio) if next_comm=="cyclictest" >> \
    /sys/kernel/debug/tracing/events/sched/sched_switch/trigger
```

We also need to create a histogram on the wakeup_latency synthetic event in order to aggregate the generated synthetic event data:

```
# echo 'hist:keys=pid,prio,lat:sort=pid,lat' >> \
    /sys/kernel/debug/tracing/events/synthetic/wakeup_latency/trigger
```

Finally, once we’ve run cyclictest to actually generate some events, we can see the output by looking at the wakeup_latency synthetic event’s hist file:

```
# cat /sys/kernel/debug/tracing/events/synthetic/wakeup_latency/hist
```

- onmax(var).save(field,...)

The ‘onmax(var).save(field,...)’ hist trigger action is invoked whenever the value of ‘var’ associated with a histogram entry exceeds the current maximum contained in that variable.

The end result is that the trace event fields specified as the `onmax.save()` params will be saved if 'var' exceeds the current maximum for that hist trigger entry. This allows context from the event that exhibited the new maximum to be saved for later reference. When the histogram is displayed, additional fields displaying the saved values will be printed.

As an example the below defines a couple of hist triggers, one for `sched_waking` and another for `sched_switch`, keyed on pid. Whenever a `sched_waking` occurs, the timestamp is saved in the entry corresponding to the current pid, and when the scheduler switches back to that pid, the timestamp difference is calculated. If the resulting latency, stored in `wakeup_lat`, exceeds the current maximum latency, the values specified in the `save()` fields are recorded:

```
# echo 'hist:keys=pid:ts0=common_timestamp.usecs \
      if comm=="cyclicttest" >> \
      /sys/kernel/debug/tracing/events/sched/sched_waking/trigger

# echo 'hist:keys=next_pid:\
      wakeup_lat=common_timestamp.usecs-$ts0:\
      onmax($wakeup_lat).save(next_comm,prev_pid,prev_prio,prev_comm) \
      if next_comm=="cyclicttest" >> \
      /sys/kernel/debug/tracing/events/sched/sched_switch/trigger
```

When the histogram is displayed, the max value and the saved values corresponding to the max are displayed following the rest of the fields:

```
# cat /sys/kernel/debug/tracing/events/sched/sched_switch/hist
{ next_pid:      2255 } hitcount:      239
  common_timestamp-ts0:      0
  max:      27
  next_comm: cyclicttest
  prev_pid:      0  prev_prio:      120  prev_comm: swapper/1

{ next_pid:      2256 } hitcount:      2355
  common_timestamp-ts0: 0
  max:      49  next_comm: cyclicttest
  prev_pid:      0  prev_prio:      120  prev_comm: swapper/0

Totals:
  Hits: 12970
  Entries: 2
  Dropped: 0
```

- `onmax(var).snapshot()`

The '`onmax(var).snapshot()`' hist trigger action is invoked whenever the value of 'var' associated with a histogram entry exceeds the current maximum contained in that variable.

The end result is that a global snapshot of the trace buffer will be saved in the `tracing/snapshot` file if 'var' exceeds the current maximum for any hist trigger entry.

Note that in this case the maximum is a global maximum for the current trace instance, which is the maximum across all buckets of the histogram. The key of the specific trace event that caused the global maximum and the global maximum itself are displayed, along with a message stating that a snapshot has been taken and where to find it. The user can

use the key information displayed to locate the corresponding bucket in the histogram for even more detail.

As an example the below defines a couple of hist triggers, one for sched_waking and another for sched_switch, keyed on pid. Whenever a sched_waking event occurs, the timestamp is saved in the entry corresponding to the current pid, and when the scheduler switches back to that pid, the timestamp difference is calculated. If the resulting latency, stored in wakeup_lat, exceeds the current maximum latency, a snapshot is taken. As part of the setup, all the scheduler events are also enabled, which are the events that will show up in the snapshot when it is taken at some point:

```
# echo 1 > /sys/kernel/debug/tracing/events/sched/enable

# echo 'hist:keys=pid:ts0=common_timestamp.usecs if comm=="cyclictest"' >>
  /sys/kernel/debug/tracing/events/sched/sched_waking/trigger

# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-$ts0:
  onmax($wakeup_lat).save(next_prio,next_comm,prev_pid,prev_prio,
  prev_comm):onmax($wakeup_lat).snapshot() if next_comm=="cyclictest"' >>
  /sys/kernel/debug/tracing/events/sched/sched_switch/trigger
```

When the histogram is displayed, for each bucket the max value and the saved values corresponding to the max are displayed following the rest of the fields.

If a snapshot was taken, there is also a message indicating that, along with the value and event that triggered the global maximum:

```
# cat /sys/kernel/debug/tracing/events/sched/sched_switch/hist

{ next_pid: 2101 } hitcount: 200 max: 52 next_prio: 120 next_comm: cyclictest
  prev_pid: 0 prev_prio: 120 prev_comm: swapper/6

{ next_pid: 2103 } hitcount: 1326 max: 572 next_prio: 19 next_comm: cyclictest
  prev_pid: 0 prev_prio: 120 prev_comm: swapper/1

{ next_pid: 2102 } hitcount: 1982 max: 74 next_prio: 19 next_comm: cyclictest
  prev_pid: 0 prev_prio: 120 prev_comm: swapper/5
```

Snapshot taken (see tracing/snapshot). Details: triggering value { on-max(\$wakeup_lat) }: 572 triggered by event with key: { next_pid: 2103 }

Totals: Hits: 3508 Entries: 3 Dropped: 0

In the above case, the event that triggered the global maximum has the key with next_pid == 2103. If you look at the bucket that has 2103 as the key, you'll find the additional values save()'d along with the local maximum for that bucket, which should be the same as the global maximum (since that was the same value that triggered the global snapshot).

And finally, looking at the snapshot data should show at or near the end the event that triggered the snapshot (in this case you can verify the timestamps between the sched_waking and sched_switch events, which should match the time displayed in the global maximum):

```
# cat /sys/kernel/debug/tracing/snapshot

<...>-2103 [005] d..3 309.873125: sched_switch: prev_
↪comm=cyclictest prev_pid=2103 prev_prio=19 prev_state=D ==> next_
↪comm=swapper/5 next_pid=0 next_prio=120
<idle>-0 [005] d.h3 309.873611: sched_waking: comm=cyclictest_
↪pid=2102 prio=19 target_cpu=005
```

```

<idle>-0      [005] dNh4    309.873613: sched_wakeup: comm=cyclicttest_
↪pid=2102 prio=19 target_cpu=005
<idle>-0      [005] d..3    309.873616: sched_switch: prev_comm=swapper/
↪5 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=cyclicttest next_
↪pid=2102 next_prio=19
<...>-2102    [005] d..3    309.873625: sched_switch: prev_
↪comm=cyclicttest prev_pid=2102 prev_prio=19 prev_state=D ==> next_
↪comm=swapper/5 next_pid=0 next_prio=120
<idle>-0      [005] d.h3    309.874624: sched_waking: comm=cyclicttest_
↪pid=2102 prio=19 target_cpu=005
<idle>-0      [005] dNh4    309.874626: sched_wakeup: comm=cyclicttest_
↪pid=2102 prio=19 target_cpu=005
<idle>-0      [005] dNh3    309.874628: sched_waking: comm=cyclicttest_
↪pid=2103 prio=19 target_cpu=005
<idle>-0      [005] dNh4    309.874630: sched_wakeup: comm=cyclicttest_
↪pid=2103 prio=19 target_cpu=005
<idle>-0      [005] d..3    309.874633: sched_switch: prev_comm=swapper/
↪5 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=cyclicttest next_
↪pid=2102 next_prio=19
<idle>-0      [004] d.h3    309.874757: sched_waking: comm=gnome-
↪terminal- pid=1699 prio=120 target_cpu=004
<idle>-0      [004] dNh4    309.874762: sched_wakeup: comm=gnome-
↪terminal- pid=1699 prio=120 target_cpu=004
<idle>-0      [004] d..3    309.874766: sched_switch: prev_comm=swapper/
↪4 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=gnome-terminal-
↪next_pid=1699 next_prio=120
gnome-terminal--1699 [004] d.h2    309.874941: sched_stat_runtime:
↪comm=gnome-terminal- pid=1699 runtime=180706 [ns] vruntime=1126870572_
↪[ns]
<idle>-0      [003] d.s4    309.874956: sched_waking: comm=rcu_sched_
↪pid=9 prio=120 target_cpu=007
<idle>-0      [003] d.s5    309.874960: sched_wake_idle_without_ipi:
↪cpu=7
<idle>-0      [003] d.s5    309.874961: sched_wakeup: comm=rcu_sched_
↪pid=9 prio=120 target_cpu=007
<idle>-0      [007] d..3    309.874963: sched_switch: prev_comm=swapper/
↪7 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=rcu_sched next_
↪pid=9 next_prio=120
rcu_sched-9    [007] d..3    309.874973: sched_stat_runtime: comm=rcu_
↪sched pid=9 runtime=13646 [ns] vruntime=22531430286 [ns]
rcu_sched-9    [007] d..3    309.874978: sched_switch: prev_comm=rcu_
↪sched prev_pid=9 prev_prio=120 prev_state=R+ ==> next_comm=swapper/7_
↪next_pid=0 next_prio=120
<...>-2102    [005] d..4    309.874994: sched_migrate_task:
↪comm=cyclicttest pid=2103 prio=19 orig_cpu=5 dest_cpu=1
<...>-2102    [005] d..4    309.875185: sched_wake_idle_without_ipi:
↪cpu=1
<idle>-0      [001] d..3    309.875200: sched_switch: prev_comm=swapper/
↪1 prev_pid=0 prev_prio=120 prev_state=S ==> next_comm=cyclicttest next_
↪pid=2103 next_prio=19

```

- `onchange(var).save(field,...)`

The `'onchange(var).save(field,...)'` hist trigger action is invoked whenever the value of `'var'` associated with a histogram entry changes.

The end result is that the trace event fields specified as the `onchange.save()` params will be saved if `'var'` changes for that hist trigger entry. This allows context from the event that changed the value to be saved for later reference. When the histogram is displayed, additional fields displaying the saved values will be printed.

- `onchange(var).snapshot()`

The `'onchange(var).snapshot()'` hist trigger action is invoked whenever the value of `'var'` associated with a histogram entry changes.

The end result is that a global snapshot of the trace buffer will be saved in the `tracing/snapshot` file if `'var'` changes for any hist trigger entry.

Note that in this case the changed value is a global variable associated with current trace instance. The key of the specific trace event that caused the value to change and the global value itself are displayed, along with a message stating that a snapshot has been taken and where to find it. The user can use the key information displayed to locate the corresponding bucket in the histogram for even more detail.

As an example the below defines a hist trigger on the `tcp_probe` event, keyed on `dport`. Whenever a `tcp_probe` event occurs, the `cwnd` field is checked against the current value stored in the `$cwnd` variable. If the value has changed, a snapshot is taken. As part of the setup, all the scheduler and tcp events are also enabled, which are the events that will show up in the snapshot when it is taken at some point:

```
# echo 1 > /sys/kernel/debug/tracing/events/sched/enable # echo 1 >
/sys/kernel/debug/tracing/events/tcp/enable
```

```
# echo 'hist:keys=dport:cwnd=snd_cwnd: onchange($cwnd).save(snd_wnd,srtt,rcv_wnd):
onchange($cwnd).snapshot()' >> /sys/kernel/debug/tracing/events/tcp/tcp_probe/trigger
```

When the histogram is displayed, for each bucket the tracked value and the saved values corresponding to that value are displayed following the rest of the fields.

If a snapshot was taken, there is also a message indicating that, along with the value and event that triggered the snapshot:

```
# cat /sys/kernel/debug/tracing/events/tcp/tcp_probe/hist

{ dport:      1521 } hitcount:      8
  changed:      10  snd_wnd:      35456  srtt:      154262  rcv_wnd:      42112
  ↳ 42112

{ dport:      80 } hitcount:      23
  changed:      10  snd_wnd:      28960  srtt:      19604  rcv_wnd:      29312
  ↳ 29312

{ dport:     9001 } hitcount:     172
  changed:      10  snd_wnd:     48384  srtt:     260444  rcv_wnd:     55168
  ↳ 55168
```

```
{ dport:      443 } hitcount:      211
  changed:      10  snd_wnd:      26960  srtt:      17379  rcv_wnd:
  ↪ 28800
```

Snapshot taken (see tracing/snapshot). Details:

```
triggering value { onchange($cwnd) }:      10
triggered by event with key: { dport:      80 }
```

Totals:

```
Hits: 414
Entries: 4
Dropped: 0
```

In the above case, the event that triggered the snapshot has the key with `dport == 80`. If you look at the bucket that has 80 as the key, you'll find the additional values `save()`'d along with the changed value for that bucket, which should be the same as the global changed value (since that was the same value that triggered the global snapshot).

And finally, looking at the snapshot data should show at or near the end the event that triggered the snapshot:

```
# cat /sys/kernel/debug/tracing/snapshot
```

```
gnome-shell-1261 [006] dN.3 49.823113: sched_stat_runtime:
↪ comm=gnome-shell pid=1261 runtime=49347 [ns] vruntime=1835730389 [ns]
kworker/u16:4-773 [003] d..3 49.823114: sched_switch: prev_
↪ comm=kworker/u16:4 prev_pid=773 prev_prio=120 prev_state=R+ ==> next_
↪ comm=kworker/3:2 next_pid=135 next_prio=120
gnome-shell-1261 [006] d..3 49.823114: sched_switch: prev_
↪ comm=gnome-shell prev_pid=1261 prev_prio=120 prev_state=R+ ==> next_
↪ comm=kworker/6:2 next_pid=387 next_prio=120
kworker/3:2-135 [003] d..3 49.823118: sched_stat_runtime:
↪ comm=kworker/3:2 pid=135 runtime=5339 [ns] vruntime=17815800388 [ns]
kworker/6:2-387 [006] d..3 49.823120: sched_stat_runtime:
↪ comm=kworker/6:2 pid=387 runtime=9594 [ns] vruntime=14589605367 [ns]
kworker/6:2-387 [006] d..3 49.823122: sched_switch: prev_
↪ comm=kworker/6:2 prev_pid=387 prev_prio=120 prev_state=R+ ==> next_
↪ comm=gnome-shell next_pid=1261 next_prio=120
kworker/3:2-135 [003] d..3 49.823123: sched_switch: prev_
↪ comm=kworker/3:2 prev_pid=135 prev_prio=120 prev_state=T ==> next_
↪ comm=swapper/3 next_pid=0 next_prio=120
<idle>-0 [004] ..s7 49.823798: tcp_probe: src=10.0.0.
↪ 10:54326 dest=23.215.104.193:80 mark=0x0 length=32 snd_nxt=0xe3ae2ff5
↪ snd_una=0xe3ae2ecd snd_cwnd=10 ssthresh=2147483647 snd_wnd=28960
↪ srtt=19604 rcv_wnd=29312
```

16.2.8 3. User space creating a trigger

Writing into `/sys/kernel/tracing/trace_marker` writes into the ftrace ring buffer. This can also act like an event, by writing into the trigger file located in `/sys/kernel/tracing/events/ftrace/print/`

Modifying `cyclictest` to write into the `trace_marker` file before it sleeps and after it wakes up, something like this:

```
static void traceputs(char *str)
{
    /* tracemark_fd is the trace_marker file descriptor */
    if (tracemark_fd < 0)
        return;
    /* write the tracemark message */
    write(tracemark_fd, str, strlen(str));
}
```

And later add something like:

```
traceputs("start");
clock_nanosleep(...);
traceputs("end");
```

We can make a histogram from this:

```
# cd /sys/kernel/tracing
# echo 'latency u64 lat' > synthetic_events
# echo 'hist:keys=common_pid:ts0=common_timestamp.usecs if buf == "start"' >
→events/ftrace/print/trigger
# echo 'hist:keys=common_pid:lat=common_timestamp.usecs-$ts0:onmatch(ftrace.
→print).latency($lat) if buf == "end"' >> events/ftrace/print/trigger
# echo 'hist:keys=lat,common_pid:sort=lat' > events/synthetic/latency/trigger
```

The above created a synthetic event called “latency” and two histograms against the `trace_marker`, one gets triggered when “start” is written into the `trace_marker` file and the other when “end” is written. If the pids match, then it will call the “latency” synthetic event with the calculated latency as its parameter. Finally, a histogram is added to the latency synthetic event to record the calculated latency along with the pid.

Now running `cyclictest` with:

```
# ./cyclictest -p80 -d0 -i250 -n -a -t --tracemark -b 1000

-p80   : run threads at priority 80
-d0    : have all threads run at the same interval
-i250  : start the interval at 250 microseconds (all threads will do this)
-n     : sleep with nanosleep
-a     : affine all threads to a separate CPU
-t     : one thread per available CPU
--tracemark : enable trace mark writing
-b 1000 : stop if any latency is greater than 1000 microseconds
```

Note, the `-b 1000` is used just to make `--tracemark` available.

Then we can see the histogram created by this with:

```
# cat events/synthetic/latency/hist
# event histogram
#
# trigger info: hist:keys=lat,common_pid:vals=hitcount:sort=lat:size=2048_
→[active]
#
{ lat:      107, common_pid:    2039 } hitcount:      1
{ lat:      122, common_pid:    2041 } hitcount:      1
{ lat:      166, common_pid:    2039 } hitcount:      1
{ lat:      174, common_pid:    2039 } hitcount:      1
{ lat:      194, common_pid:    2041 } hitcount:      1
{ lat:      196, common_pid:    2036 } hitcount:      1
{ lat:      197, common_pid:    2038 } hitcount:      1
{ lat:      198, common_pid:    2039 } hitcount:      1
{ lat:      199, common_pid:    2039 } hitcount:      1
{ lat:      200, common_pid:    2041 } hitcount:      1
{ lat:      201, common_pid:    2039 } hitcount:      2
{ lat:      202, common_pid:    2038 } hitcount:      1
{ lat:      202, common_pid:    2043 } hitcount:      1
{ lat:      203, common_pid:    2039 } hitcount:      1
{ lat:      203, common_pid:    2036 } hitcount:      1
{ lat:      203, common_pid:    2041 } hitcount:      1
{ lat:      206, common_pid:    2038 } hitcount:      2
{ lat:      207, common_pid:    2039 } hitcount:      1
{ lat:      207, common_pid:    2036 } hitcount:      1
{ lat:      208, common_pid:    2040 } hitcount:      1
{ lat:      209, common_pid:    2043 } hitcount:      1
{ lat:      210, common_pid:    2039 } hitcount:      1
{ lat:      211, common_pid:    2039 } hitcount:      4
{ lat:      212, common_pid:    2043 } hitcount:      1
{ lat:      212, common_pid:    2039 } hitcount:      2
{ lat:      213, common_pid:    2039 } hitcount:      1
{ lat:      214, common_pid:    2038 } hitcount:      1
{ lat:      214, common_pid:    2039 } hitcount:      2
{ lat:      214, common_pid:    2042 } hitcount:      1
{ lat:      215, common_pid:    2039 } hitcount:      1
{ lat:      217, common_pid:    2036 } hitcount:      1
{ lat:      217, common_pid:    2040 } hitcount:      1
{ lat:      217, common_pid:    2039 } hitcount:      1
{ lat:      218, common_pid:    2039 } hitcount:      6
{ lat:      219, common_pid:    2039 } hitcount:      9
{ lat:      220, common_pid:    2039 } hitcount:     11
{ lat:      221, common_pid:    2039 } hitcount:      5
{ lat:      221, common_pid:    2042 } hitcount:      1
{ lat:      222, common_pid:    2039 } hitcount:      7
{ lat:      223, common_pid:    2036 } hitcount:      1
{ lat:      223, common_pid:    2039 } hitcount:      3
{ lat:      224, common_pid:    2039 } hitcount:      4
```

```

{ lat:      224, common_pid:    2037 } hitcount:      1
{ lat:      224, common_pid:    2036 } hitcount:      2
{ lat:      225, common_pid:    2039 } hitcount:      5
{ lat:      225, common_pid:    2042 } hitcount:      1
{ lat:      226, common_pid:    2039 } hitcount:      7
{ lat:      226, common_pid:    2036 } hitcount:      4
{ lat:      227, common_pid:    2039 } hitcount:      6
{ lat:      227, common_pid:    2036 } hitcount:     12
{ lat:      227, common_pid:    2043 } hitcount:      1
{ lat:      228, common_pid:    2039 } hitcount:      7
{ lat:      228, common_pid:    2036 } hitcount:     14
{ lat:      229, common_pid:    2039 } hitcount:      9
{ lat:      229, common_pid:    2036 } hitcount:      8
{ lat:      229, common_pid:    2038 } hitcount:      1
{ lat:      230, common_pid:    2039 } hitcount:     11
{ lat:      230, common_pid:    2036 } hitcount:      6
{ lat:      230, common_pid:    2043 } hitcount:      1
{ lat:      230, common_pid:    2042 } hitcount:      2
{ lat:      231, common_pid:    2041 } hitcount:      1
{ lat:      231, common_pid:    2036 } hitcount:      6
{ lat:      231, common_pid:    2043 } hitcount:      1
{ lat:      231, common_pid:    2039 } hitcount:      8
{ lat:      232, common_pid:    2037 } hitcount:      1
{ lat:      232, common_pid:    2039 } hitcount:      6
{ lat:      232, common_pid:    2040 } hitcount:      2
{ lat:      232, common_pid:    2036 } hitcount:      5
{ lat:      232, common_pid:    2043 } hitcount:      1
{ lat:      233, common_pid:    2036 } hitcount:      5
{ lat:      233, common_pid:    2039 } hitcount:     11
{ lat:      234, common_pid:    2039 } hitcount:      4
{ lat:      234, common_pid:    2038 } hitcount:      2
{ lat:      234, common_pid:    2043 } hitcount:      2
{ lat:      234, common_pid:    2036 } hitcount:     11
{ lat:      234, common_pid:    2040 } hitcount:      1
{ lat:      235, common_pid:    2037 } hitcount:      2
{ lat:      235, common_pid:    2036 } hitcount:      8
{ lat:      235, common_pid:    2043 } hitcount:      2
{ lat:      235, common_pid:    2039 } hitcount:      5
{ lat:      235, common_pid:    2042 } hitcount:      2
{ lat:      235, common_pid:    2040 } hitcount:      4
{ lat:      235, common_pid:    2041 } hitcount:      1
{ lat:      236, common_pid:    2036 } hitcount:      7
{ lat:      236, common_pid:    2037 } hitcount:      1
{ lat:      236, common_pid:    2041 } hitcount:      5
{ lat:      236, common_pid:    2039 } hitcount:      3
{ lat:      236, common_pid:    2043 } hitcount:      9
{ lat:      236, common_pid:    2040 } hitcount:      7
{ lat:      237, common_pid:    2037 } hitcount:      1
{ lat:      237, common_pid:    2040 } hitcount:      1
{ lat:      237, common_pid:    2036 } hitcount:      9

```


{ lat:	237, common_pid:	2039 }	hitcount:	3
{ lat:	237, common_pid:	2043 }	hitcount:	8
{ lat:	237, common_pid:	2042 }	hitcount:	2
{ lat:	237, common_pid:	2041 }	hitcount:	2
{ lat:	238, common_pid:	2043 }	hitcount:	10
{ lat:	238, common_pid:	2040 }	hitcount:	1
{ lat:	238, common_pid:	2037 }	hitcount:	9
{ lat:	238, common_pid:	2038 }	hitcount:	1
{ lat:	238, common_pid:	2039 }	hitcount:	1
{ lat:	238, common_pid:	2042 }	hitcount:	3
{ lat:	238, common_pid:	2036 }	hitcount:	7
{ lat:	239, common_pid:	2041 }	hitcount:	1
{ lat:	239, common_pid:	2043 }	hitcount:	11
{ lat:	239, common_pid:	2037 }	hitcount:	11
{ lat:	239, common_pid:	2038 }	hitcount:	6
{ lat:	239, common_pid:	2036 }	hitcount:	7
{ lat:	239, common_pid:	2040 }	hitcount:	1
{ lat:	239, common_pid:	2042 }	hitcount:	9
{ lat:	240, common_pid:	2037 }	hitcount:	29
{ lat:	240, common_pid:	2043 }	hitcount:	15
{ lat:	240, common_pid:	2040 }	hitcount:	44
{ lat:	240, common_pid:	2039 }	hitcount:	1
{ lat:	240, common_pid:	2041 }	hitcount:	2
{ lat:	240, common_pid:	2038 }	hitcount:	1
{ lat:	240, common_pid:	2036 }	hitcount:	10
{ lat:	240, common_pid:	2042 }	hitcount:	13
{ lat:	241, common_pid:	2036 }	hitcount:	21
{ lat:	241, common_pid:	2041 }	hitcount:	36
{ lat:	241, common_pid:	2037 }	hitcount:	34
{ lat:	241, common_pid:	2042 }	hitcount:	14
{ lat:	241, common_pid:	2040 }	hitcount:	94
{ lat:	241, common_pid:	2039 }	hitcount:	12
{ lat:	241, common_pid:	2038 }	hitcount:	2
{ lat:	241, common_pid:	2043 }	hitcount:	28
{ lat:	242, common_pid:	2040 }	hitcount:	109
{ lat:	242, common_pid:	2041 }	hitcount:	506
{ lat:	242, common_pid:	2039 }	hitcount:	155
{ lat:	242, common_pid:	2042 }	hitcount:	21
{ lat:	242, common_pid:	2037 }	hitcount:	52
{ lat:	242, common_pid:	2043 }	hitcount:	21
{ lat:	242, common_pid:	2036 }	hitcount:	16
{ lat:	242, common_pid:	2038 }	hitcount:	156
{ lat:	243, common_pid:	2037 }	hitcount:	46
{ lat:	243, common_pid:	2039 }	hitcount:	40
{ lat:	243, common_pid:	2042 }	hitcount:	119
{ lat:	243, common_pid:	2041 }	hitcount:	611
{ lat:	243, common_pid:	2036 }	hitcount:	69
{ lat:	243, common_pid:	2038 }	hitcount:	784
{ lat:	243, common_pid:	2040 }	hitcount:	323
{ lat:	243, common_pid:	2043 }	hitcount:	14


```

{ lat:      244, common_pid:    2043 } hitcount:      35
{ lat:      244, common_pid:    2042 } hitcount:     305
{ lat:      244, common_pid:    2039 } hitcount:       8
{ lat:      244, common_pid:    2040 } hitcount:    4515
{ lat:      244, common_pid:    2038 } hitcount:     371
{ lat:      244, common_pid:    2037 } hitcount:      31
{ lat:      244, common_pid:    2036 } hitcount:     114
{ lat:      244, common_pid:    2041 } hitcount:    3396
{ lat:      245, common_pid:    2036 } hitcount:     700
{ lat:      245, common_pid:    2041 } hitcount:    2772
{ lat:      245, common_pid:    2037 } hitcount:     268
{ lat:      245, common_pid:    2039 } hitcount:     472
{ lat:      245, common_pid:    2038 } hitcount:    2758
{ lat:      245, common_pid:    2042 } hitcount:    3833
{ lat:      245, common_pid:    2040 } hitcount:    3105
{ lat:      245, common_pid:    2043 } hitcount:     645
{ lat:      246, common_pid:    2038 } hitcount:    3451
{ lat:      246, common_pid:    2041 } hitcount:     142
{ lat:      246, common_pid:    2037 } hitcount:    5101
{ lat:      246, common_pid:    2040 } hitcount:      68
{ lat:      246, common_pid:    2043 } hitcount:    5099
{ lat:      246, common_pid:    2039 } hitcount:    5608
{ lat:      246, common_pid:    2042 } hitcount:    3723
{ lat:      246, common_pid:    2036 } hitcount:    4738
{ lat:      247, common_pid:    2042 } hitcount:     312
{ lat:      247, common_pid:    2043 } hitcount:    2385
{ lat:      247, common_pid:    2041 } hitcount:     452
{ lat:      247, common_pid:    2038 } hitcount:     792
{ lat:      247, common_pid:    2040 } hitcount:      78
{ lat:      247, common_pid:    2036 } hitcount:    2375
{ lat:      247, common_pid:    2039 } hitcount:    1834
{ lat:      247, common_pid:    2037 } hitcount:    2655
{ lat:      248, common_pid:    2037 } hitcount:      36
{ lat:      248, common_pid:    2042 } hitcount:      11
{ lat:      248, common_pid:    2038 } hitcount:     122
{ lat:      248, common_pid:    2036 } hitcount:     135
{ lat:      248, common_pid:    2039 } hitcount:      26
{ lat:      248, common_pid:    2041 } hitcount:     503
{ lat:      248, common_pid:    2043 } hitcount:      66
{ lat:      248, common_pid:    2040 } hitcount:      46
{ lat:      249, common_pid:    2037 } hitcount:      29
{ lat:      249, common_pid:    2038 } hitcount:       1
{ lat:      249, common_pid:    2043 } hitcount:      29
{ lat:      249, common_pid:    2039 } hitcount:       8
{ lat:      249, common_pid:    2042 } hitcount:     56
{ lat:      249, common_pid:    2040 } hitcount:     27
{ lat:      249, common_pid:    2041 } hitcount:     11
{ lat:      249, common_pid:    2036 } hitcount:     27
{ lat:      250, common_pid:    2038 } hitcount:       1
{ lat:      250, common_pid:    2036 } hitcount:     30

```

{ lat:	250, common_pid:	2040 }	hitcount:	19
{ lat:	250, common_pid:	2043 }	hitcount:	22
{ lat:	250, common_pid:	2042 }	hitcount:	20
{ lat:	250, common_pid:	2041 }	hitcount:	1
{ lat:	250, common_pid:	2039 }	hitcount:	6
{ lat:	250, common_pid:	2037 }	hitcount:	48
{ lat:	251, common_pid:	2037 }	hitcount:	43
{ lat:	251, common_pid:	2039 }	hitcount:	1
{ lat:	251, common_pid:	2036 }	hitcount:	12
{ lat:	251, common_pid:	2042 }	hitcount:	2
{ lat:	251, common_pid:	2041 }	hitcount:	1
{ lat:	251, common_pid:	2043 }	hitcount:	15
{ lat:	251, common_pid:	2040 }	hitcount:	3
{ lat:	252, common_pid:	2040 }	hitcount:	1
{ lat:	252, common_pid:	2036 }	hitcount:	12
{ lat:	252, common_pid:	2037 }	hitcount:	21
{ lat:	252, common_pid:	2043 }	hitcount:	14
{ lat:	253, common_pid:	2037 }	hitcount:	21
{ lat:	253, common_pid:	2039 }	hitcount:	2
{ lat:	253, common_pid:	2036 }	hitcount:	9
{ lat:	253, common_pid:	2043 }	hitcount:	6
{ lat:	253, common_pid:	2040 }	hitcount:	1
{ lat:	254, common_pid:	2036 }	hitcount:	8
{ lat:	254, common_pid:	2043 }	hitcount:	3
{ lat:	254, common_pid:	2041 }	hitcount:	1
{ lat:	254, common_pid:	2042 }	hitcount:	1
{ lat:	254, common_pid:	2039 }	hitcount:	1
{ lat:	254, common_pid:	2037 }	hitcount:	12
{ lat:	255, common_pid:	2043 }	hitcount:	1
{ lat:	255, common_pid:	2037 }	hitcount:	2
{ lat:	255, common_pid:	2036 }	hitcount:	2
{ lat:	255, common_pid:	2039 }	hitcount:	8
{ lat:	256, common_pid:	2043 }	hitcount:	1
{ lat:	256, common_pid:	2036 }	hitcount:	4
{ lat:	256, common_pid:	2039 }	hitcount:	6
{ lat:	257, common_pid:	2039 }	hitcount:	5
{ lat:	257, common_pid:	2036 }	hitcount:	4
{ lat:	258, common_pid:	2039 }	hitcount:	5
{ lat:	258, common_pid:	2036 }	hitcount:	2
{ lat:	259, common_pid:	2036 }	hitcount:	7
{ lat:	259, common_pid:	2039 }	hitcount:	7
{ lat:	260, common_pid:	2036 }	hitcount:	8
{ lat:	260, common_pid:	2039 }	hitcount:	6
{ lat:	261, common_pid:	2036 }	hitcount:	5
{ lat:	261, common_pid:	2039 }	hitcount:	7
{ lat:	262, common_pid:	2039 }	hitcount:	5
{ lat:	262, common_pid:	2036 }	hitcount:	5
{ lat:	263, common_pid:	2039 }	hitcount:	7
{ lat:	263, common_pid:	2036 }	hitcount:	7
{ lat:	264, common_pid:	2039 }	hitcount:	9

```

{ lat:      264, common_pid:    2036 } hitcount:      9
{ lat:      265, common_pid:    2036 } hitcount:      5
{ lat:      265, common_pid:    2039 } hitcount:      1
{ lat:      266, common_pid:    2036 } hitcount:      1
{ lat:      266, common_pid:    2039 } hitcount:      3
{ lat:      267, common_pid:    2036 } hitcount:      1
{ lat:      267, common_pid:    2039 } hitcount:      3
{ lat:      268, common_pid:    2036 } hitcount:      1
{ lat:      268, common_pid:    2039 } hitcount:      6
{ lat:      269, common_pid:    2036 } hitcount:      1
{ lat:      269, common_pid:    2043 } hitcount:      1
{ lat:      269, common_pid:    2039 } hitcount:      2
{ lat:      270, common_pid:    2040 } hitcount:      1
{ lat:      270, common_pid:    2039 } hitcount:      6
{ lat:      271, common_pid:    2041 } hitcount:      1
{ lat:      271, common_pid:    2039 } hitcount:      5
{ lat:      272, common_pid:    2039 } hitcount:     10
{ lat:      273, common_pid:    2039 } hitcount:      8
{ lat:      274, common_pid:    2039 } hitcount:      2
{ lat:      275, common_pid:    2039 } hitcount:      1
{ lat:      276, common_pid:    2039 } hitcount:      2
{ lat:      276, common_pid:    2037 } hitcount:      1
{ lat:      276, common_pid:    2038 } hitcount:      1
{ lat:      277, common_pid:    2039 } hitcount:      1
{ lat:      277, common_pid:    2042 } hitcount:      1
{ lat:      278, common_pid:    2039 } hitcount:      1
{ lat:      279, common_pid:    2039 } hitcount:      4
{ lat:      279, common_pid:    2043 } hitcount:      1
{ lat:      280, common_pid:    2039 } hitcount:      3
{ lat:      283, common_pid:    2036 } hitcount:      2
{ lat:      284, common_pid:    2039 } hitcount:      1
{ lat:      284, common_pid:    2043 } hitcount:      1
{ lat:      288, common_pid:    2039 } hitcount:      1
{ lat:      289, common_pid:    2039 } hitcount:      1
{ lat:      300, common_pid:    2039 } hitcount:      1
{ lat:      384, common_pid:    2039 } hitcount:      1

```

Totals:

```

  Hits: 67625
  Entries: 278
  Dropped: 0

```

Note, the writes are around the sleep, so ideally they will all be of 250 microseconds. If you are wondering how there are several that are under 250 microseconds, that is because the way `cyclictest` works, is if one iteration comes in late, the next one will set the timer to wake up less than 250. That is, if an iteration came in 50 microseconds late, the next wake up will be at 200 microseconds.

But this could easily be done in userspace. To make this even more interesting, we can mix the histogram between events that happened in the kernel with `trace_marker`:

```
# cd /sys/kernel/tracing
# echo 'latency u64 lat' > synthetic_events
# echo 'hist:keys=pid:ts0=common_timestamp.usecs' > events/sched/sched_waking/
→ trigger
# echo 'hist:keys=common_pid:lat=common_timestamp.usecs-$ts0:onmatch(sched.
→ sched_waking).latency($lat) if buf == "end"' > events/ftrace/print/trigger
# echo 'hist:keys=lat,common_pid:sort=lat' > events/synthetic/latency/trigger
```

The difference this time is that instead of using the `trace_marker` to start the latency, the `sched_waking` event is used, matching the `common_pid` for the `trace_marker` write with the pid that is being woken by `sched_waking`.

After running `cyclicttest` again with the same parameters, we now have:

```
# cat events/synthetic/latency/hist
# event histogram
#
# trigger info: hist:keys=lat,common_pid:vals=hitcount:sort=lat:size=2048,
→ [active]
#
```

{ lat:	7,	common_pid:	2302 }	hitcount:	640
{ lat:	7,	common_pid:	2299 }	hitcount:	42
{ lat:	7,	common_pid:	2303 }	hitcount:	18
{ lat:	7,	common_pid:	2305 }	hitcount:	166
{ lat:	7,	common_pid:	2306 }	hitcount:	1
{ lat:	7,	common_pid:	2301 }	hitcount:	91
{ lat:	7,	common_pid:	2300 }	hitcount:	17
{ lat:	8,	common_pid:	2303 }	hitcount:	8296
{ lat:	8,	common_pid:	2304 }	hitcount:	6864
{ lat:	8,	common_pid:	2305 }	hitcount:	9464
{ lat:	8,	common_pid:	2301 }	hitcount:	9213
{ lat:	8,	common_pid:	2306 }	hitcount:	6246
{ lat:	8,	common_pid:	2302 }	hitcount:	8797
{ lat:	8,	common_pid:	2299 }	hitcount:	8771
{ lat:	8,	common_pid:	2300 }	hitcount:	8119
{ lat:	9,	common_pid:	2305 }	hitcount:	1519
{ lat:	9,	common_pid:	2299 }	hitcount:	2346
{ lat:	9,	common_pid:	2303 }	hitcount:	2841
{ lat:	9,	common_pid:	2301 }	hitcount:	1846
{ lat:	9,	common_pid:	2304 }	hitcount:	3861
{ lat:	9,	common_pid:	2302 }	hitcount:	1210
{ lat:	9,	common_pid:	2300 }	hitcount:	2762
{ lat:	9,	common_pid:	2306 }	hitcount:	4247
{ lat:	10,	common_pid:	2299 }	hitcount:	16
{ lat:	10,	common_pid:	2306 }	hitcount:	333
{ lat:	10,	common_pid:	2303 }	hitcount:	16
{ lat:	10,	common_pid:	2304 }	hitcount:	168
{ lat:	10,	common_pid:	2302 }	hitcount:	240
{ lat:	10,	common_pid:	2301 }	hitcount:	28
{ lat:	10,	common_pid:	2300 }	hitcount:	95

```

{ lat:      10, common_pid: 2305 } hitcount:      18
{ lat:      11, common_pid: 2303 } hitcount:         5
{ lat:      11, common_pid: 2305 } hitcount:         8
{ lat:      11, common_pid: 2306 } hitcount:      221
{ lat:      11, common_pid: 2302 } hitcount:       76
{ lat:      11, common_pid: 2304 } hitcount:       26
{ lat:      11, common_pid: 2300 } hitcount:     125
{ lat:      11, common_pid: 2299 } hitcount:         2
{ lat:      12, common_pid: 2305 } hitcount:         3
{ lat:      12, common_pid: 2300 } hitcount:         6
{ lat:      12, common_pid: 2306 } hitcount:       90
{ lat:      12, common_pid: 2302 } hitcount:         4
{ lat:      12, common_pid: 2303 } hitcount:         1
{ lat:      12, common_pid: 2304 } hitcount:     122
{ lat:      13, common_pid: 2300 } hitcount:         12
{ lat:      13, common_pid: 2301 } hitcount:          1
{ lat:      13, common_pid: 2306 } hitcount:       32
{ lat:      13, common_pid: 2302 } hitcount:         5
{ lat:      13, common_pid: 2305 } hitcount:          1
{ lat:      13, common_pid: 2303 } hitcount:          1
{ lat:      13, common_pid: 2304 } hitcount:       61
{ lat:      14, common_pid: 2303 } hitcount:          4
{ lat:      14, common_pid: 2306 } hitcount:          5
{ lat:      14, common_pid: 2305 } hitcount:          4
{ lat:      14, common_pid: 2304 } hitcount:       62
{ lat:      14, common_pid: 2302 } hitcount:       19
{ lat:      14, common_pid: 2300 } hitcount:       33
{ lat:      14, common_pid: 2299 } hitcount:          1
{ lat:      14, common_pid: 2301 } hitcount:          4
{ lat:      15, common_pid: 2305 } hitcount:          1
{ lat:      15, common_pid: 2302 } hitcount:       25
{ lat:      15, common_pid: 2300 } hitcount:       11
{ lat:      15, common_pid: 2299 } hitcount:         5
{ lat:      15, common_pid: 2301 } hitcount:          1
{ lat:      15, common_pid: 2304 } hitcount:          8
{ lat:      15, common_pid: 2303 } hitcount:          1
{ lat:      15, common_pid: 2306 } hitcount:          6
{ lat:      16, common_pid: 2302 } hitcount:       31
{ lat:      16, common_pid: 2306 } hitcount:          3
{ lat:      16, common_pid: 2300 } hitcount:          5
{ lat:      17, common_pid: 2302 } hitcount:          6
{ lat:      17, common_pid: 2303 } hitcount:          1
{ lat:      18, common_pid: 2304 } hitcount:          1
{ lat:      18, common_pid: 2302 } hitcount:          8
{ lat:      18, common_pid: 2299 } hitcount:          1
{ lat:      18, common_pid: 2301 } hitcount:          1
{ lat:      19, common_pid: 2303 } hitcount:          4
{ lat:      19, common_pid: 2304 } hitcount:          5
{ lat:      19, common_pid: 2302 } hitcount:          4
{ lat:      19, common_pid: 2299 } hitcount:          3

```

{ lat:	19, common_pid:	2306 }	hitcount:	1
{ lat:	19, common_pid:	2300 }	hitcount:	4
{ lat:	19, common_pid:	2305 }	hitcount:	5
{ lat:	20, common_pid:	2299 }	hitcount:	2
{ lat:	20, common_pid:	2302 }	hitcount:	3
{ lat:	20, common_pid:	2305 }	hitcount:	1
{ lat:	20, common_pid:	2300 }	hitcount:	2
{ lat:	20, common_pid:	2301 }	hitcount:	2
{ lat:	20, common_pid:	2303 }	hitcount:	3
{ lat:	21, common_pid:	2305 }	hitcount:	1
{ lat:	21, common_pid:	2299 }	hitcount:	5
{ lat:	21, common_pid:	2303 }	hitcount:	4
{ lat:	21, common_pid:	2302 }	hitcount:	7
{ lat:	21, common_pid:	2300 }	hitcount:	1
{ lat:	21, common_pid:	2301 }	hitcount:	5
{ lat:	21, common_pid:	2304 }	hitcount:	2
{ lat:	22, common_pid:	2302 }	hitcount:	5
{ lat:	22, common_pid:	2303 }	hitcount:	1
{ lat:	22, common_pid:	2306 }	hitcount:	3
{ lat:	22, common_pid:	2301 }	hitcount:	2
{ lat:	22, common_pid:	2300 }	hitcount:	1
{ lat:	22, common_pid:	2299 }	hitcount:	1
{ lat:	22, common_pid:	2305 }	hitcount:	1
{ lat:	22, common_pid:	2304 }	hitcount:	1
{ lat:	23, common_pid:	2299 }	hitcount:	1
{ lat:	23, common_pid:	2306 }	hitcount:	2
{ lat:	23, common_pid:	2302 }	hitcount:	6
{ lat:	24, common_pid:	2302 }	hitcount:	3
{ lat:	24, common_pid:	2300 }	hitcount:	1
{ lat:	24, common_pid:	2306 }	hitcount:	2
{ lat:	24, common_pid:	2305 }	hitcount:	1
{ lat:	24, common_pid:	2299 }	hitcount:	1
{ lat:	25, common_pid:	2300 }	hitcount:	1
{ lat:	25, common_pid:	2302 }	hitcount:	4
{ lat:	26, common_pid:	2302 }	hitcount:	2
{ lat:	27, common_pid:	2305 }	hitcount:	1
{ lat:	27, common_pid:	2300 }	hitcount:	1
{ lat:	27, common_pid:	2302 }	hitcount:	3
{ lat:	28, common_pid:	2306 }	hitcount:	1
{ lat:	28, common_pid:	2302 }	hitcount:	4
{ lat:	29, common_pid:	2302 }	hitcount:	1
{ lat:	29, common_pid:	2300 }	hitcount:	2
{ lat:	29, common_pid:	2306 }	hitcount:	1
{ lat:	29, common_pid:	2304 }	hitcount:	1
{ lat:	30, common_pid:	2302 }	hitcount:	4
{ lat:	31, common_pid:	2302 }	hitcount:	6
{ lat:	32, common_pid:	2302 }	hitcount:	1
{ lat:	33, common_pid:	2299 }	hitcount:	1
{ lat:	33, common_pid:	2302 }	hitcount:	3
{ lat:	34, common_pid:	2302 }	hitcount:	2

```

{ lat:      35, common_pid:    2302 } hitcount:      1
{ lat:      35, common_pid:    2304 } hitcount:      1
{ lat:      36, common_pid:    2302 } hitcount:      4
{ lat:      37, common_pid:    2302 } hitcount:      6
{ lat:      38, common_pid:    2302 } hitcount:      2
{ lat:      39, common_pid:    2302 } hitcount:      2
{ lat:      39, common_pid:    2304 } hitcount:      1
{ lat:      40, common_pid:    2304 } hitcount:      2
{ lat:      40, common_pid:    2302 } hitcount:      5
{ lat:      41, common_pid:    2304 } hitcount:      1
{ lat:      41, common_pid:    2302 } hitcount:      8
{ lat:      42, common_pid:    2302 } hitcount:      6
{ lat:      42, common_pid:    2304 } hitcount:      1
{ lat:      43, common_pid:    2302 } hitcount:      3
{ lat:      43, common_pid:    2304 } hitcount:      4
{ lat:      44, common_pid:    2302 } hitcount:      6
{ lat:      45, common_pid:    2302 } hitcount:      5
{ lat:      46, common_pid:    2302 } hitcount:      5
{ lat:      47, common_pid:    2302 } hitcount:      7
{ lat:      48, common_pid:    2301 } hitcount:      1
{ lat:      48, common_pid:    2302 } hitcount:      9
{ lat:      49, common_pid:    2302 } hitcount:      3
{ lat:      50, common_pid:    2302 } hitcount:      1
{ lat:      50, common_pid:    2301 } hitcount:      1
{ lat:      51, common_pid:    2302 } hitcount:      2
{ lat:      51, common_pid:    2301 } hitcount:      1
{ lat:      61, common_pid:    2302 } hitcount:      1
{ lat:     110, common_pid:    2302 } hitcount:      1

```

Totals:

```

Hits: 89565
Entries: 158
Dropped: 0

```

This doesn't tell us any information about how late cyclicttest may have woken up, but it does show us a nice histogram of how long it took from the time that cyclicttest was woken to the time it made it into user space.

HISTOGRAM DESIGN NOTES

Author Tom Zanussi <zanussi@kernel.org>

This document attempts to provide a description of how the ftrace histograms work and how the individual pieces map to the data structures used to implement them in `trace_events_hist.c` and `tracing_map.c`.

Note: All the ftrace histogram command examples assume the working directory is the ftrace /tracing directory. For example:

```
# cd /sys/kernel/debug/tracing
```

Also, the histogram output displayed for those commands will be generally be truncated - only enough to make the point is displayed.

17.1 'hist_debug' trace event files

If the kernel is compiled with `CONFIG_HIST_TRIGGERS_DEBUG` set, an event file named 'hist_debug' will appear in each event's subdirectory. This file can be read at any time and will display some of the hist trigger internals described in this document. Specific examples and output will be described in test cases below.

17.2 Basic histograms

First, basic histograms. Below is pretty much the simplest thing you can do with histograms - create one with a single key on a single event and cat the output:

```
# echo 'hist:keys=pid' >> events/sched/sched_waking/trigger

# cat events/sched/sched_waking/hist

{ pid:      18249 } hitcount:      1
{ pid:      13399 } hitcount:      1
{ pid:      17973 } hitcount:      1
{ pid:      12572 } hitcount:      1
...
{ pid:        10 } hitcount:     921
{ pid:     18255 } hitcount:     1444
{ pid:     25526 } hitcount:     2055
```

```
{ pid:      5257 } hitcount:      2055
{ pid:      27367 } hitcount:      2055
{ pid:       1728 } hitcount:      2161
```

```
Totals:
  Hits: 21305
  Entries: 183
  Dropped: 0
```

What this does is create a histogram on the `sched_waking` event using `pid` as a key and with a single value, `hitcount`, which even if not explicitly specified, exists for every histogram regardless.

The `hitcount` value is a per-bucket value that's automatically incremented on every hit for the given key, which in this case is the `pid`.

So in this histogram, there's a separate bucket for each `pid`, and each bucket contains a value for that bucket, counting the number of times `sched_waking` was called for that `pid`.

Each histogram is represented by a `hist_data` struct.

To keep track of each key and value field in the histogram, `hist_data` keeps an array of these fields named `fields[]`. The `fields[]` array is an array containing `struct hist_field` representations of each histogram `val` and `key` in the histogram (variables are also included here, but are discussed later). So for the above histogram we have one key and one value; in this case the one value is the `hitcount` value, which all histograms have, regardless of whether they define that value or not, which the above histogram does not.

Each `struct hist_field` contains a pointer to the `ftrace_event_field` from the event's `trace_event_file` along with various bits related to that such as the size, offset, type, and a `hist_field_fn_t` function, which is used to grab the field's data from the `ftrace` event buffer (in most cases - some `hist_fields` such as `hitcount` don't directly map to an event field in the trace buffer - in these cases the function implementation gets its value from somewhere else). The `flags` field indicates which type of field it is - key, value, variable, variable reference, etc., with `value` being the default.

The other important `hist_data` data structure in addition to the `fields[]` array is the `tracing_map` instance created for the histogram, which is held in the `.map` member. The `tracing_map` implements the lock-free hash table used to implement histograms (see `kernel/trace/tracing_map.h` for much more discussion about the low-level data structures implementing the `tracing_map`). For the purposes of this discussion, the `tracing_map` contains a number of buckets, each bucket corresponding to a particular `tracing_map_elt` object hashed by a given histogram key.

Below is a diagram the first part of which describes the `hist_data` and associated key and value fields for the histogram described above. As you can see, there are two fields in the `fields` array, one `val` field for the `hitcount` and one `key` field for the `pid` key.

Below that is a diagram of a run-time snapshot of what the `tracing_map` might look like for a given run. It attempts to show the relationships between the `hist_data` fields and the `tracing_map` elements for a couple hypothetical keys and values.:

```
+-----+
| hist_data |
+-----+
| .fields[] |---->| val = hitcount |-----+
```



```

| .fields |
+-----+ +-----+
| .map    |---->| map_entry |
+-----+ +-----+
          | .key    |---> 0
          +-----+
          | .val    |---> NULL
          +-----+
| map_entry |
+-----+
          | .key    |---> pid = 999
          +-----+ +-----+
          | .val    |--->| map_elt  |
          +-----+ +-----+
                  | .key    |---> full key *
                  +-----+ +-----+
                  | .fields |--->| .sum (val) |<--+
                  +-----+ +-----+
                  | 2345          |
| map_entry | +-----+
+-----+      | .offset (key) |<-----+
          | .key    |---> 0
          +-----+
          | .val    |---> NULL
          +-----+
          | map_entry |
          +-----+
          | .key    |
          +-----+ +-----+
          | .val    |--->| map_elt |
          +-----+ +-----+
| map_entry | +-----+
+-----+      | .sum (val) or |
          | .key    |
          +-----+ +-----+
          | .val    |--->| map_elt |
          +-----+ +-----+
          | .sum (val) or |
          | .offset (key) |
          +-----+
          | .key    |---> pid = 4444
          +-----+ +-----+
          | .val    |   | map_elt |
          +-----+ +-----+
                  | .key    |---> full key *
                  +-----+ +-----+
                  | .fields |--->| .sum (val) |<--+
                  +-----+ +-----+
                  | 65523          |
                  +-----+
                  | .offset (key) |<-----+
                  | 0              |
                  +-----+
                  |
                  |
                  |
                  +-----+
                  | .sum (val) or |
                  | .offset (key) |

```

```
+-----+
| .sum (val) or |
| .offset (key) |
+-----+
```

Abbreviations used in the diagrams:

```
hist_data = struct hist_trigger_data
hist_data.fields = struct hist_field
fn = hist_field_fn_t
map_entry = struct tracing_map_entry
map_elt = struct tracing_map_elt
map_elt.fields = struct tracing_map_field
```

Whenever a new event occurs and it has a hist trigger associated with it, `event_hist_trigger()` is called. `event_hist_trigger()` first deals with the key: for each subkey in the key (in the above example, there is just one subkey corresponding to `pid`), the `hist_field` that represents that subkey is retrieved from `hist_data.fields[]` and the `hist_field_fn_t fn()` associated with that field, along with the field's size and offset, is used to grab that subkey's data from the current trace record.

Once the complete key has been retrieved, it's used to look that key up in the `tracing_map`. If there's no `tracing_map_elt` associated with that key, an empty one is claimed and inserted in the map for the new key. In either case, the `tracing_map_elt` associated with that key is returned.

Once a `tracing_map_elt` available, `hist_trigger_elt_update()` is called. As the name implies, this updates the element, which basically means updating the element's fields. There's a `tracing_map_field` associated with each key and value in the histogram, and each of these correspond to the key and value `hist_fields` created when the histogram was created. `hist_trigger_elt_update()` goes through each value `hist_field` and, as for the keys, uses the `hist_field`'s `fn()` and size and offset to grab the field's value from the current trace record. Once it has that value, it simply adds that value to that field's continually-updated `tracing_map_field.sum` member. Some `hist_field fn()`s, such as for the hitcount, don't actually grab anything from the trace record (the hitcount `fn()` just increments the counter `sum` by 1), but the idea is the same.

Once all the values have been updated, `hist_trigger_elt_update()` is done and returns. Note that there are also `tracing_map_fields` for each subkey in the key, but `hist_trigger_elt_update()` doesn't look at them or update anything - those exist only for sorting, which can happen later.

17.2.1 Basic histogram test

This is a good example to try. It produces 3 value fields and 2 key fields in the output:

```
# echo 'hist:keys=common_pid,call_site.sym:values=bytes_req,bytes_alloc,
↪hitcount' >> events/kmem/kmalloc/trigger
```

To see the debug data, cat the `kmem/kmalloc`'s `'hist_debug'` file. It will show the trigger info of the histogram it corresponds to, along with the address of the `hist_data` associated with the histogram, which will become useful in later examples. It then displays the number of total `hist_fields` associated with the histogram along with a count of how many of those correspond to keys and how many correspond to values.

It then goes on to display details for each field, including the field's flags and the position of each field in the `hist_data`'s `fields[]` array, which is useful information for verifying that things internally appear correct or not, and which again will become even more useful in further examples:

```
# cat events/kmem/kmalloc/hist_debug

# event histogram
#
# trigger info: hist:keys=common_pid,call_site.sym:vals=hitcount,bytes_req,
→ bytes_alloc:sort=hitcount:size=2048 [active]
#

hist_data: 000000005e48c9a5

n_vals: 3
n_keys: 2
n_fields: 5

val fields:

  hist_data->fields[0]:
    flags:
      VAL: HIST_FIELD_FL_HITCOUNT
    type: u64
    size: 8
    is_signed: 0

  hist_data->fields[1]:
    flags:
      VAL: normal u64 value
    ftrace_event_field name: bytes_req
    type: size_t
    size: 8
    is_signed: 0

  hist_data->fields[2]:
    flags:
      VAL: normal u64 value
    ftrace_event_field name: bytes_alloc
    type: size_t
    size: 8
    is_signed: 0

key fields:

  hist_data->fields[3]:
    flags:
      HIST_FIELD_FL_KEY
    ftrace_event_field name: common_pid
    type: int
```

```

size: 8
is_signed: 1

hist_data->fields[4]:
  flags:
    HIST_FIELD_FL_KEY
  ftrace_event_field name: call_site
  type: unsigned long
  size: 8
  is_signed: 0

```

The commands below can be used to clean things up for the next test:

```
# echo '!hist:keys=common_pid,call_site.sym:values=bytes_req,bytes_alloc,
↪hitcount' >> events/kmem/kmalloc/trigger
```

17.3 Variables

Variables allow data from one hist trigger to be saved by one hist trigger and retrieved by another hist trigger. For example, a trigger on the sched_waking event can capture a timestamp for a particular pid, and later a sched_switch event that switches to that pid event can grab the timestamp and use it to calculate a time delta between the two events:

```
# echo 'hist:keys=pid:ts0=common_timestamp.usecs' >>
    events/sched/sched_waking/trigger

# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-$ts0' >>
    events/sched/sched_switch/trigger
```

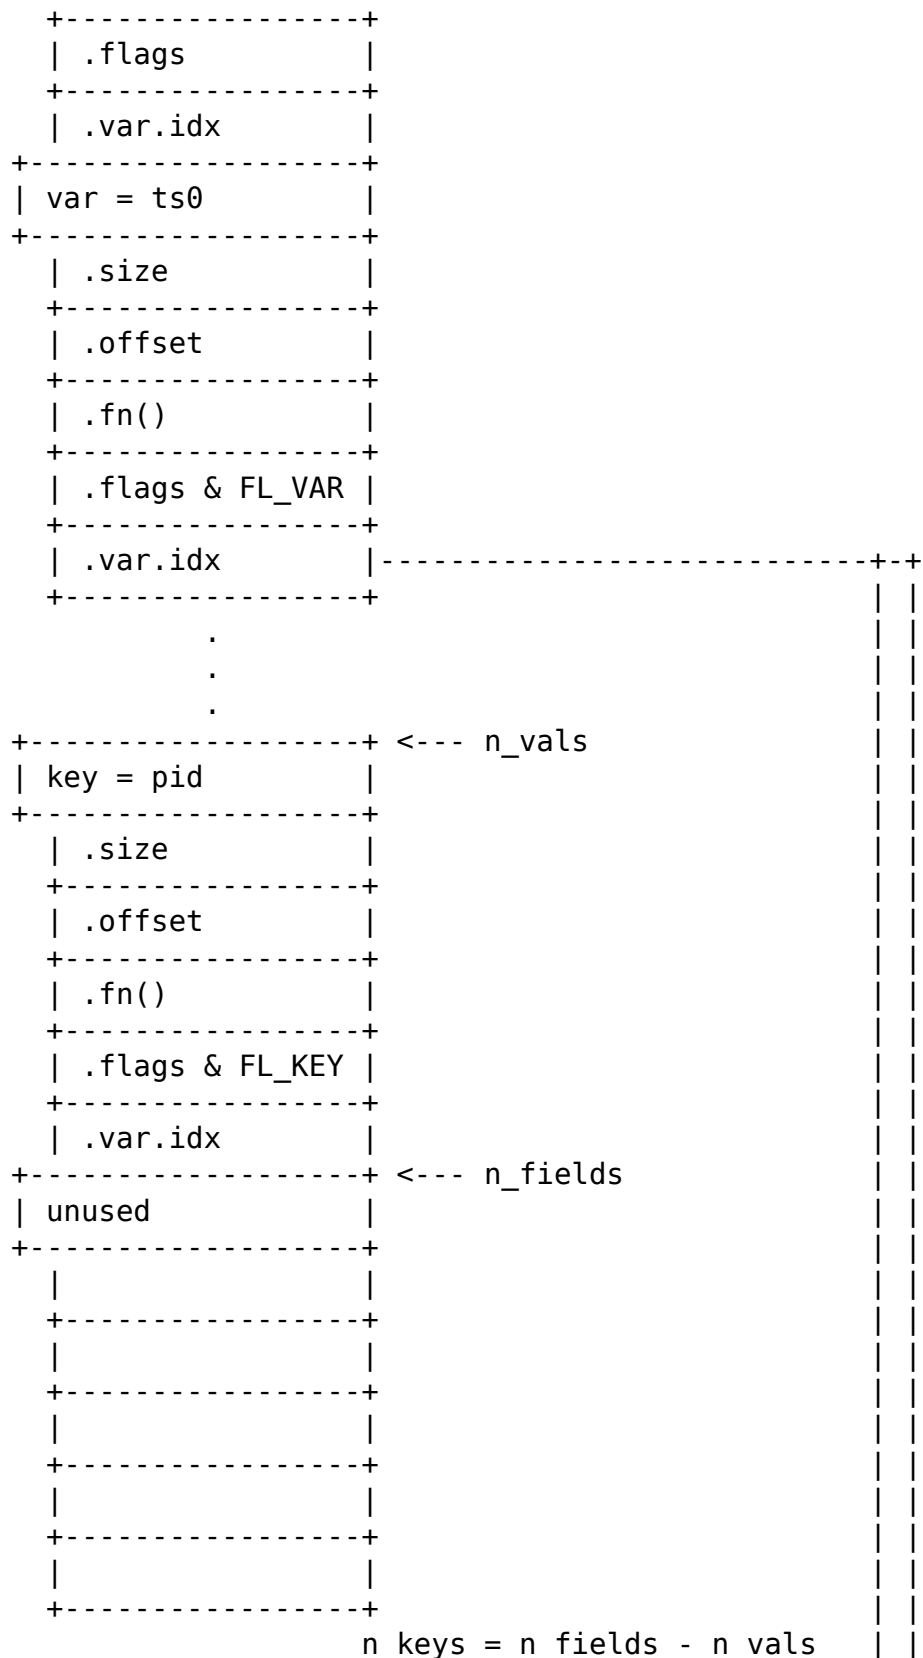
In terms of the histogram data structures, variables are implemented as another type of hist_field and for a given hist trigger are added to the hist_data.fields[] array just after all the val fields. To distinguish them from the existing key and val fields, they're given a new flag type, HIST_FIELD_FL_VAR (abbreviated FL_VAR) and they also make use of a new .var.idx field member in struct hist_field, which maps them to an index in a new map_elt.vars[] array added to the map_elt specifically designed to store and retrieve variable values. The diagram below shows those new elements and adds a new variable entry, ts0, corresponding to the ts0 variable in the sched_waking trigger above.

sched_waking histogram -----:

```

+-----+
| hist_data |<-----+
+-----+ +-----+
| .fields[] |-->| val = hitcount |
+-----+ +-----+
| .map      |   | .size      |
+-----+ +-----+
|           |   | .offset    |
|           |   +-----+
|           |   | .fn()      |
|           |   +-----+

```



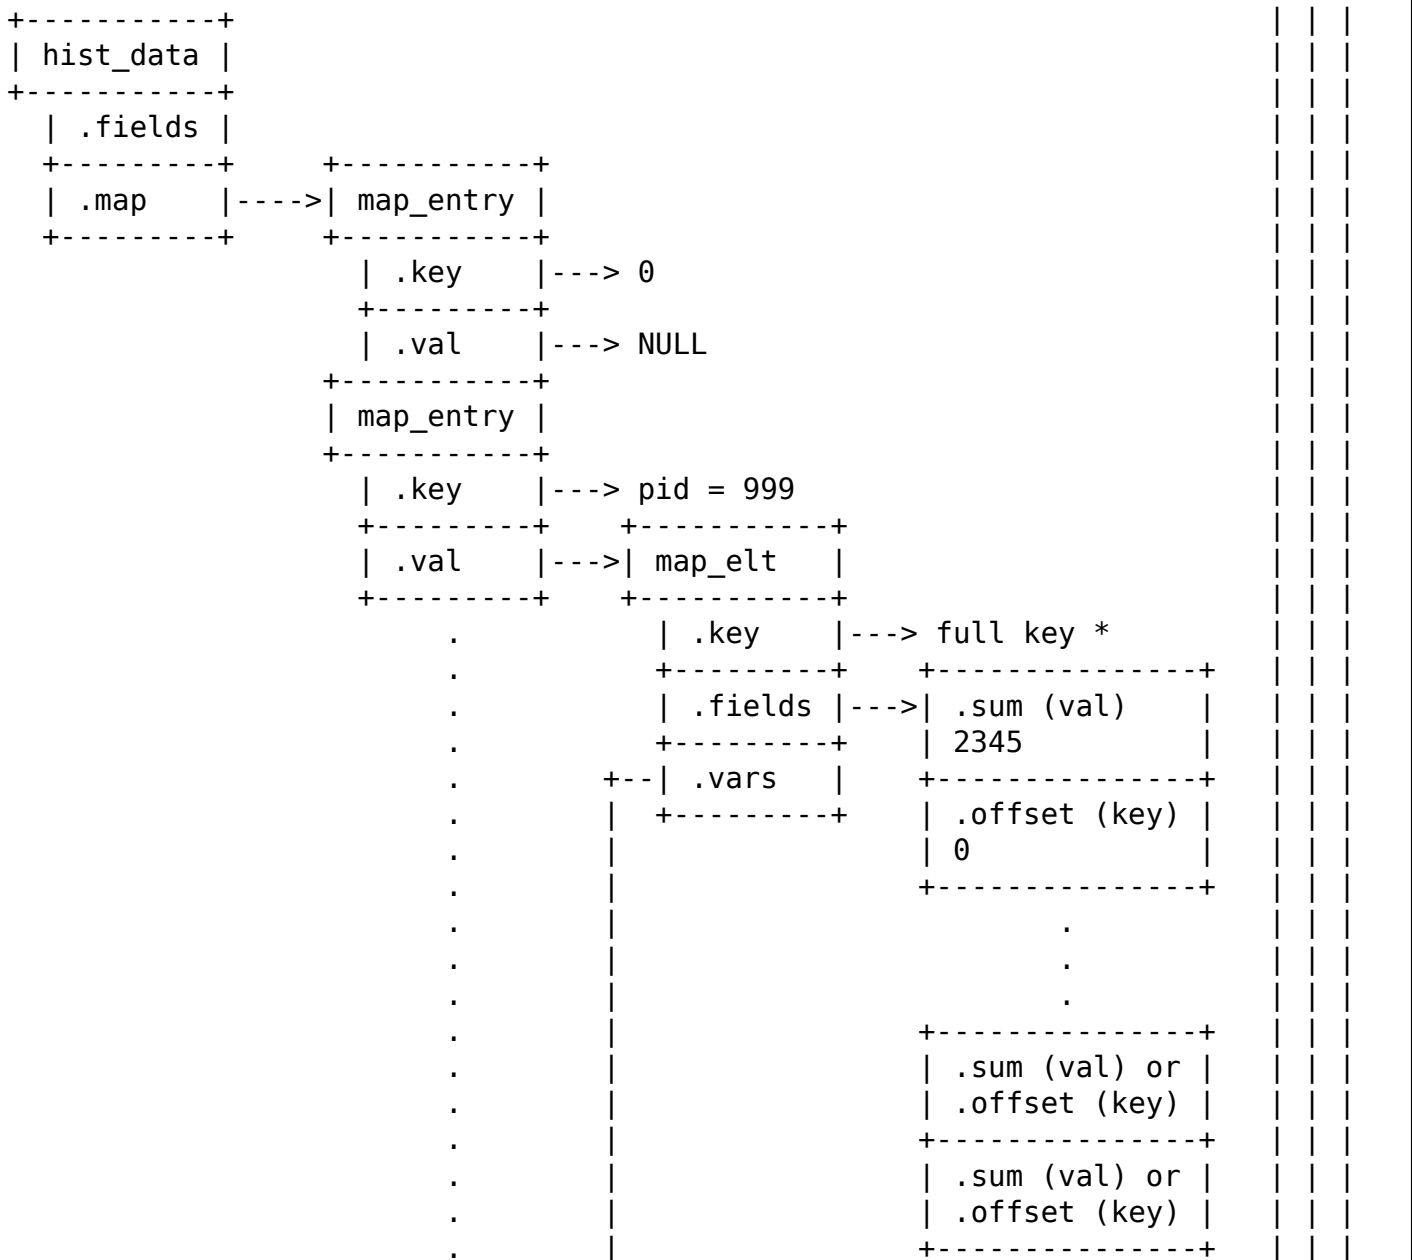
This is very similar to the basic case. In the above diagram, we can see a new .flags member has been added to the struct hist_field struct, and a new entry added to hist_data.fields representing the ts0 variable. For a normal val hist_field, .flags is just 0 (modulo

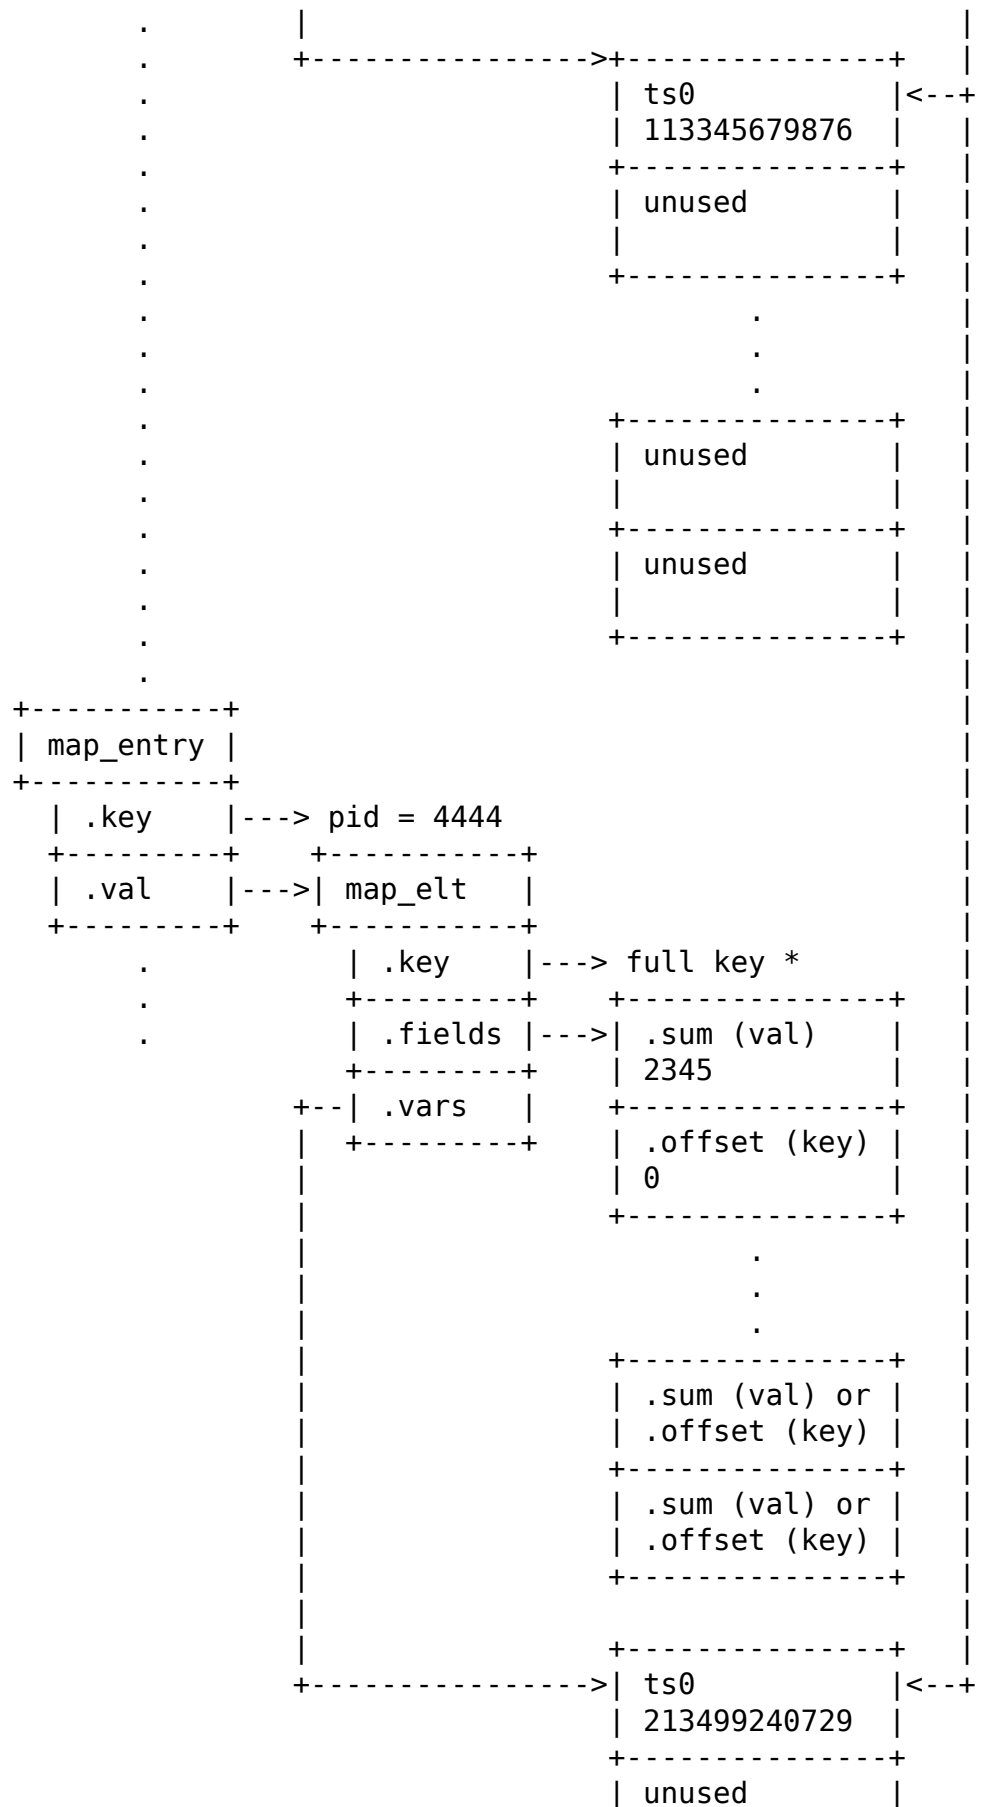
modifier flags), but if the value is defined as a variable, the .flags ||| contains a set FL_VAR bit. |||


As you can see, the ts0 entry's .var.idx member contains the index ||| into the tracing_map_elts' .vars[] array containing variable values. ||| This idx is used whenever the value of the variable is set or read. ||| The map_elt.vars idx assigned to the given variable is assigned and ||| saved in .var.idx by create_tracing_map_fields() after it calls ||| tracing_map_add_var(). |||

Below is a representation of the histogram at run-time, which ||| populates the map, along with correspondence to the above hist_data and ||| hist_field data structures. |||

The diagram attempts to show the relationship between the ||| hist_data.fields[] and the map_elt.fields[] and map_elt.vars[] with ||| the links drawn between diagrams. For each of the map_elts, you can ||| see that the .fields[] members point to the .sum or .offset of a key ||| or val and the .vars[] members point to the value of a variable. The ||| arrows between the two diagrams show the linkages between those ||| tracing_map members and the field definitions in the corresponding ||| hist_data fields[] members.:







For each used map entry, there's a `map_elt` pointing to an array of `||.vars` containing the current value of the variables associated with `||` that histogram entry. So in the above, the timestamp associated with `|| pid 999` is `113345679876`, and the timestamp variable in the same `||.var.idx` for `pid 4444` is `213499240729`. `||`

sched_switch histogram | | ----- | |

The sched_switch histogram paired with the above sched_waking | | histogram is shown below. The most important aspect of the | | sched_switch histogram is that it references a variable on the | | sched_waking histogram above. | |

The histogram diagram is very similar to the others so far displayed, but it adds variable references. You can see the normal hitcount and key fields along with a new wakeup_lat variable implemented in the same way as the sched_waking ts0 variable, but in addition there's an entry with the new FL_VAR_REF (short for HIST_FIELD_FL_VAR_REF) flag.

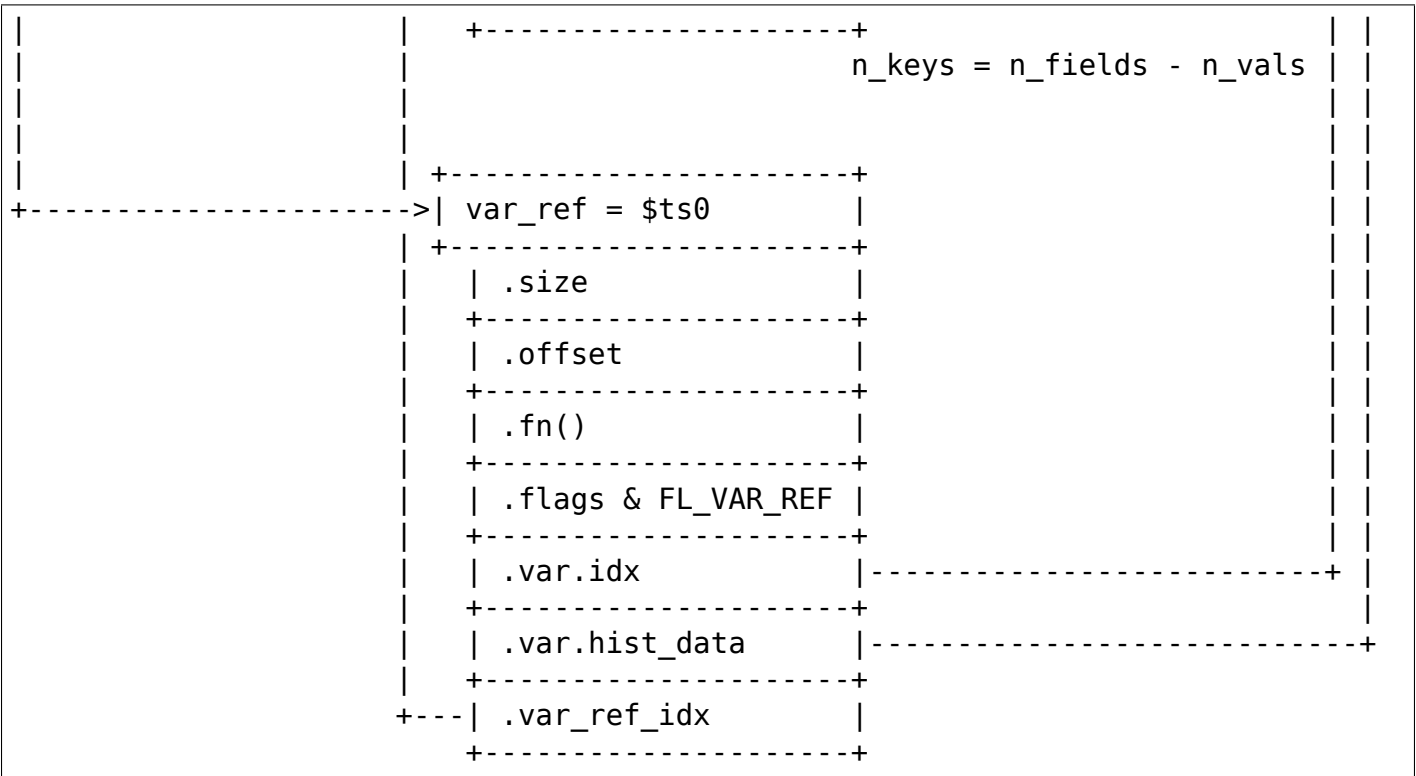
Associated with the new var_ref field are a couple of new hist_field || members, var.hist_data and var_ref_idx. For a variable reference, the || var.hist_data goes with the var.idx, which together uniquely identify || a particular variable on a particular histogram. The var_ref_idx is || just the index into the var_ref_vals[] array that caches the values of || each variable whenever a hist trigger is updated. Those resulting || values are then finally accessed by other code such as trace action || code that uses the var_ref_idx values to assign param values. ||

The diagram below describes the situation for the sched_switch || histogram referred to before:

```
# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-$ts0' >>
    events/sched/sched_switch/trigger

+-----+
| hist_data |
+-----+
+-----+ +-----+
| .fields[] | -->| val = hitcount |
+-----+ +-----+
| .map | | .size |
+-----+ +-----+
+--| .var_refs[] | | .offset |
| +-----+ +-----+
| | | .fn() |
| var_ref_vals[] +-----+
| +-----+ | .flags |
|
```

\$ts0	<---+	+-----+		
+-----+		.var.idx		
		+-----+		
+-----+		.var.hist_data		
		+-----+		
+-----+		.var_ref_idx		
		+-----+		
+-----+		var = wakeup_lat		
.		+-----+		
.		.size		
.		+-----+		
+-----+		.offset		
		+-----+		
+-----+		.fn()		
		+-----+		
+-----+		.flags & FL_VAR		
		+-----+		
		.var.idx		
		+-----+		
		.var.hist_data		
		+-----+		
		.var_ref_idx		
		+-----+		
		.		
		.		
		.		
		+-----+	<--- n_vals	
		key = pid		
		+-----+		
		.size		
		+-----+		
		.offset		
		+-----+		
		.fn()		
		+-----+		
		.flags		
		+-----+		
		.var.idx		
		+-----+	<--- n_fields	
		unused		
		+-----+		
		+-----+		
		+-----+		
		+-----+		
		+-----+		



Abbreviations used in the diagrams:

```

hist_data = struct hist_trigger_data
hist_data.fields = struct hist_field
fn = hist_field_fn_t
FL_KEY = HIST_FIELD_FL_KEY
FL_VAR = HIST_FIELD_FL_VAR
FL_VAR_REF = HIST_FIELD_FL_VAR_REF

```

When a hist trigger makes use of a variable, a new hist_field is created with flag HIST_FIELD_FL_VAR_REF. For a VAR_REF field, the var.idx and var.hist_data take the same values as the referenced variable, as well as the referenced variable's size, type, and is_signed values. The VAR_REF field's .name is set to the name of the variable it references. If a variable reference was created using the explicit system.event.\$var_ref notation, the hist_field's system and event_name variables are also set.

So, in order to handle an event for the sched_switch histogram, because we have a reference to a variable on another histogram, we need to resolve all variable references first. This is done via the resolve_var_refs() calls made from event_hist_trigger(). What this does is grabs the var_refs[] array from the hist_data representing the sched_switch histogram. For each one of those, the referenced variable's var.hist_data along with the current key is used to look up the corresponding tracing_map_elt in that histogram. Once found, the referenced variable's var.idx is used to look up the variable's value using tracing_map_read_var(elt, var.idx), which yields the value of the variable for that element, ts0 in the case above. Note that both the hist_fields representing both the variable and the variable reference have the same var.idx, so this is straightforward.

17.3.1 Variable and variable reference test

This example creates a variable on the sched_waking event, ts0, and uses it in the sched_switch trigger. The sched_switch trigger also creates its own variable, wakeup_lat, but nothing yet uses it:

```
# echo 'hist:keys=pid:ts0=common_timestamp.usecs' >> events/sched/sched_waking/  
→trigger  
  
# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-$ts0' >> events/  
→sched/sched_switch/trigger
```

Looking at the sched_waking 'hist_debug' output, in addition to the normal key and value hist_fields, in the val fields section we see a field with the HIST_FIELD_FL_VAR flag, which indicates that that field represents a variable. Note that in addition to the variable name, contained in the var.name field, it includes the var.idx, which is the index into the tracing_map_elt.vars[] array of the actual variable location. Note also that the output shows that variables live in the same part of the hist_data->fields[] array as normal values:

```
# cat events/sched/sched_waking/hist_debug  
  
# event histogram  
#  
# trigger info: hist:keys=pid:vals=hitcount:ts0=common_timestamp.  
→usecs:sort=hitcount:size=2048:clock=global [active]  
#  
  
hist_data: 000000009536f554  
  
n_vals: 2  
n_keys: 1  
n_fields: 3  
  
val fields:  
  
  hist_data->fields[0]:  
    flags:  
      VAL: HIST_FIELD_FL_HITCOUNT  
    type: u64  
    size: 8  
    is_signed: 0  
  
  hist_data->fields[1]:  
    flags:  
      HIST_FIELD_FL_VAR  
    var.name: ts0  
    var.idx (into tracing_map_elt.vars[]): 0  
    type: u64  
    size: 8  
    is_signed: 0  
  
key fields:
```

```

hist_data->fields[2]:
  flags:
    HIST_FIELD_FL_KEY
  ftrace_event_field name: pid
  type: pid_t
  size: 8
  is_signed: 1

```

Moving on to the sched_switch trigger hist_debug output, in addition to the unused wakeup_lat variable, we see a new section displaying variable references. Variable references are displayed in a separate section because in addition to being logically separate from variables and values, they actually live in a separate hist_data array, var_refs[].

In this example, the sched_switch trigger has a reference to a variable on the sched_waking trigger, \$ts0. Looking at the details, we can see that the var.hist_data value of the referenced variable matches the previously displayed sched_waking trigger, and the var.idx value matches the previously displayed var.idx value for that variable. Also displayed is the var_ref_idx value for that variable reference, which is where the value for that variable is cached for use when the trigger is invoked:

```

# cat events/sched/sched_switch/hist_debug

# event histogram
#
# trigger info: hist:keys=next_pid:vals=hitcount:wakeup_lat=common_timestamp.
→usecs-$ts0:sort=hitcount:size=2048:clock=global [active]
#

hist_data: 00000000f4ee8006

n_vals: 2
n_keys: 1
n_fields: 3

val fields:

  hist_data->fields[0]:
    flags:
      VAL: HIST_FIELD_FL_HITCOUNT
    type: u64
    size: 8
    is_signed: 0

  hist_data->fields[1]:
    flags:
      HIST_FIELD_FL_VAR
    var.name: wakeup_lat
    var.idx (into tracing_map_elt.vars[]): 0
    type: u64
    size: 0

```

```
is_signed: 0

key fields:

hist_data->fields[2]:
  flags:
    HIST_FIELD_FL_KEY
  ftrace_event_field name: next_pid
  type: pid_t
  size: 8
  is_signed: 1

variable reference fields:

hist_data->var_refs[0]:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: ts0
  var.idx (into tracing_map_elt.vars[]): 0
  var.hist_data: 000000009536f554
  var_ref_idx (into hist_data->var_refs[]): 0
  type: u64
  size: 8
  is_signed: 0
```

The commands below can be used to clean things up for the next test:

```
# echo '!hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-$ts0' >> events/
→ sched/sched_switch/trigger

# echo '!hist:keys=pid:ts0=common_timestamp.usecs' >> events/sched/sched_
→ waking/trigger
```

17.4 Actions and Handlers

Adding onto the previous example, we will now do something with that `wakeup_lat` variable, namely send it and another field as a synthetic event.

The `onmatch()` action below basically says that whenever we have a `sched_switch` event, if we have a matching `sched_waking` event, in this case if we have a `pid` in the `sched_waking` histogram that matches the `next_pid` field on this `sched_switch` event, we retrieve the variables specified in the `wakeup_latency()` trace action, and use them to generate a new `wakeup_latency` event into the trace stream.

Note that the way the trace handlers such as `wakeup_latency()` (which could equivalently be written `trace(wakeup_latency,$wakeup_lat,next_pid)`) are implemented, the parameters specified to the trace handler must be variables. In this case, `$wakeup_lat` is obviously a variable, but `next_pid` isn't, since it's just naming a field in the `sched_switch` trace event. Since this is something that almost every `trace()` and `save()` action does, a special shortcut is implemented to allow field names to be used directly in those cases. How it works is that under the covers,

a temporary variable is created for the named field, and this variable is what is actually passed to the trace handler. In the code and documentation, this type of variable is called a ‘field variable’.

Fields on other trace event’s histograms can be used as well. In that case we have to generate a new histogram and an unfortunately named ‘synthetic_field’ (the use of synthetic here has nothing to do with synthetic events) and use that special histogram field as a variable.

The diagram below illustrates the new elements described above in the context of the sched_switch histogram using the onmatch() handler and the trace() action.

First, we define the wakeup_latency synthetic event:

```
# echo 'wakeup_latency u64 lat; pid_t pid' >> synthetic_events
```

Next, the sched_waking hist trigger as before:

```
# echo 'hist:keys=pid:ts0=common_timestamp.usecs' >>
events/sched/sched_waking/trigger
```

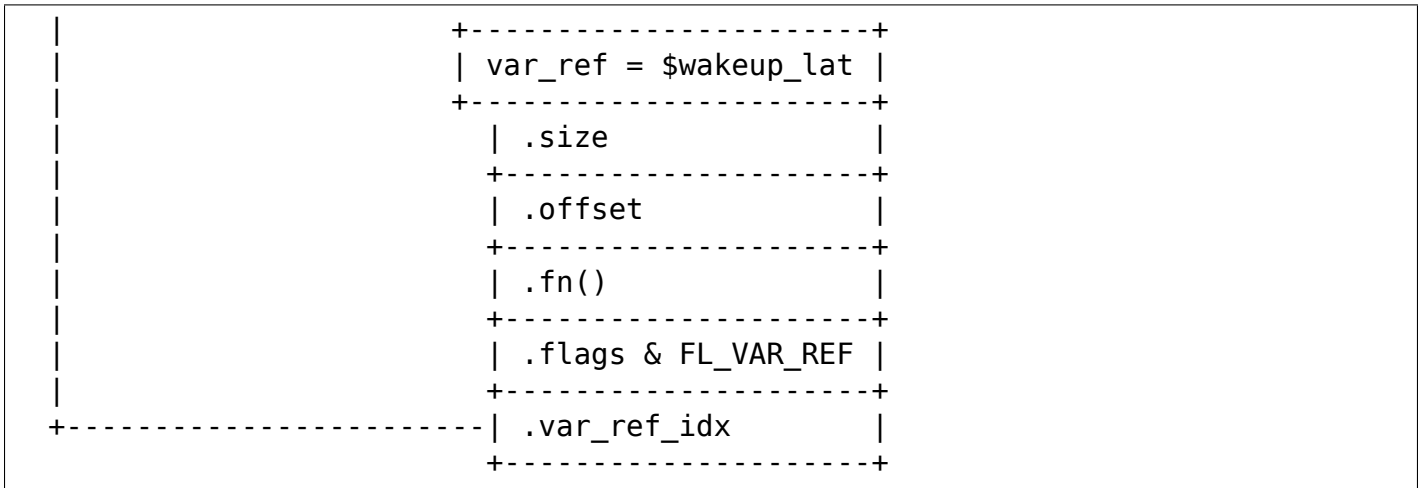
Finally, we create a hist trigger on the sched_switch event that generates a wakeup_latency() trace event. In this case we pass next_pid into the wakeup_latency synthetic event invocation, which means it will be automatically converted into a field variable:

```
# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-$ts0: \
onmatch(sched.sched_waking).wakeup_latency($wakeup_lat,next_pid)' >>
/sys/kernel/debug/tracing/events/sched/sched_switch/trigger
```

The diagram for the sched_switch event is similar to previous examples but shows the additional field_vars[] array for hist_data and shows the linkages between the field_vars and the variables and references created to implement the field variables. The details are discussed below:

```
+-----+
| hist_data |
+-----+ +-----+
| .fields[] |-->| val = hitcount |
+-----+ +-----+
| .map      | | .size      |
+-----+ +-----+
+---| .field_vars[] | | .offset |
| +-----+ +-----+
|+--| .var_refs[] | | .offset |
|| +-----+ +-----+
|| | .fn() |
|| var_ref_vals[] +-----+
|| +-----+ | .flags |
|| | $ts0 |<---+ +-----+
|| +-----+ | | .var.idx |
|| | $next_pid |<--+ +-----+
|| +-----+ | | .var.hist_data |
||+>| $wakeup_lat | | +-----+
|| +-----+ | | .var_ref_idx |
|| | | | +-----+
|| +-----+ | | | var = wakeup_lat |
```

[illegible]



As you can see, for a field variable, two `hist_fields` are created: one representing the variable, in this case `next_pid`, and one to actually get the value of the field from the trace stream, like a normal `val` field does. These are created separately from normal variable creation and are saved in the `hist_data->field_vars[]` array. See below for how these are used. In addition, a reference `hist_field` is also created, which is needed to reference the field variables such as `$next_pid` variable in the `trace()` action.

Note that `$wakeup_lat` is also a variable reference, referencing the value of the expression `common_timestamp-$ts0`, and so also needs to have a `hist` field entry representing that reference created.

When `hist_trigger_elt_update()` is called to get the normal key and value fields, it also calls `update_field_vars()`, which goes through each `field_var` created for the histogram, and available from `hist_data->field_vars` and calls `val->fn()` to get the data from the current trace record, and then uses the `var`'s `var.idx` to set the variable at the `var.idx` offset in the appropriate `tracing_map_elt`'s variable at `elt->vars[var.idx]`.

Once all the variables have been updated, `resolve_var_refs()` can be called from `event_hist_trigger()`, and not only can our `$ts0` and `$next_pid` references be resolved but the `$wakeup_lat` reference as well. At this point, the `trace()` action can simply access the values assembled in the `var_ref_vals[]` array and generate the trace event.

The same process occurs for the field variables associated with the `save()` action.

Abbreviations used in the diagram:

```

hist_data = struct hist_trigger_data
hist_data.fields = struct hist_field
field_var = struct field_var
fn = hist_field_fn_t
FL_KEY = HIST_FIELD_FL_KEY
FL_VAR = HIST_FIELD_FL_VAR
FL_VAR_REF = HIST_FIELD_FL_VAR_REF

```

17.4.1 trace() action field variable test

This example adds to the previous test example by finally making use of the `wakeup_lat` variable, but in addition also creates a couple of field variables that then are all passed to the `wakeup_latency()` trace action via the `onmatch()` handler.

First, we create the `wakeup_latency` synthetic event:

```
# echo 'wakeup_latency u64 lat; pid_t pid; char comm[16]' >> synthetic_events
```

Next, the `sched_waking` trigger from previous examples:

```
# echo 'hist:keys=pid:ts0=common_timestamp.usecs' >> events/sched/sched_waking/
↪trigger
```

Finally, as in the previous test example, we calculate and assign the wakeup latency using the `$ts0` reference from the `sched_waking` trigger to the `wakeup_lat` variable, and finally use it along with a couple `sched_switch` event fields, `next_pid` and `next_comm`, to generate a `wakeup_latency` trace event. The `next_pid` and `next_comm` event fields are automatically converted into field variables for this purpose:

```
# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-
↪$ts0:onmatch(sched.sched_waking).wakeup_latency($wakeup_lat,next_pid,next_
↪comm)' >> /sys/kernel/debug/tracing/events/sched/sched_switch/trigger
```

The `sched_waking hist_debug` output shows the same data as in the previous test example:

```
# cat events/sched/sched_waking/hist_debug

# event histogram
#
# trigger info: hist:keys=pid:vals=hitcount:ts0=common_timestamp.
↪usecs:sort=hitcount:size=2048:clock=global [active]
#

hist_data: 00000000d60ff61f

n_vals: 2
n_keys: 1
n_fields: 3

val fields:

  hist_data->fields[0]:
    flags:
      VAL: HIST_FIELD_FL_HITCOUNT
    type: u64
    size: 8
    is_signed: 0

  hist_data->fields[1]:
    flags:
      HIST_FIELD_FL_VAR
```

```
var.name: ts0
var.idx (into tracing_map_elt.vars[]): 0
type: u64
size: 8
is_signed: 0
```

key fields:

```
hist_data->fields[2]:
  flags:
    HIST_FIELD_FL_KEY
  ftrace_event_field name: pid
  type: pid_t
  size: 8
  is_signed: 1
```

The sched_switch hist_debug output shows the same key and value fields as in the previous test example - note that wakeup_lat is still in the val fields section, but that the new field variables are not there - although the field variables are variables, they're held separately in the hist_data's field_vars[] array. Although the field variables and the normal variables are located in separate places, you can see that the actual variable locations for those variables in the tracing_map_elt.vars[] do have increasing indices as expected: wakeup_lat takes the var.idx = 0 slot, while the field variables for next_pid and next_comm have values var.idx = 1, and var.idx = 2. Note also that those are the same values displayed for the variable references corresponding to those variables in the variable reference fields section. Since there are two triggers and thus two hist_data addresses, those addresses also need to be accounted for when doing the matching - you can see that the first variable refers to the 0 var.idx on the previous hist trigger (see the hist_data address associated with that trigger), while the second variable refers to the 0 var.idx on the sched_switch hist trigger, as do all the remaining variable references.

Finally, the action tracking variables section just shows the system and event name for the onmatch() handler:

```
# cat events/sched/sched_switch/hist_debug

# event histogram
#
# trigger info: hist:keys=next_pid:vals=hitcount:wakeup_lat=common_timestamp.
→usecs-$ts0:sort=hitcount:size=2048:clock=global:onmatch(sched.sched_waking).
→wakeup_latency($wakeup_lat,next_pid,next_comm) [active]
#

hist_data: 0000000008f551b7

n_vals: 2
n_keys: 1
n_fields: 3

val fields:

  hist_data->fields[0]:
```

```

    flags:
      VAL: HIST_FIELD_FL_HITCOUNT
    type: u64
    size: 8
    is_signed: 0

hist_data->fields[1]:
  flags:
    HIST_FIELD_FL_VAR
  var.name: wakeup_lat
  var.idx (into tracing_map_elt.vars[]): 0
  type: u64
  size: 0
  is_signed: 0

key fields:

hist_data->fields[2]:
  flags:
    HIST_FIELD_FL_KEY
  ftrace_event_field name: next_pid
  type: pid_t
  size: 8
  is_signed: 1

variable reference fields:

hist_data->var_refs[0]:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: ts0
  var.idx (into tracing_map_elt.vars[]): 0
  var.hist_data: 00000000d60ff61f
  var_ref_idx (into hist_data->var_refs[]): 0
  type: u64
  size: 8
  is_signed: 0

hist_data->var_refs[1]:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: wakeup_lat
  var.idx (into tracing_map_elt.vars[]): 0
  var.hist_data: 0000000008f551b7
  var_ref_idx (into hist_data->var_refs[]): 1
  type: u64
  size: 0
  is_signed: 0

hist_data->var_refs[2]:

```

```
flags:
  HIST_FIELD_FL_VAR_REF
name: next_pid
var.idx (into tracing_map_elt.vars[]): 1
var.hist_data: 0000000008f551b7
var_ref_idx (into hist_data->var_refs[]): 2
type: pid_t
size: 4
is_signed: 0

hist_data->var_refs[3]:
flags:
  HIST_FIELD_FL_VAR_REF
name: next_comm
var.idx (into tracing_map_elt.vars[]): 2
var.hist_data: 0000000008f551b7
var_ref_idx (into hist_data->var_refs[]): 3
type: char[16]
size: 256
is_signed: 0

field variables:

hist_data->field_vars[0]:

  field_vars[0].var:
  flags:
    HIST_FIELD_FL_VAR
  var.name: next_pid
  var.idx (into tracing_map_elt.vars[]): 1

  field_vars[0].val:
  ftrace_event_field name: next_pid
  type: pid_t
  size: 4
  is_signed: 1

hist_data->field_vars[1]:

  field_vars[1].var:
  flags:
    HIST_FIELD_FL_VAR
  var.name: next_comm
  var.idx (into tracing_map_elt.vars[]): 2

  field_vars[1].val:
  ftrace_event_field name: next_comm
  type: char[16]
  size: 256
  is_signed: 0
```


action tracking variables (for onmax()/onchange()/onmatch()):

```
hist_data->actions[0].match_data.event_system: sched
hist_data->actions[0].match_data.event: sched_waking
```

The commands below can be used to clean things up for the next test:

```
# echo '!hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-
->$ts0:onmatch(sched.sched_waking).wakeup_latency($wakeup_lat,next_pid,next_
->comm)' >> /sys/kernel/debug/tracing/events/sched/sched_switch/trigger

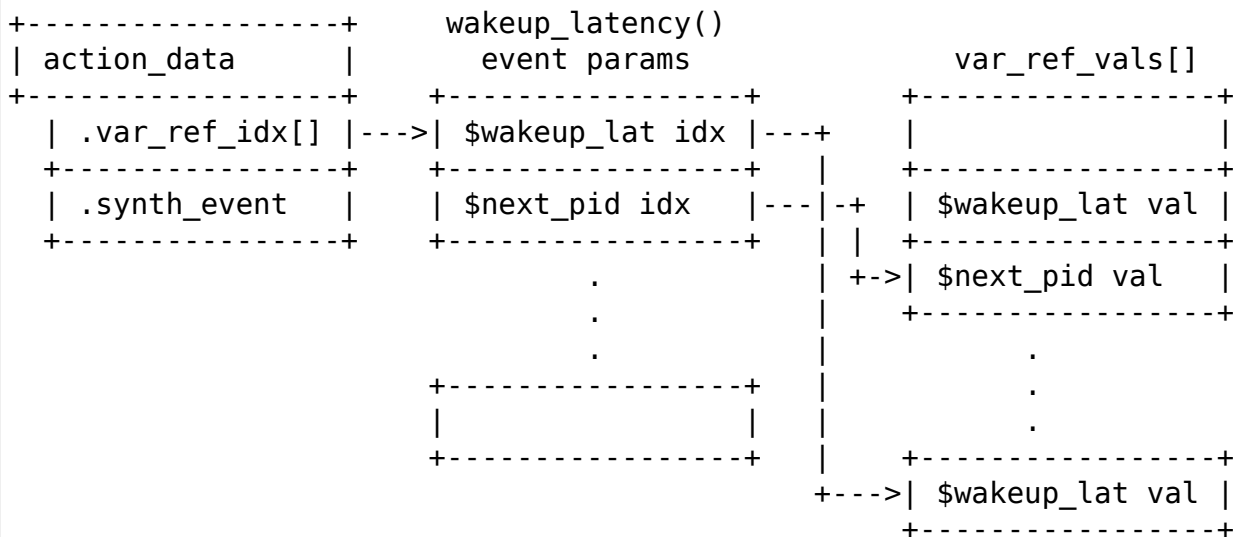
# echo '!hist:keys=pid:ts0=common_timestamp.usecs' >> events/sched/sched_
->waking/trigger

# echo '!wakeup_latency u64 lat; pid_t pid; char comm[16]' >> synthetic_events
```

17.4.2 action_data and the trace() action

As mentioned above, when the trace() action generates a synthetic event, all the parameters to the synthetic event either already are variables or are converted into variables (via field variables), and finally all those variable values are collected via references to them into a var_ref_vals[] array.

The values in the var_ref_vals[] array, however, don't necessarily follow the same ordering as the synthetic event params. To address that, struct action_data contains another array, var_ref_idx[] that maps the trace action params to the var_ref_vals[] values. Below is a diagram illustrating that for the wakeup_latency() synthetic event:



Basically, how this ends up getting used in the synthetic event probe function, trace_event_raw_event_synth(), is as follows:

```
for each field i in .synth_event
    val_idx = .var_ref_idx[i]
    val = var_ref_vals[val_idx]
```

17.4.3 action_data and the onXXX() handlers

The hist trigger onXXX() actions other than onmatch(), such as onmax() and onchange(), also make use of and internally create hidden variables. This information is contained in the action_data.track_data struct, and is also visible in the hist_debug output as will be described in the example below.

Typically, the onmax() or onchange() handlers are used in conjunction with the save() and snapshot() actions. For example:

```
# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-$ts0: \
      onmax($wakeup_lat).save(next_comm,prev_pid,prev_prio,prev_comm)' >>
      /sys/kernel/debug/tracing/events/sched/sched_switch/trigger
```

or:

```
# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-$ts0: \
      onmax($wakeup_lat).snapshot()' >>
      /sys/kernel/debug/tracing/events/sched/sched_switch/trigger
```

17.4.4 save() action field variable test

For this example, instead of generating a synthetic event, the save() action is used to save field values whenever an onmax() handler detects that a new max latency has been hit. As in the previous example, the values being saved are also field values, but in this case, are kept in a separate hist_data array named save_vars[].

As in previous test examples, we set up the sched_waking trigger:

```
# echo 'hist:keys=pid:ts0=common_timestamp.usecs' >> events/sched/sched_waking/
↪trigger
```

In this case, however, we set up the sched_switch trigger to save some sched_switch field values whenever we hit a new maximum latency. For both the onmax() handler and save() action, variables will be created, which we can use the hist_debug files to examine:

```
# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-$ts0:onmax(
↪$wakeup_lat).save(next_comm,prev_pid,prev_prio,prev_comm)' >> events/sched/
↪sched_switch/trigger
```

The sched_waking hist_debug output shows the same data as in the previous test examples:

```
# cat events/sched/sched_waking/hist_debug
#
# trigger info: hist:keys=pid:vals=hitcount:ts0=common_timestamp.
↪usecs:sort=hitcount:size=2048:clock=global [active]
#
```

```

hist_data: 00000000e6290f48

n_vals: 2
n_keys: 1
n_fields: 3

val fields:

  hist_data->fields[0]:
    flags:
      VAL: HIST_FIELD_FL_HITCOUNT
    type: u64
    size: 8
    is_signed: 0

  hist_data->fields[1]:
    flags:
      HIST_FIELD_FL_VAR
    var.name: ts0
    var.idx (into tracing_map_elt.vars[]): 0
    type: u64
    size: 8
    is_signed: 0

key fields:

  hist_data->fields[2]:
    flags:
      HIST_FIELD_FL_KEY
    ftrace_event_field name: pid
    type: pid_t
    size: 8
    is_signed: 1

```

The output of the `sched_switch` trigger shows the same val and key values as before, but also shows a couple new sections.

First, the action tracking variables section now shows the `actions[].track_data` information describing the special tracking variables and references used to track, in this case, the running maximum value. The `actions[].track_data.var_ref` member contains the reference to the variable being tracked, in this case the `$wakeup_lat` variable. In order to perform the `onmax()` handler function, there also needs to be a variable that tracks the current maximum by getting updated whenever a new maximum is hit. In this case, we can see that an auto-generated variable named `'__max'` has been created and is visible in the `actions[].track_data.track_var` variable.

Finally, in the new 'save action variables' section, we can see that the 4 params to the `save()` function have resulted in 4 field variables being created for the purposes of saving the values of the named fields when the max is hit. These variables are kept in a separate `save_vars[]` array off of `hist_data`, so are displayed in a separate section:

```
# cat events/sched/sched_switch/hist_debug

# event histogram
#
# trigger info: hist:keys=next_pid:vals=hitcount:wakeup_lat=common_timestamp.
→usecs-$ts0:sort=hitcount:size=2048:clock=global:onmax($wakeup_lat).save(next_
→comm,prev_pid,prev_prio,prev_comm) [active]
#

hist_data: 0000000057bcd28d

n_vals: 2
n_keys: 1
n_fields: 3

val fields:

  hist_data->fields[0]:
    flags:
      VAL: HIST_FIELD_FL_HITCOUNT
    type: u64
    size: 8
    is_signed: 0

  hist_data->fields[1]:
    flags:
      HIST_FIELD_FL_VAR
    var.name: wakeup_lat
    var.idx (into tracing_map_elt.vars[]): 0
    type: u64
    size: 0
    is_signed: 0

key fields:

  hist_data->fields[2]:
    flags:
      HIST_FIELD_FL_KEY
    ftrace_event_field name: next_pid
    type: pid_t
    size: 8
    is_signed: 1

variable reference fields:

  hist_data->var_refs[0]:
    flags:
      HIST_FIELD_FL_VAR_REF
    name: ts0
    var.idx (into tracing_map_elt.vars[]): 0
```

```

var.hist_data: 00000000e6290f48
var_ref_idx (into hist_data->var_refs[]): 0
type: u64
size: 8
is_signed: 0

hist_data->var_refs[1]:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: wakeup_lat
  var.idx (into tracing_map_elt.vars[]): 0
  var.hist_data: 0000000057bcd28d
  var_ref_idx (into hist_data->var_refs[]): 1
  type: u64
  size: 0
  is_signed: 0

action tracking variables (for onmax()/onchange()/onmatch()):

hist_data->actions[0].track_data.var_ref:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: wakeup_lat
  var.idx (into tracing_map_elt.vars[]): 0
  var.hist_data: 0000000057bcd28d
  var_ref_idx (into hist_data->var_refs[]): 1
  type: u64
  size: 0
  is_signed: 0

hist_data->actions[0].track_data.track_var:
  flags:
    HIST_FIELD_FL_VAR
  var.name: __max
  var.idx (into tracing_map_elt.vars[]): 1
  type: u64
  size: 8
  is_signed: 0

save action variables (save() params):

hist_data->save_vars[0]:

  save_vars[0].var:
  flags:
    HIST_FIELD_FL_VAR
  var.name: next_comm
  var.idx (into tracing_map_elt.vars[]): 2

  save_vars[0].val:

```

```
ftrace_event_field name: next_comm
type: char[16]
size: 256
is_signed: 0

hist_data->save_vars[1]:

save_vars[1].var:
flags:
    HIST_FIELD_FL_VAR
var.name: prev_pid
var.idx (into tracing_map_elt.vars[]): 3

save_vars[1].val:
ftrace_event_field name: prev_pid
type: pid_t
size: 4
is_signed: 1

hist_data->save_vars[2]:

save_vars[2].var:
flags:
    HIST_FIELD_FL_VAR
var.name: prev_prio
var.idx (into tracing_map_elt.vars[]): 4

save_vars[2].val:
ftrace_event_field name: prev_prio
type: int
size: 4
is_signed: 1

hist_data->save_vars[3]:

save_vars[3].var:
flags:
    HIST_FIELD_FL_VAR
var.name: prev_comm
var.idx (into tracing_map_elt.vars[]): 5

save_vars[3].val:
ftrace_event_field name: prev_comm
type: char[16]
size: 256
is_signed: 0
```

The commands below can be used to clean things up for the next test:

```
# echo '!hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-$ts0:onmax(
↪$wakeup_lat).save(next_comm,prev_pid,prev_prio,prev_comm)' >> events/sched/
↪sched_switch/trigger
```

```
# echo '!hist:keys=pid:ts0=common_timestamp.usecs' >> events/sched/sched_
↳waking/trigger
```

17.5 A couple special cases

While the above covers the basics of the histogram internals, there are a couple of special cases that should be discussed, since they tend to create even more confusion. Those are field variables on other histograms, and aliases, both described below through example tests using the `hist_debug` files.

17.5.1 Test of field variables on other histograms

This example is similar to the previous examples, but in this case, the `sched_switch` trigger references a `hist` trigger field on another event, namely the `sched_waking` event. In order to accomplish this, a field variable is created for the other event, but since an existing histogram can't be used, as existing histograms are immutable, a new histogram with a matching variable is created and used, and we'll see that reflected in the `hist_debug` output shown below.

First, we create the `wakeup_latency` synthetic event. Note the addition of the `prio` field:

```
# echo 'wakeup_latency u64 lat; pid_t pid; int prio' >> synthetic_events
```

As in previous test examples, we set up the `sched_waking` trigger:

```
# echo 'hist:keys=pid:ts0=common_timestamp.usecs' >> events/sched/sched_waking/
↳trigger
```

Here we set up a `hist` trigger on `sched_switch` to send a `wakeup_latency` event using an on-match handler naming the `sched_waking` event. Note that the third param being passed to the `wakeup_latency()` is `prio`, which is a field name that needs to have a field variable created for it. There isn't however any `prio` field on the `sched_switch` event so it would seem that it wouldn't be possible to create a field variable for it. The matching `sched_waking` event does have a `prio` field, so it should be possible to make use of it for this purpose. The problem with that is that it's not currently possible to define a new variable on an existing histogram, so it's not possible to add a new `prio` field variable to the existing `sched_waking` histogram. It is however possible to create an additional new 'matching' `sched_waking` histogram for the same event, meaning that it uses the same key and filters, and define the new `prio` field variable on that.

Here's the `sched_switch` trigger:

```
# echo 'hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-
↳$ts0:onmatch(sched.sched_waking).wakeup_latency($wakeup_lat,next_pid,prio)' >
↳events/sched/sched_switch/trigger
```

And here's the output of the `hist_debug` information for the `sched_waking` `hist` trigger. Note that there are two histograms displayed in the output: the first is the normal `sched_waking` histogram we've seen in the previous examples, and the second is the special histogram we created to provide the `prio` field variable.

Looking at the second histogram below, we see a variable with the name `synthetic_prio`. This is the field variable created for the `prio` field on that `sched_waking` histogram:

```
# cat events/sched/sched_waking/hist_debug

# event histogram
#
# trigger info: hist:keys=pid:vals=hitcount:ts0=common_timestamp.
→usecs:sort=hitcount:size=2048:clock=global [active]
#

hist_data: 00000000349570e4

n_vals: 2
n_keys: 1
n_fields: 3

val fields:

  hist_data->fields[0]:
    flags:
      VAL: HIST_FIELD_FL_HITCOUNT
    type: u64
    size: 8
    is_signed: 0

  hist_data->fields[1]:
    flags:
      HIST_FIELD_FL_VAR
    var.name: ts0
    var.idx (into tracing_map_elt.vars[]): 0
    type: u64
    size: 8
    is_signed: 0

key fields:

  hist_data->fields[2]:
    flags:
      HIST_FIELD_FL_KEY
    ftrace_event_field name: pid
    type: pid_t
    size: 8
    is_signed: 1

# event histogram
#
# trigger info: hist:keys=pid:vals=hitcount:synthetic_
→prio=prio:sort=hitcount:size=2048 [active]
#
```



```

hist_data: 000000006920cf38

n_vals: 2
n_keys: 1
n_fields: 3

val fields:

  hist_data->fields[0]:
    flags:
      VAL: HIST_FIELD_FL_HITCOUNT
    type: u64
    size: 8
    is_signed: 0

  hist_data->fields[1]:
    flags:
      HIST_FIELD_FL_VAR
    ftrace_event_field name: prio
    var.name: synthetic_prio
    var.idx (into tracing_map_elt.vars[]): 0
    type: int
    size: 4
    is_signed: 1

key fields:

  hist_data->fields[2]:
    flags:
      HIST_FIELD_FL_KEY
    ftrace_event_field name: pid
    type: pid_t
    size: 8
    is_signed: 1

```

Looking at the sched_switch histogram below, we can see a reference to the synthetic_prio variable on sched_waking, and looking at the associated hist_data address we see that it is indeed associated with the new histogram. Note also that the other references are to a normal variable, wakeup_lat, and to a normal field variable, next_pid, the details of which are in the field variables section:

```

# cat events/sched/sched_switch/hist_debug

# event histogram
#
# trigger info: hist:keys=next_pid:vals=hitcount:wakeup_lat=common_timestamp.
→usecs-$ts0:sort=hitcount:size=2048:clock=global:onmatch(sched.sched_waking).
→wakeup_latency($wakeup_lat,next_pid,prio) [active]
#

```

```
hist_data: 00000000a73b67df

n_vals: 2
n_keys: 1
n_fields: 3

val fields:

  hist_data->fields[0]:
    flags:
      VAL: HIST_FIELD_FL_HITCOUNT
    type: u64
    size: 8
    is_signed: 0

  hist_data->fields[1]:
    flags:
      HIST_FIELD_FL_VAR
    var.name: wakeup_lat
    var.idx (into tracing_map_elt.vars[]): 0
    type: u64
    size: 0
    is_signed: 0

key fields:

  hist_data->fields[2]:
    flags:
      HIST_FIELD_FL_KEY
    ftrace_event_field name: next_pid
    type: pid_t
    size: 8
    is_signed: 1

variable reference fields:

  hist_data->var_refs[0]:
    flags:
      HIST_FIELD_FL_VAR_REF
    name: ts0
    var.idx (into tracing_map_elt.vars[]): 0
    var.hist_data: 00000000349570e4
    var_ref_idx (into hist_data->var_refs[]): 0
    type: u64
    size: 8
    is_signed: 0

  hist_data->var_refs[1]:
    flags:
      HIST_FIELD_FL_VAR_REF
```

```

name: wakeup_lat
var.idx (into tracing_map_elt.vars[]): 0
var.hist_data: 00000000a73b67df
var_ref_idx (into hist_data->var_refs[]): 1
type: u64
size: 0
is_signed: 0

hist_data->var_refs[2]:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: next_pid
  var.idx (into tracing_map_elt.vars[]): 1
  var.hist_data: 00000000a73b67df
  var_ref_idx (into hist_data->var_refs[]): 2
  type: pid_t
  size: 4
  is_signed: 0

hist_data->var_refs[3]:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: synthetic_prio
  var.idx (into tracing_map_elt.vars[]): 0
  var.hist_data: 000000006920cf38
  var_ref_idx (into hist_data->var_refs[]): 3
  type: int
  size: 4
  is_signed: 1

field variables:

hist_data->field_vars[0]:

  field_vars[0].var:
  flags:
    HIST_FIELD_FL_VAR
  var.name: next_pid
  var.idx (into tracing_map_elt.vars[]): 1

  field_vars[0].val:
  ftrace_event_field name: next_pid
  type: pid_t
  size: 4
  is_signed: 1

action tracking variables (for onmax()/onchange()/onmatch()):

hist_data->actions[0].match_data.event_system: sched
hist_data->actions[0].match_data.event: sched_waking

```

The commands below can be used to clean things up for the next test:

```
# echo '!hist:keys=next_pid:wakeup_lat=common_timestamp.usecs-  
→$ts0:onmatch(sched.sched_waking).wakeup_latency($wakeup_lat,next_pid,prio)' >  
→ events/sched/sched_switch/trigger  
  
# echo '!hist:keys=pid:ts0=common_timestamp.usecs' >> events/sched/sched_  
→waking/trigger  
  
# echo '!wakeup_latency u64 lat; pid_t pid; int prio' >> synthetic_events
```

17.5.2 Alias test

This example is very similar to previous examples, but demonstrates the alias flag.

First, we create the wakeup_latency synthetic event:

```
# echo 'wakeup_latency u64 lat; pid_t pid; char comm[16]' >> synthetic_events
```

Next, we create a sched_waking trigger similar to previous examples, but in this case we save the pid in the waking_pid variable:

```
# echo 'hist:keys=pid:waking_pid=pid:ts0=common_timestamp.usecs' >> events/  
→sched/sched_waking/trigger
```

For the sched_switch trigger, instead of using \$waking_pid directly in the wakeup_latency synthetic event invocation, we create an alias of \$waking_pid named \$woken_pid, and use that in the synthetic event invocation instead:

```
# echo 'hist:keys=next_pid:woken_pid=$waking_pid:wakeup_lat=common_timestamp.  
→usecs-$ts0:onmatch(sched.sched_waking).wakeup_latency($wakeup_lat,$woken_pid,  
→next_comm)' >> events/sched/sched_switch/trigger
```

Looking at the sched_waking hist_debug output, in addition to the normal fields, we can see the waking_pid variable:

```
# cat events/sched/sched_waking/hist_debug  
  
# event histogram  
#  
# trigger info: hist:keys=pid:vals=hitcount:waking_pid=pid,ts0=common_  
→timestamp.usecs:sort=hitcount:size=2048:clock=global [active]  
#  
  
hist_data: 00000000a250528c  
  
n_vals: 3  
n_keys: 1  
n_fields: 4  
  
val fields:
```

```

hist_data->fields[0]:
  flags:
    VAL: HIST_FIELD_FL_HITCOUNT
  type: u64
  size: 8
  is_signed: 0

hist_data->fields[1]:
  flags:
    HIST_FIELD_FL_VAR
  ftrace_event_field name: pid
  var.name: waking_pid
  var.idx (into tracing_map_elt.vars[]): 0
  type: pid_t
  size: 4
  is_signed: 1

hist_data->fields[2]:
  flags:
    HIST_FIELD_FL_VAR
  var.name: ts0
  var.idx (into tracing_map_elt.vars[]): 1
  type: u64
  size: 8
  is_signed: 0

key fields:

hist_data->fields[3]:
  flags:
    HIST_FIELD_FL_KEY
  ftrace_event_field name: pid
  type: pid_t
  size: 8
  is_signed: 1

```

The sched_switch hist_debug output shows that a variable named woken_pid has been created but that it also has the HIST_FIELD_FL_ALIAS flag set. It also has the HIST_FIELD_FL_VAR flag set, which is why it appears in the val field section.

Despite that implementation detail, an alias variable is actually more like a variable reference; in fact it can be thought of as a reference to a reference. The implementation copies the var_ref->fn() from the variable reference being referenced, in this case, the waking_pid fn(), which is hist_field_var_ref() and makes that the fn() of the alias. The hist_field_var_ref() fn() requires the var_ref_idx of the variable reference it's using, so waking_pid's var_ref_idx is also copied to the alias. The end result is that when the value of alias is retrieved, in the end it just does the same thing the original reference would have done and retrieves the same value from the var_ref_vals[] array. You can verify this in the output by noting that the var_ref_idx of the alias, in this case woken_pid, is the same as the var_ref_idx of the reference, waking_pid, in the variable reference fields section.

Additionally, once it gets that value, since it is also a variable, it then saves that value into its

var.idx. So the var.idx of the woken_pid alias is 0, which it fills with the value from var_ref_idx 0 when its fn() is called to update itself. You'll also notice that there's a woken_pid var_ref in the variable refs section. That is the reference to the woken_pid alias variable, and you can see that it retrieves the value from the same var.idx as the woken_pid alias, 0, and then in turn saves that value in its own var_ref_idx slot, 3, and the value at this position is finally what gets assigned to the \$woken_pid slot in the trace event invocation:

```
# cat events/sched/sched_switch/hist_debug

# event histogram
#
# trigger info: hist:keys=next_pid:vals=hitcount:woken_pid=$waking_pid,wakeup_
→lat=common_timestamp.usecs-
→$ts0:sort=hitcount:size=2048:clock=global:onmatch(sched.sched_waking).wakeup_
→latency($wakeup_lat,$woken_pid,next_comm) [active]
#

hist_data: 0000000055d65ed0

n_vals: 3
n_keys: 1
n_fields: 4

val fields:

hist_data->fields[0]:
  flags:
    VAL: HIST_FIELD_FL_HITCOUNT
  type: u64
  size: 8
  is_signed: 0

hist_data->fields[1]:
  flags:
    HIST_FIELD_FL_VAR
    HIST_FIELD_FL_ALIAS
  var.name: woken_pid
  var.idx (into tracing_map_elt.vars[]): 0
  var_ref_idx (into hist_data->var_refs[]): 0
  type: pid_t
  size: 4
  is_signed: 1

hist_data->fields[2]:
  flags:
    HIST_FIELD_FL_VAR
  var.name: wakeup_lat
  var.idx (into tracing_map_elt.vars[]): 1
  type: u64
  size: 0
  is_signed: 0
```

key fields:

```
hist_data->fields[3]:
  flags:
    HIST_FIELD_FL_KEY
  ftrace_event_field name: next_pid
  type: pid_t
  size: 8
  is_signed: 1
```

variable reference fields:

```
hist_data->var_refs[0]:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: waking_pid
  var.idx (into tracing_map_elt.vars[]): 0
  var.hist_data: 00000000a250528c
  var_ref_idx (into hist_data->var_refs[]): 0
  type: pid_t
  size: 4
  is_signed: 1
```

```
hist_data->var_refs[1]:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: ts0
  var.idx (into tracing_map_elt.vars[]): 1
  var.hist_data: 00000000a250528c
  var_ref_idx (into hist_data->var_refs[]): 1
  type: u64
  size: 8
  is_signed: 0
```

```
hist_data->var_refs[2]:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: wakeup_lat
  var.idx (into tracing_map_elt.vars[]): 1
  var.hist_data: 0000000055d65ed0
  var_ref_idx (into hist_data->var_refs[]): 2
  type: u64
  size: 0
  is_signed: 0
```

```
hist_data->var_refs[3]:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: woken_pid
```

```
var.idx (into tracing_map_elt.vars[]): 0
var.hist_data: 0000000055d65ed0
var_ref_idx (into hist_data->var_refs[]): 3
type: pid_t
size: 4
is_signed: 1

hist_data->var_refs[4]:
  flags:
    HIST_FIELD_FL_VAR_REF
  name: next_comm
  var.idx (into tracing_map_elt.vars[]): 2
  var.hist_data: 0000000055d65ed0
  var_ref_idx (into hist_data->var_refs[]): 4
  type: char[16]
  size: 256
  is_signed: 0

field variables:

hist_data->field_vars[0]:

  field_vars[0].var:
  flags:
    HIST_FIELD_FL_VAR
  var.name: next_comm
  var.idx (into tracing_map_elt.vars[]): 2

  field_vars[0].val:
  ftrace_event_field name: next_comm
  type: char[16]
  size: 256
  is_signed: 0

action tracking variables (for onmax()/onchange()/onmatch()):

hist_data->actions[0].match_data.event_system: sched
hist_data->actions[0].match_data.event: sched_waking
```

The commands below can be used to clean things up for the next test:

```
# echo '!hist:keys=next_pid:woken_pid=$waking_pid:wakeup_lat=common_timestamp.
↪usecs-$ts0:onmatch(sched.sched_waking).wakeup_latency($wakeup_lat,$woken_pid,
↪next_comm)' >> events/sched/sched_switch/trigger

# echo '!hist:keys=pid:ts0=common_timestamp.usecs' >> events/sched/sched_
↪waking/trigger

# echo '!wakeup_latency u64 lat; pid_t pid; char comm[16]' >> synthetic_events
```


BOOT-TIME TRACING

Author Masami Hiramatsu <mhiramat@kernel.org>

18.1 Overview

Boot-time tracing allows users to trace boot-time process including device initialization with full features of ftrace including per-event filter and actions, histograms, kprobe-events and synthetic-events, and trace instances. Since kernel command line is not enough to control these complex features, this uses bootconfig file to describe tracing feature programming.

18.2 Options in the Boot Config

Here is the list of available options list for boot time tracing in boot config file¹. All options are under “ftrace.” or “kernel.” prefix. See kernel parameters for the options which starts with “kernel.” prefix².

18.2.1 Ftrace Global Options

Ftrace global options have “kernel.” prefix in boot config, which means these options are passed as a part of kernel legacy command line.

kernel.tp_printk Output trace-event data on printk buffer too.

kernel.dump_on_oops [= MODE] Dump ftrace on Oops. If MODE = 1 or omitted, dump trace buffer on all CPUs. If MODE = 2, dump a buffer on a CPU which kicks Oops.

kernel.traceoff_on_warning Stop tracing if WARN_ON() occurs.

kernel.fgraph_max_depth = MAX_DEPTH Set MAX_DEPTH to maximum depth of fgraph tracer.

kernel.fgraph_filters = FILTER[, FILTER2...] Add fgraph tracing function filters.

kernel.fgraph_notraces = FILTER[, FILTER2...] Add fgraph non-tracing function filters.

¹ See Documentation/admin-guide/bootconfig.rst

² See Documentation/admin-guide/kernel-parameters.rst

18.2.2 Ftrace Per-instance Options

These options can be used for each instance including global ftrace node.

ftrace.[instance.INSTANCE.]options = OPT1[, OPT2[...]] Enable given ftrace options.

ftrace.[instance.INSTANCE.]tracing_on = 0|1 Enable/Disable tracing on this instance when starting boot-time tracing. (you can enable it by the “traceon” event trigger action)

ftrace.[instance.INSTANCE.]trace_clock = CLOCK Set given CLOCK to ftrace’s trace_clock.

ftrace.[instance.INSTANCE.]buffer_size = SIZE Configure ftrace buffer size to SIZE. You can use “KB” or “MB” for that SIZE.

ftrace.[instance.INSTANCE.]alloc_snapshot Allocate snapshot buffer.

ftrace.[instance.INSTANCE.]cpumask = CPUMASK Set CPUMASK as trace cpu-mask.

ftrace.[instance.INSTANCE.]events = EVENT[, EVENT2[...]] Enable given events on boot. You can use a wild card in EVENT.

ftrace.[instance.INSTANCE.]tracer = TRACER Set TRACER to current tracer on boot. (e.g. function)

ftrace.[instance.INSTANCE.]ftrace.filters This will take an array of tracing function filter rules.

ftrace.[instance.INSTANCE.]ftrace.notraces This will take an array of NON-tracing function filter rules.

18.2.3 Ftrace Per-Event Options

These options are setting per-event options.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.enable Enable GROUP:EVENT tracing.

ftrace.[instance.INSTANCE.]event.GROUP.enable Enable all event tracing within GROUP.

ftrace.[instance.INSTANCE.]event.enable Enable all event tracing.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.filter = FILTER Set FILTER rule to the GROUP:EVENT.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.actions = ACTION[, ACTION2[...]]
Set ACTIONS to the GROUP:EVENT.

ftrace.[instance.INSTANCE.]event.kprobes.EVENT.probes = PROBE[, PROBE2[...]]
Defines new kprobe event based on PROBEs. It is able to define multiple probes on one event, but those must have same type of arguments. This option is available only for the event which group name is “kprobes”.

ftrace.[instance.INSTANCE.]event.synthetic.EVENT.fields = FIELD[, FIELD2[...]]
Defines new synthetic event with FIELDS. Each field should be “type varname”.

Note that kprobe and synthetic event definitions can be written under instance node, but those are also visible from other instances. So please take care for event name conflict.

18.2.4 Ftrace Histogram Options

Since it is too long to write a histogram action as a string for per-event action option, there are tree-style options under per-event 'hist' subkey for the histogram actions. For the detail of the each parameter, please read the event histogram document ([Event Histograms](#))

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]keys = KEY1[, KEY2[...]]

Set histogram key parameters. (Mandatory) The 'N' is a digit string for the multiple histogram. You can omit it if there is one histogram on the event.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]values = VAL1[, VAL2[...]]

Set histogram value parameters.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]sort = SORT1[, SORT2[...]]

Set histogram sort parameter options.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]size = NR_ENTRIES Set histogram size (number of entries).

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]name = NAME Set histogram name.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]var.VARIABLE = EXPR

Define a new VARIABLE by EXPR expression.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]<pause|continue|clear>

Set histogram control parameter. You can set one of them.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]onmatch.[M.]event = GROUP.EVENT

Set histogram 'onmatch' handler matching event parameter. The 'M' is a digit string for the multiple 'onmatch' handler. You can omit it if there is one 'onmatch' handler on this histogram.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]onmatch.[M.]trace = EVENT[, ARG1...]

Set histogram 'trace' action for 'onmatch'. EVENT must be a synthetic event name, and ARG1... are parameters for that event. Mandatory if 'onmatch.event' option is set.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]onmax.[M.]var = VAR Set

histogram 'onmax' handler variable parameter.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]onchange.[M.]var = VAR

Set histogram 'onchange' handler variable parameter.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]<onmax|onchange>.[M.]save = A

Set histogram 'save' action parameters for 'onmax' or 'onchange' handler. This option or below 'snapshot' option is mandatory if 'onmax.var' or 'onchange.var' option is set.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.[N.]<onmax|onchange>.[M.]snapsho

Set histogram 'snapshot' action for 'onmax' or 'onchange' handler. This option or above 'save' option is mandatory if 'onmax.var' or 'onchange.var' option is set.

ftrace.[instance.INSTANCE.]event.GROUP.EVENT.hist.filter = FILTER_EXPR Set his-

istogram filter expression. You don't need 'if' in the FILTER_EXPR.

Note that this 'hist' option can conflict with the per-event 'actions' option if the 'actions' option has a histogram action.

18.3 When to Start

All boot-time tracing options starting with `ftrace` will be enabled at the end of `core_initcall`. This means you can trace the events from `postcore_initcall`. Most of the subsystems and architecture dependent drivers will be initialized after that (`arch_initcall` or `subsys_initcall`). Thus, you can trace those with boot-time tracing. If you want to trace events before `core_initcall`, you can use the options starting with `kernel`. Some of them will be enabled earlier than the `initcall` processing (for example, `kernel.ftrace=function` and `kernel.trace_event` will start before the `initcall`.)

18.4 Examples

For example, to add filter and actions for each event, define `kprobe` events, and synthetic events with histogram, write a boot config like below:

```
ftrace.event {
    task.task_newtask {
        filter = "pid < 128"
        enable
    }
    kprobes.vfs_read {
        probes = "vfs_read $arg1 $arg2"
        filter = "common_pid < 200"
        enable
    }
    synthetic.initcall_latency {
        fields = "unsigned long func", "u64 lat"
        hist {
            keys = func.sym, lat
            values = lat
            sort = lat
        }
    }
    initcall.initcall_start.hist {
        keys = func
        var.ts0 = common_timestamp.usecs
    }
    initcall.initcall_finish.hist {
        keys = func
        var.lat = common_timestamp.usecs - $ts0
        onmatch {
            event = initcall.initcall_start
            trace = initcall_latency, func, $lat
        }
    }
}
```

Also, boot-time tracing supports “instance” node, which allows us to run several tracers for different purpose at once. For example, one tracer is for tracing functions starting with “`user_`”, and others tracing “`kernel_`” functions, you can write boot config as below:

```
ftrace.instance {
    foo {
        tracer = "function"
        ftrace.filters = "user_*"
    }
    bar {
        tracer = "function"
        ftrace.filters = "kernel_*"
    }
}
```

The instance node also accepts event nodes so that each instance can customize its event tracing.

With the trigger action and kprobes, you can trace function-graph while a function is called. For example, this will trace all function calls in the `pci_proc_init()`:

```
ftrace {
    tracing_on = 0
    tracer = function_graph
    event.kprobes {
        start_event {
            probes = "pci_proc_init"
            actions = "traceon"
        }
        end_event {
            probes = "pci_proc_init%return"
            actions = "traceoff"
        }
    }
}
```

This boot-time tracing also supports ftrace kernel parameters via boot config. For example, following kernel parameters:

```
trace_options=sym-addr trace_event=initcall:* tp_printk trace_buf_size=1M
↪ ftrace=function ftrace_filter="vfs*"
```

This can be written in boot config like below:

```
kernel {
    trace_options = sym-addr
    trace_event = "initcall:*"
    tp_printk
    trace_buf_size = 1M
    ftrace = function
    ftrace_filter = "vfs*"
}
```

Note that parameters start with “kernel” prefix instead of “ftrace”.

HARDWARE LATENCY DETECTOR

19.1 Introduction

The tracer `hwlat_detector` is a special purpose tracer that is used to detect large system latencies induced by the behavior of certain underlying hardware or firmware, independent of Linux itself. The code was developed originally to detect SMIs (System Management Interrupts) on x86 systems, however there is nothing x86 specific about this patchset. It was originally written for use by the “RT” patch since the Real Time kernel is highly latency sensitive.

SMIs are not serviced by the Linux kernel, which means that it does not even know that they are occurring. SMIs are instead set up by BIOS code and are serviced by BIOS code, usually for “critical” events such as management of thermal sensors and fans. Sometimes though, SMIs are used for other tasks and those tasks can spend an inordinate amount of time in the handler (sometimes measured in milliseconds). Obviously this is a problem if you are trying to keep event service latencies down in the microsecond range.

The hardware latency detector works by hogging one of the cpus for configurable amounts of time (with interrupts disabled), polling the CPU Time Stamp Counter for some period, then looking for gaps in the TSC data. Any gap indicates a time when the polling was interrupted and since the interrupts are disabled, the only thing that could do that would be an SMI or other hardware hiccup (or an NMI, but those can be tracked).

Note that the `hwlat` detector should *NEVER* be used in a production environment. It is intended to be run manually to determine if the hardware platform has a problem with long system firmware service routines.

19.2 Usage

Write the ASCII text “`hwlat`” into the `current_tracer` file of the tracing system (mounted at `/sys/kernel/tracing` or `/sys/kernel/tracing`). It is possible to redefine the threshold in microseconds (us) above which latency spikes will be taken into account.

Example:

```
# echo hwlat > /sys/kernel/tracing/current_tracer
# echo 100 > /sys/kernel/tracing/tracing_thresh
```

The `/sys/kernel/tracing/hwlat_detector` interface contains the following files:

- **width - time period to sample with CPUs held (usecs)** must be less than the total window size (enforced)

- `window` - total period of sampling, width being inside (usecs)

By default the width is set to 500,000 and window to 1,000,000, meaning that for every 1,000,000 usecs (1s) the hwlat detector will spin for 500,000 usecs (0.5s). If `tracing_thresh` contains zero when hwlat tracer is enabled, it will change to a default of 10 usecs. If any latencies that exceed the threshold is observed then the data will be written to the tracing ring buffer.

The minimum sleep time between periods is 1 millisecond. Even if width is less than 1 millisecond apart from window, to allow the system to not be totally starved.

If `tracing_thresh` was zero when hwlat detector was started, it will be set back to zero if another tracer is loaded. Note, the last value in `tracing_thresh` that hwlat detector had will be saved and this value will be restored in `tracing_thresh` if it is still zero when hwlat detector is started again.

The following tracing directory files are used by the `hwlat_detector`:

in `/sys/kernel/tracing`:

- `tracing_threshold` - minimum latency value to be considered (usecs)
- `tracing_max_latency` - maximum hardware latency actually observed (usecs)
- `tracing_cpumask` - the CPUs to move the hwlat thread across
- `hwlat_detector/width` - specified amount of time to spin within window (usecs)
- `hwlat_detector/window` - amount of time between (width) runs (usecs)
- `hwlat_detector/mode` - the thread mode

By default, one hwlat detector's kernel thread will migrate across each CPU specified in `cpumask` at the beginning of a new window, in a round-robin fashion. This behavior can be changed by changing the thread mode, the available options are:

- `none`: do not force migration
- `round-robin`: migrate across each CPU specified in `cpumask` [default]
- `per-cpu`: create one thread for each cpu in `tracing_cpumask`

OSNOISE TRACER

In the context of high-performance computing (HPC), the Operating System Noise (*osnoise*) refers to the interference experienced by an application due to activities inside the operating system. In the context of Linux, NMIs, IRQs, SoftIRQs, and any other system thread can cause noise to the system. Moreover, hardware-related jobs can also cause noise, for example, via SMIs.

`hwlat_detector` is one of the tools used to identify the most complex source of noise: *hardware noise*.

In a nutshell, the `hwlat_detector` creates a thread that runs periodically for a given period. At the beginning of a period, the thread disables interrupt and starts sampling. While running, the `hwlatd` thread reads the time in a loop. As interrupts are disabled, threads, IRQs, and SoftIRQs cannot interfere with the `hwlatd` thread. Hence, the cause of any gap between two different reads of the time roots either on NMI or in the hardware itself. At the end of the period, `hwlatd` enables interrupts and reports the max observed gap between the reads. It also prints a NMI occurrence counter. If the output does not report NMI executions, the user can conclude that the hardware is the culprit for the latency. The `hwlat` detects the NMI execution by observing the entry and exit of a NMI.

The `osnoise` tracer leverages the `hwlat_detector` by running a similar loop with preemption, SoftIRQs and IRQs enabled, thus allowing all the sources of *osnoise* during its execution. Using the same approach of `hwlat`, `osnoise` takes note of the entry and exit point of any source of interferences, increasing a per-cpu interference counter. The `osnoise` tracer also saves an interference counter for each source of interference. The interference counter for NMI, IRQs, SoftIRQs, and threads is increased anytime the tool observes these interferences' entry events. When a noise happens without any interference from the operating system level, the hardware noise counter increases, pointing to a hardware-related noise. In this way, `osnoise` can account for any source of interference. At the end of the period, the `osnoise` tracer prints the sum of all noise, the max single noise, the percentage of CPU available for the thread, and the counters for the noise sources.

20.1 Usage

Write the ASCII text “osnoise” into the `current_tracer` file of the tracing system (generally mounted at `/sys/kernel/tracing`).

For example:

```
[root@f32 ~]# cd /sys/kernel/tracing/
[root@f32 tracing]# echo osnoise > current_tracer
```

It is possible to follow the trace by reading the trace file:

```
[root@f32 tracing]# cat trace
# tracer: osnoise
#
#          _-----=> irqs-off
#         /_-----=> need-resched
#        |/_-----=> hardirq/softirq
#       ||/_--=> preempt-depth
#
#          MAX
#
#          || /
# SINGLE      Interference counters:
#          ||||
#          CPU NOISE +-----+
#          TASK-PID   CPU# ||||  TIMESTAMP    IN US     NOISE    % OF
# AVAILABLE  IN US   HW  NMI   IRQ   SIRQ  THREAD  IN US
#
#           |  |      |    |    |      |      |
#           |  |      |    |    |      |      |
#           <...>-859 [000] .... 81.637220: 1000000    190  99.98100
#           9      18    0   1007    18      1
#           <...>-860 [001] .... 81.638154: 1000000    656  99.93440
#           74     23    0   1006    16      3
#           <...>-861 [002] .... 81.638193: 1000000   5675  99.43250
#           202     6    0   1013    25     21
#           <...>-862 [003] .... 81.638242: 1000000    125  99.98750
#           45     1    0   1011    23      0
#           <...>-863 [004] .... 81.638260: 1000000   1721  99.82790
#           168     7    0   1002    49     41
#           <...>-864 [005] .... 81.638286: 1000000    263  99.97370
#           57     6    0   1006    26      2
#           <...>-865 [006] .... 81.638302: 1000000    109  99.98910
#           21     3    0   1006    18      1
#           <...>-866 [007] .... 81.638326: 1000000   7816  99.21840
#           107     8    0   1016    39     19
```

In addition to the regular trace fields (from TASK-PID to TIMESTAMP), the tracer prints a message at the end of each period for each CPU that is running an osnoise/ thread. The osnoise specific fields report:

- The `RUNTIME IN US` reports the amount of time in microseconds that the `osnoise` thread kept looping reading the time.
- The `NOISE IN US` reports the sum of noise in microseconds observed by the `osnoise` tracer

during the associated runtime.

- The % OF CPU AVAILABLE reports the percentage of CPU available for the osnoise thread during the runtime window.
- The MAX SINGLE NOISE IN US reports the maximum single noise observed during the runtime window.
- The Interference counters display how many each of the respective interference happened during the runtime window.

Note that the example above shows a high number of HW noise samples. The reason being is that this sample was taken on a virtual machine, and the host interference is detected as a hardware interference.

20.2 Tracer options

The tracer has a set of options inside the osnoise directory, they are:

- osnoise/cpus: CPUs at which a osnoise thread will execute.
- osnoise/period_us: the period of the osnoise thread.
- osnoise/runtime_us: how long an osnoise thread will look for noise.
- osnoise/stop_tracing_us: stop the system tracing if a single noise higher than the configured value happens. Writing 0 disables this option.
- osnoise/stop_tracing_total_us: stop the system tracing if total noise higher than the configured value happens. Writing 0 disables this option.
- tracing_threshold: the minimum delta between two time() reads to be considered as noise, in us. When set to 0, the default value will be used, which is currently 5 us.

20.3 Additional Tracing

In addition to the tracer, a set of tracepoints were added to facilitate the identification of the osnoise source.

- osnoise:sample_threshold: printed anytime a noise is higher than the configurable tolerance_ns.
- osnoise:nmi_noise: noise from NMI, including the duration.
- osnoise:irq_noise: noise from an IRQ, including the duration.
- osnoise:softirq_noise: noise from a SoftIRQ, including the duration.
- osnoise:thread_noise: noise from a thread, including the duration.

Note that all the values are *net values*. For example, if while osnoise is running, another thread preempts the osnoise thread, it will start a thread_noise duration at the start. Then, an IRQ takes place, preempting the thread_noise, starting a irq_noise. When the IRQ ends its execution, it will compute its duration, and this duration will be subtracted from the thread_noise, in such a way as to avoid the double accounting of the IRQ execution. This logic is valid for all sources of noise.

Here is one example of the usage of these tracepoints:

```
osnoise/8-961      [008] d.h.  5789.857532: irq_noise: local_timer:236 start_
↪5789.857529929 duration 1845 ns
osnoise/8-961      [008] dNh.  5789.858408: irq_noise: local_timer:236 start_
↪5789.858404871 duration 2848 ns
migration/8-54     [008] d...  5789.858413: thread_noise: migration/8:54_
↪start 5789.858409300 duration 3068 ns
osnoise/8-961      [008] ....  5789.858413: sample_threshold: start 5789.
↪858404555 duration 8812 ns interferences 2
```

In this example, a noise sample of 8 microseconds was reported in the last line, pointing to two interferences. Looking backward in the trace, the two previous entries were about the migration thread running after a timer IRQ execution. The first event is not part of the noise because it took place one millisecond before.

It is worth noticing that the sum of the duration reported in the tracepoints is smaller than eight us reported in the `sample_threshold`. The reason roots in the overhead of the entry and exit code that happens before and after any interference execution. This justifies the dual approach: measuring thread and tracing.

TIMERLAT TRACER

The timerlat tracer aims to help the preemptive kernel developers to find sources of wakeup latencies of real-time threads. Like cyclictst, the tracer sets a periodic timer that wakes up a thread. The thread then computes a *wakeup latency* value as the difference between the *current time* and the *absolute time* that the timer was set to expire. The main goal of timerlat is tracing in such a way to help kernel developers.

21.1 Usage

Write the ASCII text “timerlat” into the `current_tracer` file of the tracing system (generally mounted at `/sys/kernel/tracing`).

For example:

```
[root@f32 ~]# cd /sys/kernel/tracing/
[root@f32 tracing]# echo timerlat > current tracer
```

It is possible to follow the trace by reading the trace trace file:

```
[root@f32 tracing]# cat trace
# tracer: timerlat
#
#          _-----=> irqsoft
#         /_-----=> need-resched
#        |/_-----=> hardirq/softirq
#       ||/_--=> preempt-depth
#      ||/_
#     ||/_
#    ||/_
#   ||/_
#  ||/_
# ||/_
#||/_
#|||_|
#|||_| ACTIVATION
#|||_| ID CONTEXT
# TASK-PID CPU# |||_| TIMESTAMP ID CONTEXT
# LATENCY
# | | | | |
# | | | | |
# | | | | |
# <idle>-0 [000] d.h1 54.029328: #1 context irq timer_
↪ latency 932 ns
# <...>-867 [000] .... 54.029339: #1 context thread timer_
↪ latency 11700 ns
# <idle>-0 [001] dNh1 54.029346: #1 context irq timer_
↪ latency 2833 ns
# <...>-868 [001] .... 54.029353: #1 context thread timer_
↪ latency 9820 ns
```

```

<idle>-0      [000] d.h1    54.030328: #2    context    irq timer_
→latency      769 ns
<...>-867     [000] ....    54.030330: #2    context thread timer_
→latency      3070 ns
<idle>-0      [001] d.h1    54.030344: #2    context    irq timer_
→latency      935 ns
<...>-868     [001] ....    54.030347: #2    context thread timer_
→latency      4351 ns

```

The tracer creates a per-cpu kernel thread with real-time priority that prints two lines at every activation. The first is the *timer latency* observed at the *hardirq* context before the activation of the thread. The second is the *timer latency* observed by the thread. The ACTIVATION ID field serves to relate the *irq* execution to its respective *thread* execution.

The *irq/thread* splitting is important to clarify in which context the unexpected high value is coming from. The *irq* context can be delayed by hardware-related actions, such as SMIs, NMIs, IRQs, or by thread masking interrupts. Once the timer happens, the delay can also be influenced by blocking caused by threads. For example, by postponing the scheduler execution via `preempt_disable()`, scheduler execution, or masking interrupts. Threads can also be delayed by the interference from other threads and IRQs.

21.2 Tracer options

The timerlat tracer is built on top of osnoise tracer. So its configuration is also done in the osnoise/ config directory. The timerlat configs are:

- `cpus`: CPUs at which a timerlat thread will execute.
- `timerlat_period_us`: the period of the timerlat thread.
- `stop_tracing_us`: stop the system tracing if a timer latency at the *irq* context higher than the configured value happens. Writing 0 disables this option.
- `stop_tracing_total_us`: stop the system tracing if a timer latency at the *thread* context is higher than the configured value happens. Writing 0 disables this option.
- `print_stack`: save the stack of the IRQ occurrence. The stack is printed after the *thread context* event, or at the IRQ handler if `stop_tracing_us` is hit.

21.3 timerlat and osnoise

The timerlat can also take advantage of the osnoise: traceevents. For example:

```

[root@f32 ~]# cd /sys/kernel/tracing/
[root@f32 tracing]# echo timerlat > current_tracer
[root@f32 tracing]# echo 1 > events/osnoise/enable
[root@f32 tracing]# echo 25 > osnoise/stop_tracing_total_us
[root@f32 tracing]# tail -10 trace
      cc1-87882    [005] d..h...    548.771078: #402268 context    irq timer_
→latency      13585 ns
      cc1-87882    [005] dNLh1..    548.771082: irq_noise: local_timer:236_
→start 548.771077442 duration 7597 ns


```

```
cc1-87882    [005] dNLh2..    548.771099: irq_noise: qxl:21 start 548.
↪771085017 duration 7139 ns
cc1-87882    [005] d...3..    548.771102: thread_noise:      cc1:87882_
↪start 548.771078243 duration 9909 ns
timerlat/5-1035    [005] .....    548.771104: #402268 context thread timer_
↪latency      39960 ns
```

In this case, the root cause of the timer latency does not point to a single cause but to multiple ones. Firstly, the timer IRQ was delayed for 13 us, which may point to a long IRQ disabled section (see IRQ stacktrace section). Then the timer interrupt that wakes up the timerlat thread took 7597 ns, and the qxl:21 device IRQ took 7139 ns. Finally, the cc1 thread noise took 9909 ns of time before the context switch. Such pieces of evidence are useful for the developer to use other tracing methods to figure out how to debug and optimize the system.

It is worth mentioning that the *duration* values reported by the osnoise: events are *net* values. For example, the thread_noise does not include the duration of the overhead caused by the IRQ execution (which indeed accounted for 12736 ns). But the values reported by the timerlat tracer (timerlat latency) are *gross* values.

The art below illustrates a CPU timeline and how the timerlat tracer observes it at the top and the osnoise: events at the bottom. Each “-” in the timelines means circa 1 us, and the time moves ==>:



21.4 IRQ stacktrace

The `osnoise/print_stack` option is helpful for the cases in which a thread noise causes the major factor for the timer latency, because of preempt or irq disabled. For example:

```
[root@f32 tracing]# echo 500 > osnoise/stop_tracing_total_us
[root@f32 tracing]# echo 500 > osnoise/print_stack
[root@f32 tracing]# echo timerlat > current_tracer
[root@f32 tracing]# tail -21 per_cpu/cpu7/trace
    insmod-1026    [007] dN.h1..    200.201948: irq_noise: local_timer:236_
→start 200.201939376 duration 7872 ns
    insmod-1026    [007] d..h1..    200.202587: #29800 context    irq timer_
→latency    1616 ns
    insmod-1026    [007] dN.h2..    200.202598: irq_noise: local_timer:236_
→start 200.202586162 duration 11855 ns
    insmod-1026    [007] dN.h3..    200.202947: irq_noise: local_timer:236_
→start 200.202939174 duration 7318 ns
    insmod-1026    [007] d...3..    200.203444: thread_noise:    insmod:1026_
→start 200.202586933 duration 838681 ns
    timerlat/7-1001 [007] .....    200.203445: #29800 context thread timer_
→latency    859978 ns
    timerlat/7-1001 [007] ....1..    200.203446: <stack trace>
=> timerlat_irq
=> __hrtimer_run_queues
=> hrtimer_interrupt
=> __sysvec_apic_timer_interrupt
=> asm_call_irq_on_stack
=> sysvec_apic_timer_interrupt
=> asm_sysvec_apic_timer_interrupt
=> delay_tsc
=> dummy_load_1ms_pd_init
=> do_one_initcall
=> do_init_module
=> __do_sys_finit_module
=> do_syscall_64
=> entry_SYSCALL_64_after_hwframe
```

In this case, it is possible to see that the thread added the highest contribution to the *timer latency* and the stack trace, saved during the timerlat IRQ handler, points to a function named `dummy_load_1ms_pd_init`, which had the following code (on purpose):

```
static int __init dummy_load_1ms_pd_init(void)
{
    preempt_disable();
    mdelay(1);
    preempt_enable();
    return 0;
}
```


INTEL(R) TRACE HUB (TH)

22.1 Overview

Intel(R) Trace Hub (TH) is a set of hardware blocks that produce, switch and output trace data from multiple hardware and software sources over several types of trace output ports encoded in System Trace Protocol (MIPI STPv2) and is intended to perform full system debugging. For more information on the hardware, see Intel(R) Trace Hub developer's manual [1].

It consists of trace sources, trace destinations (outputs) and a switch (Global Trace Hub, GTH). These devices are placed on a bus of their own ("intel_th"), where they can be discovered and configured via sysfs attributes.

Currently, the following Intel TH subdevices (blocks) are supported:

- Software Trace Hub (STH), trace source, which is a System Trace Module (STM) device,
- Memory Storage Unit (MSU), trace output, which allows storing trace hub output in system memory,
- Parallel Trace Interface output (PTI), trace output to an external debug host via a PTI port,
- Global Trace Hub (GTH), which is a switch and a central component of Intel(R) Trace Hub architecture.

Common attributes for output devices are described in Documentation/ABI/testing/sysfs-bus-intel_th-output-devices, the most notable of them is "active", which enables or disables trace output into that particular output device.

GTH allows directing different STP masters into different output ports via its "masters" attribute group. More detailed GTH interface description is at Documentation/ABI/testing/sysfs-bus-intel_th-devices-gth.

STH registers an stm class device, through which it provides interface to userspace and kernel-space software trace sources. See [System Trace Module](#) for more information on that.

MSU can be configured to collect trace data into a system memory buffer, which can later on be read from its device nodes via read() or mmap() interface and directed to a "software sink" driver that will consume the data and/or relay it further.

On the whole, Intel(R) Trace Hub does not require any special userspace software to function; everything can be configured, started and collected via sysfs attributes, and device nodes.

[1] <https://software.intel.com/sites/default/files/managed/d3/3c/intel-th-developer-manual.pdf>

22.2 Bus and Subdevices

For each Intel TH device in the system a bus of its own is created and assigned an id number that reflects the order in which TH devices were enumerated. All TH subdevices (devices on intel_th bus) begin with this id: 0-gth, 0-msc0, 0-msc1, 0-pti, 0-sth, which is followed by device's name and an optional index.

Output devices also get a device node in /dev/intel_thN, where N is the Intel TH device id. For example, MSU's memory buffers, when allocated, are accessible via /dev/intel_th0/msc{0,1}.

22.3 Quick example

figure out which GTH port is the first memory controller:

```
$ cat /sys/bus/intel_th/devices/0-msc0/port
0
```

looks like it's port 0, configure master 33 to send data to port 0:

```
$ echo 0 > /sys/bus/intel_th/devices/0-gth/masters/33
```

allocate a 2-windowed multiblock buffer on the first memory # controller, each with 64 pages:

```
$ echo multi > /sys/bus/intel_th/devices/0-msc0/mode
$ echo 64,64 > /sys/bus/intel_th/devices/0-msc0/nr_pages
```

enable wrapping for this controller, too:

```
$ echo 1 > /sys/bus/intel_th/devices/0-msc0/wrap
```

and enable tracing into this port:

```
$ echo 1 > /sys/bus/intel_th/devices/0-msc0/active
```

.. send data to master 33, see [System Trace Module](#) for more details .. # .. wait for traces to pile up .. # .. and stop the trace:

```
$ echo 0 > /sys/bus/intel_th/devices/0-msc0/active
```

and now you can collect the trace from the device node:

```
$ cat /dev/intel_th0/msc0 > my_stp_trace
```

22.4 Host Debugger Mode

It is possible to configure the Trace Hub and control its trace capture from a remote debug host, which should be connected via one of the hardware debugging interfaces, which will then be used to both control Intel Trace Hub and transfer its trace data to the debug host.

The driver needs to be told that such an arrangement is taking place so that it does not touch any capture/port configuration and avoids conflicting with the debug host's configuration accesses. The only activity that the driver will perform in this mode is collecting software traces to the Software Trace Hub (an `stm` class device). The user is still responsible for setting up adequate master/channel mappings that the decoder on the receiving end would recognize.

In order to enable the host mode, set the `'host_mode'` parameter of the `'intel_th'` kernel module to `'y'`. None of the virtual output devices will show up on the `intel_th` bus. Also, trace configuration and capture controlling attribute groups of the `'gth'` device will not be exposed. The `'sth'` device will operate as usual.

22.5 Software Sinks

The Memory Storage Unit (MSU) driver provides an in-kernel API for drivers to register themselves as software sinks for the trace data. Such drivers can further export the data via other devices, such as USB device controllers or network cards.

The API has two main parts::

- notifying the software sink that a particular window is full, and “locking” that window, that is, making it unavailable for the trace collection; when this happens, the MSU driver will automatically switch to the next window in the buffer if it is unlocked, or stop the trace capture if it's not;
- tracking the “locked” state of windows and providing a way for the software sink driver to notify the MSU driver when a window is unlocked and can be used again to collect trace data.

An example sink driver, `msu-sink` illustrates the implementation of a software sink. Functionally, it simply unlocks windows as soon as they are full, keeping the MSU running in a circular buffer mode. Unlike the “multi” mode, it will fill out all the windows in the buffer as opposed to just the first one. It can be enabled by writing “sink” to the “mode” file (assuming `msu-sink.ko` is loaded).

LOCKLESS RING BUFFER DESIGN

Copyright 2009 Red Hat Inc.

Author Steven Rostedt <srostedt@redhat.com>

License The GNU Free Documentation License, Version 1.2 (dual licensed under the GPL v2)

Reviewers Mathieu Desnoyers, Huang Ying, Hidetoshi Seto, and Frederic Weisbecker.

Written for: 2.6.31

23.1 Terminology used in this Document

tail

- where new writes happen in the ring buffer.

head

- where new reads happen in the ring buffer.

producer

- the task that writes into the ring buffer (same as writer)

writer

- same as producer

consumer

- the task that reads from the buffer (same as reader)

reader

- same as consumer.

reader_page

- A page outside the ring buffer used solely (for the most part) by the reader.

head_page

- a pointer to the page that the reader will use next

tail_page

- a pointer to the page that will be written to next

commit_page

- a pointer to the page with the last finished non-nested write.

cmpxchg

- hardware-assisted atomic transaction that performs the following:

```
A = B if previous A == C

R = cmpxchg(A, C, B) is saying that we replace A with B if and only
    if current A is equal to C, and we put the old (current)
    A into R

R gets the previous A regardless if A is updated with B or not.
```

To see if the update was successful a compare of `R == C` may be used.

23.2 The Generic Ring Buffer

The ring buffer can be used in either an overwrite mode or in producer/consumer mode.

Producer/consumer mode is where if the producer were to fill up the buffer before the consumer could free up anything, the producer will stop writing to the buffer. This will lose most recent events.

Overwrite mode is where if the producer were to fill up the buffer before the consumer could free up anything, the producer will overwrite the older data. This will lose the oldest events.

No two writers can write at the same time (on the same per-cpu buffer), but a writer may interrupt another writer, but it must finish writing before the previous writer may continue. This is very important to the algorithm. The writers act like a “stack”. The way interrupts works enforces this behavior:

```
writer1 start
  <preempted> writer2 start
    <preempted> writer3 start
      writer3 finishes
    writer2 finishes
  writer1 finishes
```

This is very much like a writer being preempted by an interrupt and the interrupt doing a write as well.

Readers can happen at any time. But no two readers may run at the same time, nor can a reader preempt/interrupt another reader. A reader cannot preempt/interrupt a writer, but it may read/consume from the buffer at the same time as a writer is writing, but the reader must be on another processor to do so. A reader may read on its own processor and can be preempted by a writer.

A writer can preempt a reader, but a reader cannot preempt a writer. But a reader can read the buffer at the same time (on another processor) as a writer.

The ring buffer is made up of a list of pages held together by a linked list.

At initialization a reader page is allocated for the reader that is not part of the ring buffer.

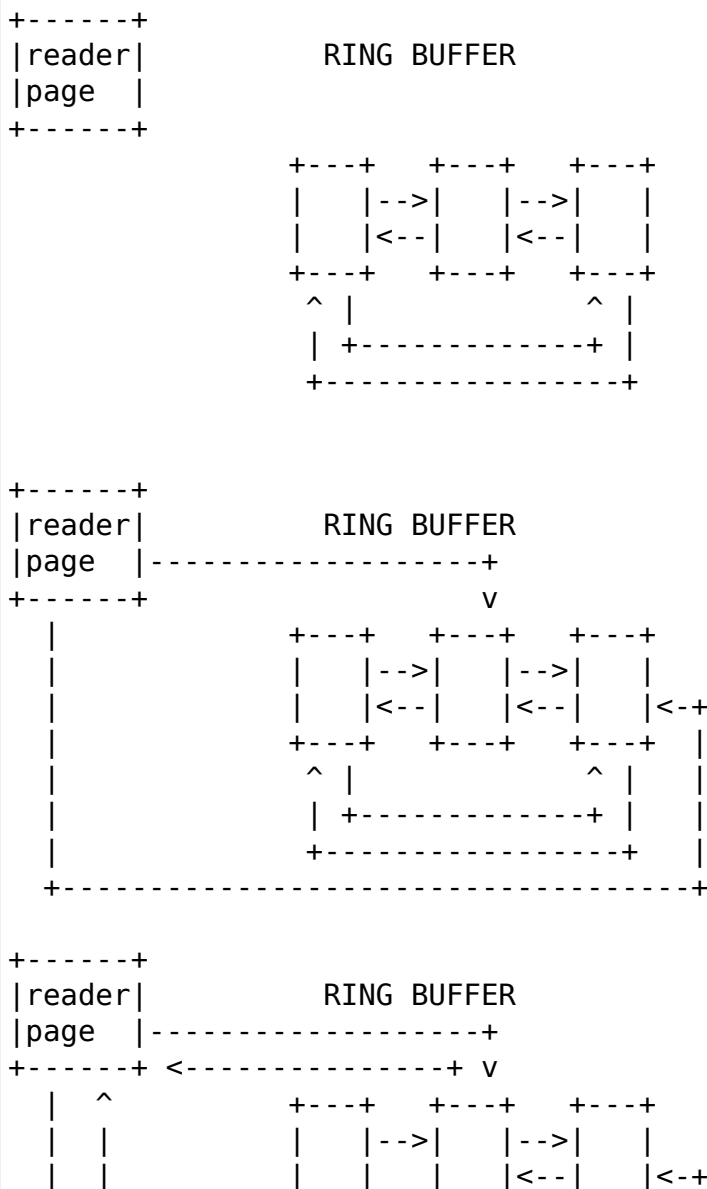
The `head_page`, `tail_page` and `commit_page` are all initialized to point to the same page.

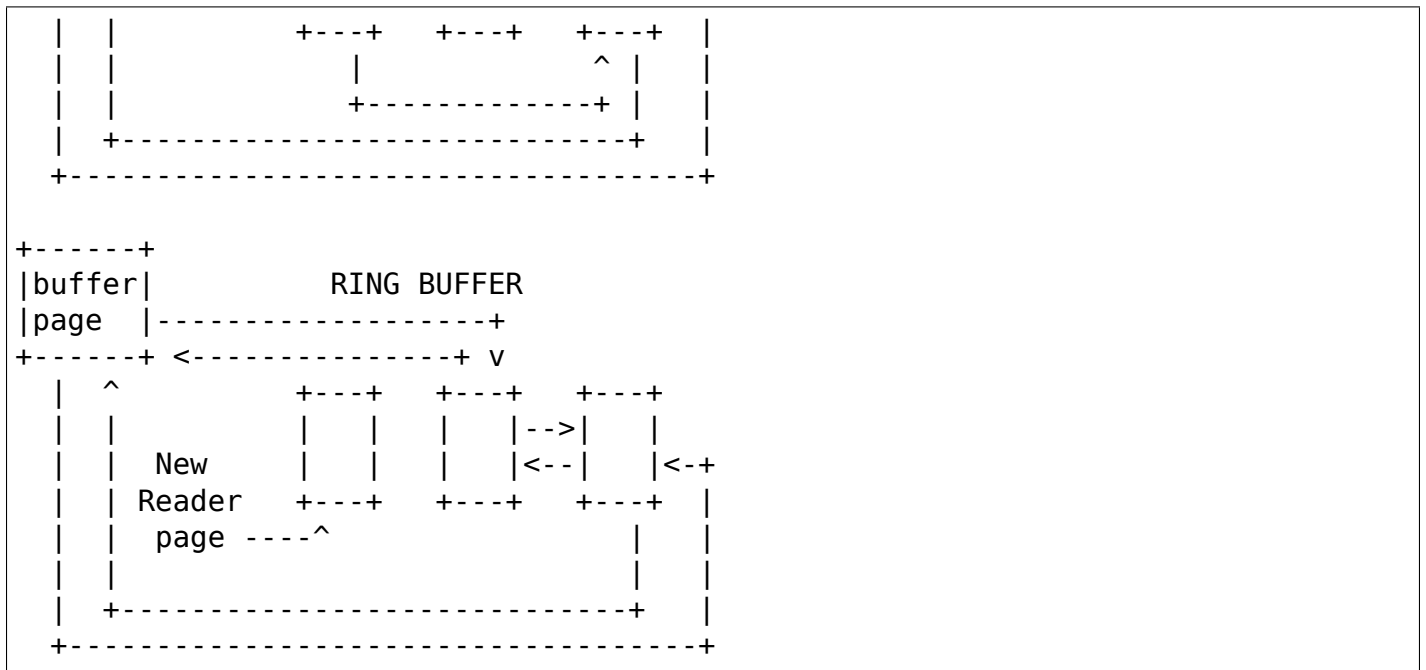
The reader page is initialized to have its next pointer pointing to the head page, and its previous pointer pointing to a page before the head page.

The reader has its own page to use. At start up time, this page is allocated but is not attached to the list. When the reader wants to read from the buffer, if its page is empty (like it is on start-up), it will swap its page with the `head_page`. The old reader page will become part of the ring buffer and the `head_page` will be removed. The page after the inserted page (old reader_page) will become the new head page.

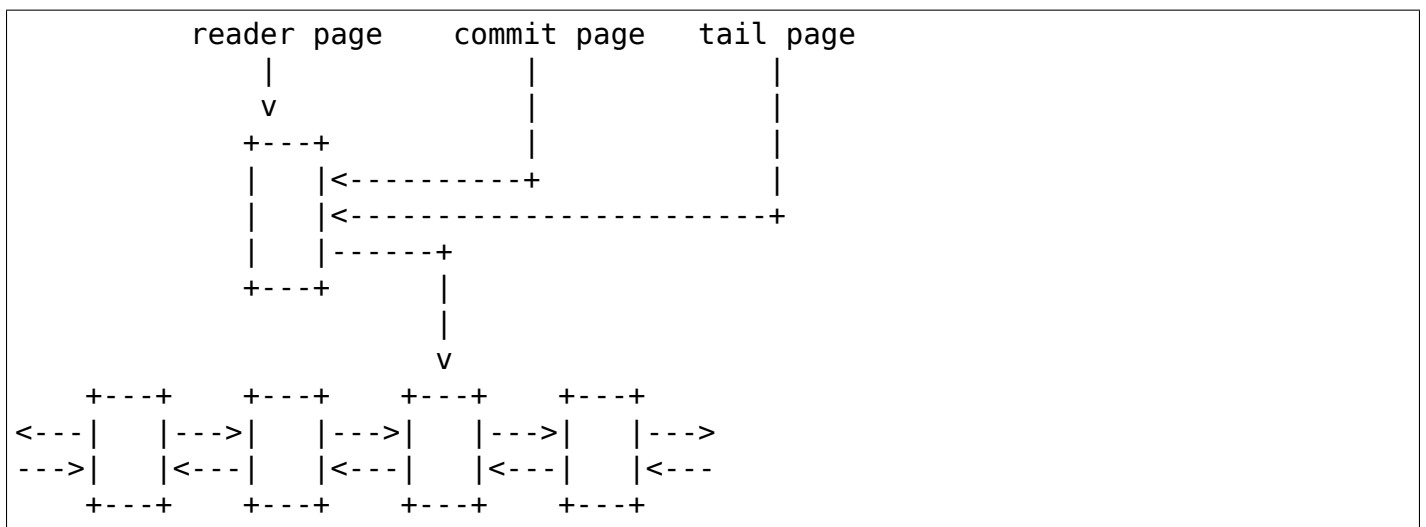
Once the new page is given to the reader, the reader could do what it wants with it, as long as a writer has left that page.

A sample of how the reader page is swapped: Note this does not show the head page in the buffer, it is for demonstrating a swap only.





It is possible that the page swapped is the commit page and the tail page, if what is in the ring buffer is less than what is held in a buffer page.



This case is still valid for this algorithm. When the writer leaves the page, it simply goes into the ring buffer since the reader page still points to the next location in the ring buffer.

The main pointers:

reader page

- The page used solely by the reader and is not part of the ring buffer (may be swapped in)

head page

- the next page in the ring buffer that will be swapped with the reader page.

tail page

- the page where the next write will take place.

commit page

- the page that last finished a write.

The commit page only is updated by the outermost writer in the writer stack. A writer that preempts another writer will not move the commit page.

When data is written into the ring buffer, a position is reserved in the ring buffer and passed back to the writer. When the writer is finished writing data into that position, it commits the write.

Another write (or a read) may take place at anytime during this transaction. If another write happens it must finish before continuing with the previous write.

Write reserve:

```

Buffer page
+-----+
|written |
+-----+ <--- given back to writer (current commit)
|reserved|
+-----+ <--- tail pointer
| empty  |
+-----+
```

Write commit:

```

Buffer page
+-----+
|written |
+-----+
|written |
+-----+ <--- next position for write (current commit)
| empty  |
+-----+
```

If a write happens after the first reserve:

```

Buffer page
+-----+
|written |
+-----+ <-- current commit
|reserved|
+-----+ <--- given back to second writer
|reserved|
+-----+ <--- tail pointer
```

After second writer commits::

```

Buffer page
+-----+
|written |
+-----+ <--(last full commit)
```

```

|reserved |
+-----+
|pending  |
|commit   |
+-----+ <--- tail pointer

```

When the first writer commits::

```

    Buffer page
+-----+
|written  |
+-----+
|written  |
+-----+
|written  |
+-----+ <--(last full commit and tail pointer)

```

The commit pointer points to the last write location that was committed without preempting another write. When a write that preempted another write is committed, it only becomes a pending commit and will not be a full commit until all writes have been committed.

The commit page points to the page that has the last full commit. The tail page points to the page with the last write (before committing).

The tail page is always equal to or after the commit page. It may be several pages ahead. If the tail page catches up to the commit page then no more writes may take place (regardless of the mode of the ring buffer: overwrite and produce/consumer).

The order of pages is:

```

head page
commit page
tail page

```

Possible scenario:

```

                tail page
            commit page |
            |             |
            v             v
+---+ +---+ +---+ +---+
<---| |--->| |--->| |--->| |--->
--->| |<---| |<---| |<---| |<---
+---+ +---+ +---+ +---+

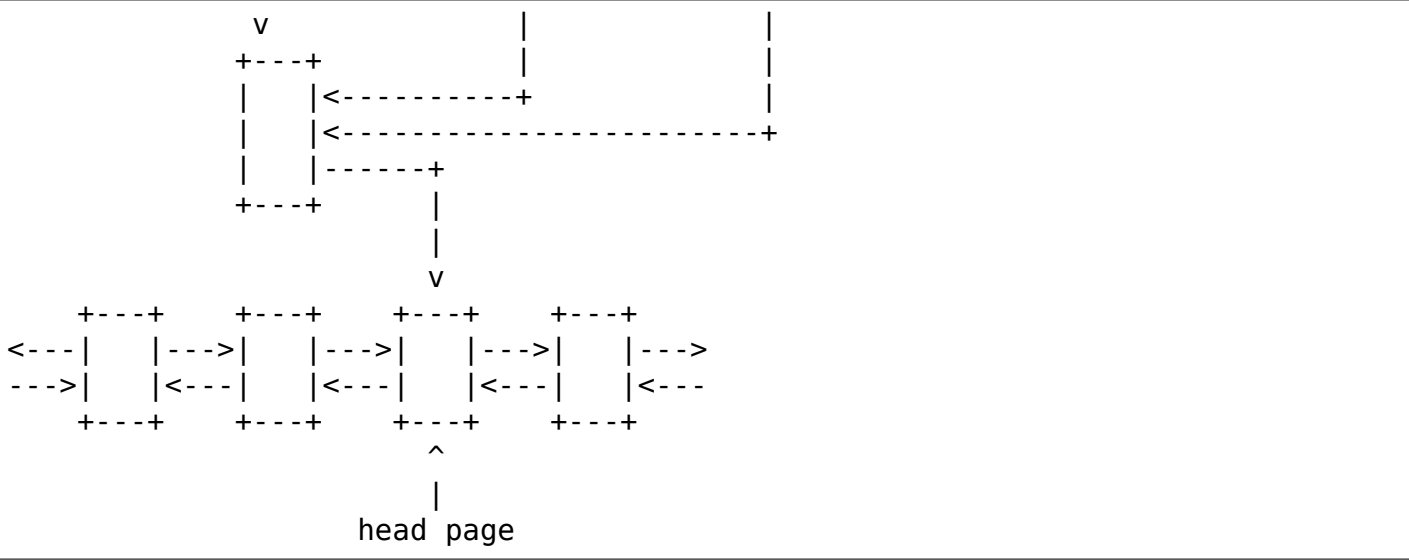
```

There is a special case that the head page is after either the commit page and possibly the tail page. That is when the commit (and tail) page has been swapped with the reader page. This is because the head page is always part of the ring buffer, but the reader page is not. Whenever there has been less than a full page that has been committed inside the ring buffer, and a reader swaps out a page, it will be swapping out the commit page.

```

    reader page    commit page    tail page
        |           |           |

```

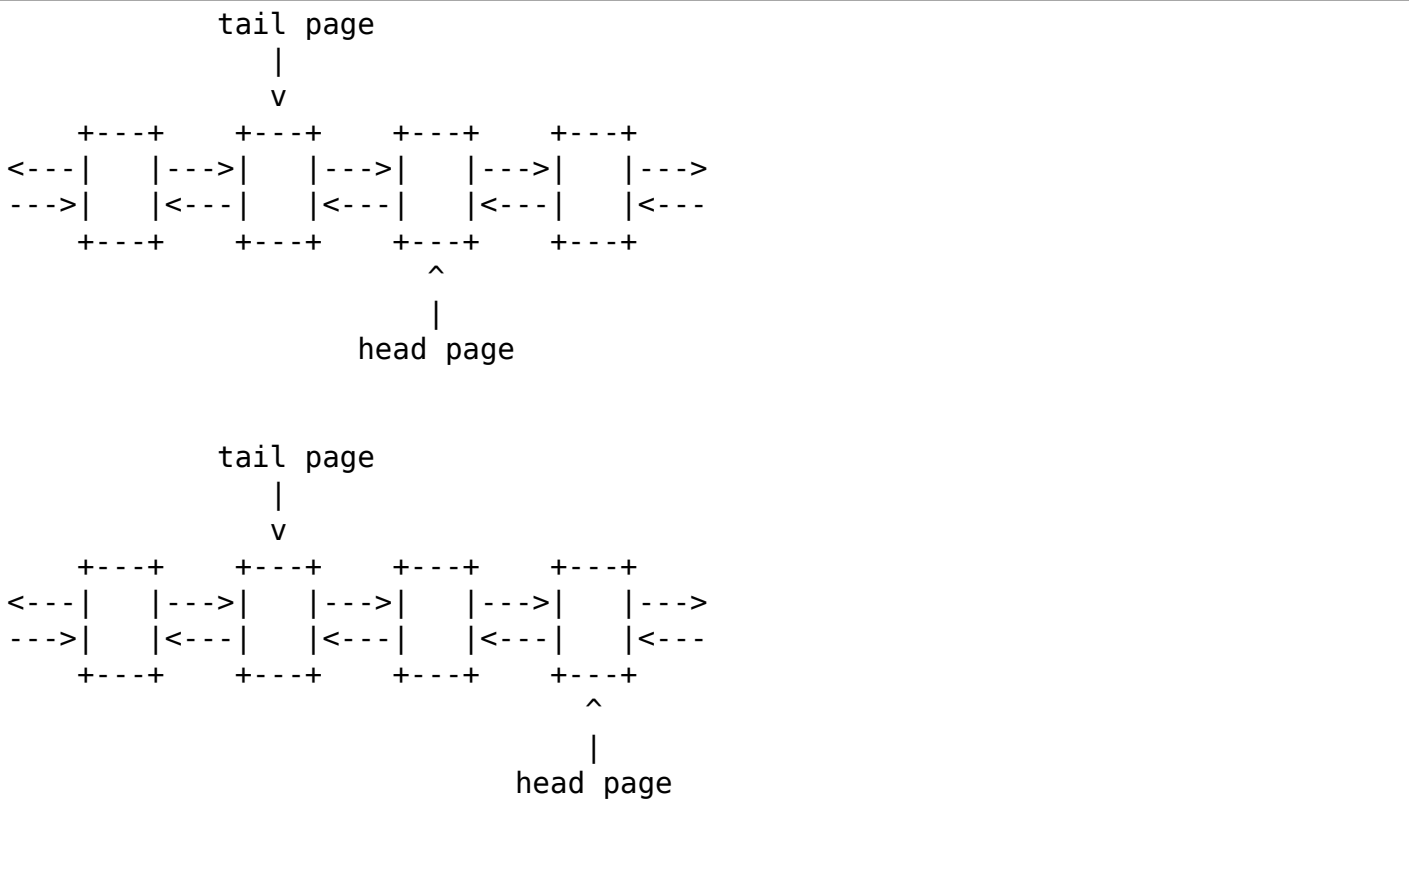



In this case, the head page will not move when the tail and commit move back into the ring buffer.

The reader cannot swap a page into the ring buffer if the commit page is still on that page. If the read meets the last commit (real commit not pending or reserved), then there is nothing more to read. The buffer is considered empty until another full commit finishes.

When the tail meets the head page, if the buffer is in overwrite mode, the head page will be pushed ahead one. If the buffer is in producer/consumer mode, the write will fail.

Overwrite mode:





Note, the reader page will still point to the previous head page. But when a swap takes place, it will use the most recent head page.

23.3 Making the Ring Buffer Lockless:

The main idea behind the lockless algorithm is to combine the moving of the head_page pointer with the swapping of pages with the reader. State flags are placed inside the pointer to the page. To do this, each page must be aligned in memory by 4 bytes. This will allow the 2 least significant bits of the address to be used as flags, since they will always be zero for the address. To get the address, simply mask out the flags:

```
MASK = ~3
```

```
address & MASK
```


Two flags will be kept by these two bits:

HEADER

- the page being pointed to is a head page


UPDATE

- the page being pointed to is being updated by a writer and was or is about to be a head page.



The above pointer “-H->” would have the HEADER flag set. That is the next page is the next page to be swapped out by the reader. This pointer means the next page is the head page.

When the tail page meets the head pointer, it will use `cmpxchg` to change the pointer to the UPDATE state:




“-U->” represents a pointer in the UPDATE state.


Any access to the reader will need to take some sort of lock to serialize the readers. But the writers will never take a lock to write to the ring buffer. This means we only need to worry about a single reader, and writes only preempt in “stack” formation.


When the reader tries to swap the page with the ring buffer, it will also use `cmpxchg`. If the flag bit in the pointer to the head page does not have the HEADER flag set, the compare will fail and the reader will need to look for the new head page and try again. Note, the flags UPDATE and HEADER are never set at the same time.

The reader swaps the reader page as follows:




The reader sets the reader page next pointer as HEADER to the page after the head page:

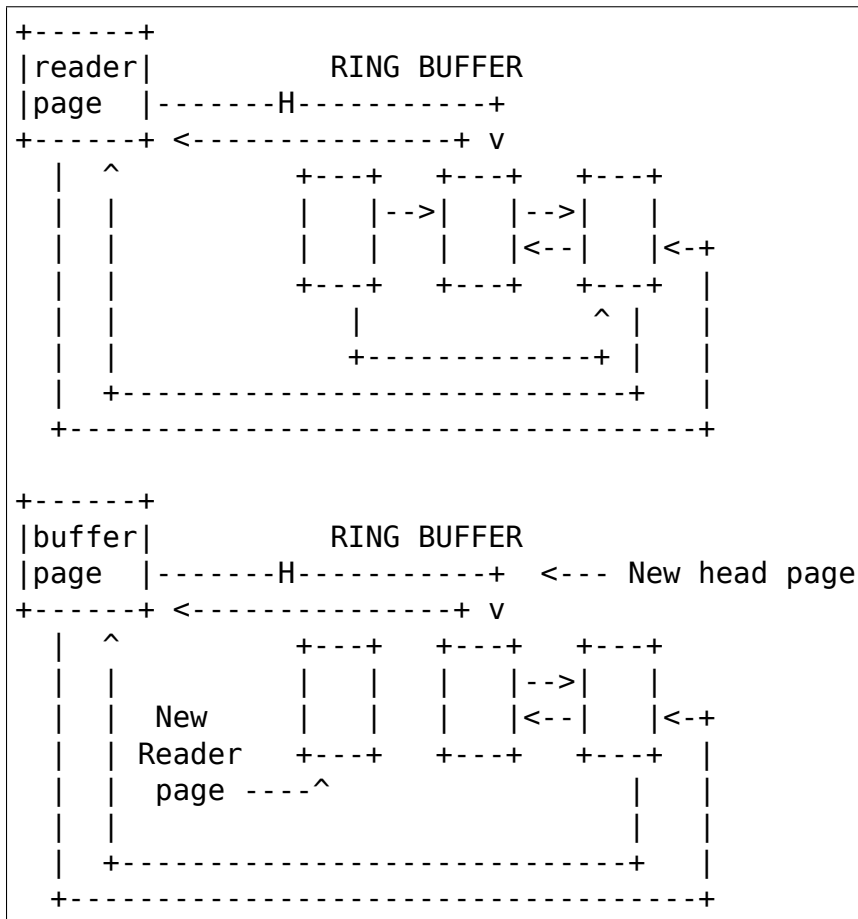




It does a `cmpxchg` with the pointer to the previous head page to make it point to the reader page. Note that the new pointer does not have the HEADER flag set. This action atomically moves the head page forward:




After the new head page is set, the previous pointer of the head page is updated to the reader page:



Another important point: The page that the reader page points back to by its previous pointer (the one that now points to the new head page) never points back to the reader page. That is because the reader page is not part of the ring buffer. Traversing the ring buffer via the next pointers will always stay in the ring buffer. Traversing the ring buffer via the prev pointers may not.


Note, the way to determine a reader page is simply by examining the previous pointer of the page. If the next pointer of the previous page does not point back to the original page, then the original page is a reader page:




The way the head page moves forward:


When the tail page meets the head page and the buffer is in overwrite mode and more writes take place, the head page must be moved forward before the writer may move the tail page. The way this is done is that the writer performs a `cmpxchg` to convert the pointer to the head page from the `HEADER` flag to have the `UPDATE` flag set. Once this is done, the reader will not be able to swap the head page from the buffer, nor will it be able to move the head page, until the writer is finished with the move.

This eliminates any races that the reader can have on the writer. The reader must spin, and this is why the reader cannot preempt the writer:




The following page will be made into the new head page:






After the new head page has been set, we can set the old head page pointer back to NORMAL:

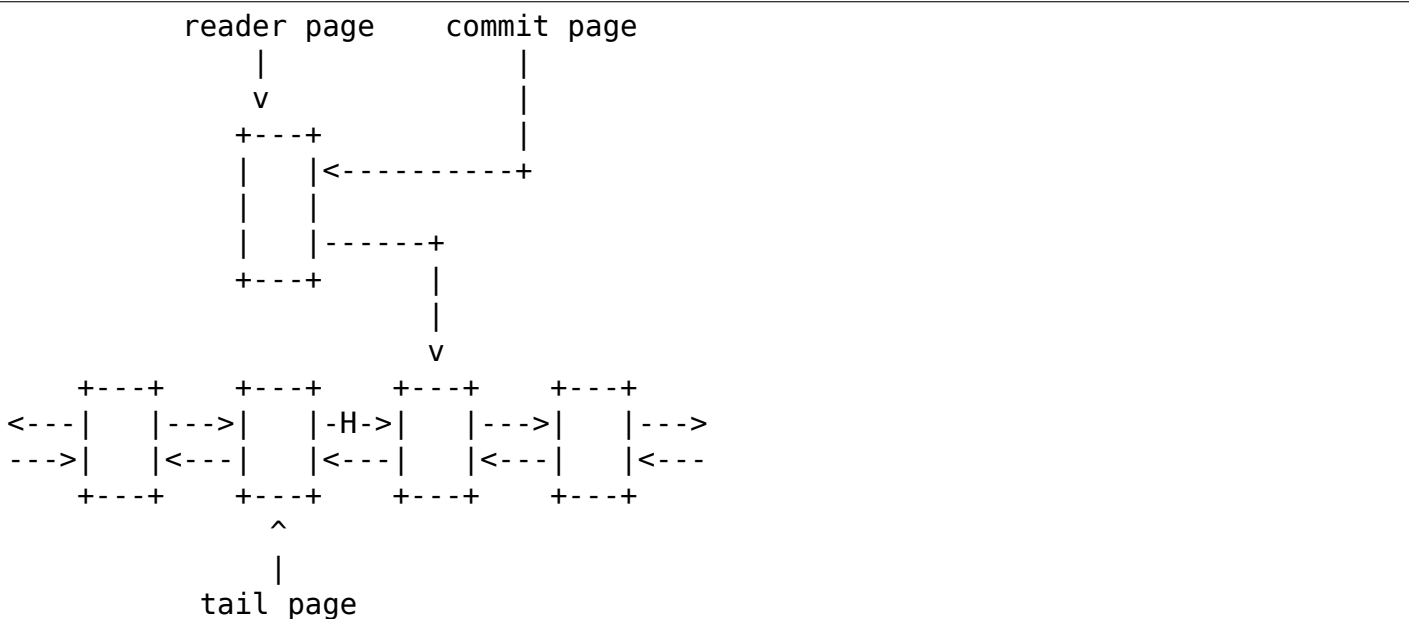


After the head page has been moved, the tail page may now move forward:



The above are the trivial updates. Now for the more complex scenarios.

As stated before, if enough writes preempt the first write, the tail page may make it all the way around the buffer and meet the commit page. At this time, we must start dropping writes (usually with some kind of warning to the user). But what happens if the commit was still on the reader page? The commit page is not part of the ring buffer. The tail page must account for this:



If the tail page were to simply push the head page forward, the commit when leaving the reader

page would not be pointing to the correct page.

The solution to this is to test if the commit page is on the reader page before pushing the head page. If it is, then it can be assumed that the tail page wrapped the buffer, and we must drop new writes.

This is not a race condition, because the commit page can only be moved by the outermost writer (the writer that was preempted). This means that the commit will not move while a writer is moving the tail page. The reader cannot swap the reader page if it is also being used as the commit page. The reader can simply check that the commit is off the reader page. Once the commit page leaves the reader page it will never go back on it unless a reader does another swap with the buffer page that is also the commit page.


23.4 Nested writes

In the pushing forward of the tail page we must first push forward the head page if the head page is the next page. If the head page is not the next page, the tail page is simply updated with a `cpxchg`.

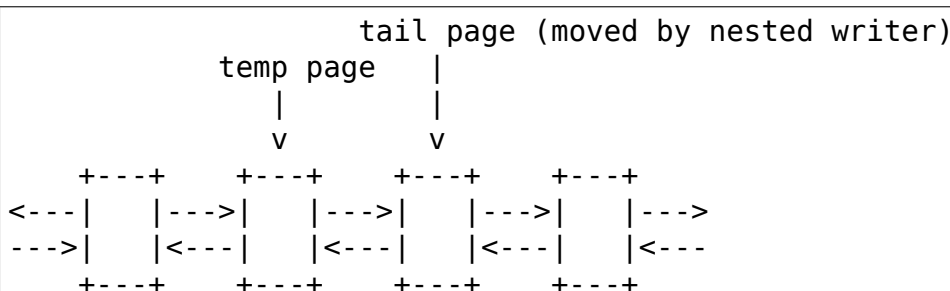
Only writers move the tail page. This must be done atomically to protect against nested writers:

```
temp_page = tail_page
next_page = temp_page->next
cpxchg(tail_page, temp_page, next_page)
```

The above will update the tail page if it is still pointing to the expected page. If this fails, a nested write pushed it forward, the current write does not need to push it:




Nested write comes in and moves the tail page forward:




The above would fail the `cpxchg`, but since the tail page has already been moved forward, the writer will just try again to reserve storage on the new tail page.

But the moving of the head page is a bit more complex:




The write converts the head page pointer to UPDATE:




But if a nested writer preempts here, it will see that the next page is a head page, but it is also nested. It will detect that it is nested and will save that information. The detection is the fact that it sees the UPDATE flag instead of a HEADER or NORMAL pointer.


The nested writer will set the new head page pointer:



But it will not reset the update back to normal. Only the writer that converted a pointer from HEAD to UPDATE will convert it back to NORMAL:



After the nested writer finishes, the outermost writer will convert the UPDATE pointer to NORMAL:



```

--->|   |<---|   |<---|   |<---|   |<---
    +---+   +---+   +---+   +---+

```

It can be even more complex if several nested writes came in and moved the tail page ahead several pages:

(first writer)

```

      tail page
      |
      v
+---+ +---+ +---+ +---+
<---|   |--->|   |-H->|   |--->|   |--->
--->|   |<---|   |<---|   |<---|   |<---
    +---+   +---+   +---+   +---+

```

The write converts the head page pointer to UPDATE:

```

      tail page
      |
      v
+---+ +---+ +---+ +---+
<---|   |--->|   |-U->|   |--->|   |--->
--->|   |<---|   |<---|   |<---|   |<---
    +---+   +---+   +---+   +---+

```

Next writer comes in, and sees the update and sets up the new head page:

(second writer)

```

      tail page
      |
      v
+---+ +---+ +---+ +---+
<---|   |--->|   |-U->|   |-H->|   |--->
--->|   |<---|   |<---|   |<---|   |<---
    +---+   +---+   +---+   +---+

```

The nested writer moves the tail page forward. But does not set the old update page to NORMAL because it is not the outermost writer:


```

      tail page
      |
      v
+---+ +---+ +---+ +---+
<---|   |--->|   |-U->|   |-H->|   |--->
--->|   |<---|   |<---|   |<---|   |<---
    +---+   +---+   +---+   +---+

```


Another writer preempts and sees the page after the tail page is a head page. It changes it from HEAD to UPDATE:

(third writer)




The writer will move the head page forward:

(third writer)




But now that the third writer did change the HEAD flag to UPDATE it will convert it to normal:

(third writer)



Then it will move the tail page, and return back to the second writer:


(second writer)



The second writer will fail to move the tail page because it was already moved, so it will try again and add its data to the new tail page. It will return to the first writer:


(first writer)

tail page




The first writer cannot know atomically if the tail page moved while it updates the HEAD page. It will then update the head page to what it thinks is the new head page:

(first writer)




Since the `cmpxchg` returns the old value of the pointer the first writer will see it succeeded in updating the pointer from NORMAL to HEAD. But as we can see, this is not good enough. It must also check to see if the tail page is either where it use to be or on the next page:

(first writer)




If `tail page != A` and `tail page != B`, then it must reset the pointer back to NORMAL. The fact that it only needs to worry about nested writers means that it only needs to check this after setting the HEAD page:

(first writer)



Now the writer can update the head page. This is also why the head page must remain in UPDATE and only reset by the outermost writer. This prevents the reader from seeing the incorrect head page:

(first writer)



SYSTEM TRACE MODULE

System Trace Module (STM) is a device described in MIPI STP specs as STP trace stream generator. STP (System Trace Protocol) is a trace protocol multiplexing data from multiple trace sources, each one of which is assigned a unique pair of master and channel. While some of these masters and channels are statically allocated to certain hardware trace sources, others are available to software. Software trace sources are usually free to pick for themselves any master/channel combination from this pool.

On the receiving end of this STP stream (the decoder side), trace sources can only be identified by master/channel combination, so in order for the decoder to be able to make sense of the trace that involves multiple trace sources, it needs to be able to map those master/channel pairs to the trace sources that it understands.

For instance, it is helpful to know that syslog messages come on master 7 channel 15, while arbitrary user applications can use masters 48 to 63 and channels 0 to 127.

To solve this mapping problem, stm class provides a policy management mechanism via configs, that allows defining rules that map string identifiers to ranges of masters and channels. If these rules (policy) are consistent with what decoder expects, it will be able to properly process the trace data.

This policy is a tree structure containing rules (policy_node) that have a name (string identifier) and a range of masters and channels associated with it, located in “stp-policy” subsystem directory in configs. The topmost directory’s name (the policy) is formatted as the STM device name to which this policy applies and an arbitrary string identifier separated by a stop. From the example above, a rule may look like this:

```
$ ls /config/stp-policy/dummy_stm.my-policy/user
channels masters
$ cat /config/stp-policy/dummy_stm.my-policy/user/masters
48 63
$ cat /config/stp-policy/dummy_stm.my-policy/user/channels
0 127
```

which means that the master allocation pool for this rule consists of masters 48 through 63 and channel allocation pool has channels 0 through 127 in it. Now, any producer (trace source) identifying itself with “user” identification string will be allocated a master and channel from within these ranges.

These rules can be nested, for example, one can define a rule “dummy” under “user” directory from the example above and this new rule will be used for trace sources with the id string of “user/dummy”.

Trace sources have to open the stm class device's node and write their trace data into its file descriptor.

In order to find an appropriate policy node for a given trace source, several mechanisms can be used. First, a trace source can explicitly identify itself by calling an `STP_POLICY_ID_SET` ioctl on the character device's file descriptor, providing their id string, before they write any data there. Secondly, if they chose not to perform the explicit identification (because you may not want to patch existing software to do this), they can just start writing the data, at which point the stm core will try to find a policy node with the name matching the task's name (e.g., "syslogd") and if one exists, it will be used. Thirdly, if the task name can't be found among the policy nodes, the catch-all entry "default" will be used, if it exists. This entry also needs to be created and configured by the system administrator or whatever tools are taking care of the policy configuration. Finally, if all the above steps failed, the `write()` to an stm file descriptor will return a error (EINVAL).

Previously, if no policy nodes were found for a trace source, the stm class would silently fall back to allocating the first available contiguous range of master/channels from the beginning of the device's master/channel range. The new requirement for a policy node to exist will help programmers and sysadmins identify gaps in configuration and have better control over the un-identified sources.

Some STM devices may allow direct mapping of the channel mmio regions to userspace for zero-copy writing. One mappable page (in terms of mmu) will usually contain multiple channels' mmios, so the user will need to allocate that many channels to themselves (via the aforementioned ioctl() call) to be able to do this. That is, if your stm device's channel mmio region is 64 bytes and hardware page size is 4096 bytes, after a successful `STP_POLICY_ID_SET` ioctl() call with `width==64`, you should be able to `mmap()` one page on this file descriptor and obtain direct access to an mmio region for 64 channels.

Examples of STM devices are Intel(R) Trace Hub [1] and Coresight STM [2].

24.1 stm_source

For kernel-based trace sources, there is "stm_source" device class. Devices of this class can be connected and disconnected to/from stm devices at runtime via a sysfs attribute called "stm_source_link" by writing the name of the desired stm device there, for example:

```
$ echo dummy_stm.0 > /sys/class/stm_source/console/stm_source_link
```

For examples on how to use stm_source interface in the kernel, refer to stm_console, stm_heartbeat or stm_fttrace drivers.

Each stm_source device will need to assume a master and a range of channels, depending on how many channels it requires. These are allocated for the device according to the policy configuration. If there's a node in the root of the policy directory that matches the stm_source device's name (for example, "console"), this node will be used to allocate master and channel numbers. If there's no such policy node, the stm core will use the catch-all entry "default", if one exists. If neither policy nodes exist, the `write()` to stm_source_link will return an error.

24.2 stm_console

One implementation of this interface also used in the example above is the “stm_console” driver, which basically provides a one-way console for kernel messages over an stm device.

To configure the master/channel pair that will be assigned to this console in the STP stream, create a “console” policy entry (see the beginning of this text on how to do that). When initialized, it will consume one channel.

24.3 stm_fttrace

This is another “stm_source” device, once the stm_fttrace has been linked with an stm device, and if “function” tracer is enabled, function address and parent function address which Ftrace subsystem would store into ring buffer will be exported via the stm device at the same time.

Currently only Ftrace “function” tracer is supported.

- [1] <https://software.intel.com/sites/default/files/managed/d3/3c/intel-th-developer-manual.pdf>
- [2] <http://infocenter.arm.com/help/index.jsp?topic=/com.arm.doc.ddi0444b/index.html>

MIPI SYS-T OVER STP

The MIPI SyS-T protocol driver can be used with STM class devices to generate standardized trace stream. Aside from being a standard, it provides better trace source identification and timestamp correlation.

In order to use the MIPI SyS-T protocol driver with your STM device, first, you'll need CONFIG_STM_PROTO_SYS_T.

Now, you can select which protocol driver you want to use when you create a policy for your STM device, by specifying it in the policy name:

```
# mkdir /config/stp-policy/dummy_stm.0:p_sys-t.my-policy/
```

In other words, the policy name format is extended like this:

`<device_name>:<protocol_name>.<policy_name>`

With Intel TH, therefore it can look like "0-sth:p_sys-t.my-policy".

If the protocol name is omitted, the STM class will chose whichever protocol driver was loaded first.

You can also double check that everything is working as expected by

```
# cat /config/stp-policy/dummy_stm.0:p_sys-t.my-policy/protocol p_sys-t
```

Now, with the MIPI SyS-T protocol driver, each policy node in the configs gets a few additional attributes, which determine per-source parameters specific to the protocol:

```
# mkdir /config/stp-policy/dummy_stm.0:p_sys-t.my-policy/default # ls /config/stp-policy/dummy_stm.0:p_sys-t.my-policy/default channels clocksync_interval do_len masters ts_interval uuid
```

The most important one here is the "uuid", which determines the UUID that will be used to tag all data coming from this source. It is automatically generated when a new node is created, but it is likely that you would want to change it.

do_len switches on/off the additional "payload length" field in the MIPI SyS-T message header. It is off by default as the STP already marks message boundaries.

ts_interval and clocksync_interval determine how much time in milliseconds can pass before we need to include a protocol (not transport, aka STP) timestamp in a message header or send a CLOCKSNC packet, respectively.

See Documentation/ABI/testing/configfs-stp-policy-p_sys-t for more details.

- [1] <https://www.mipi.org/specifications/sys-t>

CORESIGHT - ARM HARDWARE TRACE

26.1 Coresight - HW Assisted Tracing on ARM

Author Mathieu Poirier <mathieu.poirier@linaro.org>

Date September 11th, 2014

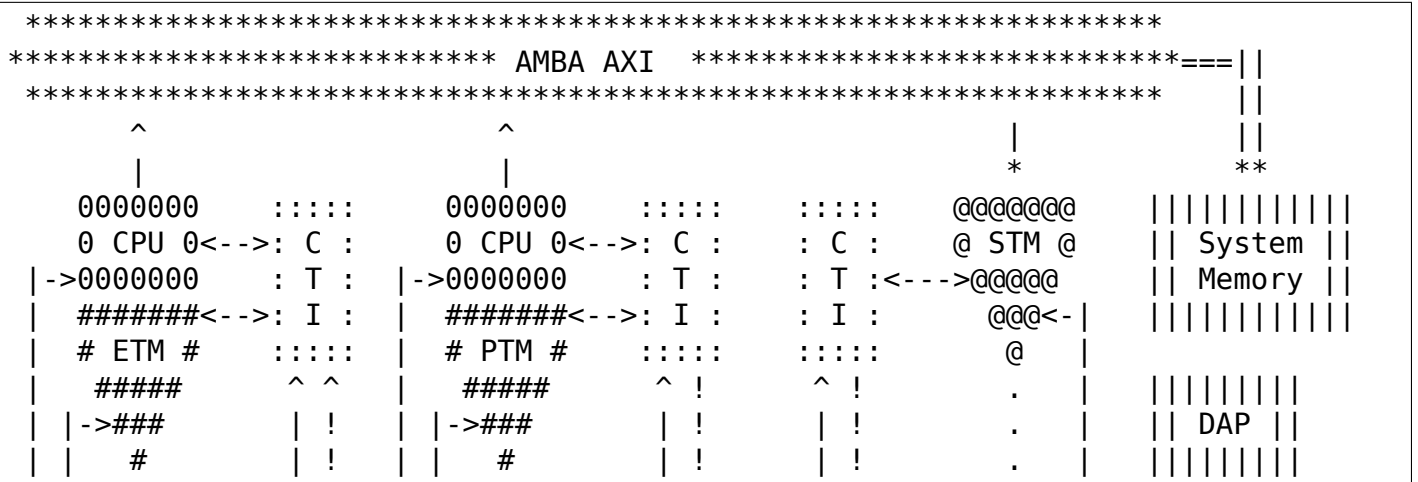
26.1.1 Introduction

Coresight is an umbrella of technologies allowing for the debugging of ARM based SoC. It includes solutions for JTAG and HW assisted tracing. This document is concerned with the latter.

HW assisted tracing is becoming increasingly useful when dealing with systems that have many SoCs and other components like GPU and DMA engines. ARM has developed a HW assisted tracing solution by means of different components, each being added to a design at synthesis time to cater to specific tracing needs. Components are generally categorised as source, link and sinks and are (usually) discovered using the AMBA bus.

“Sources” generate a compressed stream representing the processor instruction path based on tracing scenarios as configured by users. From there the stream flows through the coresight system (via ATB bus) using links that are connecting the emanating source to a sink(s). Sinks serve as endpoints to the coresight implementation, either storing the compressed stream in a memory buffer or creating an interface to the outside world where data can be transferred to a host without fear of filling up the onboard coresight memory buffer.

At typical coresight system would look like this:



The coresight framework provides a central point to represent, configure and manage coresight devices on a platform. This first implementation centers on the basic tracing functionality, enabling components such ETM/PTM, funnel, replicator, TMC, TPIU and ETB. Future work will enable more intricate IP blocks such as STM and CTI.

26.1.2 Acronyms and Classification

Acronyms:

PTM: Program Trace Macrocell

ETM: Embedded Trace Macrocell

STM: System trace Macrocell

ETB: Embedded Trace Buffer

ITM: Instrumentation Trace Macrocell

TPIU: Trace Port Interface Unit

TMC-ETR: Trace Memory Controller, configured as Embedded Trace Router

TMC-ETF: Trace Memory Controller, configured as Embedded Trace FIFO

CTI: Cross Trigger Interface

Classification:

Source: ETMv3.x ETMv4, PTMv1.0, PTMv1.1, STM, STM500, ITM

Link: Funnel, replicator (intelligent or not), TMC-ETR

Sinks: ETBv1.0, ETB1.1, TPIU, TMC-ETF

Misc: CTI

26.1.3 Device Tree Bindings

See Documentation/devicetree/bindings/arm/coresight.txt for details.

As of this writing drivers for ITM, STMs and CTIs are not provided but are expected to be added as the solution matures.

26.1.4 Framework and implementation

The coresight framework provides a central point to represent, configure and manage coresight devices on a platform. Any coresight compliant device can register with the framework for as long as they use the right APIs:

```
struct coresight_device *coresight_register(struct coresight_desc *desc);
```

```
void coresight_unregister(struct coresight_device *csdev);
```

The registering function is taking a `struct coresight_desc *desc` and register the device with the core framework. The unregister function takes a reference to a `struct coresight_device *csdev` obtained at registration time.

If everything goes well during the registration process the new devices will show up under `/sys/bus/coresight/devices`, as shown here for a TC2 platform:

```
root:~# ls /sys/bus/coresight/devices/
replicator 20030000.tpiu    2201c000.ptm  2203c000.etm  2203e000.etm
20010000.etb      20040000.funnel  2201d000.ptm  2203d000.etm
root:~#
```

The functions take a struct `coresight_device`, which looks like this:

```
struct coresight_desc {
    enum coresight_dev_type type;
    struct coresight_dev_subtype subtype;
    const struct coresight_ops *ops;
    struct coresight_platform_data *pdata;
    struct device *dev;
    const struct attribute_group **groups;
};
```

The “`coresight_dev_type`” identifies what the device is, i.e, source link or sink while the “`coresight_dev_subtype`” will characterise that type further.

The struct `coresight_ops` is mandatory and will tell the framework how to perform base operations related to the components, each component having a different set of requirement. For that struct `coresight_ops_sink`, struct `coresight_ops_link` and struct `coresight_ops_source` have been provided.

The next field `struct coresight_platform_data *pdata` is acquired by calling `of_get_coresight_platform_data()`, as part of the driver’s `_probe` routine and struct device `*dev` gets the device reference embedded in the `amba_device`:

```
static int etm_probe(struct amba_device *adev, const struct amba_id *id)
{
    ...
    ...
    drvdata->dev = &adev->dev;
    ...
}
```

Specific class of device (source, link, or sink) have generic operations that can be performed on them (see struct `coresight_ops`). The `**groups` is a list of sysfs entries pertaining to operations specific to that component only. “Implementation defined” customisations are expected to be accessed and controlled using those entries.

26.1.5 Device Naming scheme

The devices that appear on the “coresight” bus were named the same as their parent devices, i.e, the real devices that appears on AMBA bus or the platform bus. Thus the names were based on the Linux Open Firmware layer naming convention, which follows the base physical address of the device followed by the device type. e.g:

```
root:~# ls /sys/bus/coresight/devices/
20010000.etf  20040000.funnel  20100000.stm  22040000.etm
22140000.etm  230c0000.funnel  23240000.etm  20030000.tpiu
```



```
20070000.etr  20120000.replicator  220c0000.funnel
23040000.etm  23140000.etm          23340000.etm
```

However, with the introduction of ACPI support, the names of the real devices are a bit cryptic and non-obvious. Thus, a new naming scheme was introduced to use more generic names based on the type of the device. The following rules apply:

- 1) Devices that are bound to CPUs, are named based on the CPU logical number.

e.g, ETM bound to CPU0 is named "etm0"
- 2) All other devices follow a pattern, "<device_type_prefix>N", where :
 - <device_type_prefix> - A prefix specific to the type of the device
 - N - a sequential number assigned based on the order of probing.
e.g, tmc_etf0, tmc_etr0, funnel0, funnel1

Thus, with the new scheme the devices could appear as

```
root:~# ls /sys/bus/coresight/devices/
etm0      etm1      etm2      etm3  etm4      etm5      funnel0
funnel1   funnel2   replicator0  stm0  tmc_etf0  tmc_etr0  tpiu0
```

Some of the examples below might refer to old naming scheme and some to the newer scheme, to give a confirmation that what you see on your system is not unexpected. One must use the "names" as they appear on the system under specified locations.

26.1.6 Topology Representation

Each CoreSight component has a connections directory which will contain links to other CoreSight components. This allows the user to explore the trace topology and for larger systems, determine the most appropriate sink for a given source. The connection information can also be used to establish which CTI devices are connected to a given component. This directory contains a `nr_links` attribute detailing the number of links in the directory.

For an ETM source, in this case etm0 on a Juno platform, a typical arrangement will be:

```
linaro-developer:~# ls -l /sys/bus/coresight/devices/etm0/connections
<file details>  cti_cpu0 -> ../../../../23020000.cti/cti_cpu0
<file details>  nr_links
<file details>  out:0 -> ../../../../230c0000.funnel/funnel2
```

Following the out port to funnel2:

```
linaro-developer:~# ls -l /sys/bus/coresight/devices/funnel2/connections
<file details> in:0 -> ../../../../23040000.etm/etm0
<file details> in:1 -> ../../../../23140000.etm/etm3
<file details> in:2 -> ../../../../23240000.etm/etm4
```

```
<file details> in:3 -> ../../../23340000.etm/etm5
<file details> nr_links
<file details> out:0 -> ../../../20040000.funnel/funnel0
```

And again to funnel0:

```
linaro-developer:~# ls -l /sys/bus/coresight/devices/funnel0/connections
<file details> in:0 -> ../../../220c0000.funnel/funnel1
<file details> in:1 -> ../../../230c0000.funnel/funnel2
<file details> nr_links
<file details> out:0 -> ../../../20010000.etf/tmc_etf0
```

Finding the first sink tmc_etf0. This can be used to collect data as a sink, or as a link to propagate further along the chain:

```
linaro-developer:~# ls -l /sys/bus/coresight/devices/tmc_etf0/connections
<file details> cti_sys0 -> ../../../20020000.cti/cti_sys0
<file details> in:0 -> ../../../20040000.funnel/funnel0
<file details> nr_links
<file details> out:0 -> ../../../20150000.funnel/funnel4
```

via funnel4:

```
linaro-developer:~# ls -l /sys/bus/coresight/devices/funnel4/connections
<file details> in:0 -> ../../../20010000.etf/tmc_etf0
<file details> in:1 -> ../../../20140000.etf/tmc_etf1
<file details> nr_links
<file details> out:0 -> ../../../20120000.replicator/replicator0
```

and a replicator0:

```
linaro-developer:~# ls -l /sys/bus/coresight/devices/replicator0/connections
<file details> in:0 -> ../../../20150000.funnel/funnel4
<file details> nr_links
<file details> out:0 -> ../../../20030000.tpiu/tpiu0
<file details> out:1 -> ../../../20070000.etr/tmc_etr0
```

Arriving at the final sink in the chain, tmc_etr0:

```
linaro-developer:~# ls -l /sys/bus/coresight/devices/tmc_etr0/connections
<file details> cti_sys0 -> ../../../20020000.cti/cti_sys0
<file details> in:0 -> ../../../20120000.replicator/replicator0
<file details> nr_links
```

As described below, when using sysfs it is sufficient to enable a sink and a source for successful trace. The framework will correctly enable all intermediate links as required.

Note: cti_sys0 appears in two of the connections lists above. CTIs can connect to multiple devices and are arranged in a star topology via the CTM. See (*CoreSight Embedded Cross Trigger (CTI & CTM)*).⁴ for further details. Looking at this device we see 4 connections:

⁴ *CoreSight Embedded Cross Trigger (CTI & CTM)*.

```
linaro-developer:~# ls -l /sys/bus/coresight/devices/cti_sys0/connections
<file details> nr_links
<file details> stm0 -> ../../../../20100000.stm/stm0
<file details> tmc_etf0 -> ../../../../20010000.etf/tmc_etf0
<file details> tmc_etr0 -> ../../../../20070000.etr/tmc_etr0
<file details> tpiu0 -> ../../../../20030000.tpiu/tpiu0
```

26.1.7 How to use the tracer modules

There are two ways to use the Coresight framework:

1. using the perf cmd line tools.
2. interacting directly with the Coresight devices using the sysFS interface.

Preference is given to the former as using the sysFS interface requires a deep understanding of the Coresight HW. The following sections provide details on using both methods.

1) Using the sysFS interface:

Before trace collection can start, a coresight sink needs to be identified. There is no limit on the amount of sinks (nor sources) that can be enabled at any given moment. As a generic operation, all device pertaining to the sink class will have an “active” entry in sysfs:

```
root:/sys/bus/coresight/devices# ls
replicator  20030000.tpiu  2201c000.ptm  2203c000.etm  2203e000.etm
20010000.etb  20040000.funnel  2201d000.ptm  2203d000.etm
root:/sys/bus/coresight/devices# ls 20010000.etb
enable_sink  status  trigger_cntr
root:/sys/bus/coresight/devices# echo 1 > 20010000.etb/enable_sink
root:/sys/bus/coresight/devices# cat 20010000.etb/enable_sink
1
root:/sys/bus/coresight/devices#
```

At boot time the current etm3x driver will configure the first address comparator with “_stext” and “_etext”, essentially tracing any instruction that falls within that range. As such “enabling” a source will immediately trigger a trace capture:

```
root:/sys/bus/coresight/devices# echo 1 > 2201c000.ptm/enable_source
root:/sys/bus/coresight/devices# cat 2201c000.ptm/enable_source
1
root:/sys/bus/coresight/devices# cat 20010000.etb/status
Depth:          0x2000
Status:          0x1
RAM read ptr:    0x0
RAM wrt ptr:     0x19d3  <----- The write pointer is moving
Trigger cnt:     0x0
Control:         0x1
Flush status:    0x0
Flush ctrl:      0x2001
root:/sys/bus/coresight/devices#
```

Trace collection is stopped the same way:

```
root:/sys/bus/coresight/devices# echo 0 > 2201c000.ptm/enable_source
root:/sys/bus/coresight/devices#
```

The content of the ETB buffer can be harvested directly from /dev:

```
root:/sys/bus/coresight/devices# dd if=/dev/20010000.etb \
of=~/.cstrace.bin
64+0 records in
64+0 records out
32768 bytes (33 kB) copied, 0.00125258 s, 26.2 MB/s
root:/sys/bus/coresight/devices#
```

The file cstrace.bin can be decompressed using “ptm2human”, DS-5 or Trace32.

Following is a DS-5 output of an experimental loop that increments a variable up to a certain value. The example is simple and yet provides a glimpse of the wealth of possibilities that coresight provides.

Info	Tracing enabled					
Instruction	106378866	0x8026B53C	E52DE004	false	PUSH	↳
↳ {lr}						
Instruction	0	0x8026B540	E24DD00C	false	SUB	sp, sp,
↳ #0xc						
Instruction	0	0x8026B544	E3A03000	false	MOV	r3, #0
Instruction	0	0x8026B548	E58D3004	false	STR	r3,
↳ [sp, #4]						
Instruction	0	0x8026B54C	E59D3004	false	LDR	r3,
↳ [sp, #4]						
Instruction	0	0x8026B550	E3530004	false	CMP	r3, #4
Instruction	0	0x8026B554	E2833001	false	ADD	r3, r3,
↳ #1						
Instruction	0	0x8026B558	E58D3004	false	STR	r3,
↳ [sp, #4]						
Instruction	0	0x8026B55C	DAFFFFFFFA	true	BLE	{pc}-
↳ 0x10 ; 0x8026b54c						
Timestamp	Timestamp: 17106715833					
Instruction	319	0x8026B54C	E59D3004	false	LDR	r3,
↳ [sp, #4]						
Instruction	0	0x8026B550	E3530004	false	CMP	r3, #4
Instruction	0	0x8026B554	E2833001	false	ADD	r3, r3,
↳ #1						
Instruction	0	0x8026B558	E58D3004	false	STR	r3,
↳ [sp, #4]						
Instruction	0	0x8026B55C	DAFFFFFFFA	true	BLE	{pc}-
↳ 0x10 ; 0x8026b54c						
Instruction	9	0x8026B54C	E59D3004	false	LDR	r3,
↳ [sp, #4]						
Instruction	0	0x8026B550	E3530004	false	CMP	r3, #4
Instruction	0	0x8026B554	E2833001	false	ADD	r3, r3,
↳ #1						
Instruction	0	0x8026B558	E58D3004	false	STR	r3,
↳ [sp, #4]						

Instruction	0	0x8026B55C	DAFFFFFFFA	true	BLE	{pc}-
→0x10 ; 0x8026b54c						
Instruction	7	0x8026B54C	E59D3004	false	LDR	r3,
→[sp,#4]						
Instruction	0	0x8026B550	E3530004	false	CMP	r3,#4
Instruction	0	0x8026B554	E2833001	false	ADD	r3,r3,
→#1						
Instruction	0	0x8026B558	E58D3004	false	STR	r3,
→[sp,#4]						
Instruction	0	0x8026B55C	DAFFFFFFFA	true	BLE	{pc}-
→0x10 ; 0x8026b54c						
Instruction	7	0x8026B54C	E59D3004	false	LDR	r3,
→[sp,#4]						
Instruction	0	0x8026B550	E3530004	false	CMP	r3,#4
Instruction	0	0x8026B554	E2833001	false	ADD	r3,r3,
→#1						
Instruction	0	0x8026B558	E58D3004	false	STR	r3,
→[sp,#4]						
Instruction	0	0x8026B55C	DAFFFFFFFA	true	BLE	{pc}-
→0x10 ; 0x8026b54c						
Instruction	10	0x8026B54C	E59D3004	false	LDR	r3,
→[sp,#4]						
Instruction	0	0x8026B550	E3530004	false	CMP	r3,#4
Instruction	0	0x8026B554	E2833001	false	ADD	r3,r3,
→#1						
Instruction	0	0x8026B558	E58D3004	false	STR	r3,
→[sp,#4]						
Instruction	0	0x8026B55C	DAFFFFFFFA	true	BLE	{pc}-
→0x10 ; 0x8026b54c						
Instruction	6	0x8026B560	EE1D3F30	false	MRC	p15,
→#0x0,r3,c13,c0,#1						
Instruction	0	0x8026B564	E1A0100D	false	MOV	r1,sp
Instruction	0	0x8026B568	E3C12D7F	false	BIC	r2,r1,
→#0x1fc0						
Instruction	0	0x8026B56C	E3C2203F	false	BIC	r2,r2,
→#0x3f						
Instruction	0	0x8026B570	E59D1004	false	LDR	r1,
→[sp,#4]						
Instruction	0	0x8026B574	E59F0010	false	LDR	r0,
→[pc,#16] ; [0x8026B58C] = 0x80550368						
Instruction	0	0x8026B578	E592200C	false	LDR	r2,
→[r2,#0xc]						
Instruction	0	0x8026B57C	E59221D0	false	LDR	r2,
→[r2,#0x1d0]						
Instruction	0	0x8026B580	EB07A4CF	true	BL	{pc}
→+0x1e9344 ; 0x804548c4						
Info Tracing enabled						
Instruction	13570831	0x8026B584	E28DD00C	false	ADD	sp,sp,#0xc
→ sp,sp,#0xc						
Instruction	0	0x8026B588	E8BD8000	true	LDM	sp!,
→{pc}						

Timestamp

Timestamp: 17107041535

2) Using perf framework:

Coresight tracers are represented using the Perf framework's Performance Monitoring Unit (PMU) abstraction. As such the perf framework takes charge of controlling when tracing gets enabled based on when the process of interest is scheduled. When configured in a system, Coresight PMUs will be listed when queried by the perf command line tool:

```
linaro@linaro-nano:~$ ./perf list pmu
```

List of pre-defined events (to be used in -e):

```
cs_etm// [Kernel PMU event]
```

```
linaro@linaro-nano:~$
```

Regardless of the number of tracers available in a system (usually equal to the amount of processor cores), the "cs_etm" PMU will be listed only once.

A Coresight PMU works the same way as any other PMU, i.e the name of the PMU is listed along with configuration options within forward slashes '/'. Since a Coresight system will typically have more than one sink, the name of the sink to work with needs to be specified as an event option. On newer kernels the available sinks are listed in sysFS under (\$SYSFS)/bus/event_source/devices/cs_etm/sinks/:

```
root@localhost:/sys/bus/event_source/devices/cs_etm/sinks# ls
tmc_etf0  tmc_etr0  tpiu0
```

On older kernels, this may need to be found from the list of coresight devices, available under (\$SYSFS)/bus/coresight/devices/:

```
root:~# ls /sys/bus/coresight/devices/
etm0      etm1      etm2      etm3  etm4      etm5      funnel0
funnel1  funnel2  replicator0  stm0  tmc_etf0  tmc_etr0  tpiu0
root@linaro-nano:~# perf record -e cs_etm/@tmc_etr0/u --per-thread program
```

As mentioned above in section "Device Naming scheme", the names of the devices could look different from what is used in the example above. One must use the device names as it appears under the sysFS.

The syntax within the forward slashes '/' is important. The '@' character tells the parser that a sink is about to be specified and that this is the sink to use for the trace session.

More information on the above and other example on how to use Coresight with the perf tools can be found in the "HOWTO.md" file of the openCSD gitHub repository³.

2.1) AutoFDO analysis using the perf tools:

perf can be used to record and analyze trace of programs.

Execution can be recorded using 'perf record' with the cs_etm event, specifying the name of the sink to record to, e.g:

```
perf record -e cs_etm/@tmc_etr0/u --per-thread
```

³ <https://github.com/Linaro/perf-opencsd>

The ‘perf report’ and ‘perf script’ commands can be used to analyze execution, synthesizing instruction and branch events from the instruction trace. ‘perf inject’ can be used to replace the trace data with the synthesized events. The -itrace option controls the type and frequency of synthesized events (see perf documentation).

Note that only 64-bit programs are currently supported - further work is required to support instruction decode of 32-bit Arm programs.

2.2) Tracing PID

The kernel can be built to write the PID value into the PE ContextID registers. For a kernel running at EL1, the PID is stored in CONTEXTIDR_EL1. A PE may implement Arm Virtualization Host Extensions (VHE), which the kernel can run at EL2 as a virtualisation host; in this case, the PID value is stored in CONTEXTIDR_EL2.

perf provides PMU formats that program the ETM to insert these values into the trace data; the PMU formats are defined as below:

“contextid1”: Available on both EL1 kernel and EL2 kernel. When the kernel is running at EL1, “contextid1” enables the PID tracing; when the kernel is running at EL2, this enables tracing the PID of guest applications.

“contextid2”: Only usable when the kernel is running at EL2. When selected, enables PID tracing on EL2 kernel.

“contextid”: Will be an alias for the option that enables PID tracing. I.e, contextid == contextid1, on EL1 kernel. contextid == contextid2, on EL2 kernel.

perf will always enable PID tracing at the relevant EL, this is accomplished by automatically enable the “contextid” config - but for EL2 it is possible to make specific adjustments using configs “contextid1” and “contextid2”, E.g. if a user wants to trace PIDs for both host and guest, the two configs “contextid1” and “contextid2” can be set at the same time:

```
perf record -e cs_etm/contextid1,contextid2/u - vm
```

26.1.8 Generating coverage files for Feedback Directed Optimization: AutoFDO

‘perf inject’ accepts the -itrace option in which case tracing data is removed and replaced with the synthesized events. e.g.

```
perf inject --itrace --strip -i perf.data -o perf.data.new
```

Below is an example of using ARM ETM for autoFDO. It requires autofdo (<https://github.com/google/autofdo>) and gcc version 5. The bubble sort example is from the AutoFDO tutorial (<https://gcc.gnu.org/wiki/AutoFDO/Tutorial>).

```
$ gcc-5 -O3 sort.c -o sort
$ taskset -c 2 ./sort
Bubble sorting array of 30000 elements
5910 ms

$ perf record -e cs_etm/@tmc_etr0/u --per-thread taskset -c 2 ./sort
Bubble sorting array of 30000 elements
12543 ms
```



```
[ perf record: Woken up 35 times to write data ]
[ perf record: Captured and wrote 69.640 MB perf.data ]

$ perf inject -i perf.data -o inj.data --itrace=il64 --strip
$ create_gcov --binary=./sort --profile=inj.data --gcov=sort.gcov -gcov_
→version=1
$ gcc-5 -O3 -fauto-profile=sort.gcov sort.c -o sort_autofdo
$ taskset -c 2 ./sort_autofdo
Bubble sorting array of 30000 elements
5806 ms
```

26.1.9 How to use the STM module

Using the System Trace Macrocell module is the same as the tracers - the only difference is that clients are driving the trace capture rather than the program flow through the code.

As with any other CoreSight component, specifics about the STM tracer can be found in sysfs with more information on each entry being found in¹:

```
root@genericarmv8:~# ls /sys/bus/coresight/devices/stm0
enable_source  hwevent_select  port_enable      subsystem        uevent
hwevent_enable mgmt             port_select      traceid
```

Like any other source a sink needs to be identified and the STM enabled before being used:

```
root@genericarmv8:~# echo 1 > /sys/bus/coresight/devices/tmc_etf0/enable_sink
root@genericarmv8:~# echo 1 > /sys/bus/coresight/devices/stm0/enable_source
```

From there user space applications can request and use channels using the devfs interface provided for that purpose by the generic STM API:

```
root@genericarmv8:~# ls -l /dev/stm0
crw----- 1 root    root      10,  61 Jan  3 18:11 /dev/stm0
root@genericarmv8:~#
```

Details on how to use the generic STM API can be found here: - [System Trace Module](#)².

26.1.10 The CTI & CTM Modules

The CTI (Cross Trigger Interface) provides a set of trigger signals between individual CTIs and components, and can propagate these between all CTIs via channels on the CTM (Cross Trigger Matrix).

A separate documentation file is provided to explain the use of these devices. ([CoreSight Embedded Cross Trigger \(CTI & CTM\)](#)).[?]

¹ Documentation/ABI/testing/sysfs-bus-coresight-devices-stm

² [System Trace Module](#)

26.1.11 CoreSight System Configuration

CoreSight components can be complex devices with many programming options. Furthermore, components can be programmed to interact with each other across the complete system.

A CoreSight System Configuration manager is provided to allow these complex programming configurations to be selected and used easily from perf and sysfs.

See the separate document for further information. (*CoreSight System Configuration Manager*)⁵.

26.2 CoreSight System Configuration Manager

Author Mike Leach <mike.leach@linaro.org>

Date October 2020

26.2.1 Introduction

The CoreSight System Configuration manager is an API that allows the programming of the CoreSight system with pre-defined configurations that can then be easily enabled from sysfs or perf.

Many CoreSight components can be programmed in complex ways - especially ETMs. In addition, components can interact across the CoreSight system, often via the cross trigger components such as CTI and CTM. These system settings can be defined and enabled as named configurations.

26.2.2 Basic Concepts

This section introduces the basic concepts of a CoreSight system configuration.

Features

A feature is a named set of programming for a CoreSight device. The programming is device dependent, and can be defined in terms of absolute register values, resource usage and parameter values.

The feature is defined using a descriptor. This descriptor is used to load onto a matching device, either when the feature is loaded into the system, or when the CoreSight device is registered with the configuration manager.

The load process involves interpreting the descriptor into a set of register accesses in the driver - the resource usage and parameter descriptions translated into appropriate register accesses. This interpretation makes it easy and efficient for the feature to be programmed onto the device when required.

The feature will not be active on the device until the feature is enabled, and the device itself is enabled. When the device is enabled then enabled features will be programmed into the device hardware.

⁵ *CoreSight System Configuration Manager*

A feature is enabled as part of a configuration being enabled on the system.

Parameter Value

A parameter value is a named value that may be set by the user prior to the feature being enabled that can adjust the behaviour of the operation programmed by the feature.

For example, this could be a count value in a programmed operation that repeats at a given rate. When the feature is enabled then the current value of the parameter is used in programming the device.

The feature descriptor defines a default value for a parameter, which is used if the user does not supply a new value.

Users can update parameter values using the configs API for the CoreSight system - which is described below.

The current value of the parameter is loaded into the device when the feature is enabled on that device.

Configurations

A configuration defines a set of features that are to be used in a trace session where the configuration is selected. For any trace session only one configuration may be selected.

The features defined may be on any type of device that is registered to support system configuration. A configuration may select features to be enabled on a class of devices - i.e. any ETMv4, or specific devices, e.g. a specific CTI on the system.

As with the feature, a descriptor is used to define the configuration. This will define the features that must be enabled as part of the configuration as well as any preset values that can be used to override default parameter values.

Preset Values

Preset values are easily selectable sets of parameter values for the features that the configuration uses. The number of values in a single preset set, equals the sum of parameter values in the features used by the configuration.

e.g. a configuration consists of 3 features, one has 2 parameters, one has a single parameter, and another has no parameters. A single preset set will therefore have 3 values.

Presets are optionally defined by the configuration, up to 15 can be defined. If no preset is selected, then the parameter values defined in the feature are used as normal.

Operation

The following steps take place in the operation of a configuration.

- 1) In this example, the configuration is 'autofdo', which has an associated feature 'strobing' that works on ETMv4 CoreSight Devices.
- 2) The configuration is enabled. For example 'perf' may select the configuration as part of its command line:

```
perf record -e cs_etm/autofdo/ myapp
```

which will enable the 'autofdo' configuration.

- 3) perf starts tracing on the system. As each ETMv4 that perf uses for trace is enabled, the configuration manager will check if the ETMv4 has a feature that relates to the currently active configuration. In this case 'strobing' is enabled & programmed into the ETMv4.
- 4) When the ETMv4 is disabled, any registers marked as needing to be saved will be read back.
- 5) At the end of the perf session, the configuration will be disabled.

26.2.3 Viewing Configurations and Features

The set of configurations and features that are currently loaded into the system can be viewed using the configs API.

Mount configs as normal and the 'cs-syscfg' subsystem will appear:

```
$ ls /config
cs-syscfg  stp-policy
```

This has two sub-directories:

```
$ cd cs-syscfg/
$ ls
configurations  features
```

The system has the configuration 'autofdo' built in. It may be examined as follows:

```
$ cd configurations/
$ ls
autofdo
$ cd autofdo/
$ ls
description  feature_refs  preset1  preset3  preset5  preset7  preset9
enable       preset        preset2  preset4  preset6  preset8
$ cat description
Setup ETMs with strobing for autofdo
$ cat feature_refs
strobing
```

Each preset declared has a 'preset<n>' subdirectory declared. The values for the preset can be examined:

```
$ cat preset1/values
strobing.window = 0x1388 strobing.period = 0x2
$ cat preset2/values
strobing.window = 0x1388 strobing.period = 0x4
```

The 'enable' and 'preset' files allow the control of a configuration when using CoreSight with sysfs.

The features referenced by the configuration can be examined in the features directory:

```
$ cd ../../features/strobing/
$ ls
description matches nr_params params
$ cat description
Generate periodic trace capture windows.
parameter 'window': a number of CPU cycles (W)
parameter 'period': trace enabled for W cycles every period x W cycles
$ cat matches
SRC_ETMV4
$ cat nr_params
2
```

Move to the params directory to examine and adjust parameters:

```
cd params
$ ls
period window
$ cd period
$ ls
value
$ cat value
0x2710
# echo 15000 > value
# cat value
0x3a98
```

Parameters adjusted in this way are reflected in all device instances that have loaded the feature.

26.2.4 Using Configurations in perf

The configurations loaded into the CoreSight configuration management are also declared in the perf 'cs_etm' event infrastructure so that they can be selected when running trace under perf:

```
$ ls /sys/devices/cs_etm
cpu0  cpu2  events  nr_addr_filters  power  subsystem  uevent
cpu1  cpu3  format  perf_event_mux_interval_ms  sinks  type
```

The key directory here is 'events' - a generic perf directory which allows selection on the perf command line. As with the sinks entries, this provides a hash of the configuration name.

The entry in the 'events' directory uses perfs built in syntax generator to substitute the syntax for the name when evaluating the command:

```
$ ls events/  
autofdo  
$ cat events/autofdo  
configid=0xa7c3dddd
```

The 'autofdo' configuration may be selected on the perf command line:

```
$ perf record -e cs_etm/autofdo/u --per-thread <application>
```

A preset to override the current parameter values can also be selected:

```
$ perf record -e cs_etm/autofdo,preset=1/u --per-thread <application>
```

When configurations are selected in this way, then the trace sink used is automatically selected.

26.2.5 Using Configurations in sysfs

Coresight can be controlled using sysfs. When this is in use then a configuration can be made active for the devices that are used in the sysfs session.

In a configuration there are 'enable' and 'preset' files.

To enable a configuration for use with sysfs:

```
$ cd configurations/autofdo  
$ echo 1 > enable
```

This will then use any default parameter values in the features - which can be adjusted as described above.

To use a preset<n> set of parameter values:

```
$ echo 3 > preset
```

This will select preset3 for the configuration. The valid values for preset are 0 - to deselect presets, and any value of <n> where a preset<n> sub-directory is present.

Note that the active sysfs configuration is a global parameter, therefore only a single configuration can be active for sysfs at any one time. Attempting to enable a second configuration will result in an error. Additionally, attempting to disable the configuration while in use will also result in an error.

The use of the active configuration by sysfs is independent of the configuration used in perf.

26.2.6 Creating and Loading Custom Configurations

Custom configurations and / or features can be dynamically loaded into the system by using a loadable module.

An example of a custom configuration is found in `./samples/coresight`.

This creates a new configuration that uses the existing built in strobing feature, but provides a different set of presets.

When the module is loaded, then the configuration appears in the `configs` file system and is selectable in the same way as the built in configuration described above.

Configurations can use previously loaded features. The system will ensure that it is not possible to unload a feature that is currently in use, by enforcing the unload order as the strict reverse of the load order.

26.3 Coresight CPU Debug Module

Author Leo Yan <leo.yan@linaro.org>

Date April 5th, 2017

26.3.1 Introduction

Coresight CPU debug module is defined in ARMv8-a architecture reference manual (ARM DDI 0487A.k) Chapter 'Part H: External debug', the CPU can integrate debug module and it is mainly used for two modes: self-hosted debug and external debug. Usually the external debug mode is well known as the external debugger connects with SoC from JTAG port; on the other hand the program can explore debugging method which rely on self-hosted debug mode, this document is to focus on this part.

The debug module provides sample-based profiling extension, which can be used to sample CPU program counter, secure state and exception level, etc; usually every CPU has one dedicated debug module to be connected. Based on self-hosted debug mechanism, Linux kernel can access these related registers from mmio region when the kernel panic happens. The callback notifier for kernel panic will dump related registers for every CPU; finally this is good for assistant analysis for panic.

26.3.2 Implementation

- During driver registration, it uses `EDDEVID` and `EDDEVID1` - two device ID registers to decide if sample-based profiling is implemented or not. On some platforms this hardware feature is fully or partially implemented; and if this feature is not supported then registration will fail.
- At the time this documentation was written, the debug driver mainly relies on information gathered by the kernel panic callback notifier from three sampling registers: `EDPCSR`, `EDVIDSR` and `EDCIDS`: from `EDPCSR` we can get program counter; `EDVIDSR` has information for secure state, exception level, bit width, etc; `EDCIDS` is context ID value which contains the sampled value of `CONTEXTIDR_EL1`.

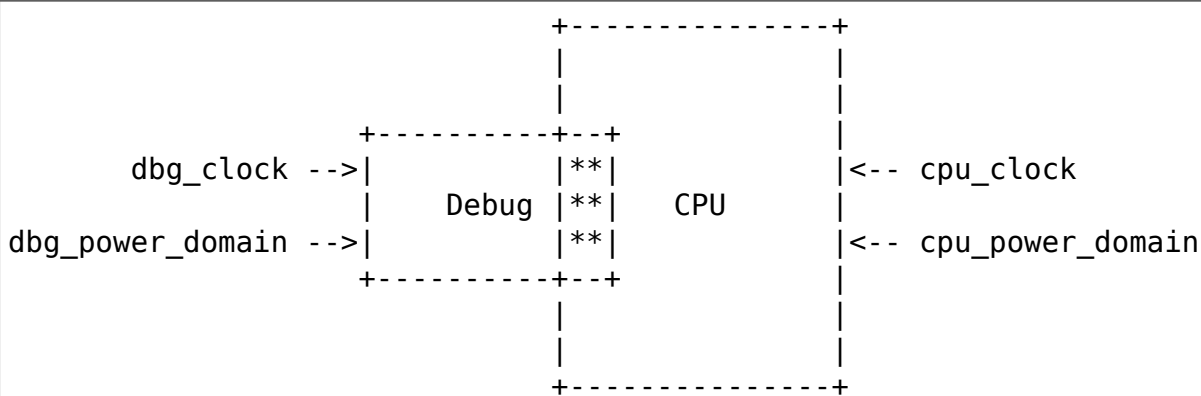
- The driver supports a CPU running in either AArch64 or AArch32 mode. The registers naming convention is a bit different between them, AArch64 uses 'ED' for register prefix (ARM DDI 0487A.k, chapter H9.1) and AArch32 uses 'DBG' as prefix (ARM DDI 0487A.k, chapter G5.1). The driver is unified to use AArch64 naming convention.
- ARMv8-a (ARM DDI 0487A.k) and ARMv7-a (ARM DDI 0406C.b) have different register bits definition. So the driver consolidates two difference:

If PCSROffset=0b0000, on ARMv8-a the feature of EDPCSR is not implemented; but ARMv7-a defines "PCSR samples are offset by a value that depends on the instruction set state". For ARMv7-a, the driver checks furthermore if CPU runs with ARM or thumb instruction set and calibrate PCSR value, the detailed description for offset is in ARMv7-a ARM (ARM DDI 0406C.b) chapter C11.11.34 "DBGPCSR, Program Counter Sampling Register".

If PCSROffset=0b0010, ARMv8-a defines "EDPCSR implemented, and samples have no offset applied and do not sample the instruction set state in AArch32 state". So on ARMv8 if EDDEVID1.PCSROffset is 0b0010 and the CPU operates in AArch32 state, EDPCSR is not sampled; when the CPU operates in AArch64 state EDPCSR is sampled and no offset are applied.

26.3.3 Clock and power domain

Before accessing debug registers, we should ensure the clock and power domain have been enabled properly. In ARMv8-a ARM (ARM DDI 0487A.k) chapter 'H9.1 Debug registers', the debug registers are spread into two domains: the debug domain and the CPU domain.



For debug domain, the user uses DT binding "clocks" and "power-domains" to specify the corresponding clock source and power supply for the debug logic. The driver calls the `pm_runtime_{put|get}` operations as needed to handle the debug power domain.

For CPU domain, the different SoC designs have different power management schemes and finally this heavily impacts external debug module. So we can divide into below cases:

- On systems with a sane power controller which can behave correctly with respect to CPU power domain, the CPU power domain can be controlled by register EDPRCR in driver. The driver firstly writes bit EDPRCR.COREPURQ to power up the CPU, and then writes bit EDPRCR.CORENPDRQ for emulation of CPU power down. As result, this can ensure the CPU power domain is powered on properly during the period when access debug related registers;

- Some designs will power down an entire cluster if all CPUs on the cluster are powered down - including the parts of the debug registers that should remain powered in the debug power domain. The bits in EDPRCR are not respected in these cases, so these designs do not support debug over power down in the way that the CoreSight / Debug designers anticipated. This means that even checking EDPRSR has the potential to cause a bus hang if the target register is unpowered.

In this case, accessing to the debug registers while they are not powered is a recipe for disaster; so we need preventing CPU low power states at boot time or when user enable module at the run time. Please see chapter “How to use the module” for detailed usage info for this.

26.3.4 Device Tree Bindings

See Documentation/devicetree/bindings/arm/coresight-cpu-debug.txt for details.

26.3.5 How to use the module

If you want to enable debugging functionality at boot time, you can add “core-sight_cpu_debug.enable=1” to the kernel command line parameter.

The driver also can work as module, so can enable the debugging when insmod module:

```
# insmod coresight_cpu_debug.ko debug=1
```

When boot time or insmod module you have not enabled the debugging, the driver uses the debugfs file system to provide a knob to dynamically enable or disable debugging:

To enable it, write a ‘1’ into /sys/kernel/debug/coresight_cpu_debug/enable:

```
# echo 1 > /sys/kernel/debug/coresight_cpu_debug/enable
```

To disable it, write a ‘0’ into /sys/kernel/debug/coresight_cpu_debug/enable:

```
# echo 0 > /sys/kernel/debug/coresight_cpu_debug/enable
```

As explained in chapter “Clock and power domain”, if you are working on one platform which has idle states to power off debug logic and the power controller cannot work well for the request from EDPRCR, then you should firstly constraint CPU idle states before enable CPU debugging feature; so can ensure the accessing to debug logic.

If you want to limit idle states at boot time, you can use “nohlt” or “cpuidle.off=1” in the kernel command line.

At the runtime you can disable idle states with below methods:

It is possible to disable CPU idle states by way of the PM QoS subsystem, more specifically by using the “/dev/cpu_dma_latency” interface (see Documentation/power/pm_qos_interface.rst for more details). As specified in the PM QoS documentation the requested parameter will stay in effect until the file descriptor is released. For example:

```
# exec 3<> /dev/cpu_dma_latency; echo 0 >&3
...
Do some work...
```



```
...
# exec 3<>-
```

The same can also be done from an application program.

Disable specific CPU's specific idle state from cpuidle sysfs (see Documentation/admin-guide/pm/cpuidle.rst):

```
# echo 1 > /sys/devices/system/cpu/cpu$cpu/cpuidle/state$state/disable
```

26.3.6 Output format

Here is an example of the debugging output format:

```
ARM external debug module:
coresight-cpu-debug 850000.debug: CPU[0]:
coresight-cpu-debug 850000.debug:  EDPRSR:  00000001 (Power:On DLK:Unlock)
coresight-cpu-debug 850000.debug:  EDPCSR:  handle_IPI+0x174/0x1d8
coresight-cpu-debug 850000.debug:  EDCIDSR:  00000000
coresight-cpu-debug 850000.debug:  EDVIDSR:  90000000 (State:Non-secure_
↳Mode:EL1/0 Width:64bits VMID:0)
coresight-cpu-debug 852000.debug: CPU[1]:
coresight-cpu-debug 852000.debug:  EDPRSR:  00000001 (Power:On DLK:Unlock)
coresight-cpu-debug 852000.debug:  EDPCSR:  debug_notifier_call+0x23c/0x358
coresight-cpu-debug 852000.debug:  EDCIDSR:  00000000
coresight-cpu-debug 852000.debug:  EDVIDSR:  90000000 (State:Non-secure_
↳Mode:EL1/0 Width:64bits VMID:0)
```

26.4 CoreSight Embedded Cross Trigger (CTI & CTM).

Author Mike Leach <mike.leach@linaro.org>

Date November 2019

26.4.1 Hardware Description

The CoreSight Cross Trigger Interface (CTI) is a hardware device that takes individual input and output hardware signals known as triggers to and from devices and interconnects them via the Cross Trigger Matrix (CTM) to other devices via numbered channels, in order to propagate events between devices.

e.g.:

```
00000000 in_trigs : :::::
0 C 0----->: : +=====>(other CTI channel IO)
0 P 0<-----: : v
0 U 0 out_trigs : : Channels ***** : :::::
00000000 : CTI :<=====>*CTM*<====>: CTI :---+
##### in_trigs : : (id 0-3) ***** : ::::: v
```

```
# ETM #----->:      :      ^      #####
#      #<-----:      :      +---# ETR #
##### out_trigs ::::      #####
```

The CTI driver enables the programming of the CTI to attach triggers to channels. When an input trigger becomes active, the attached channel will become active. Any output trigger attached to that channel will also become active. The active channel is propagated to other CTIs via the CTM, activating connected output triggers there, unless filtered by the CTI channel gate.

It is also possible to activate a channel using system software directly programming registers in the CTI.

The CTIs are registered by the system to be associated with CPUs and/or other CoreSight devices on the trace data path. When these devices are enabled the attached CTIs will also be enabled. By default/on power up the CTIs have no programmed trigger/channel attachments, so will not affect the system until explicitly programmed.

The hardware trigger connections between CTIs and devices is implementation defined, unless the CPU/ETM combination is a v8 architecture, in which case the connections have an architecturally defined standard layout.

The hardware trigger signals can also be connected to non-CoreSight devices (e.g. UART), or be propagated off chip as hardware IO lines.

All the CTI devices are associated with a CTM. On many systems there will be a single effective CTM (one CTM, or multiple CTMs all interconnected), but it is possible that systems can have nets of CTIs+CTM that are not interconnected by a CTM to each other. On these systems a CTM index is declared to associate CTI devices that are interconnected via a given CTM.

26.4.2 Sysfs files and directories

The CTI devices appear on the existing CoreSight bus alongside the other CoreSight devices:

```
>$ ls /sys/bus/coresight/devices
cti_cpu0  cti_cpu2  cti_sys0  etm0  etm2  funnel0  replicator0  tmc_etr0
cti_cpu1  cti_cpu3  cti_sys1  etm1  etm3  funnel1  tmc_etf0    tpiu0
```

The `cti_cpu<N>` named CTIs are associated with a CPU, and any ETM used by that core. The `cti_sys<N>` CTIs are general system infrastructure CTIs that can be associated with other CoreSight devices, or other system hardware capable of generating or using trigger signals.:

```
>$ ls /sys/bus/coresight/devices/etm0/cti_cpu0
channels  ctmid  enable  nr_trigger_cons  mgmt  power  powered  regs
connections  subsystem  triggers0  triggers1  uevent
```

Key file items are:-

- `enable`: enables/disables the CTI. Read to determine current state. If this shows as enabled (1), but `powered` shows unpowered (0), then the `enable` indicates a request to enabled when the device is powered.
- `ctmid`: associated CTM - only relevant if system has multiple CTI+CTM clusters that are not interconnected.

- `nr_trigger_cons` : total connections - triggers<N> directories.
- `powered` : Read to determine if the CTI is currently powered.

Sub-directories:-

- `triggers<N>`: contains list of triggers for an individual connection.
- `channels`: Contains the channel API - CTI main programming interface.
- `regs`: Gives access to the raw programmable CTI regs.
- `mgmt`: the standard CoreSight management registers.
- `connections`: Links to connected *CoreSight* devices. The number of links can be 0 to `nr_trigger_cons`. Actual number given by `nr_links` in this directory.

triggers<N> directories

Individual trigger connection information. This describes trigger signals for CoreSight and non-CoreSight connections.

Each triggers directory has a set of parameters describing the triggers for the connection.

- `name` : name of connection
- `in_signals` : input trigger signal indexes used in this connection.
- `in_types` : functional types for in signals.
- `out_signals` : output trigger signals for this connection.
- `out_types` : functional types for out signals.

e.g:

```
>$ ls ./cti_cpu0/triggers0/
in_signals in_types name out_signals out_types
>$ cat ./cti_cpu0/triggers0/name
cpu0
>$ cat ./cti_cpu0/triggers0/out_signals
0-2
>$ cat ./cti_cpu0/triggers0/out_types
pe_edbgreq pe_dbgrestart pe_ctiirq
>$ cat ./cti_cpu0/triggers0/in_signals
0-1
>$ cat ./cti_cpu0/triggers0/in_types
pe_dbgtrigger pe_pmuiirq
```

If a connection has zero signals in either the 'in' or 'out' triggers then those parameters will be omitted.

Channels API Directory

This provides an easy way to attach triggers to channels, without needing the multiple register operations that are required if manipulating the 'regs' sub-directory elements directly.

A number of files provide this API:

```
>$ ls ./cti_sys0/channels/  
chan_clear          chan_inuse          chan_xtrigs_out      trigin_attach  
chan_free           chan_pulse          chan_xtrigs_reset    trigin_detach  
chan_gate_disable   chan_set            chan_xtrigs_sel       trigout_attach  
chan_gate_enable    chan_xtrigs_in      trig_filter_enable    trigout_detach  
trigout_filtered
```

Most access to these elements take the form:

```
echo <chan> [<trigger>] > /<device_path>/<operation>
```

where the optional <trigger> is only needed for trigXX_attach | detach operations.

e.g.:

```
>$ echo 0 1 > ./cti_sys0/channels/trigout_attach  
>$ echo 0 > ./cti_sys0/channels/chan_set
```

Attaches trigout(1) to channel(0), then activates channel(0) generating a set state on cti_sys0.trigout(1)

API operations

- `trigin_attach`, `trigout_attach`: Attach a channel to a trigger signal.
- `trigin_detach`, `trigout_detach`: Detach a channel from a trigger signal.
- `chan_set`: Set the channel - the set state will be propagated around the CTM to other connected devices.
- `chan_clear`: Clear the channel.
- `chan_pulse`: Set the channel for a single CoreSight clock cycle.
- `chan_gate_enable`: Write operation sets the CTI gate to propagate (enable) the channel to other devices. This operation takes a channel number. CTI gate is enabled for all channels by default at power up. Read to list the currently enabled channels on the gate.
- `chan_gate_disable`: Write channel number to disable gate for that channel.
- `chan_inuse`: Show the current channels attached to any signal
- `chan_free`: Show channels with no attached signals.
- `chan_xtrigs_sel`: write a channel number to select a channel to view, read to show the selected channel number.
- `chan_xtrigs_in`: Read to show the input triggers attached to the selected view channel.
- `chan_xtrigs_out`: Read to show the output triggers attached to the selected view channel.
- `trig_filter_enable`: Defaults to enabled, disable to allow potentially dangerous output signals to be set.

- `trigout_filtered`: Trigger out signals that are prevented from being set if filtering `trig_filter_enable` is enabled. One use is to prevent accidental EDBGREQ signals stopping a core.
- `chan_xtrigs_reset`: Write 1 to clear all channel / trigger programming. Resets device hardware to default state.

The example below attaches input trigger index 1 to channel 2, and output trigger index 6 to the same channel. It then examines the state of the channel / trigger connections using the appropriate sysfs attributes.

The settings mean that if either input trigger 1, or channel 2 go active then trigger out 6 will go active. We then enable the CTI, and use the software channel control to activate channel 2. We see the active channel on the `choutstatus` register and the active signal on the `trigoutstatus` register. Finally clearing the channel removes this.

e.g.:

```
.../cti_sys0/channels# echo 2 1 > trigin_attach
.../cti_sys0/channels# echo 2 6 > trigout_attach
.../cti_sys0/channels# cat chan_free
0-1,3
.../cti_sys0/channels# cat chan_inuse
2
.../cti_sys0/channels# echo 2 > chan_xtrigs_sel
.../cti_sys0/channels# cat chan_xtrigs_trigin
1
.../cti_sys0/channels# cat chan_xtrigs_trigout
6
.../cti_sys0/# echo 1 > enable
.../cti_sys0/channels# echo 2 > chan_set
.../cti_sys0/channels# cat ../regs/choutstatus
0x4
.../cti_sys0/channels# cat ../regs/trigoutstatus
0x40
.../cti_sys0/channels# echo 2 > chan_clear
.../cti_sys0/channels# cat ../regs/trigoutstatus
0x0
.../cti_sys0/channels# cat ../regs/choutstatus
0x0
```

26.5 ETMv4 sysfs linux driver programming reference.

Author Mike Leach <mike.leach@linaro.org>

Date October 11th, 2019

Supplement to existing ETMv4 driver documentation.

26.5.1 Sysfs files and directories

Root: /sys/bus/coresight/devices/etm<N>

The following paragraphs explain the association between sysfs files and the ETMv4 registers that they effect. Note the register names are given without the ‘TRC’ prefix.

File mode (rw)

Trace Registers {CONFIGR + others}

Notes Bit select trace features. See ‘mode’ section below. Bits in this will cause equivalent programming of trace config and other registers to enable the features requested.

Syntax & eg echo bitfield > mode

bitfield up to 32 bits setting trace features.

Example \$> echo 0x012 > mode

File reset (wo)

Trace Registers All

Notes Reset all programming to trace nothing / no logic programmed.

Syntax echo 1 > reset

File enable_source (wo)

Trace Registers PRGCTLR, All hardware regs.

Notes

- > 0 : Programs up the hardware with the current values held in the driver and enables trace.
- = 0 : disable trace hardware.

Syntax echo 1 > enable_source

File cpu (ro)

Trace Registers None.

Notes CPU ID that this ETM is attached to.

Example \$> cat cpu

\$> 0

File addr_idx (rw)

Trace Registers None.

Notes Virtual register to index address comparator and range features. Set index for first of the pair in a range.

Syntax `echo idx > addr_idx`

Where `idx < nr_addr_cmp x 2`

File `addr_range (rw)`

Trace Registers `ACVR[idx, idx+1], VIIECTLR`

Notes Pair of addresses for a range selected by `addr_idx`. Include / exclude according to the optional parameter, or if omitted uses the current 'mode' setting. Select comparator range in control register. Error if index is odd value.

Depends `mode, addr_idx`

Syntax `echo addr1 addr2 [exclude] > addr_range`

Where `addr1` and `addr2` define the range and `addr1 < addr2`.

Optional exclude value:-

- 0 for include
- 1 for exclude.

Example `$> echo 0x0000 0x2000 0 > addr_range`

File `addr_single (rw)`

Trace Registers `ACVR[idx]`

Notes Set a single address comparator according to `addr_idx`. This is used if the address comparator is used as part of event generation logic etc.

Depends `addr_idx`

Syntax `echo addr1 > addr_single`

File `addr_start (rw)`

Trace Registers `ACVR[idx], VISSCTLR`

Notes Set a trace start address comparator according to `addr_idx`. Select comparator in control register.

Depends `addr_idx`

Syntax `echo addr1 > addr_start`

File `addr_stop (rw)`

Trace Registers `ACVR[idx], VISSCTLR`

Notes Set a trace stop address comparator according to `addr_idx`. Select comparator in control register.

Depends addr_idx

Syntax echo addr1 > addr_stop

File addr_context (rw)

Trace Registers ACATR[idx,{6:4}]

Notes Link context ID comparator to address comparator addr_idx

Depends addr_idx

Syntax echo ctxt_idx > addr_context

Where ctxt_idx is the index of the linked context id / vmid comparator.

File addr_ctxtype (rw)

Trace Registers ACATR[idx,{3:2}]

Notes Input value string. Set type for linked context ID comparator

Depends addr_idx

Syntax echo type > addr_ctxtype

Type one of {all, vmid, ctxid, none}

Example \$> echo ctxid > addr_ctxtype

File addr_exlevel_s_ns (rw)

Trace Registers ACATR[idx,{14:8}]

Notes Set the ELx secure and non-secure matching bits for the selected address comparator

Depends addr_idx

Syntax echo val > addr_exlevel_s_ns

val is a 7 bit value for exception levels to exclude. Input value shifted to correct bits in register.

Example \$> echo 0x4F > addr_exlevel_s_ns

File addr_instdatatype (rw)

Trace Registers ACATR[idx,{1:0}]

Notes Set the comparator address type for matching. Driver only supports setting instruction address type.

Depends addr_idx

File addr_cmp_view (ro)

Trace Registers ACVR[idx, idx+1], ACATR[idx], VIIECTLR

Notes Read the currently selected address comparator. If part of address range then display both addresses.

Depends addr_idx

Syntax cat addr_cmp_view

Example

```
$> cat addr_cmp_view
addr_cmp[0] range 0x0 0xffffffffffffffff include ctrl(0x4b00)
```

File nr_addr_cmp (ro)

Trace Registers From IDR4

Notes Number of address comparator pairs

File sshot_idx (rw)

Trace Registers None

Notes Select single shot register set.

File sshot_ctrl (rw)

Trace Registers SSCCR[idx]

Notes Access a single shot comparator control register.

Depends sshot_idx

Syntax echo val > sshot_ctrl

Writes val into the selected control register.

File sshot_status (ro)

Trace Registers SSCSR[idx]

Notes Read a single shot comparator status register

Depends sshot_idx

Syntax cat sshot_status

Read status.

Example \$> cat sshot_status

```
0x1
```

File sshot_pe_ctrl (rw)

Trace Registers SSPCICR[idx]

Notes Access a single shot PE comparator input control register.

Depends sshot_idx

Syntax echo val > sshot_pe_ctrl

Writes val into the selected control register.

File ns_exlevel_vinst (rw)

Trace Registers VICTLR{23:20}

Notes Program non-secure exception level filters. Set / clear NS exception filter bits. Setting '1' excludes trace from the exception level.

Syntax echo bitfield > ns_exlevel_viinst

Where bitfield contains bits to set clear for EL0 to EL2

Example %> echo 0x4 > ns_exlevel_viinst

Excludes EL2 NS trace.

File vinst_pe_cmp_start_stop (rw)

Trace Registers VIPCSSCTLR

Notes Access PE start stop comparator input control registers

File bb_ctrl (rw)

Trace Registers BBCTLR

Notes Define ranges that Branch Broadcast will operate in. Default (0x0) is all addresses.

Depends BB enabled.

File cyc_threshold (rw)

Trace Registers CCCTLR

Notes Set the threshold for which cycle counts will be emitted. Error if attempt to set below minimum defined in IDR3, masked to width of valid bits.

Depends CC enabled.

File syncfreq (rw)

Trace Registers SYNCPR

Notes Set trace synchronisation period. Power of 2 value, 0 (off) or 8-20. Driver defaults to 12 (every 4096 bytes).

File cntr_idx (rw)

Trace Registers none

Notes Select the counter to access

Syntax echo idx > cntr_idx

Where idx < nr_cntr

File cntr_ctrl (rw)

Trace Registers CNTCTLR[idx]

Notes Set counter control value.

Depends cntr_idx

Syntax echo val > cntr_ctrl

Where val is per ETMv4 spec.

File cntrldvr (rw)

Trace Registers CNTRLDVR[idx]

Notes Set counter reload value.

Depends cntr_idx

Syntax echo val > cntrldvr

Where val is per ETMv4 spec.

File nr_cntr (ro)

Trace Registers From IDR5

Notes Number of counters implemented.

File ctxid_idx (rw)

Trace Registers None

Notes Select the context ID comparator to access

Syntax echo idx > ctxid_idx

Where idx < numcidc

File ctxid_pid (rw)

Trace Registers CIDCVR[idx]

Notes Set the context ID comparator value

Depends ctxid_idx

File `ctxid_masks` (rw)

Trace Registers CIDCCTLR0, CIDCCTLR1, CIDCVR<0-7>

Notes Pair of values to set the byte masks for 1-8 context ID comparators. Automatically clears masked bytes to 0 in CID value registers.

Syntax `echo m3m2m1m0 [m7m6m5m4] > ctxid_masks`

32 bit values made up of mask bytes, where mN represents a byte mask value for Context ID comparator N.

Second value not required on systems that have fewer than 4 context ID comparators

File `numcidc` (ro)

Trace Registers From IDR4

Notes Number of Context ID comparators

File `vmid_idx` (rw)

Trace Registers None

Notes Select the VM ID comparator to access.

Syntax `echo idx > vmid_idx`

Where `idx < numvmidc`

File `vmid_val` (rw)

Trace Registers VMIDCVR[idx]

Notes Set the VM ID comparator value

Depends `vmid_idx`

File `vmid_masks` (rw)

Trace Registers VMIDCCTLR0, VMIDCCTLR1, VMIDCVR<0-7>

Notes Pair of values to set the byte masks for 1-8 VM ID comparators. Automatically clears masked bytes to 0 in VMID value registers.

Syntax `echo m3m2m1m0 [m7m6m5m4] > vmid_masks`

Where mN represents a byte mask value for VMID comparator N. Second value not required on systems that have fewer than 4 VMID comparators.

File `numvmidc` (ro)

Trace Registers From IDR4

Notes Number of VMID comparators

File `res_idx` (rw)

Trace Registers None.

Notes Select the resource selector control to access. Must be 2 or higher as selectors 0 and 1 are hardwired.

Syntax `echo idx > res_idx`

Where $2 \leq \text{idx} < \text{nr_resource} \times 2$

File `res_ctrl` (rw)

Trace Registers `RSCTLR[idx]`

Notes Set resource selector control value. Value per ETMv4 spec.

Depends `res_idx`

Syntax `echo val > res_cntr`

Where `val` is per ETMv4 spec.

File `nr_resource` (ro)

Trace Registers From `IDR4`

Notes Number of resource selector pairs

File `event` (rw)

Trace Registers `EVENTCTRL0R`

Notes Set up to 4 implemented event fields.

Syntax `echo ev3ev2ev1ev0 > event`

Where `evN` is an 8 bit event field. Up to 4 event fields make up the 32-bit input value. Number of valid fields is implementation dependent, defined in `IDR0`.

File `event_instren` (rw)

Trace Registers `EVENTCTRL1R`

Notes Choose events which insert event packets into trace stream.

Depends `EVENTCTRL0R`

Syntax `echo bitfield > event_instren`

Where `bitfield` is up to 4 bits according to number of event fields.

File `event_ts` (rw)

Trace Registers `TSCTLR`

Notes Set the event that will generate timestamp requests.

Depends TS activated

Syntax echo evfield > event_ts

Where evfield is an 8 bit event selector.

File seq_idx (rw)

Trace Registers None

Notes Sequencer event register select - 0 to 2

File seq_state (rw)

Trace Registers SEQSTR

Notes Sequencer current state - 0 to 3.

File seq_event (rw)

Trace Registers SEQEVRR[idx]

Notes State transition event registers

Depends seq_idx

Syntax echo evBevF > seq_event

Where evBevF is a 16 bit value made up of two event selectors,

- evB : back
 - evF : forwards.
-

File seq_reset_event (rw)

Trace Registers SEQRSTEVR

Notes Sequencer reset event

Syntax echo evfield > seq_reset_event

Where evfield is an 8 bit event selector.

File nrseqstate (ro)

Trace Registers From IDR5

Notes Number of sequencer states (0 or 4)

File nr_pe_cmp (ro)

Trace Registers From IDR4

Notes Number of PE comparator inputs

File nr_ext_inp (ro)

Trace Registers From IDR5

Notes Number of external inputs

File nr_ss_cmp (ro)

Trace Registers From IDR4

Notes Number of Single Shot control registers

Note: When programming any address comparator the driver will tag the comparator with a type used - i.e. RANGE, SINGLE, START, STOP. Once this tag is set, then only the values can be changed using the same sysfs file / type used to program it.

Thus:

```
% echo 0 > addr_idx          ; select address comparator 0
% echo 0x1000 0x5000 0 > addr_range ; set address range on comparators 0, 1.
% echo 0x2000 > addr_start      ; error as comparator 0 is a range comparator
% echo 2 > addr_idx             ; select address comparator 2
% echo 0x2000 > addr_start      ; this is OK as comparator 2 is unused.
% echo 0x3000 > addr_stop       ; error as comparator 2 set as start address.
% echo 2 > addr_idx             ; select address comparator 3
% echo 0x3000 > addr_stop       ; this is OK
```

To remove programming on all the comparators (and all the other hardware) use the reset parameter:

```
% echo 1 > reset
```

26.5.2 The 'mode' sysfs parameter.

This is a bitfield selection parameter that sets the overall trace mode for the ETM. The table below describes the bits, using the defines from the driver source file, along with a description of the feature these represent. Many features are optional and therefore dependent on implementation in the hardware.

Bit assignments shown below:-

bit (0): ETM_MODE_EXCLUDE

description: This is the default value for the include / exclude function when setting address ranges. Set 1 for exclude range. When the mode parameter is set this value is applied to the currently indexed address range.

bit (4): ETM_MODE_BB

description: Set to enable branch broadcast if supported in hardware [IDR0].

bit (5): ETMv4_MODE_CYCACC

description: Set to enable cycle accurate trace if supported [IDR0].

bit (6): ETMv4_MODE_CTXID

description: Set to enable context ID tracing if supported in hardware [IDR2].

bit (7): ETM_MODE_VMID

description: Set to enable virtual machine ID tracing if supported [IDR2].

bit (11): ETMv4_MODE_TIMESTAMP

description: Set to enable timestamp generation if supported [IDR0].

bit (12): ETM_MODE_RETURNSTACK

description: Set to enable trace return stack use if supported [IDR0].

bit (13-14): ETM_MODE_QELEM(val)

description: 'val' determines level of Q element support enabled if implemented by the ETM [IDR0]

bit (19): ETM_MODE_ATB_TRIGGER

description: Set to enable the ATBTRIGGER bit in the event control register [EVENTCTLR1] if supported [IDR5].

bit (20): ETM_MODE_LPOVERRIDE

description: Set to enable the LPOVERRIDE bit in the event control register [EVENTCTLR1], if supported [IDR5].

bit (21): ETM_MODE_ISTALL_EN

description: Set to enable the ISTALL bit in the stall control register [STALLCTLR]

bit (23): ETM_MODE_INSTPRIO

description: Set to enable the INSTPRIORITY bit in the stall control register [STALLCTLR] , if supported [IDR0].

bit (24): ETM_MODE_NOOVERFLOW

description: Set to enable the NOOVERFLOW bit in the stall control register [STALLCTLR], if supported [IDR3].

bit (25): ETM_MODE_TRACE_RESET

description: Set to enable the TRCRESET bit in the viewinst control register [VICTLR] , if supported [IDR3].

bit (26): ETM_MODE_TRACE_ERR

description: Set to enable the TRCCTRL bit in the viewinst control register [VICTLR].

bit (27): ETM_MODE_VIEWINST_STARTSTOP

description: Set the initial state value of the ViewInst start / stop logic in the viewinst control register [VICTLR]

bit (30): ETM_MODE_EXCL_KERN

description: Set default trace setup to exclude kernel mode trace (see note a)

bit (31): ETM_MODE_EXCL_USER

description: Set default trace setup to exclude user space trace (see note a)

Note a) On startup the ETM is programmed to trace the complete address space using address range comparator 0. 'mode' bits 30 / 31 modify this setting to set EL exclude bits for NS state in either user space (EL0) or kernel space (EL1) in the address range comparator. (the default setting excludes all secure EL, and NS EL2)

Once the reset parameter has been used, and/or custom programming has been implemented - using these bits will result in the EL bits for address comparator 0 being set in the same way.

Note b) Bits 2-3, 8-10, 15-16, 18, 22, control features that only work with data trace. As A-profile data trace is architecturally prohibited in ETMv4, these have been omitted here. Possible uses could be where a kernel has support for control of R or M profile infrastructure as part of a heterogeneous system.

Bits 17, 28-29 are unused.

26.6 Trace Buffer Extension (TRBE).

Author Anshuman Khandual <anshuman.khandual@arm.com>

Date November 2020

26.6.1 Hardware Description

Trace Buffer Extension (TRBE) is a percpu hardware which captures in system memory, CPU traces generated from a corresponding percpu tracing unit. This gets plugged in as a coresight sink device because the corresponding trace generators (ETE), are plugged in as source device.

The TRBE is not compliant to CoreSight architecture specifications, but is driven via the CoreSight driver framework to support the ETE (which is CoreSight compliant) integration.

26.6.2 Sysfs files and directories

The TRBE devices appear on the existing coresight bus alongside the other coresight devices:

```
>$ ls /sys/bus/coresight/devices
trbe0 trbe1 trbe2 trbe3
```

The trbe<N> named TRBEs are associated with a CPU.:

```
>$ ls /sys/bus/coresight/devices/trbe0/
align flag
```

Key file items are:-

- align: TRBE write pointer alignment
- flag: TRBE updates memory with access and dirty flags

USER_EVENTS: USER-BASED EVENT TRACING

Author Beau Belgrave

27.1 Overview

User based trace events allow user processes to create events and trace data that can be viewed via existing tools, such as ftrace and perf. To enable this feature, build your kernel with CONFIG_USER_EVENTS=y.

Programs can view status of the events via /sys/kernel/debug/tracing/user_events_status and can both register and write data out via /sys/kernel/debug/tracing/user_events_data.

Programs can also use /sys/kernel/debug/tracing/dynamic_events to register and delete user based events via the u: prefix. The format of the command to dynamic_events is the same as the ioctl with the u: prefix applied.

Typically programs will register a set of events that they wish to expose to tools that can read trace_events (such as ftrace and perf). The registration process gives back two ints to the program for each event. The first int is the status index. This index describes which byte in the /sys/kernel/debug/tracing/user_events_status file represents this event. The second int is the write index. This index describes the data when a write() or writev() is called on the /sys/kernel/debug/tracing/user_events_data file.

The structures referenced in this document are contained with the /include/uapi/linux/user_events.h file in the source tree.

NOTE: Both user_events_status and user_events_data are under the tracefs filesystem and may be mounted at different paths than above.

27.2 Registering

Registering within a user process is done via ioctl() out to the /sys/kernel/debug/tracing/user_events_data file. The command to issue is DIAG_IOCSREG.

This command takes a struct user_reg as an argument:

```
struct user_reg {
    u32 size;
    u64 name_args;
    u32 status_index;
```

```
    u32 write_index;
};
```

The struct `user_reg` requires two inputs, the first is the size of the structure to ensure forward and backward compatibility. The second is the command string to issue for registering. Upon success two outputs are set, the status index and the write index.

User based events show up under `tracefs` like any other event under the subsystem named “`user_events`”. This means tools that wish to attach to the events need to use `/sys/kernel/debug/tracing/events/user_events/[name]/enable` or `perf record -e user_events:[name]` when attaching/recording.

NOTE: *The `write_index` returned is only valid for the FD that was used*

27.2.1 Command Format

The command string format is as follows:

```
name[:FLAG1[,FLAG2...]] [Field1[;Field2...]]
```

27.2.2 Supported Flags

None yet

27.2.3 Field Format

```
type name [size]
```

Basic types are supported (`__data_loc`, `u32`, `u64`, `int`, `char`, `char[20]`, etc). User programs are encouraged to use clearly sized types like `u32`.

NOTE: *Long is not supported since size can vary between user and kernel.*

The size is only valid for types that start with a struct prefix. This allows user programs to describe custom structs out to tools, if required.

For example, a struct in C that looks like this:

```
struct mytype {
    char data[20];
};
```

Would be represented by the following field:

```
struct mytype myname 20
```

27.3 Deleting

Deleting an event from within a user process is done via `ioctl()` out to the `/sys/kernel/debug/tracing/user_events_data` file. The command to issue is `DIAG_IOCSDDEL`.

This command only requires a single string specifying the event to delete by its name. Delete will only succeed if there are no references left to the event (in both user and kernel space). User programs should use a separate file to request deletes than the one used for registration due to this.

27.4 Status

When tools attach/record user based events the status of the event is updated in realtime. This allows user programs to only incur the cost of the `write()` or `writew()` calls when something is actively attached to the event.

User programs call `mmap()` on `/sys/kernel/debug/tracing/user_events_status` to check the status for each event that is registered. The byte to check in the file is given back after the register `ioctl()` via `user_reg.status_index`. Currently the size of `user_events_status` is a single page, however, custom kernel configurations can change this size to allow more user based events. In all cases the size of the file is a multiple of a page size.

For example, if the register `ioctl()` gives back a `status_index` of 3 you would check byte 3 of the returned `mmap` data to see if anything is attached to that event.

Administrators can easily check the status of all registered events by reading the `user_events_status` file directly via a terminal. The output is as follows:

```
Byte:Name [# Comments]
...

Active: ActiveCount
Busy: BusyCount
Max: MaxCount
```

For example, on a system that has a single event the output looks like this:

```
1:test

Active: 1
Busy: 0
Max: 4096
```

If a user enables the user event via `ftrace`, the output would change to this:

```
1:test # Used by ftrace

Active: 1
Busy: 1
Max: 4096
```

NOTE: A status index of 0 will never be returned. This allows user programs to have an index that can be used on error cases.

27.4.1 Status Bits

The byte being checked will be non-zero if anything is attached. Programs can check specific bits in the byte to see what mechanism has been attached.

The following values are defined to aid in checking what has been attached:

EVENT_STATUS_FTRACE - Bit set if ftrace has been attached (Bit 0).

EVENT_STATUS_PERF - Bit set if perf has been attached (Bit 1).

27.5 Writing Data

After registering an event the same fd that was used to register can be used to write an entry for that event. The write_index returned must be at the start of the data, then the remaining data is treated as the payload of the event.

For example, if write_index returned was 1 and I wanted to write out an int payload of the event. Then the data would have to be 8 bytes (2 ints) in size, with the first 4 bytes being equal to 1 and the last 4 bytes being equal to the value I want as the payload.

In memory this would look like this:

```
int index;  
int payload;
```

User programs might have well known structs that they wish to use to emit out as payloads. In those cases writev() can be used, with the first vector being the index and the following vector(s) being the actual event payload.

For example, if I have a struct like this:

```
struct payload {  
    int src;  
    int dst;  
    int flags;  
};
```

It's advised for user programs to do the following:

```
struct iovec io[2];  
struct payload e;  
  
io[0].iov_base = &write_index;  
io[0].iov_len = sizeof(write_index);  
io[1].iov_base = &e;  
io[1].iov_len = sizeof(e);  
  
writev(fd, (const struct iovec*)io, 2);
```

NOTE: *The `write_index` is not emitted out into the trace being recorded.*

27.6 Example Code

See sample code in `samples/user_events`.

C

coresight_register (*C function*), [303](#)
coresight_unregister (*C function*), [303](#)

D

disable_fprobe (*C function*), [90](#)

E

enable_fprobe (*C function*), [90](#)

F

fprobe (*C struct*), [89](#)

R

register_fprobe (*C function*), [90](#)
register_fprobe_ips (*C function*), [91](#)
register_fprobe_syms (*C function*), [91](#)

U

unregister_fprobe (*C function*), [91](#)